

# Review: Image Filtering

# Image Filtering

- Point Operation
- Neighborhood Operation -> Filtering-> **Correlation & Convolution**

$$(f * I)(x, y) = \sum_{i,j=-\infty}^{\infty} f(i, j)I(x + i, y + j) \quad (f * I)(x, y) = \sum_{i,j=-\infty}^{\infty} f(i, j)I(x - i, y - j)$$

# Image Filtering

- Point Operation
- Neighborhood Operation -> Filtering-> **Correlation & Convolution**

$$(f * I)(x, y) = \sum_{i,j=-\infty}^{\infty} f(i, j)I(x + i, y + j)$$

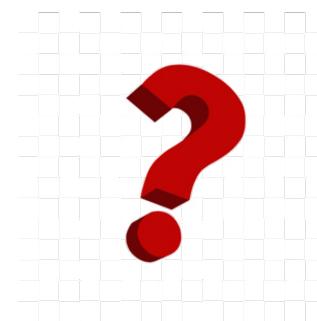
$$(f * I)(x, y) = \sum_{i,j=-\infty}^{\infty} f(i, j)I(x - i, y - j)$$

0	0	0	0	0
0	0	0	0	0
0	0	1	0	0
0	0	0	0	0
0	0	0	0	0

Image

a	b	c
d	e	f
g	h	i

Kernel  
[filter]



# Image Filtering

- Neighborhood Operation -> Filtering-> **Correlation** & **Convolution**

$$(f * I)(x, y) = \sum_{i,j=-\infty}^{\infty} f(i, j)I(x + i, y + j)$$

$$(f * I)(x, y) = \sum_{i,j=-\infty}^{\infty} f(i, j)I(x - i, y - j)$$

0	0	0	0	0
0	0	0	0	0
0	0	1	0	0
0	0	0	0	0
0	0	0	0	0

Image

a	b	c
d	e	f
g	h	i

Kernel  
[filter]

i	h	g
f	e	d
c	b	a

Correlation

a	b	c
d	e	f
g	h	i

Convolution

# Image Filtering

- Neighborhood Operation -> Filtering-> **Correlation** & **Convolution**
- Separable filters [Cost & Property Understanding]

Gaussian  $\frac{1}{16}$

1	2	1
2	4	2
1	2	1

Sobel

1	0	-1
2	0	-2
1	0	-1

1	2	1
0	0	0
-1	-2	-1

Box [Average]

$\frac{1}{9}$

1	1	1
1	1	1
1	1	1

Prewitt

1	0	-1
1	0	-1
1	0	-1

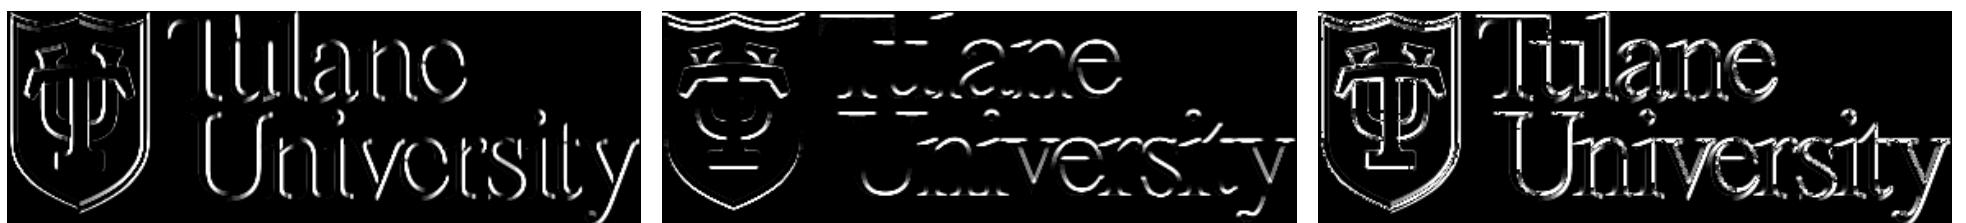
1	1	1
0	0	0
-1	-1	-1



Sobel



Prewitt



OpenCV has a built-in function **cv2.filter2D()** to convolve a kernel with an image.

<https://theailearner.com/tag/cv2-filter2d/>

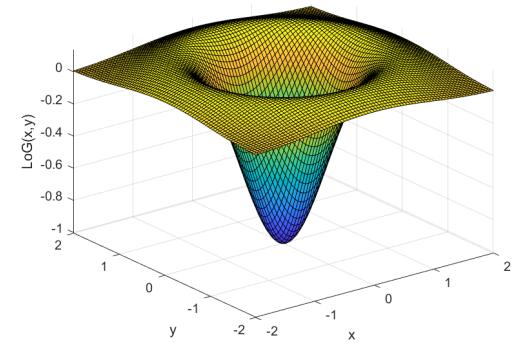


Laplacian Filter [2<sup>nd</sup>-order Derivative]

0	1	0
1	-4	1
0	1	0

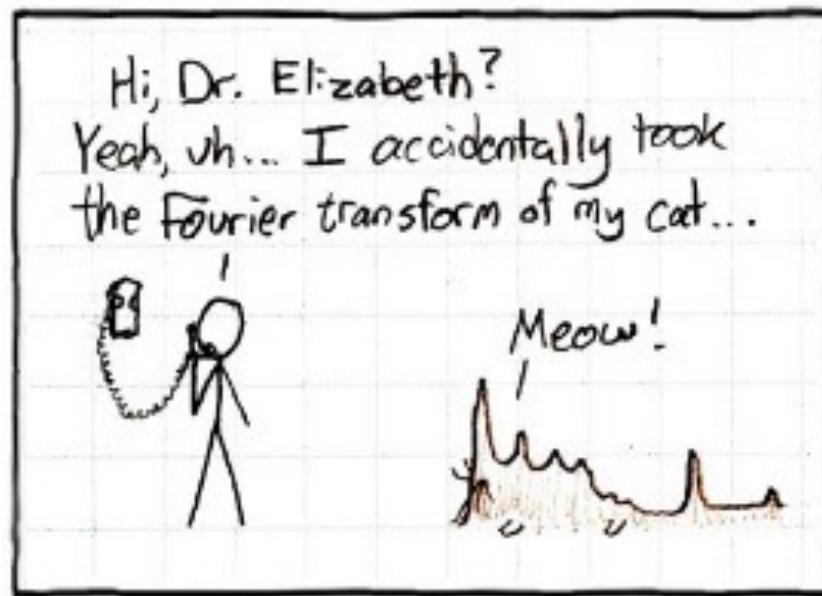


1	1	1
1	-8	1
1	1	1

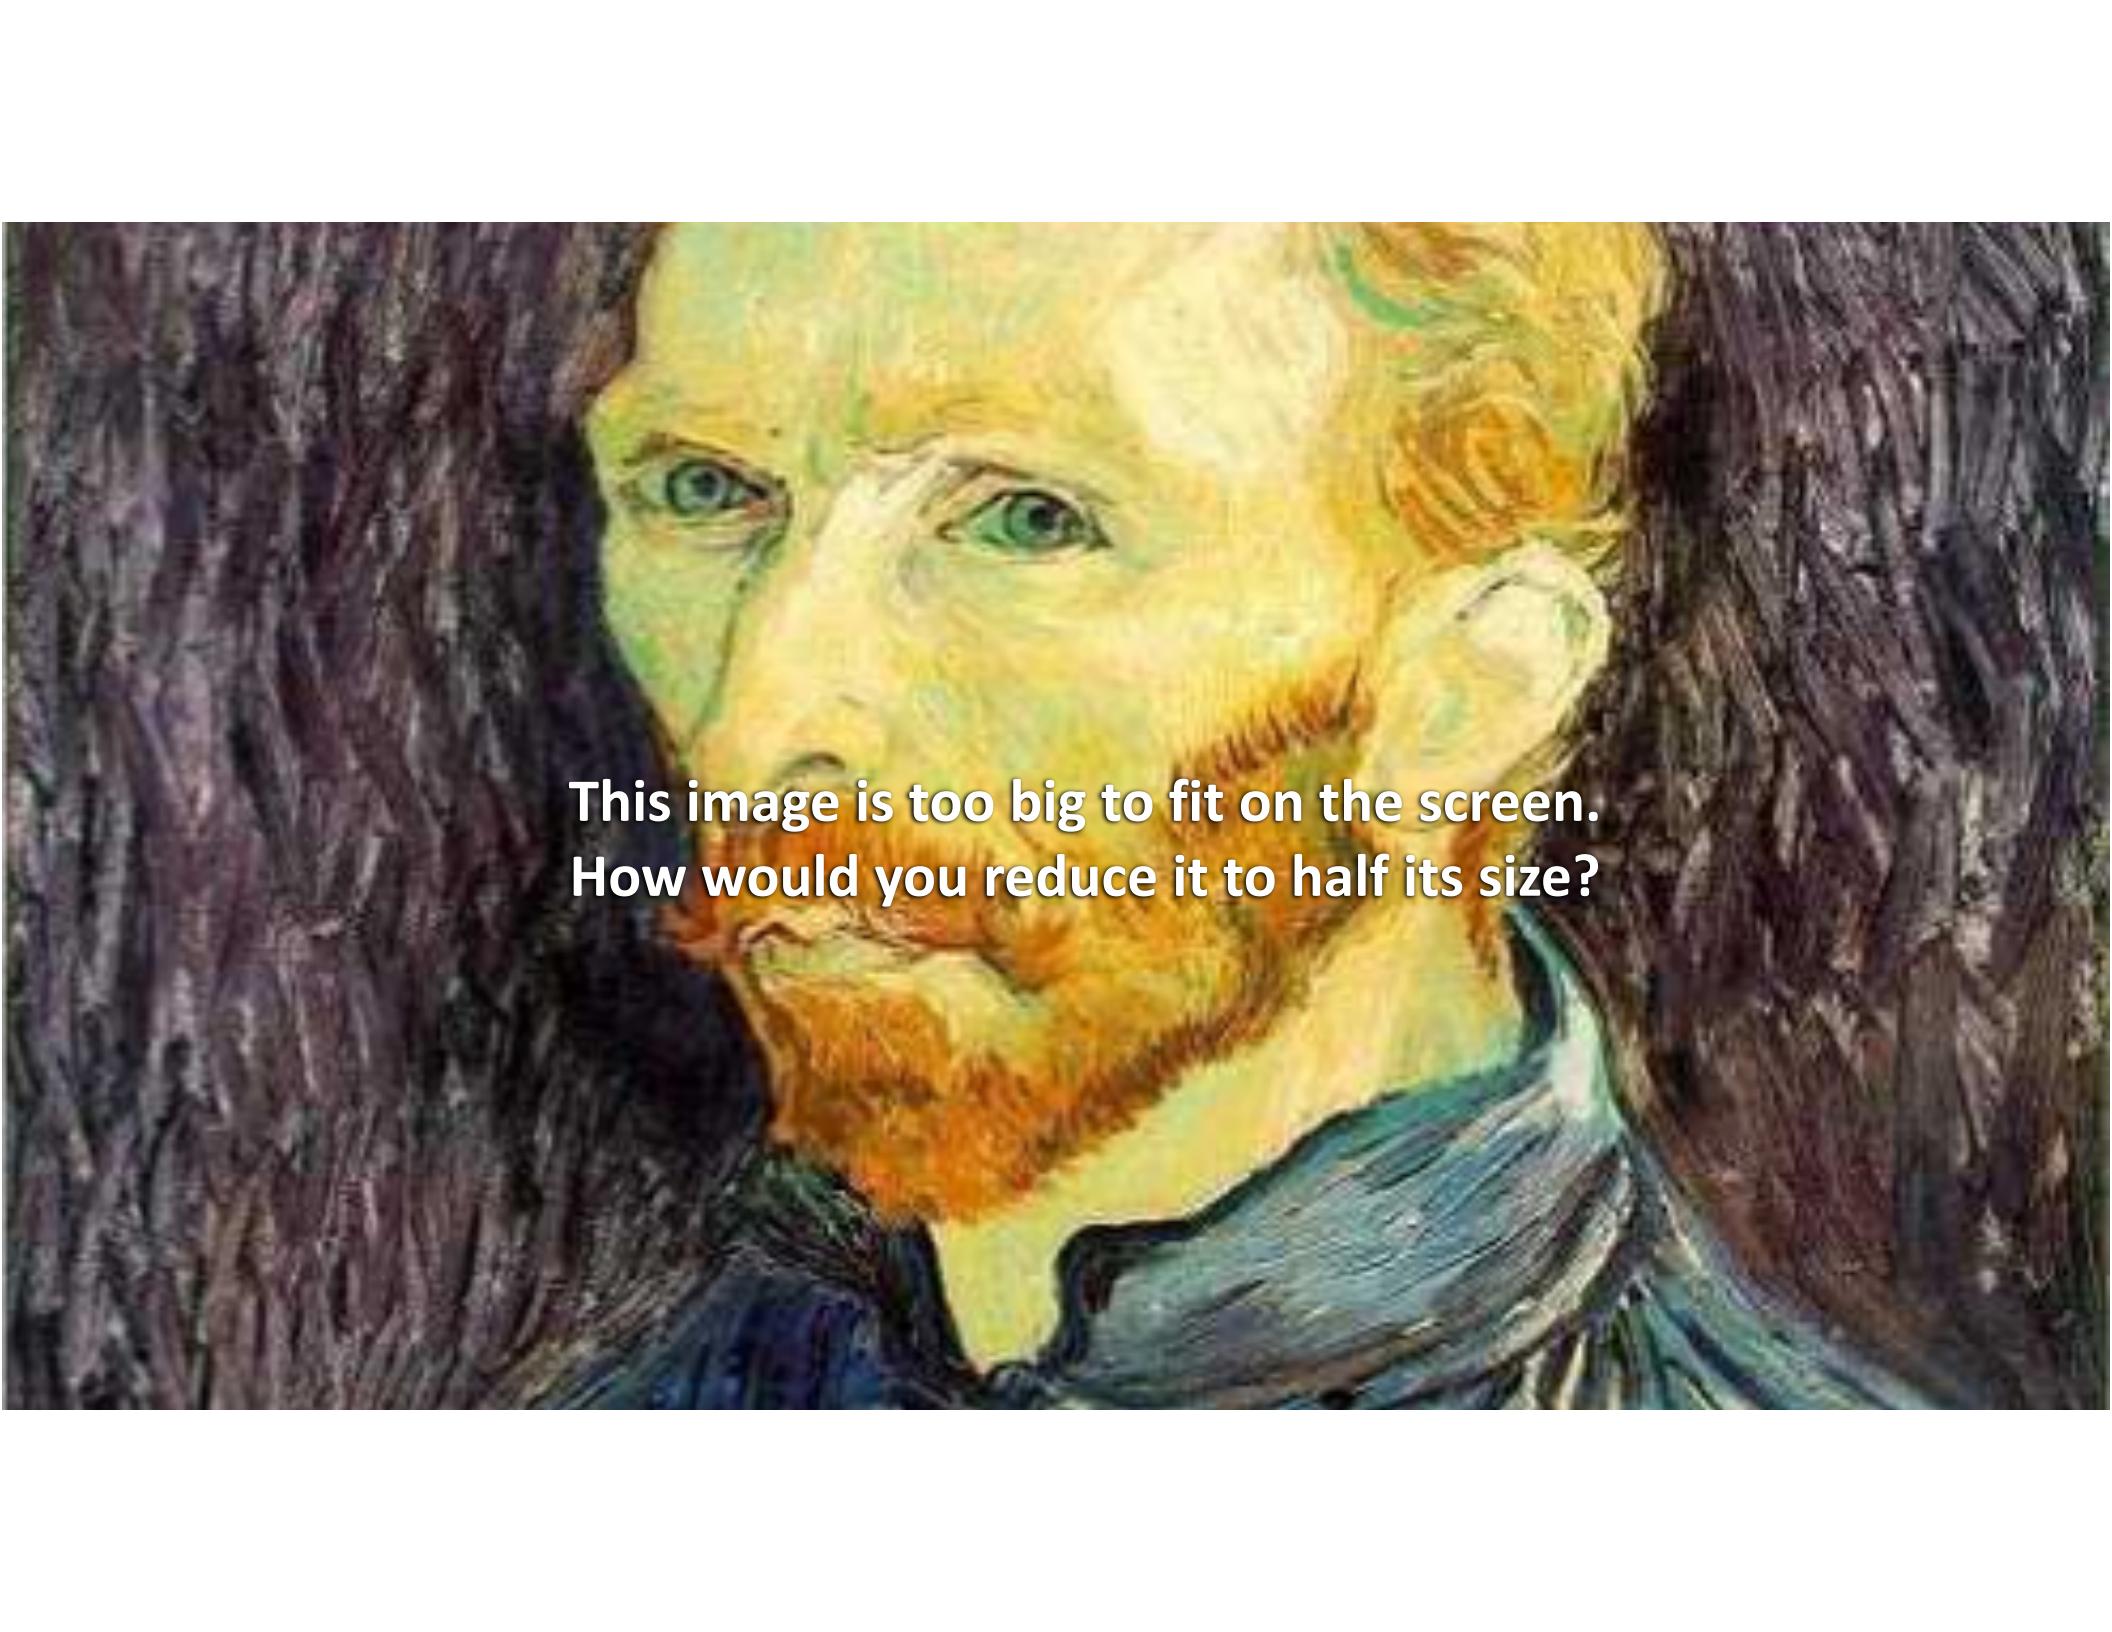


<https://homepages.inf.ed.ac.uk/rbf/HIPR2/log.htm>

# Image pyramids and frequency domain

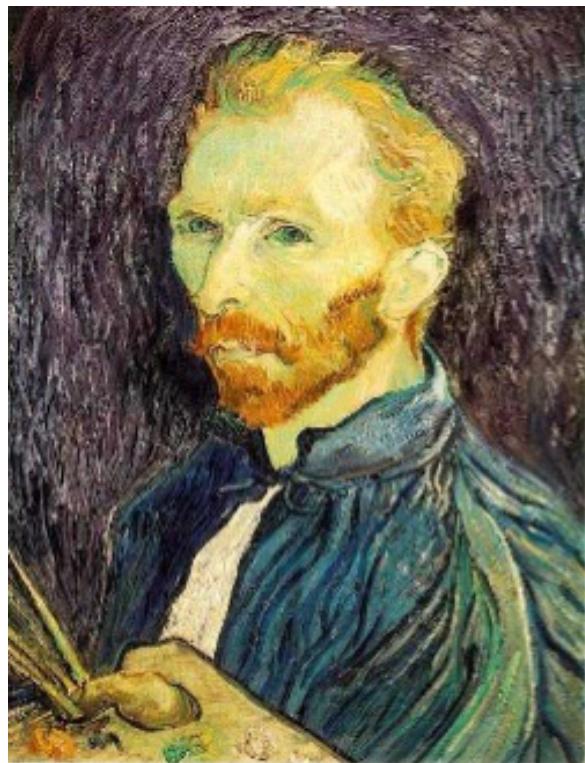


# Image downsampling



This image is too big to fit on the screen.  
How would you reduce it to half its size?

# Naïve image downsampling



1/2

Throw away half the rows and columns

delete even rows  
delete even columns



1/4

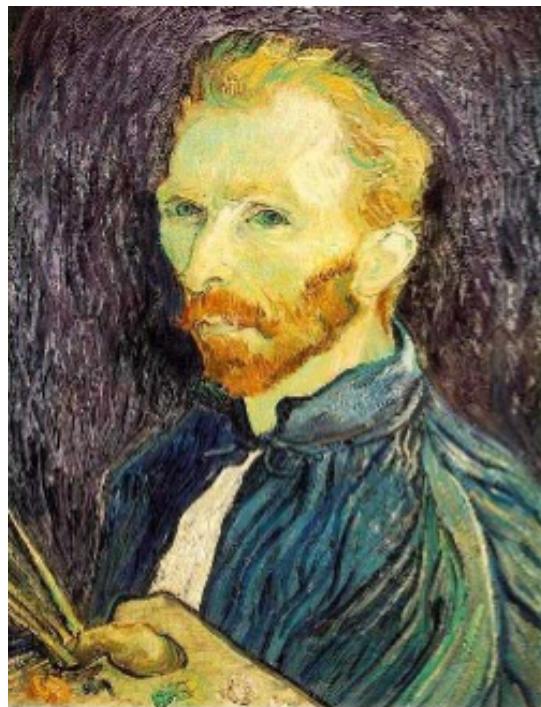
delete even rows  
delete even columns



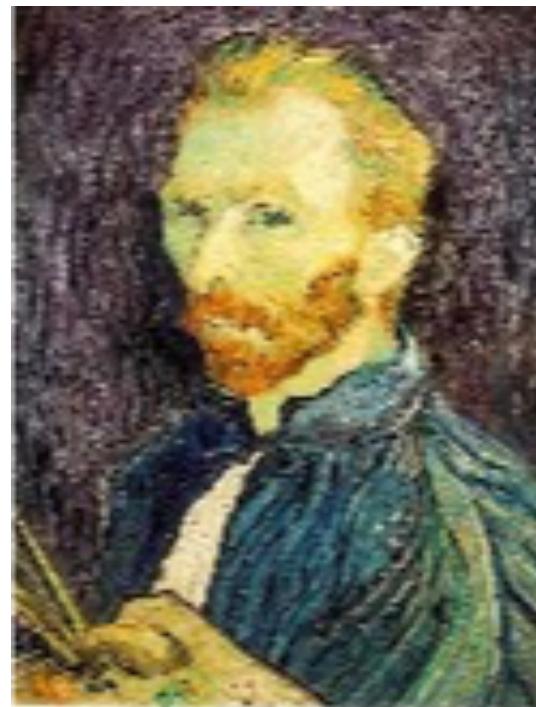
1/8

What is the problem with this approach?

# Naïve image downsampling



1/2



1/4 (2x zoom)



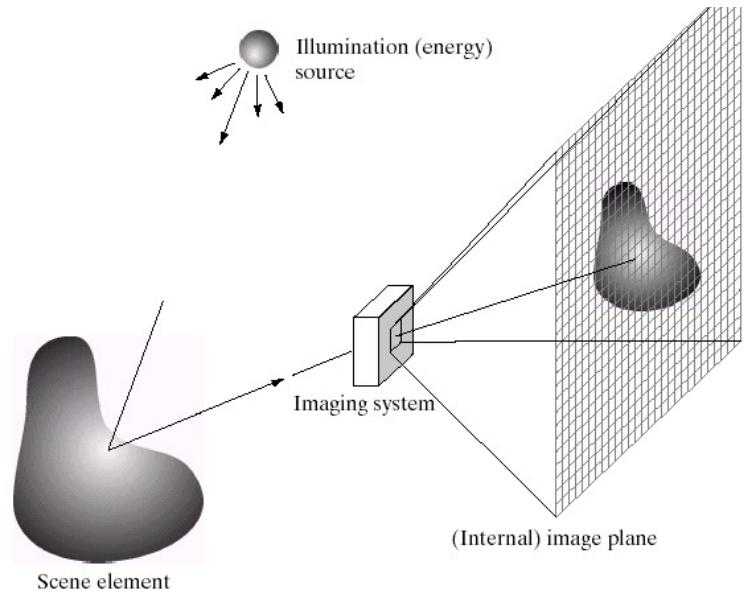
1/8 (4x zoom)

Why is the 1/8 image so pixelated (and do you know what this effect is called)?

# Aliasing

Aliasing occurs when an oscilloscope does not sample the signal fast enough to construct an accurate waveform record. The signal frequency is misidentified, and the waveforms displayed on an oscilloscope become indistinguishable. Aliasing is basically a form of undersampling.

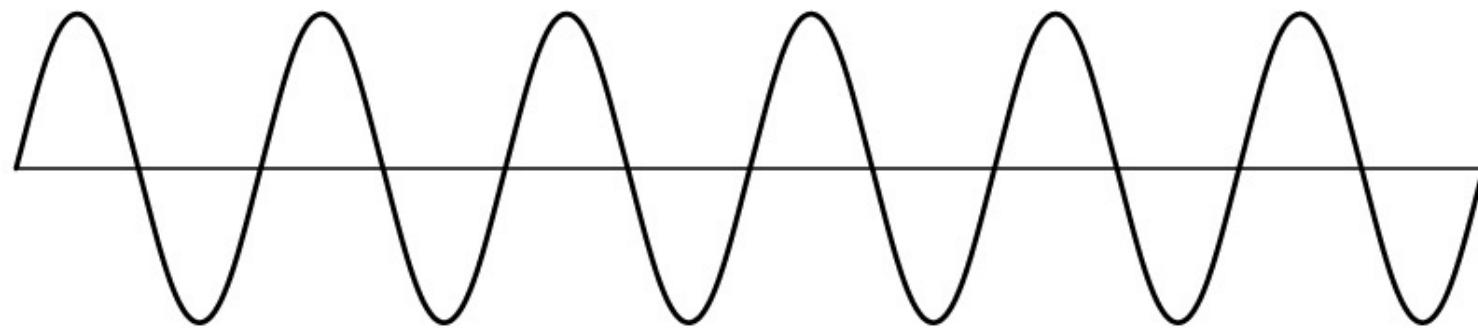
# Reminder



Images are a *discrete*, or *sampled*, representation of a *continuous* world

# Sampling

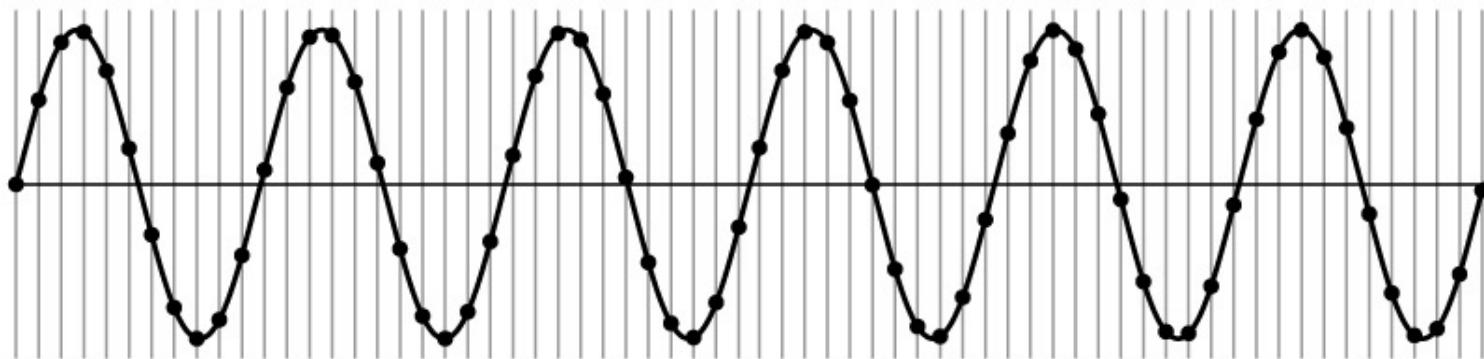
Very simple example: a sine wave



How would you discretize this signal?

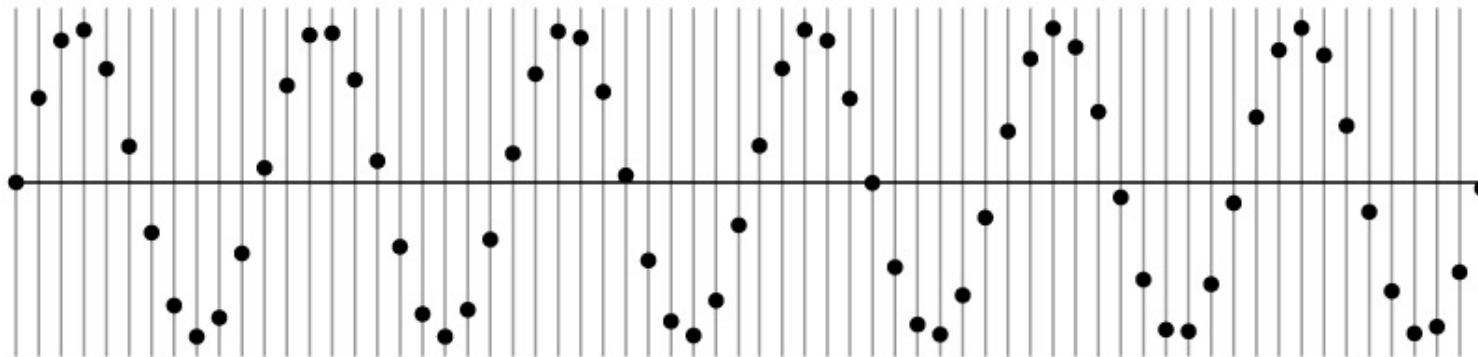
# Sampling

Very simple example: a sine wave



# Sampling

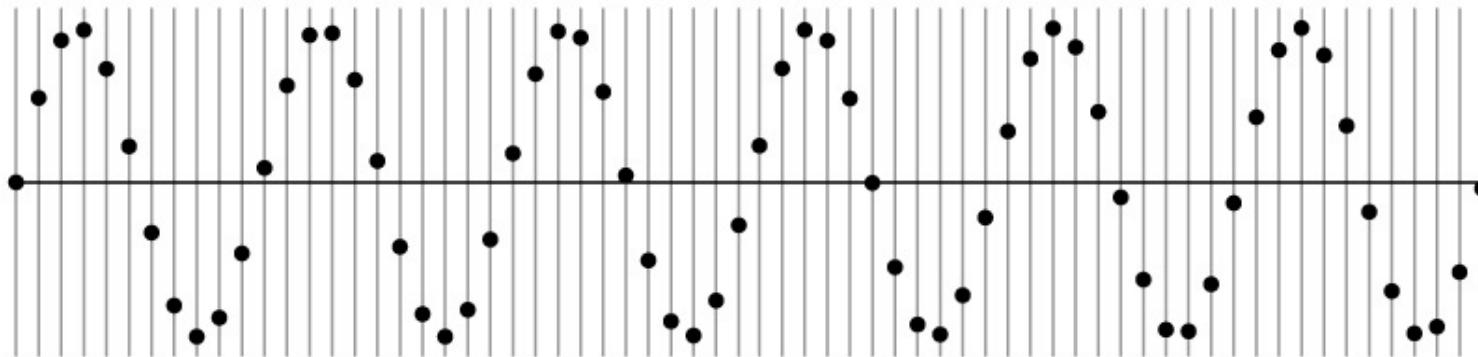
Very simple example: a sine wave



How many samples should I take?  
Can I take as *many* samples as I want?

# Sampling

Very simple example: a sine wave

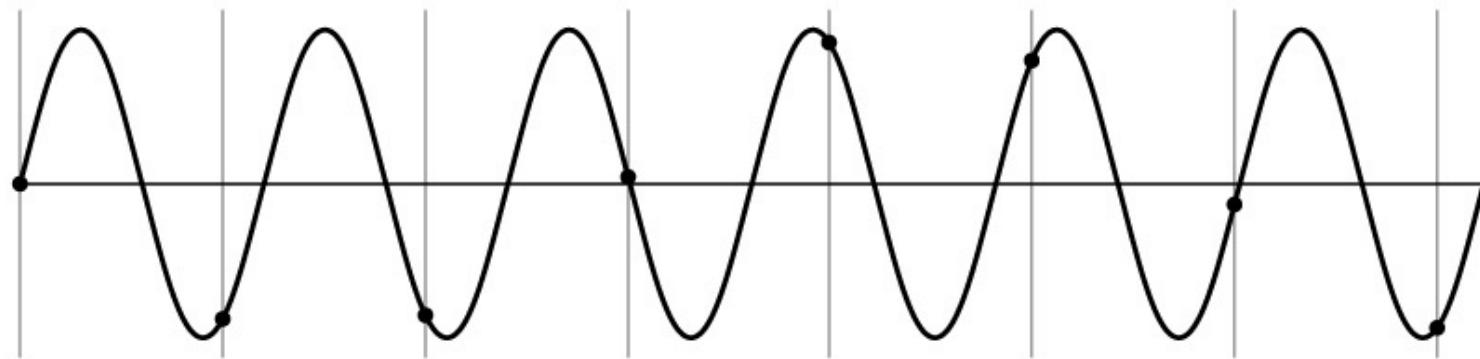


How many samples should I take?

Can I take as *few* samples as I want?

# Undersampling

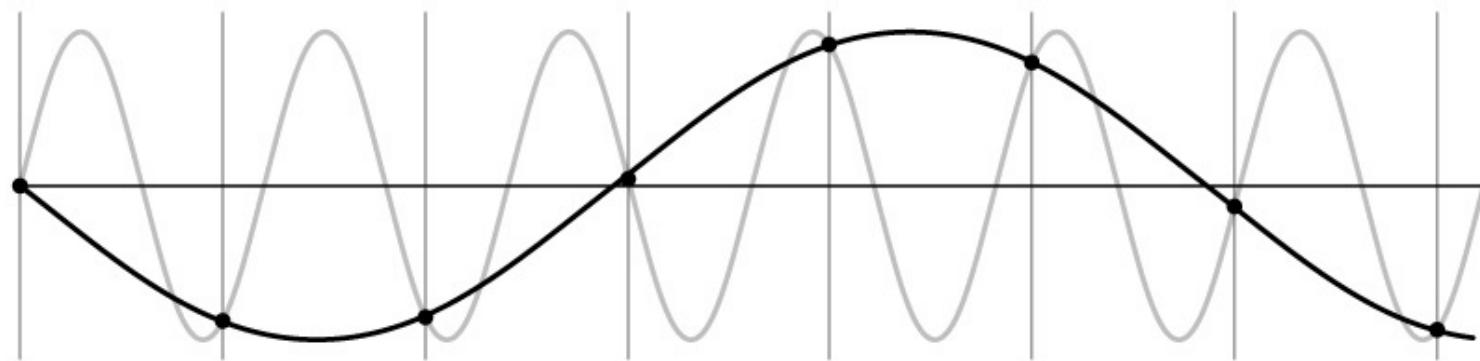
Very simple example: a sine wave



Unsurprising effect: information is lost.

# Undersampling

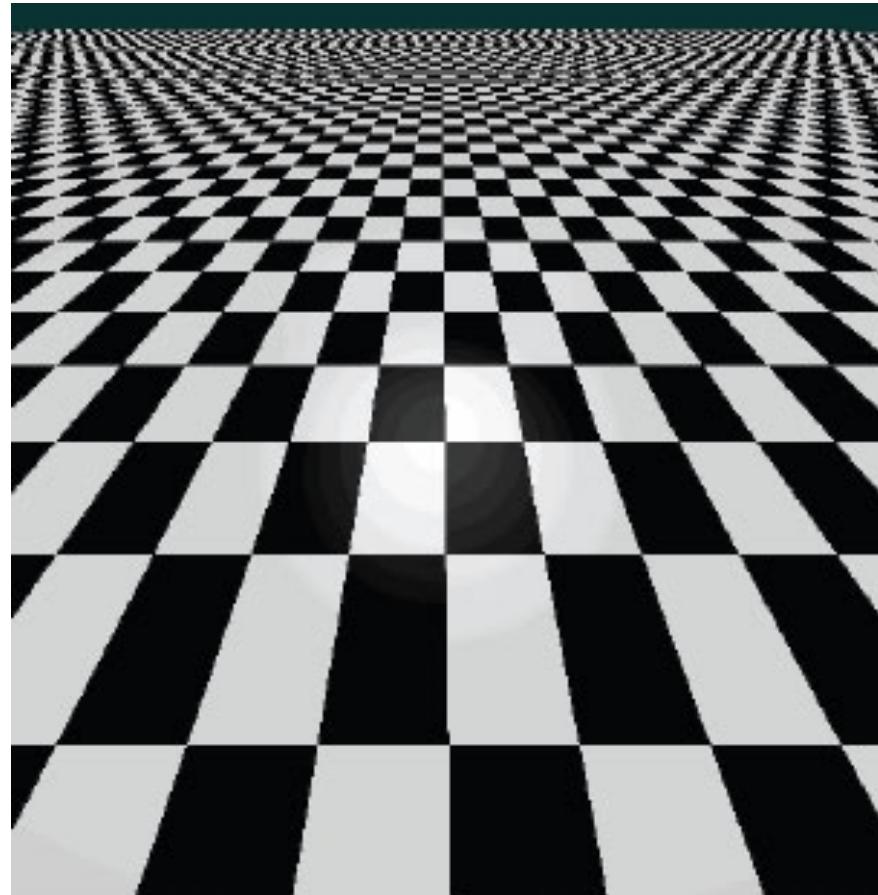
Very simple example: a sine wave



**Unsurprising effect:** information is lost.

**Surprising effect:** can confuse the signal with one of *lower* frequency.

# Aliasing in textures



# Aliasing in photographs

This is also known as “**moire**” pattern.



# Anti-aliasing

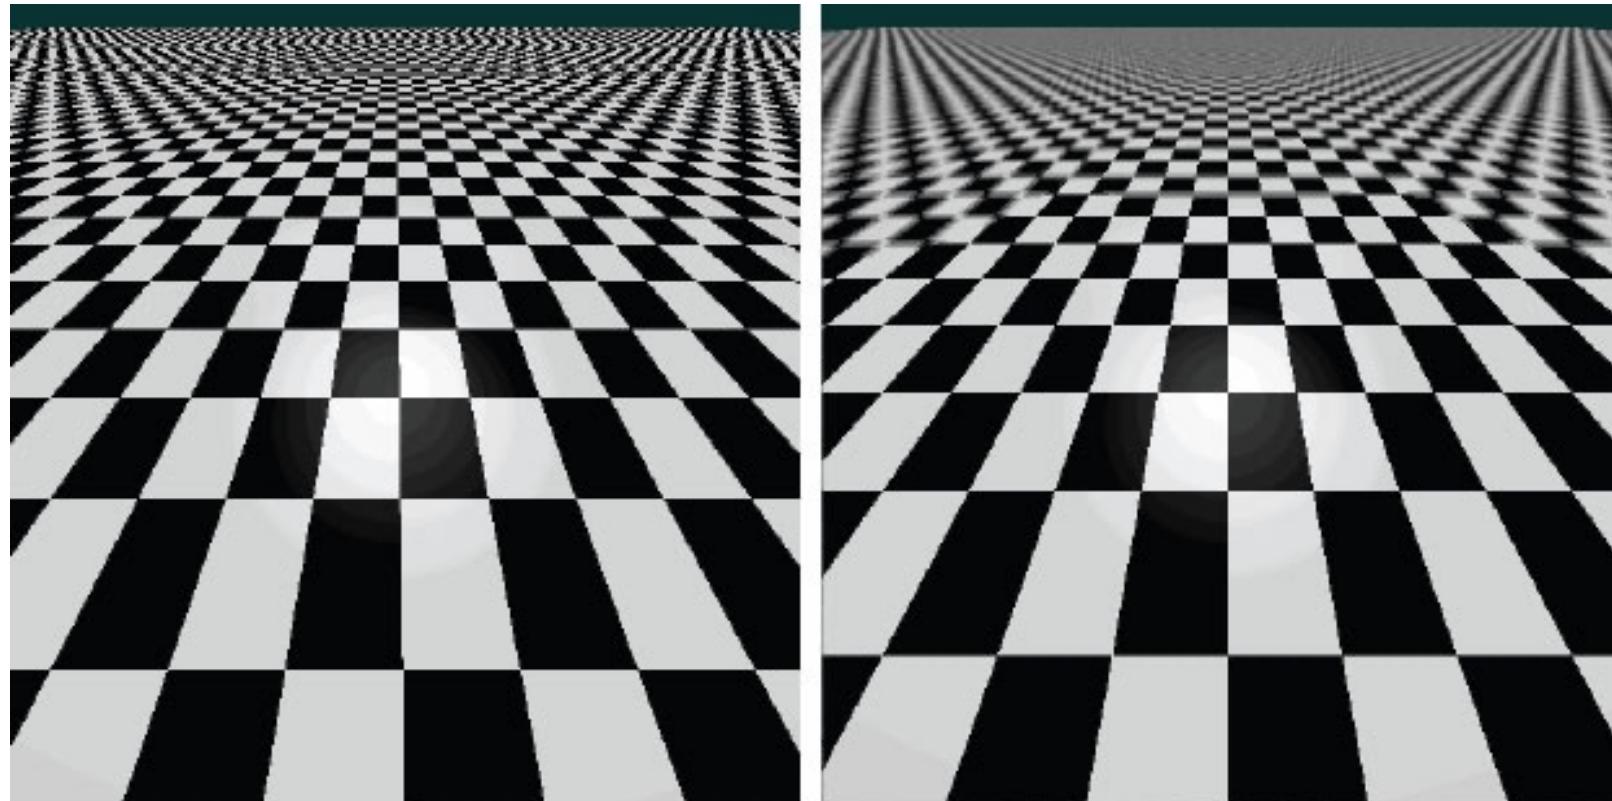
How would you deal with aliasing?

# Anti-aliasing

How would you deal with aliasing?

**Approach 1:** Oversample the signal

# Anti-aliasing in textures



aliasing artifacts

anti-aliasing by oversampling

# Anti-aliasing

How would you deal with aliasing?

**Approach 1:** Oversample the signal

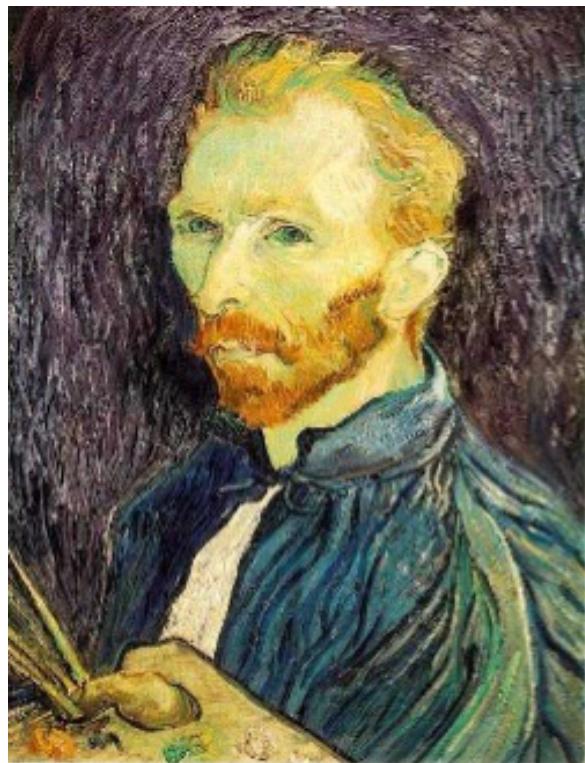
**Approach 2:** Smooth the signal

- Remove some of the detail effects that cause aliasing.
- Lose information, but better than aliasing artifacts.

How would you smooth a signal?

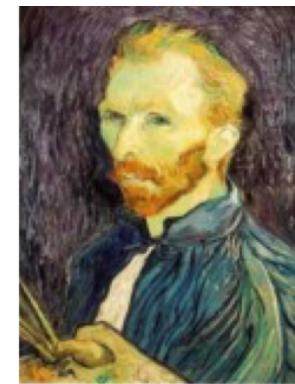
# Better image downsampling

Apply a smoothing filter first, then throw away half the rows and columns



1/2

Gaussian filter  
delete even rows  
delete even columns



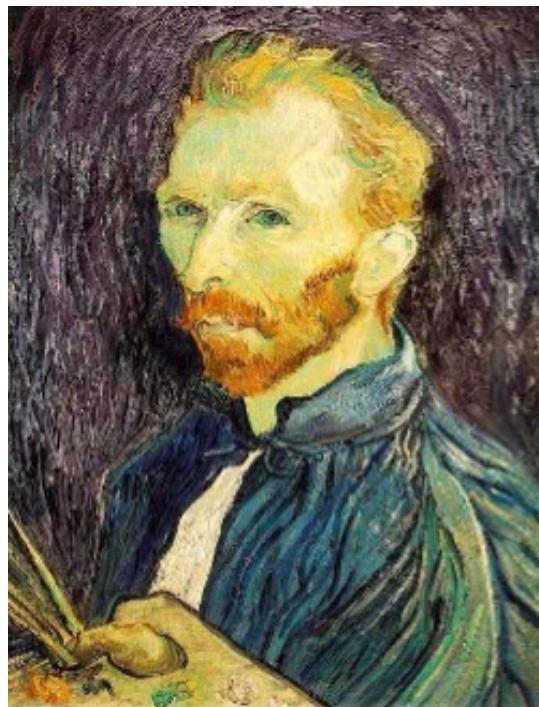
1/4

Gaussian filter  
delete even rows  
delete even columns

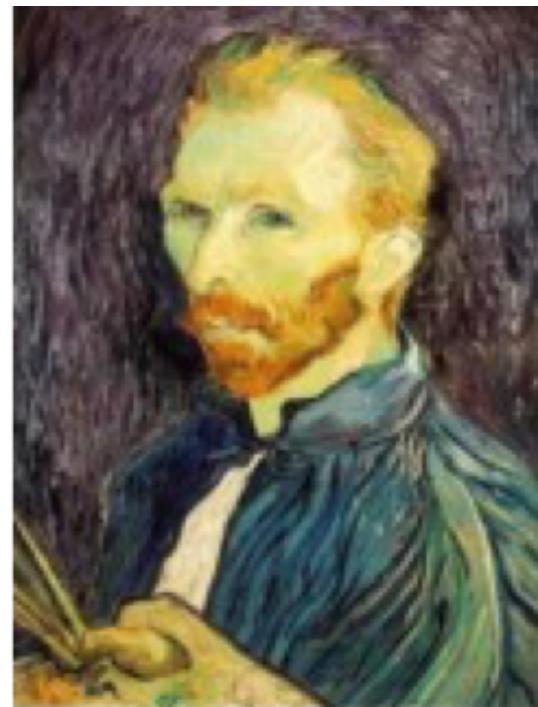


1/8

# Better image downsampling



1/2

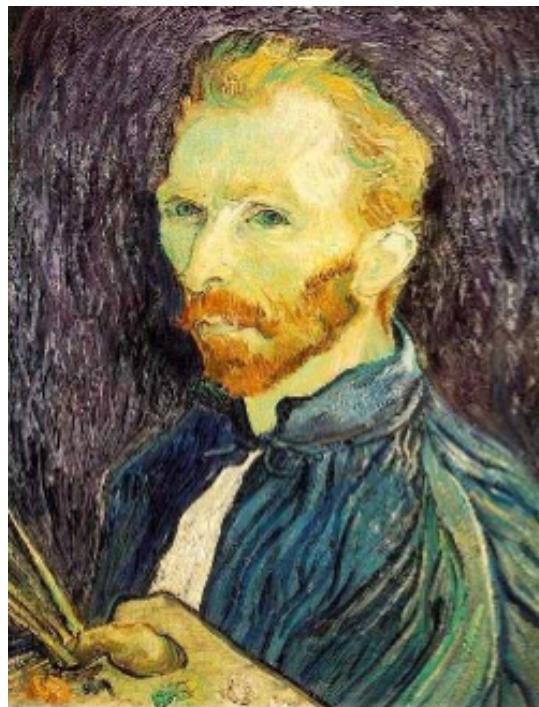


1/4 (2x zoom)

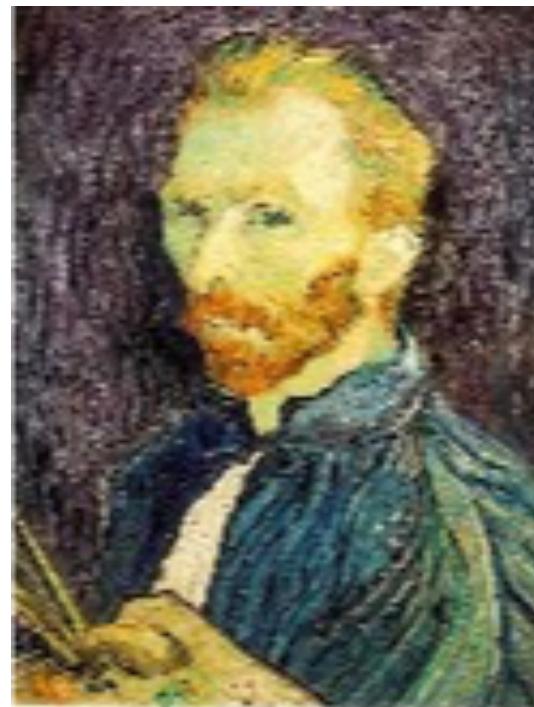


1/8 (4x zoom)

# Naïve image downsampling



1/2



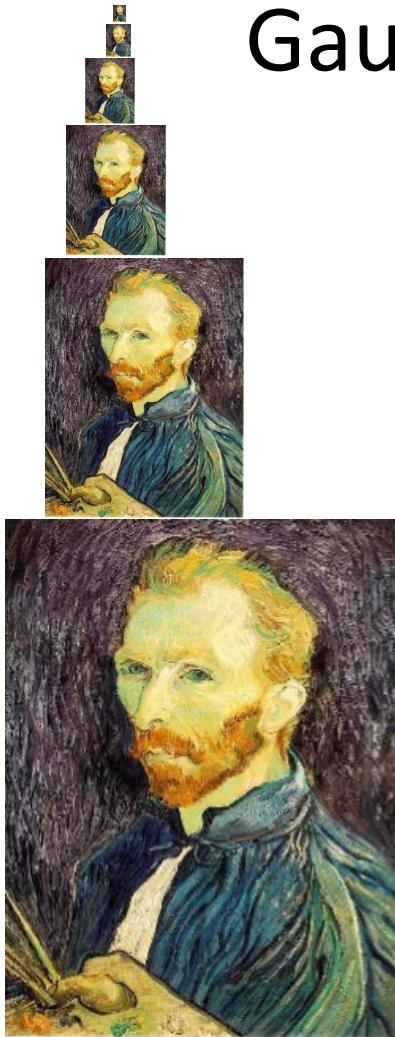
1/4 (2x zoom)



1/8 (4x zoom)

# Gaussian image pyramid

# Gaussian image pyramid

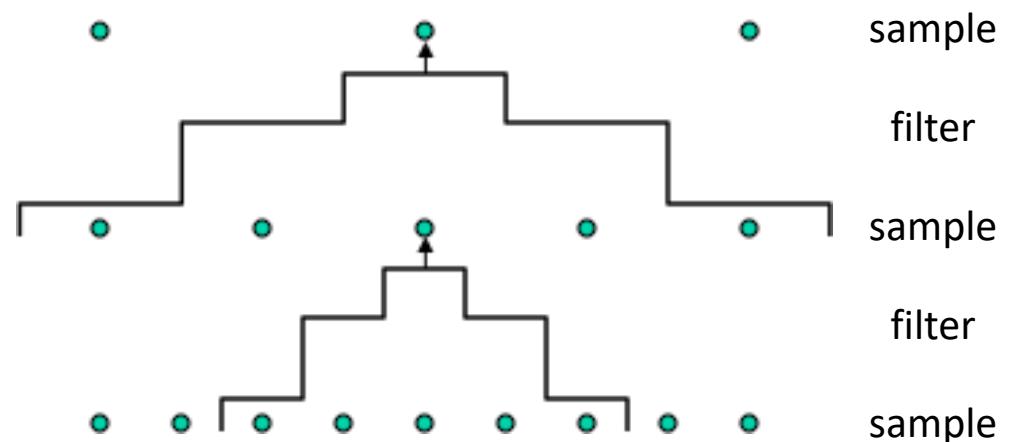


The name of this sequence of subsampled images

# Constructing a Gaussian pyramid

## Algorithm

```
repeat:  
    filter  
    subsample  
until min resolution reached
```

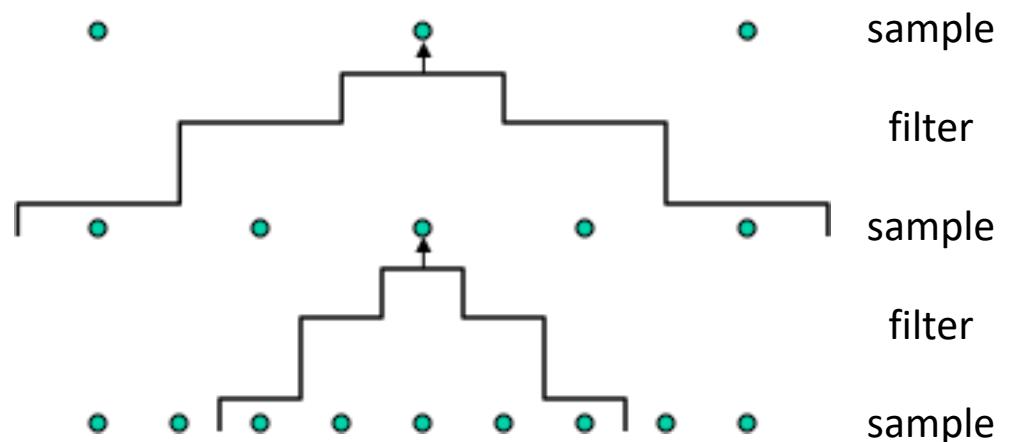


**Question:** How much bigger than the original image is the whole pyramid?

# Constructing a Gaussian pyramid

## Algorithm

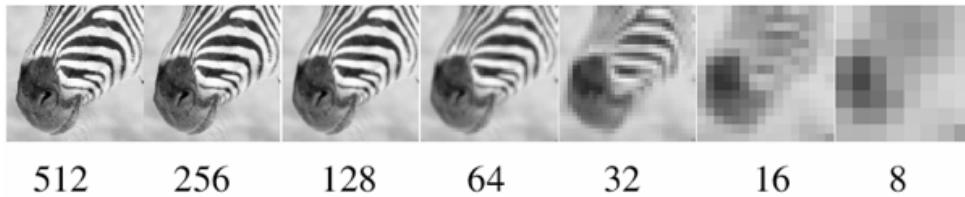
```
repeat:  
    filter  
    subsample  
until min resolution reached
```



**Question:** How much bigger than the original image is the whole pyramid?

**Answer:** Just  $4/3$  times the size of the original image! (How did I come up with this number?)

# Some properties of the Gaussian pyramid



**What happens to the details of the image?**

- They get smoothed out as we move to higher levels.

**What is preserved at the higher levels?**

- Mostly large uniform regions in the original image.

**How would you reconstruct the original image from the image at the upper level?**

- That's not possible.

# Blurring is lossy



level 0



level 1 (before downsampling)

What does the residual look like?

# Blurring is lossy



level 0



level 1 (before downsampling)

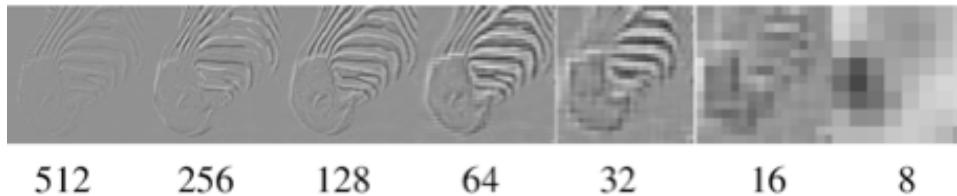


residual

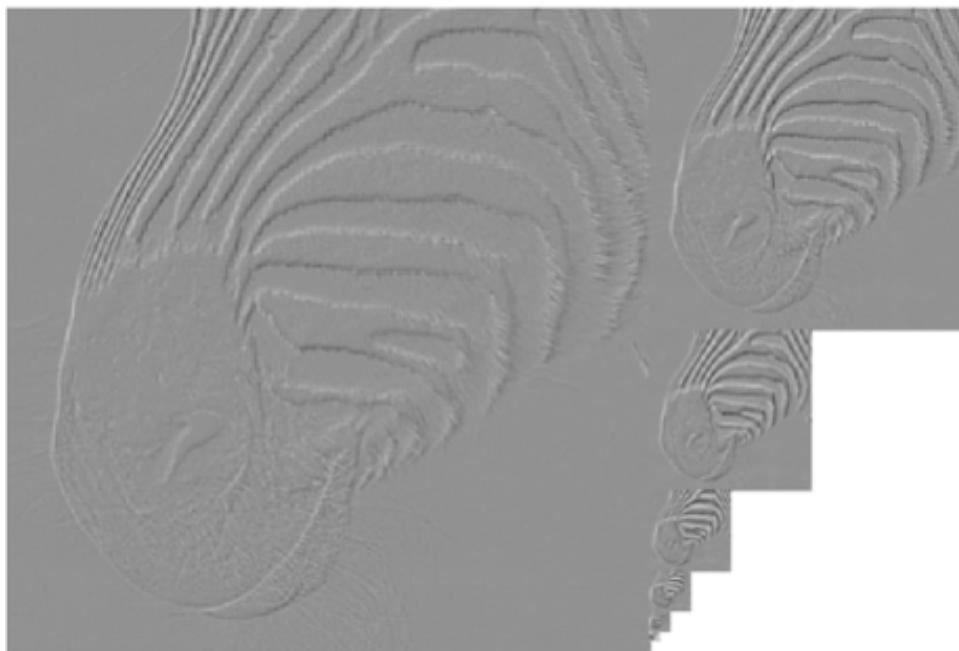
Can we make a pyramid that is lossless?

# Laplacian image pyramid

# Laplacian image pyramid



At each level, retain the residuals instead of the blurred images themselves.



Can we reconstruct the original image using the pyramid?

# Laplacian image pyramid



512

256

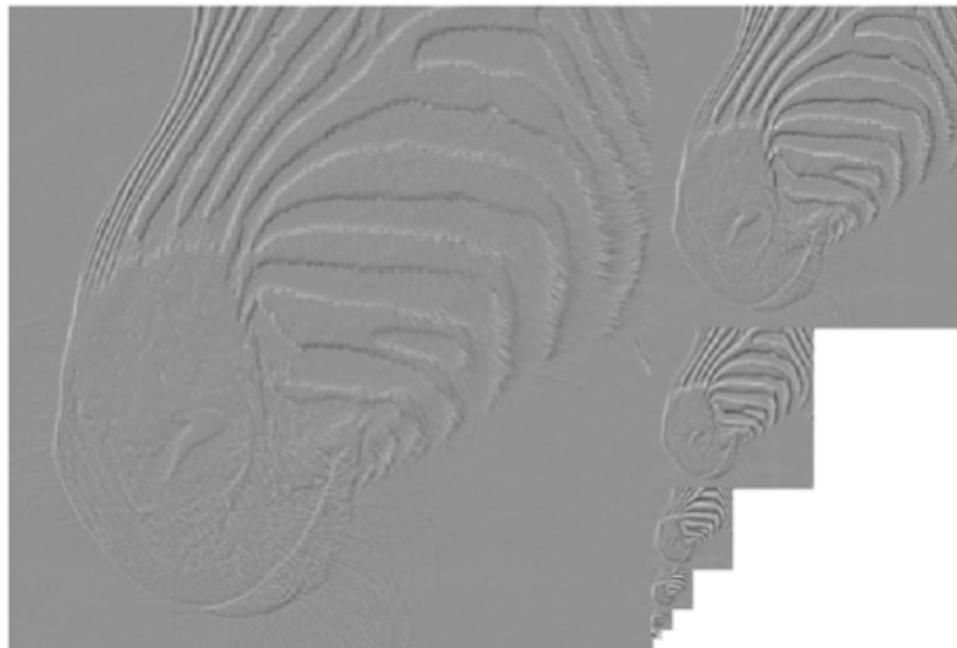
128

64

32

16

8



At each level, retain the residuals instead of the blurred images themselves.

Can we reconstruct the original image using the pyramid?

- Yes we can!

[Deep residual learning for image recognition](#)

K He, X Zhang, S Ren, J Sun - Proceedings of the IEEE ..., 2016 - openaccess.thecvf.com

Deeper neural networks are more difficult to train. We present a residual learning framework to ease the training of networks that are substantially deeper than those used previously. We .

[☆ Save](#) [⤒ Cite](#) [Cited by 131288](#) [Related articles](#) [⤓](#)

# Let's start by looking at just one level



level 0



level 1 (upsampled)

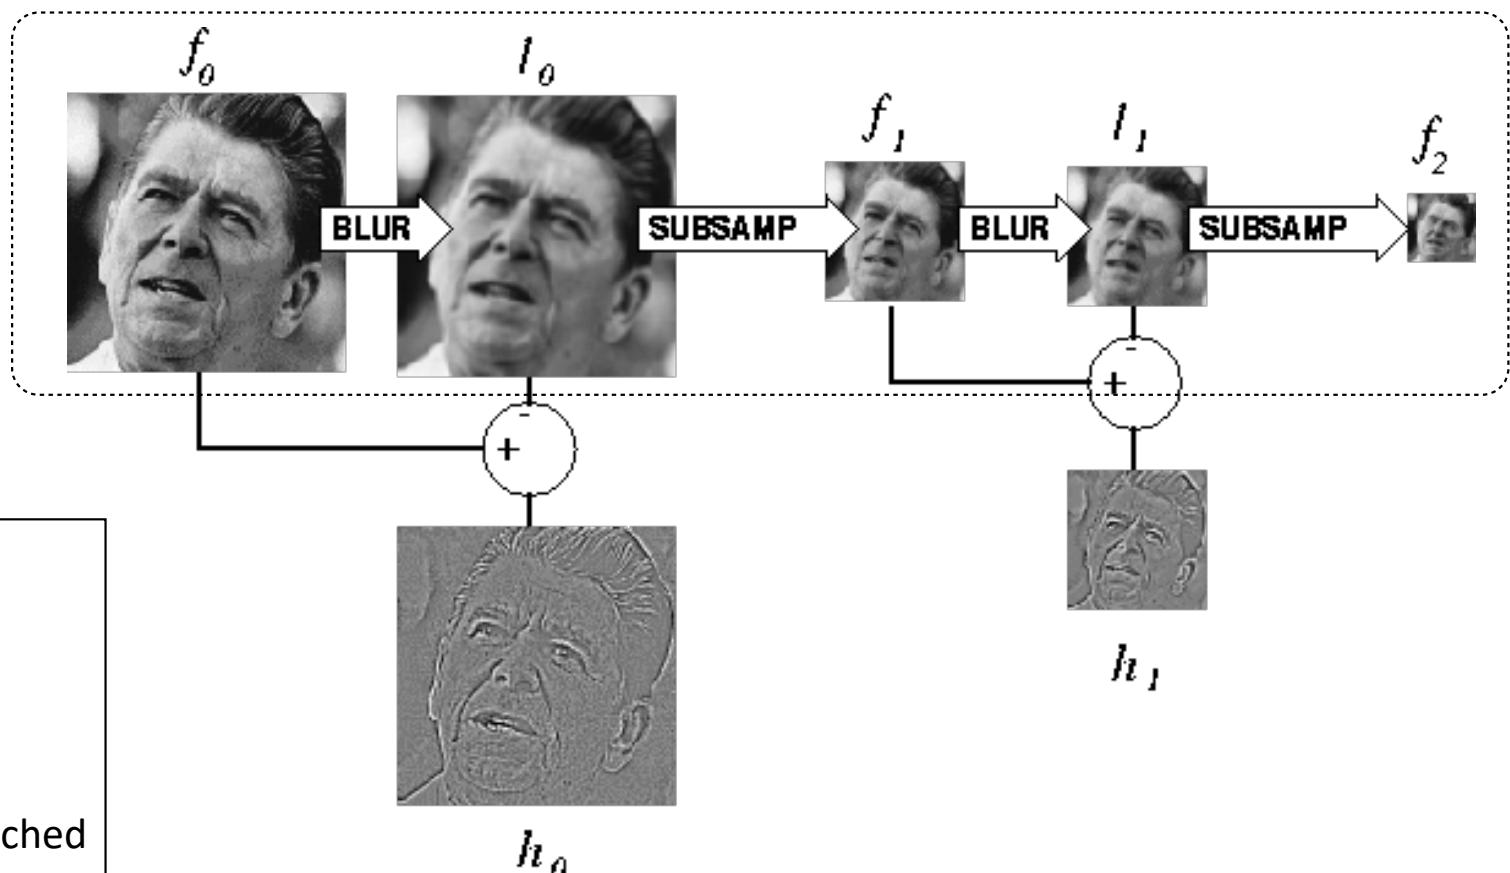


residual

Does this mean we need to store both residuals and the blurred copies of the original?

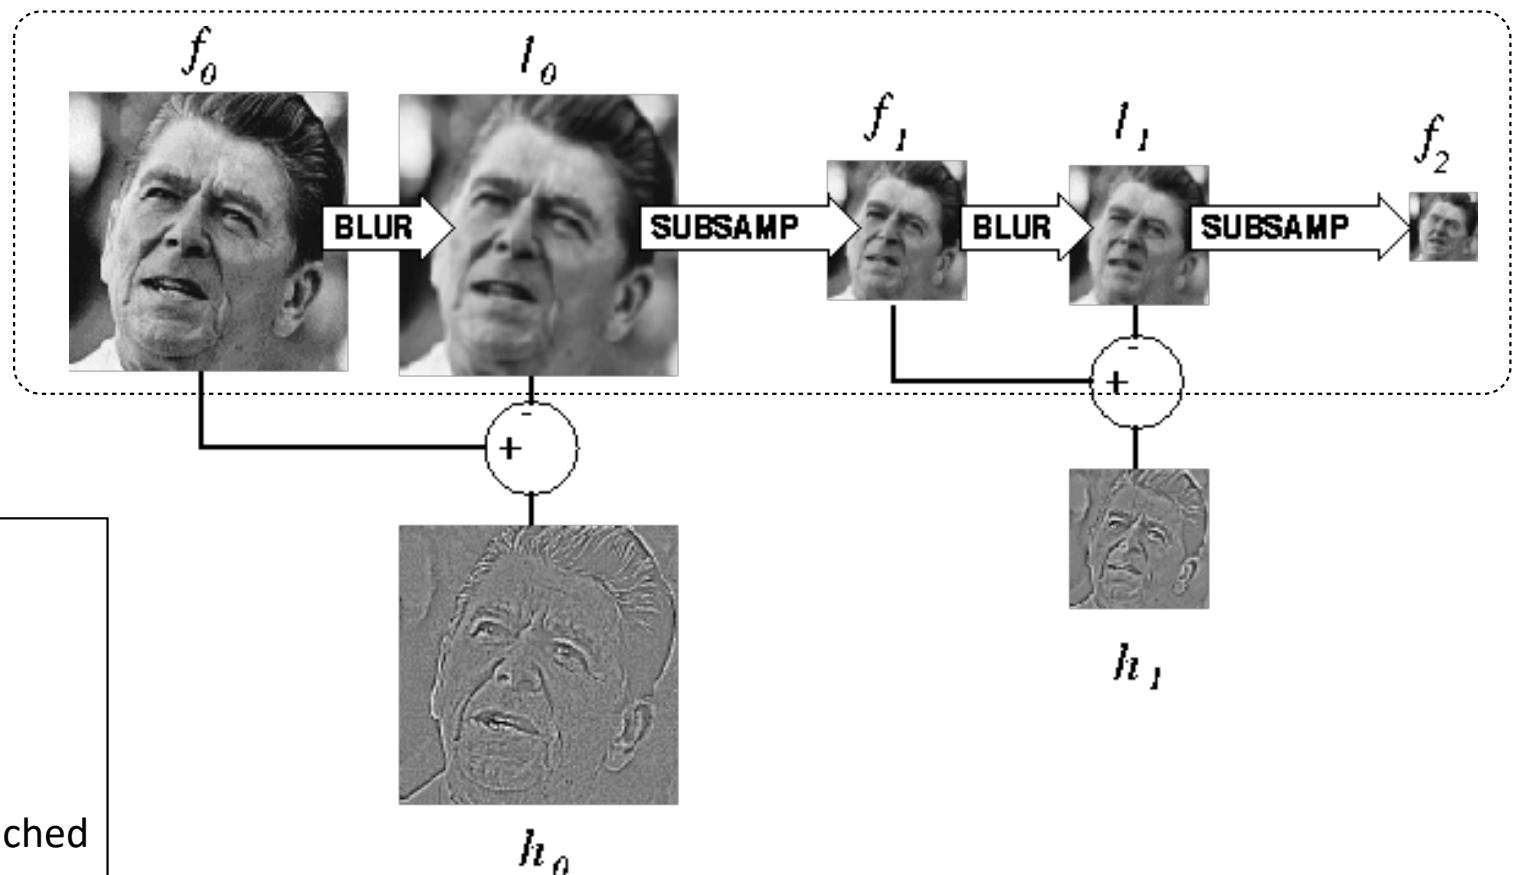
# Constructing a Laplacian pyramid

What is this part?

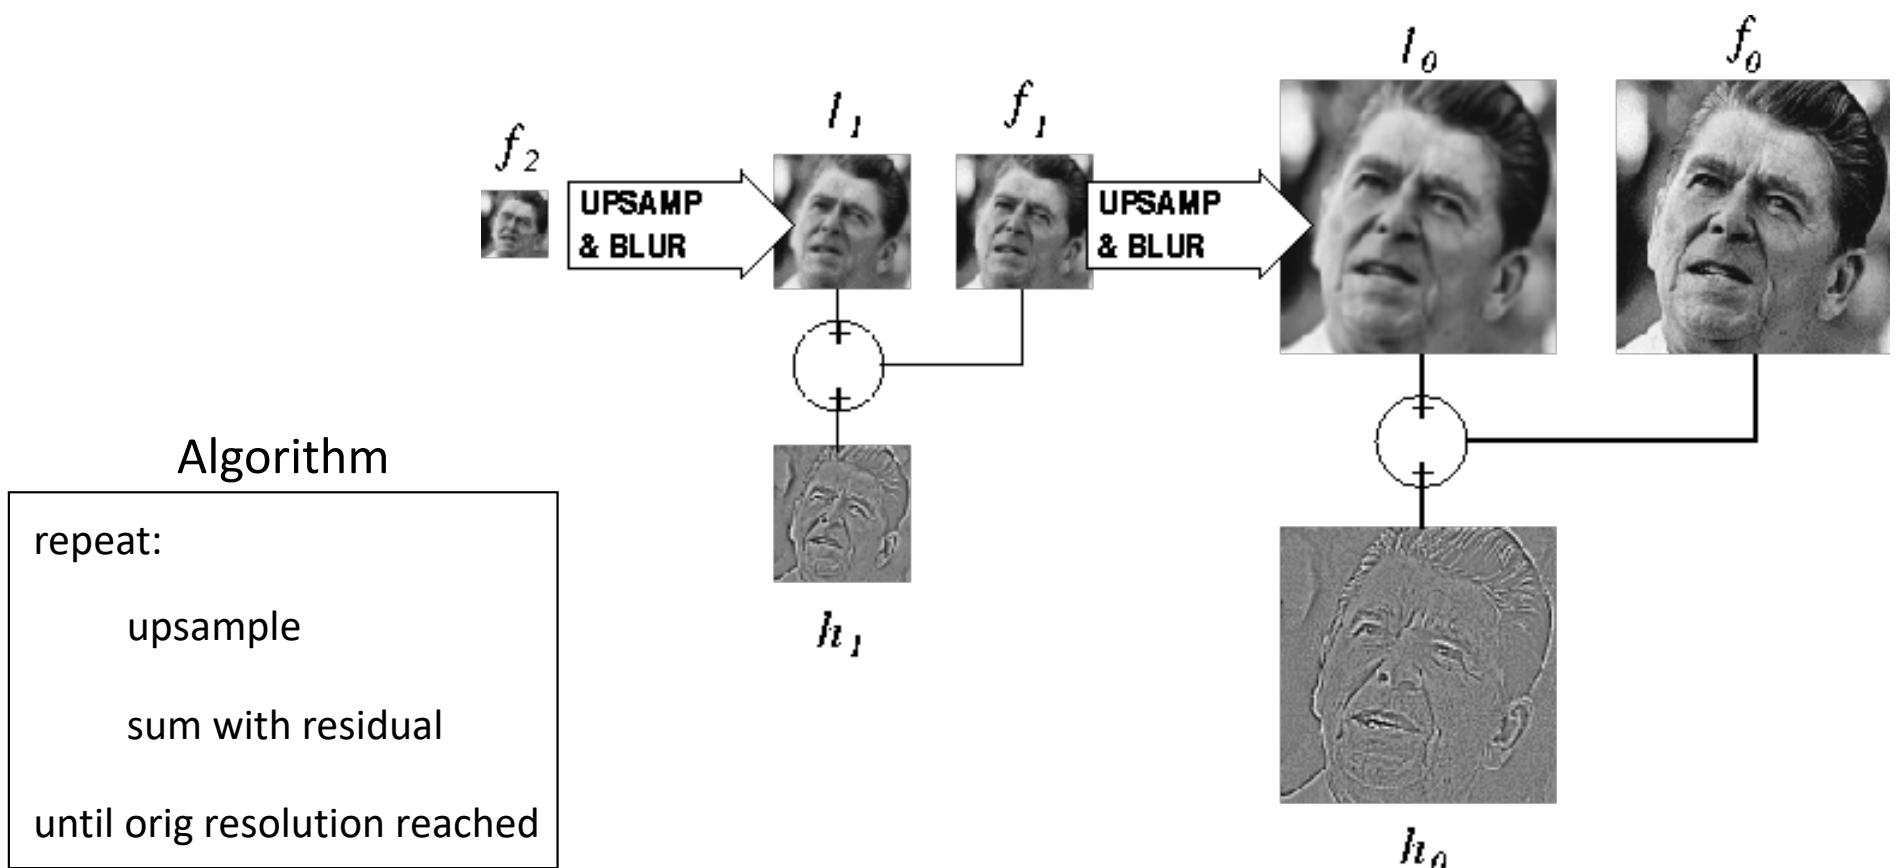


# Constructing a Laplacian pyramid

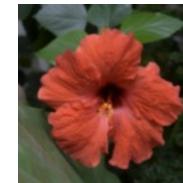
It's a Gaussian pyramid.



# Reconstructing the original image

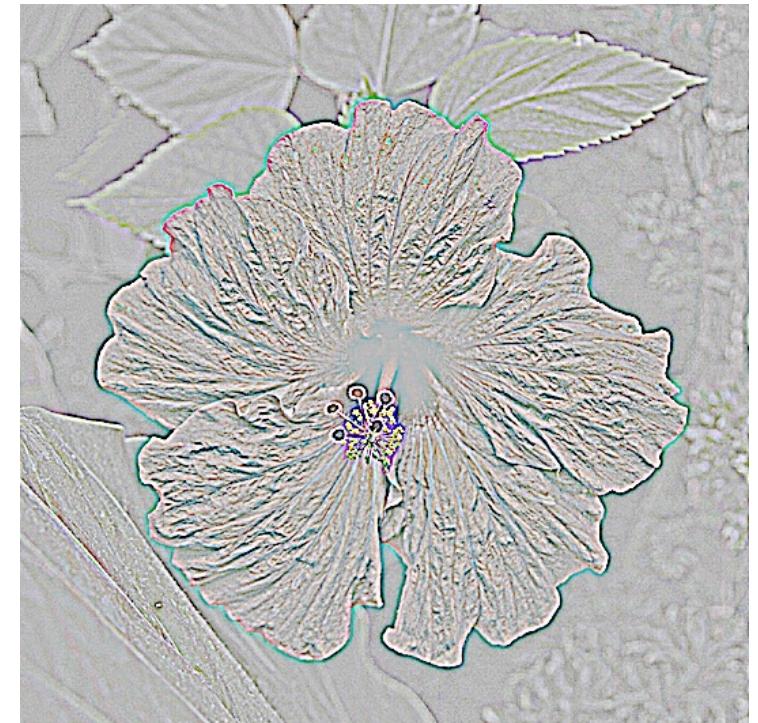
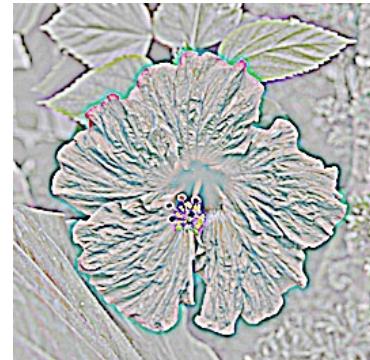
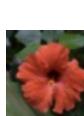


# Gaussian vs Laplacian Pyramid



Shown in opposite  
order for space.

Which one takes  
more space to store?



# What are image pyramids used for?

image compression



multi-scale texture mapping

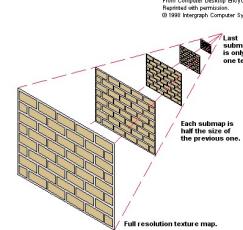
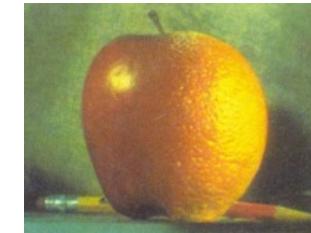
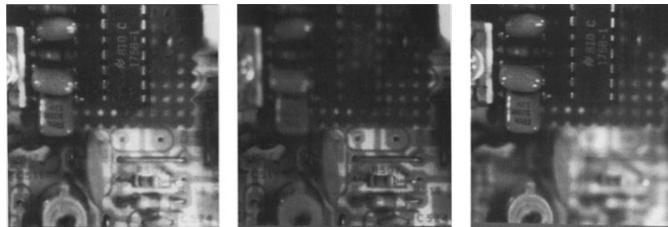


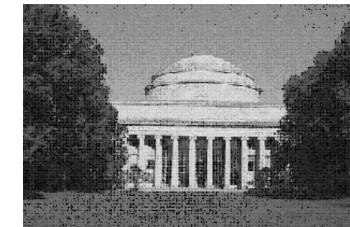
image blending



focal stack compositing



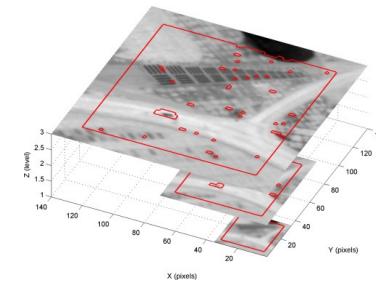
denoising



multi-scale detection



multi-scale registration



# Fourier series

# Basic building block

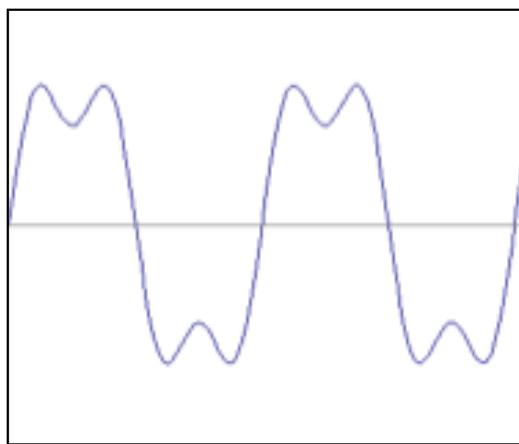
$$A \sin(\omega x + \phi)$$

The diagram illustrates the components of a sinusoidal function. The expression  $A \sin(\omega x + \phi)$  is shown at the top. Five arrows point from labels below to specific parts of the expression: 'amplitude' points to  $A$ , 'sinusoid' points to  $\sin$ , 'angular frequency' points to  $\omega$ , 'variable' points to  $x$ , and 'phase' points to  $\phi$ .

Fourier's claim: Add enough of these to get any periodic signal you want!

# Examples

How would you generate this function?



=

?

+

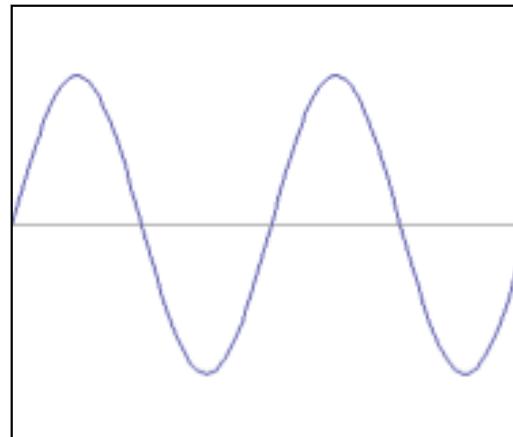
?

# Examples

How would you generate this function?



=



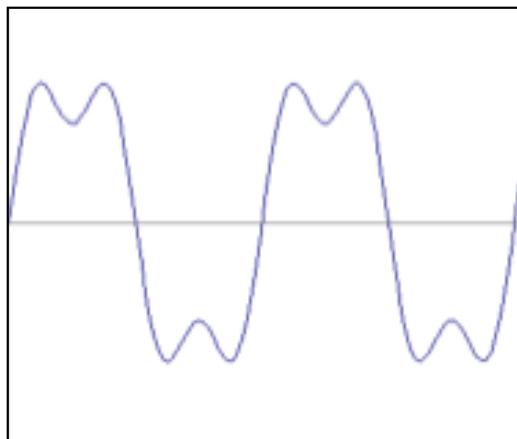
+

?

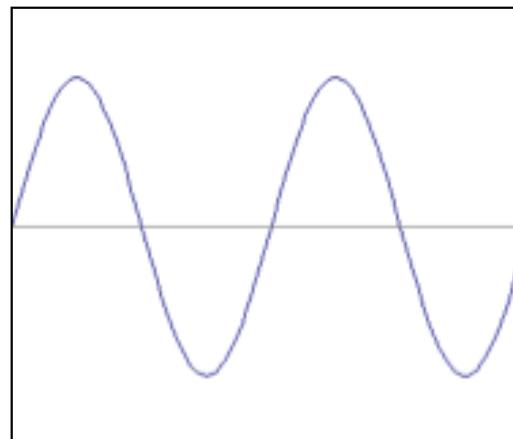
$$\sin(2\pi x)$$

# Examples

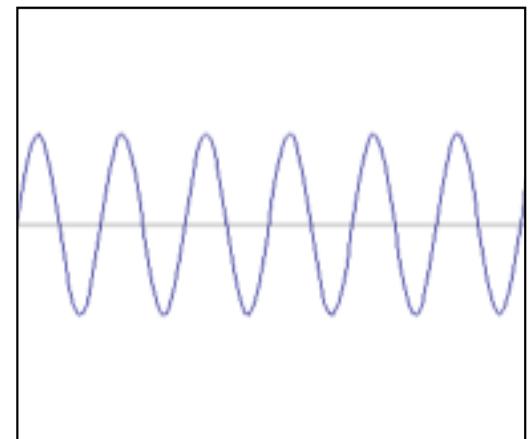
How would you generate this function?



=



+



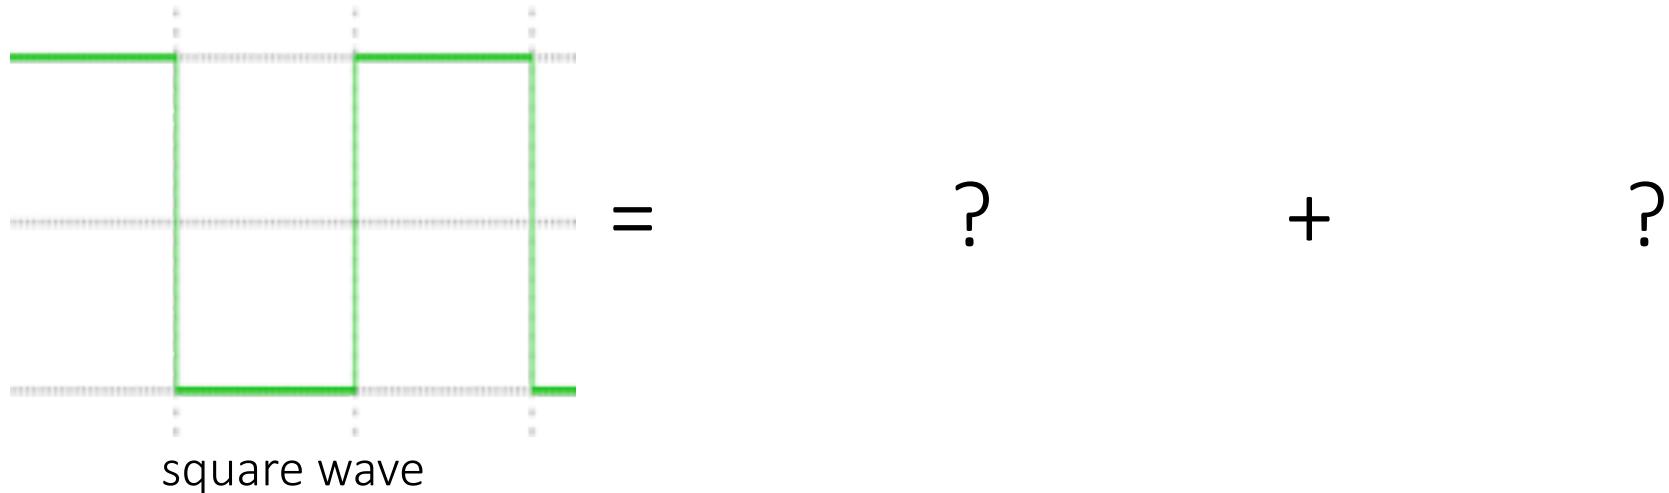
$$f(x) = \sin(2\pi x) + \frac{1}{3} \sin(2\pi 3x)$$

$$\sin(2\pi x)$$

$$\frac{1}{3} \sin(2\pi 3x)$$

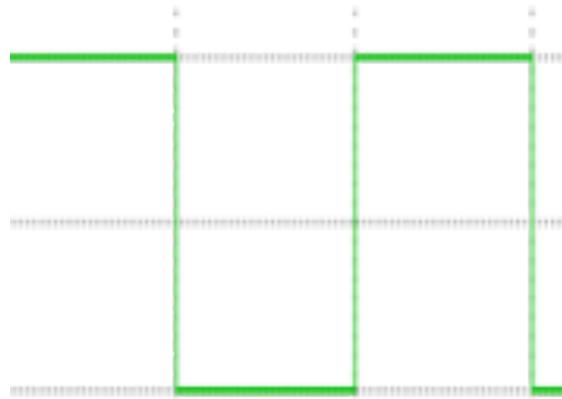
# Examples

How would you generate this function?



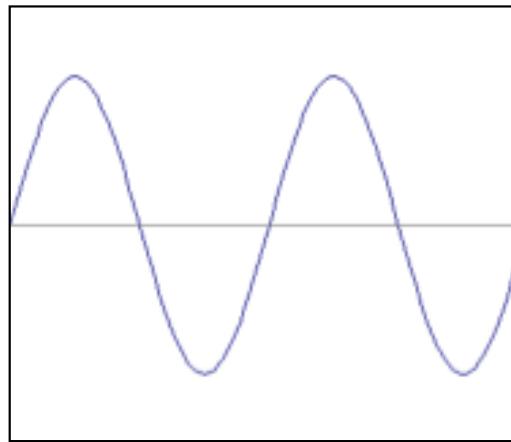
# Examples

How would you generate this function?

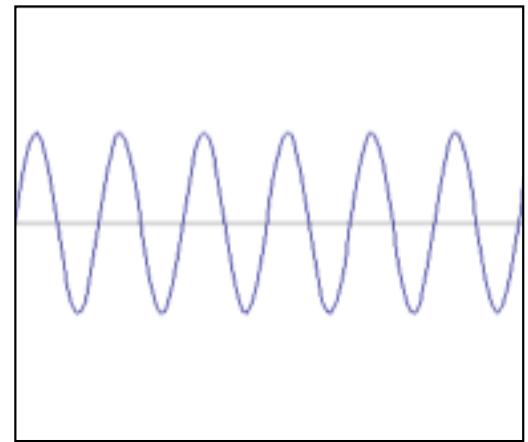


square wave

$\approx$



+

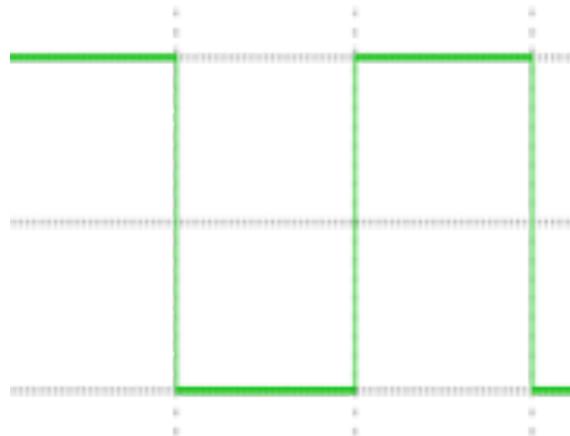


$=$

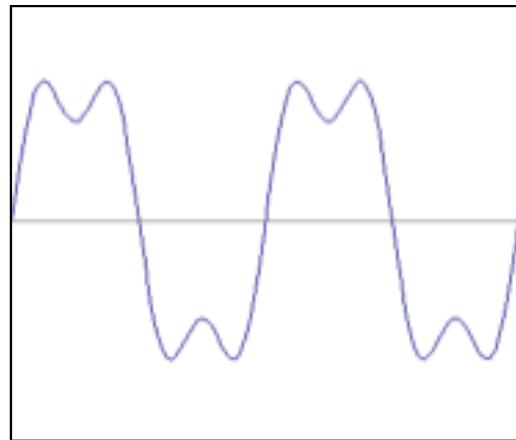


# Examples

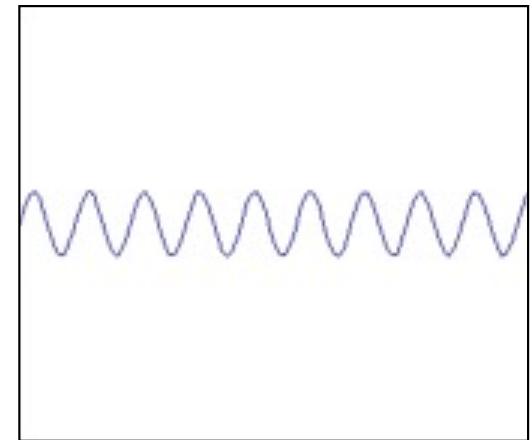
How would you generate this function?



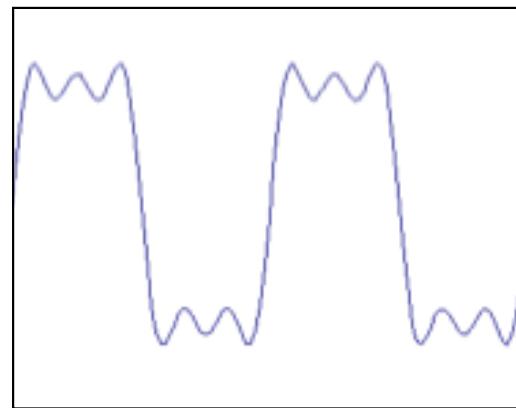
$\approx$



+

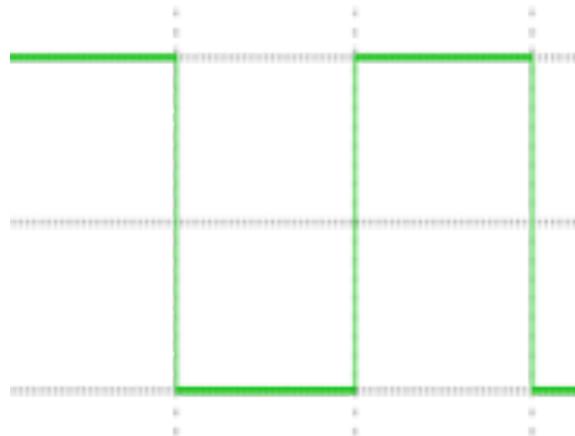


=



# Examples

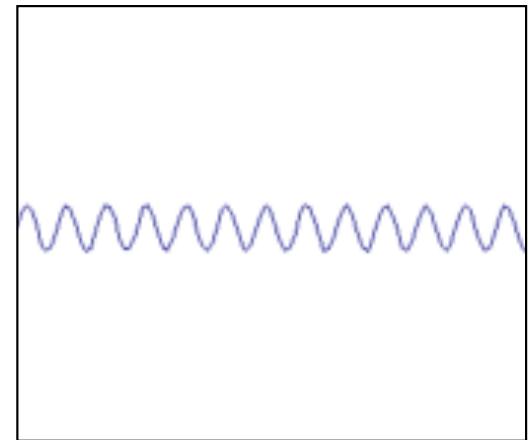
How would you generate this function?



$\approx$



+

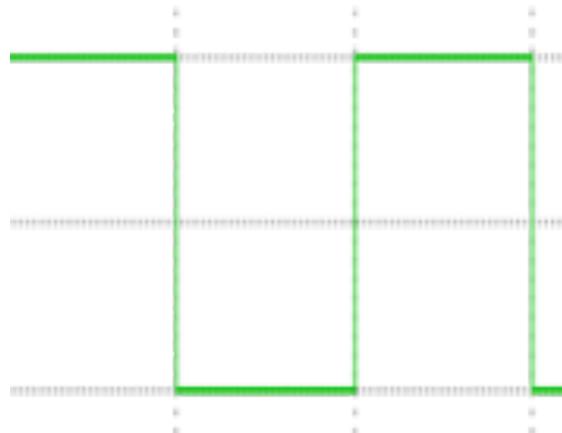


=



# Examples

How would you generate this function?

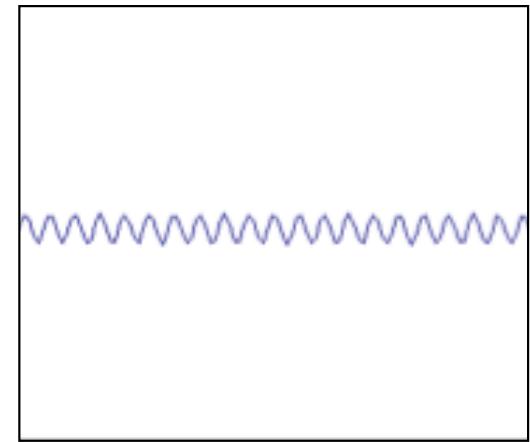


square wave

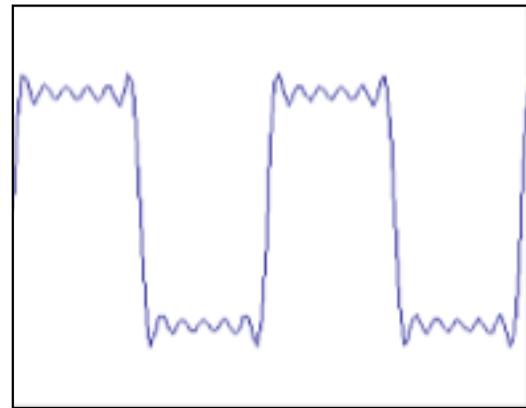
$\approx$



+

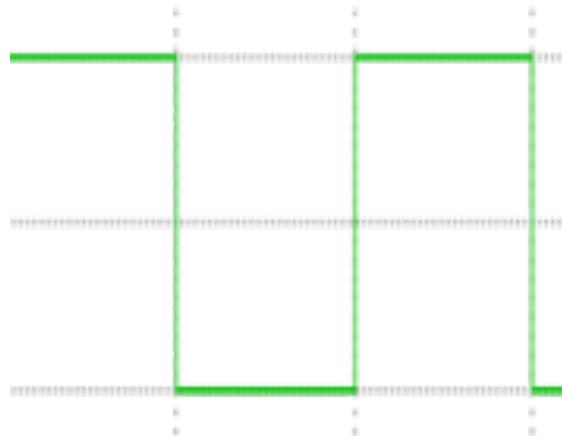


=



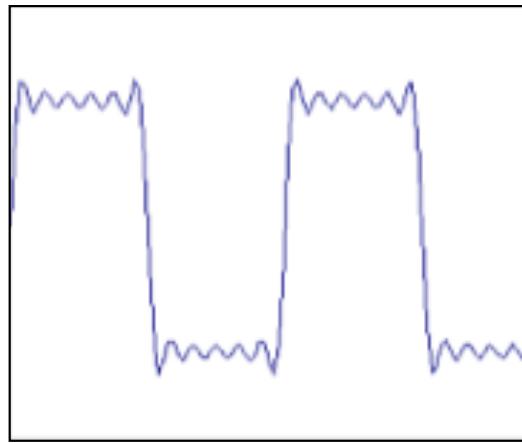
# Examples

How would you generate this function?

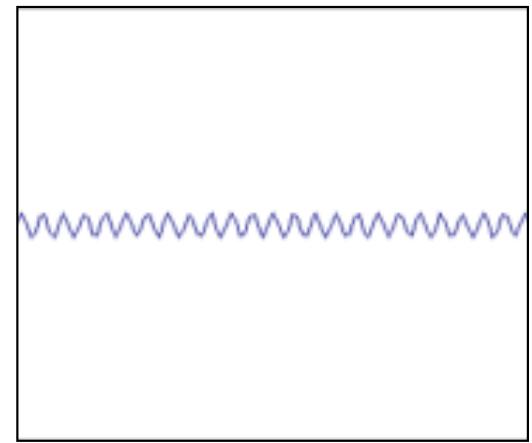


square wave

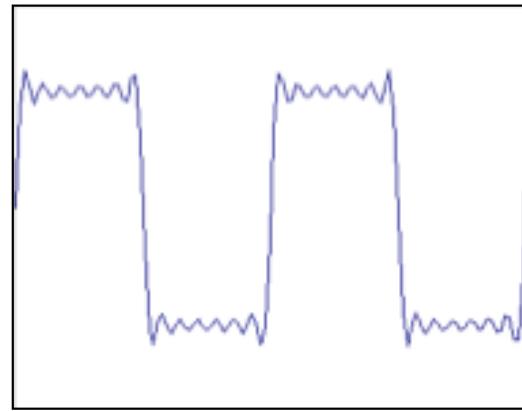
$\approx$



+

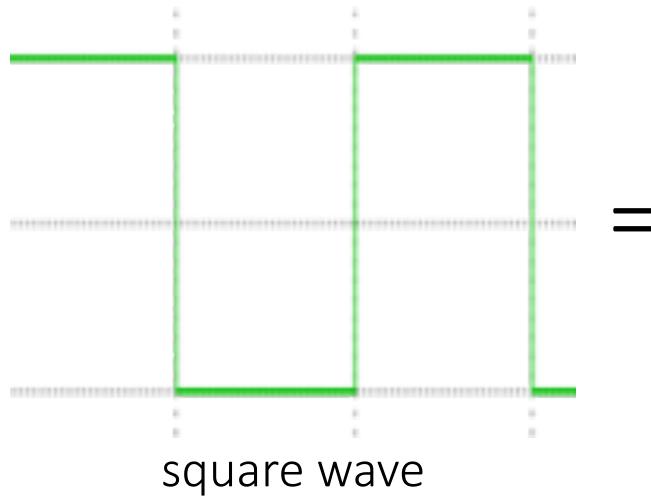


$=$



How would you express  
this mathematically?

# Examples



=

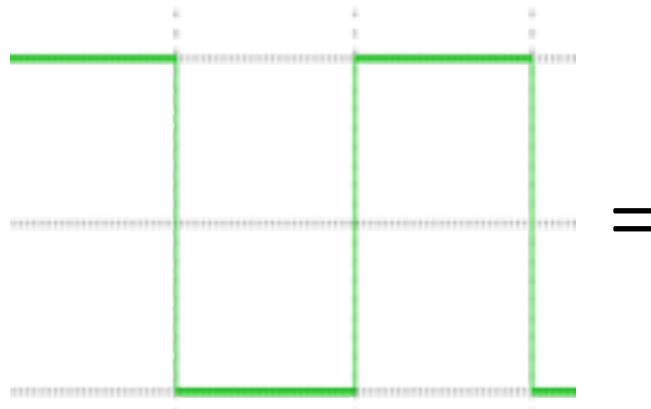
$$A \sum_{k=1}^{\infty} \frac{1}{k} \sin(2\pi kx)$$

square wave

infinite sum of sine waves

How would you visualize this in the frequency domain?

# Examples



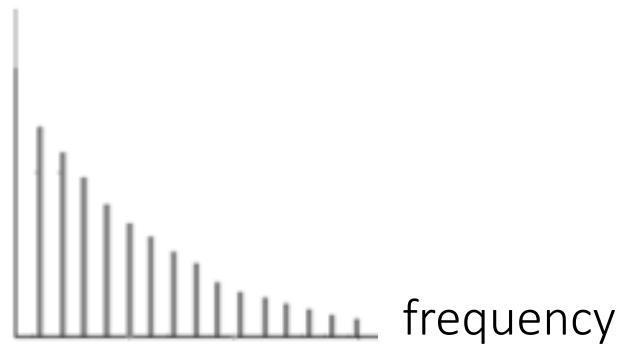
square wave

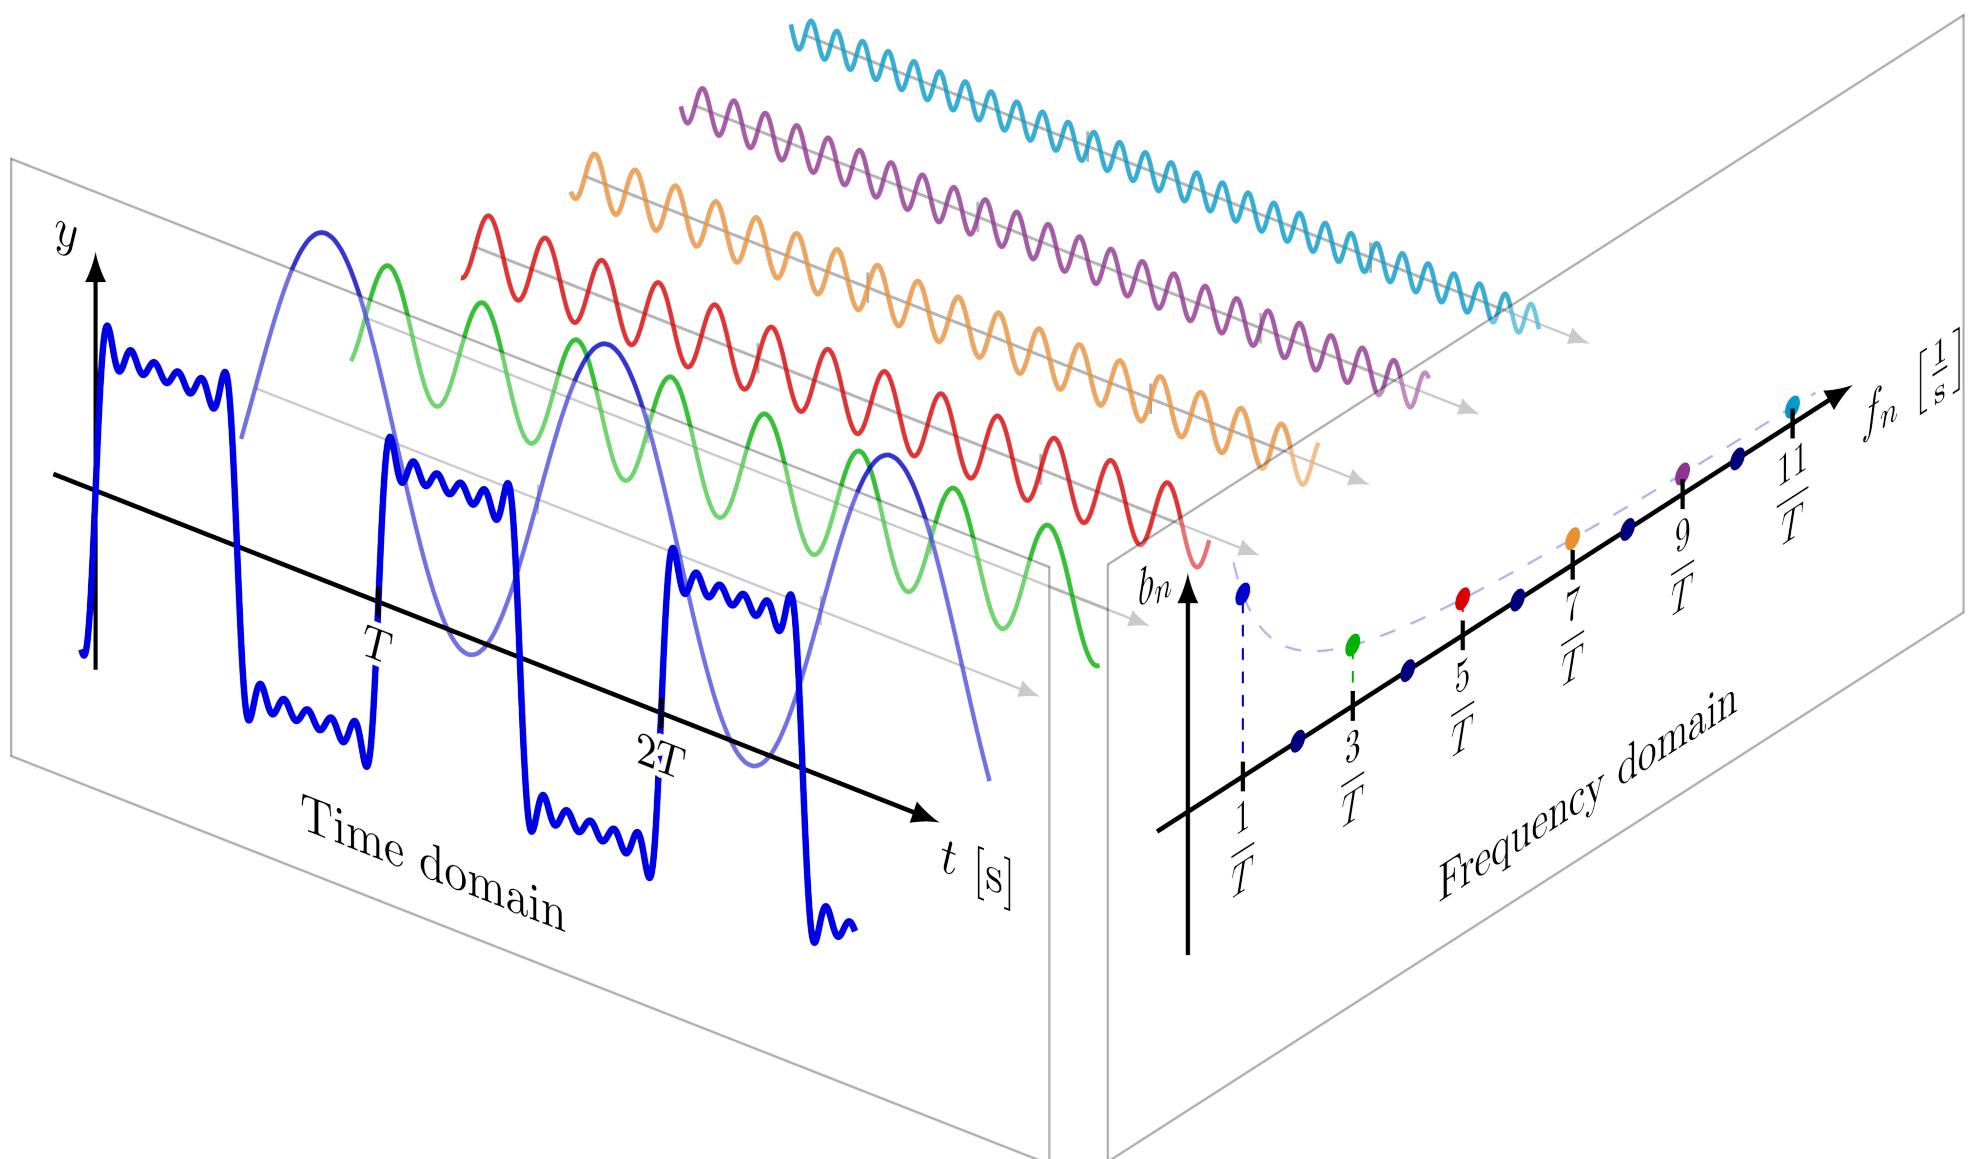
=

$$A \sum_{k=1}^{\infty} \frac{1}{k} \sin(2\pi kx)$$

infinite sum of sine waves

magnitude





# Fourier transform

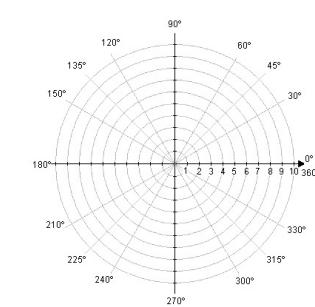
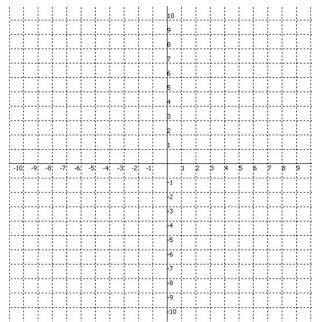
# Recalling some basics

Complex numbers have two parts:

rectangular  
coordinates

$$R + jI$$

real      imaginary



Alternative reparameterization:

polar  
coordinates

$$r(\cos \theta + j \sin \theta)$$

polar transform

$$\theta = \tan^{-1}\left(\frac{I}{R}\right) \quad r = \sqrt{R^2 + I^2}$$

or  
equivalently

$$re^{j\theta}$$

Euler's formula

$$e^{j\theta} = \cos \theta + j \sin \theta$$

exponential  
form

This will help us understand the Fourier transform equations

# Fourier transform

The connection to the ‘summation of sine waves’ idea

$$F(k) = \frac{1}{N} \sum_{x=0}^{N-1} f(x) e^{-j2\pi kx/N}$$

$$f(x) = \sum_{k=0}^{N-1} F(k) e^{j2\pi kx/N}$$

Euler's formula  
 $e^{j\theta} = \cos \theta + j \sin \theta$

sum over frequencies

$$f(x) = \sum_{k=0}^{N-1} F(k) \left\{ \cos(2\pi kx/N) + j \sin(2\pi kx/N) \right\}$$

scaling parameter

wave components

# 2D Fourier transform

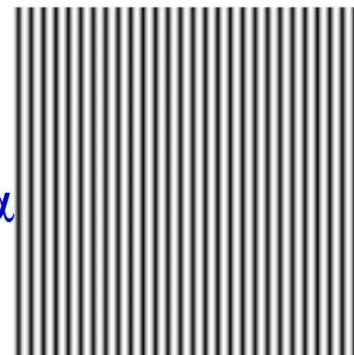
$$F[k, l] = \frac{1}{MN} \sum_{m=0}^{M-1} \sum_{n=0}^{N-1} f[m, n] e^{-j2\pi \left( \frac{k}{M}m + \frac{l}{N}n \right)}$$

Inverse 2D DFT

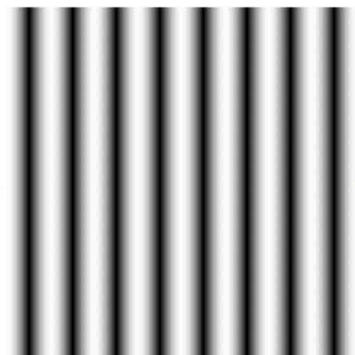
$$f[m, n] = \sum_{k=0}^{M-1} \sum_{l=0}^{N-1} F[k, l] e^{j2\pi \left( \frac{k}{M}m + \frac{l}{N}n \right)}$$



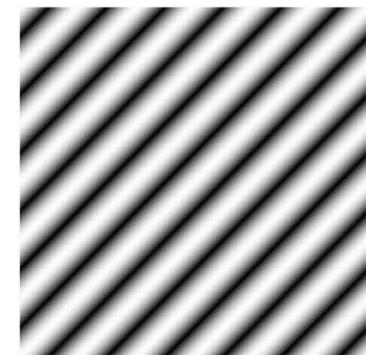
$$= \alpha$$



$$+ \beta$$

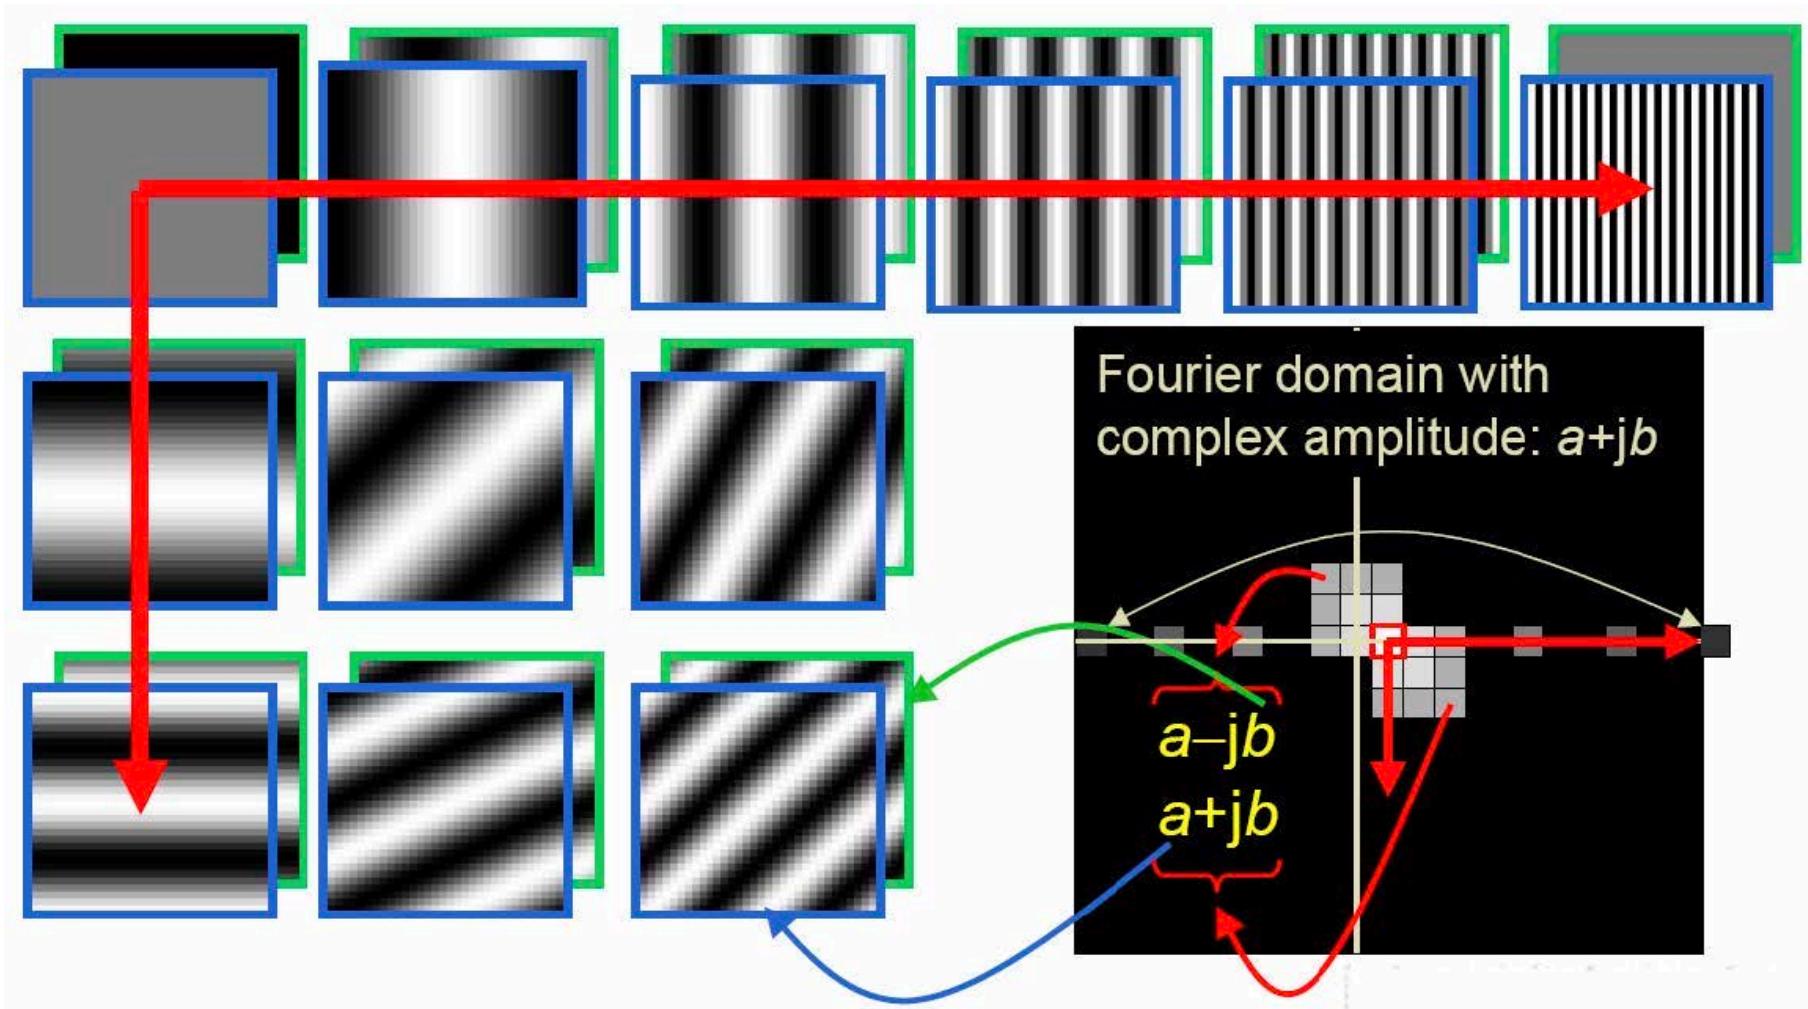


$$+ \gamma$$



$$+ \dots$$

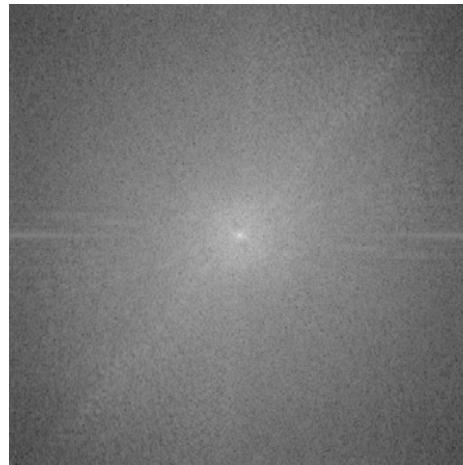
.....



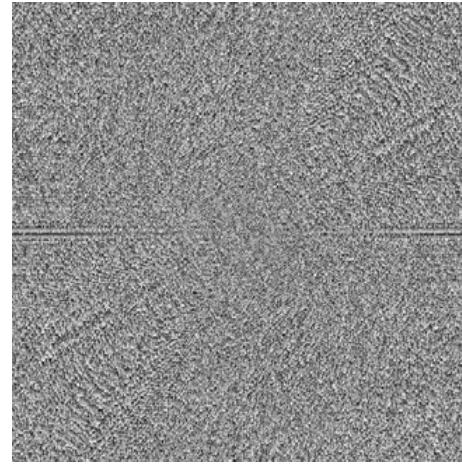
# Fourier transforms of natural images



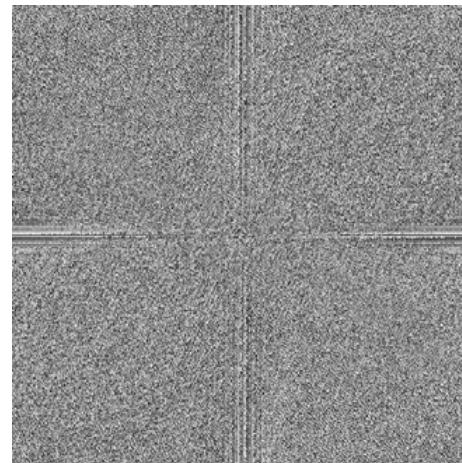
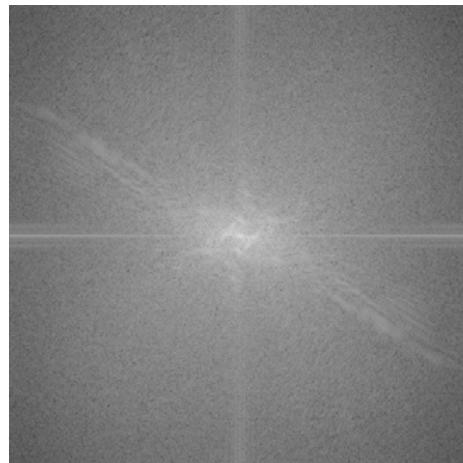
original



amplitude



phase

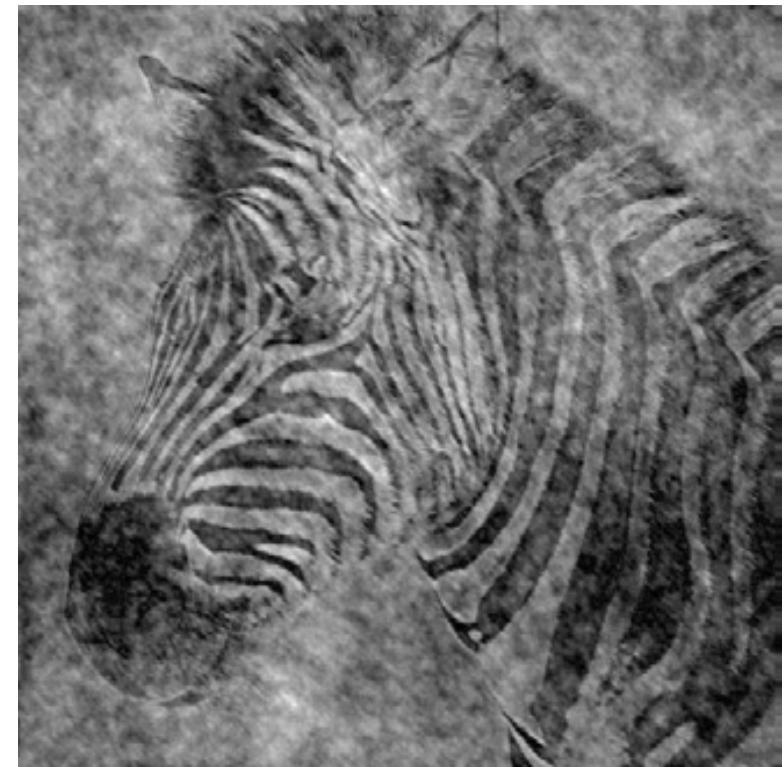


# Fourier transforms of natural images

... • • • • • • • • • • •



cheetah phase with zebra amplitude



zebra phase with cheetah amplitude

# The convolution theorem

The Fourier transform of the convolution of two functions is the product of their Fourier transforms:

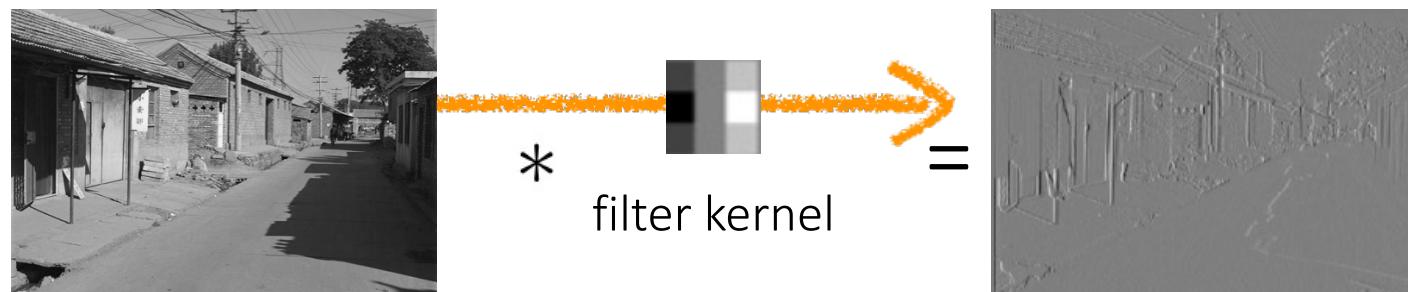
$$\mathcal{F}\{g * h\} = \mathcal{F}\{g\}\mathcal{F}\{h\}$$

The inverse Fourier transform of the product of two Fourier transforms is the convolution of the two inverse Fourier transforms:

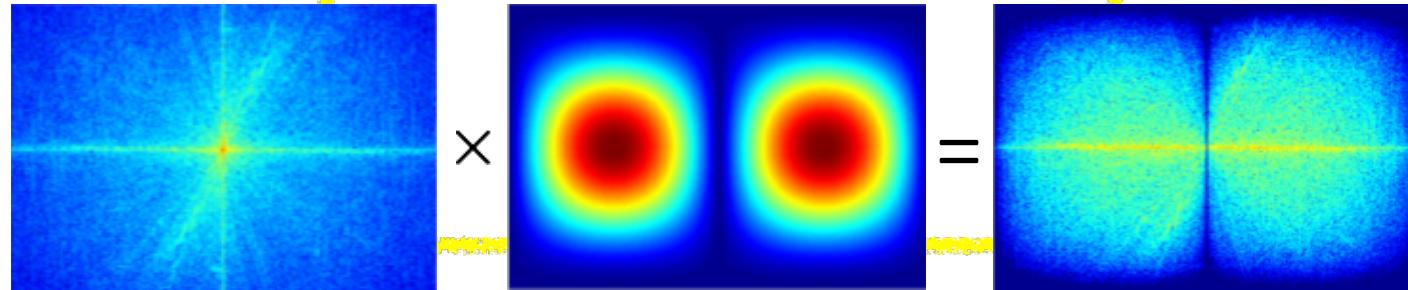
$$\mathcal{F}^{-1}\{gh\} = \mathcal{F}^{-1}\{g\} * \mathcal{F}^{-1}\{h\}$$

Convolution in spatial domain is equivalent to multiplication in frequency domain!

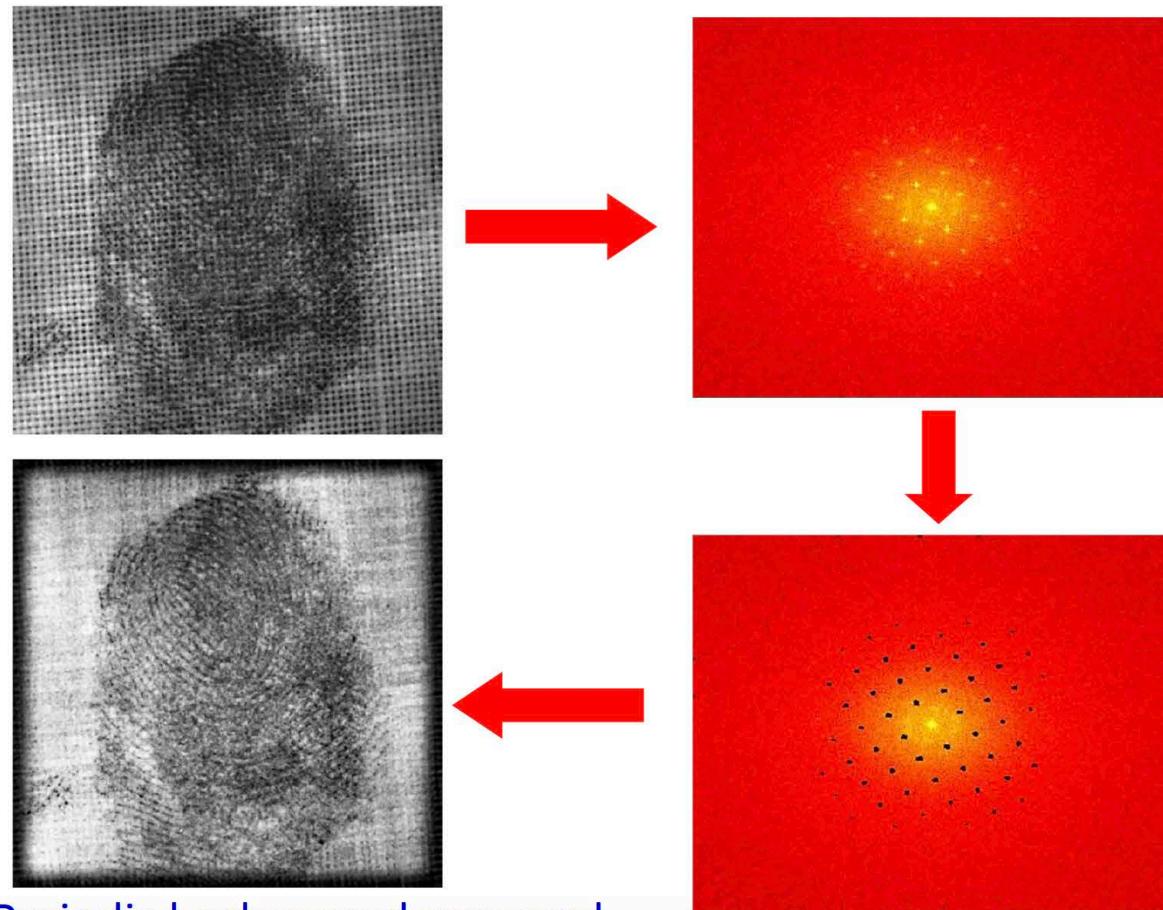
# Spatial domain filtering



Fourier transform



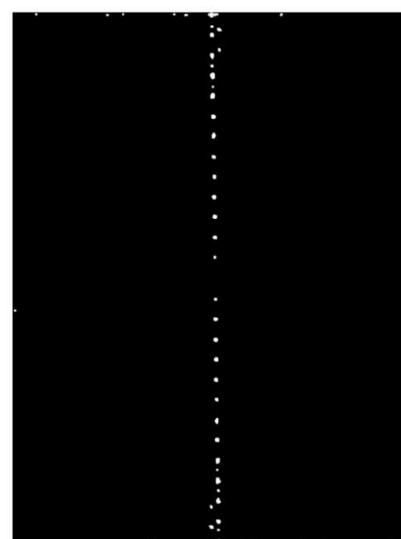
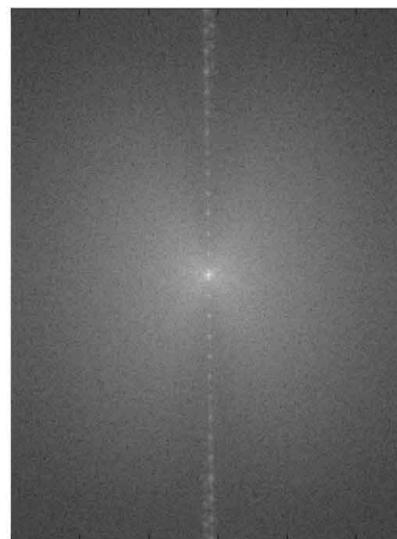
# Frequency domain filtering



Periodic background removed

remove  
peaks

Lunar orbital image (1966)



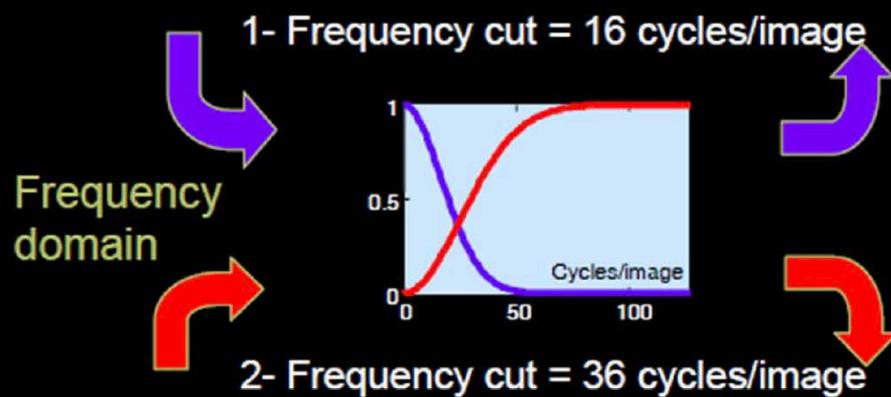
remove  
peaks

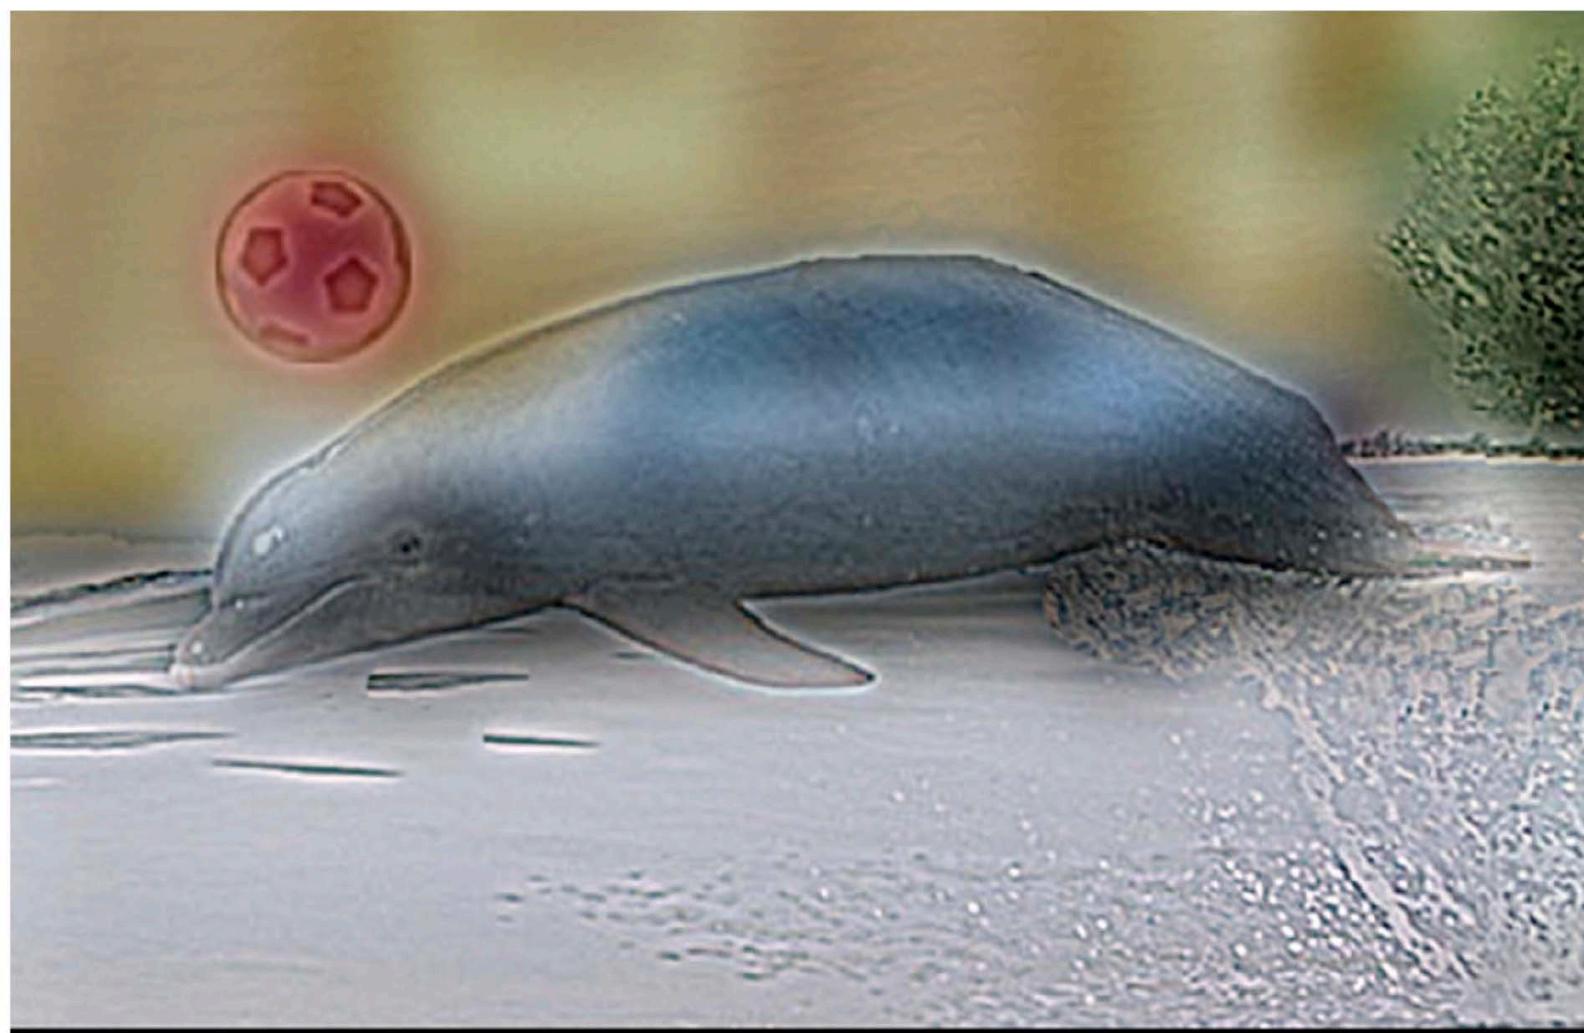
join lines  
removed

# Perception of hybrid images



SIGGRAPH2006





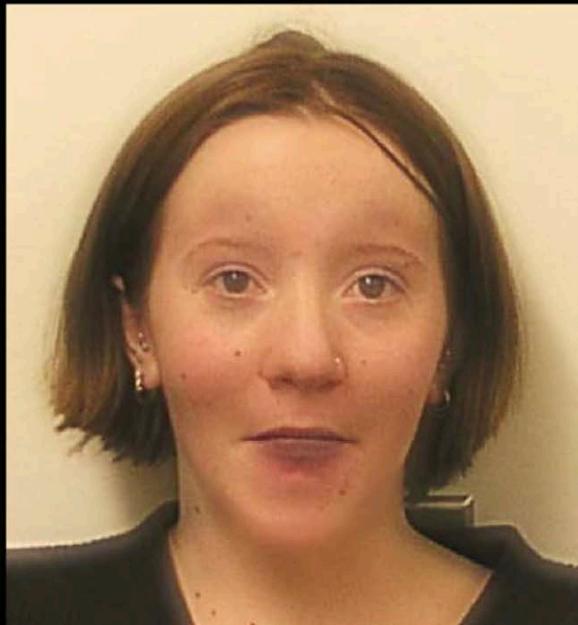
© 2006 Antonio Torralba and Aude Oliva

# Changing expression



Sad

Surprised



Aude Oliva & Antonio Torralba & Philippe G Schyns, SIGGRAPH 2006