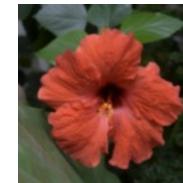


Review: Image Pyramid & Fourier Transform

Gaussian vs Laplacian Pyramid



Shown in opposite
order for space.

Which one takes
more space to store?

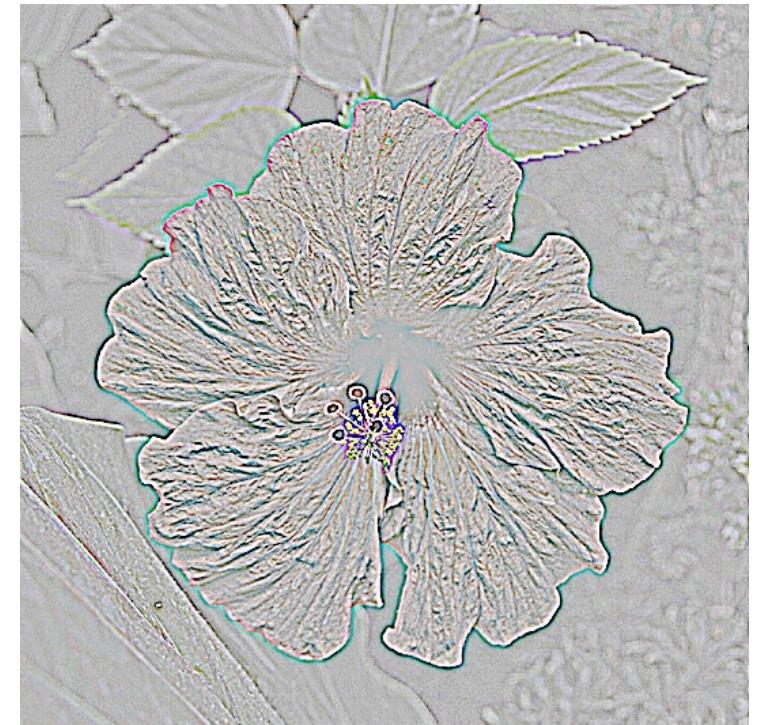
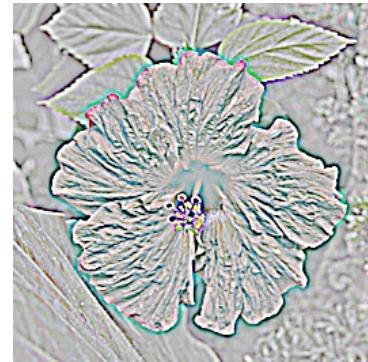
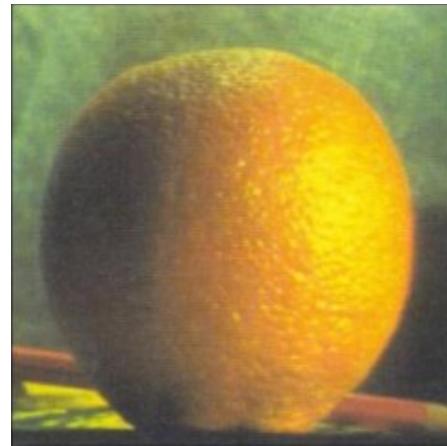
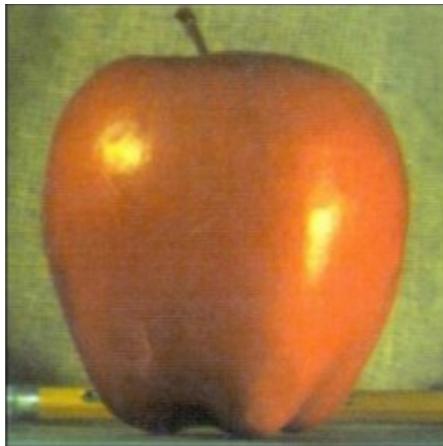
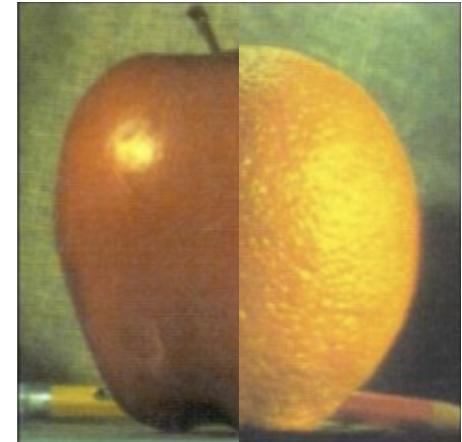


Image Blending

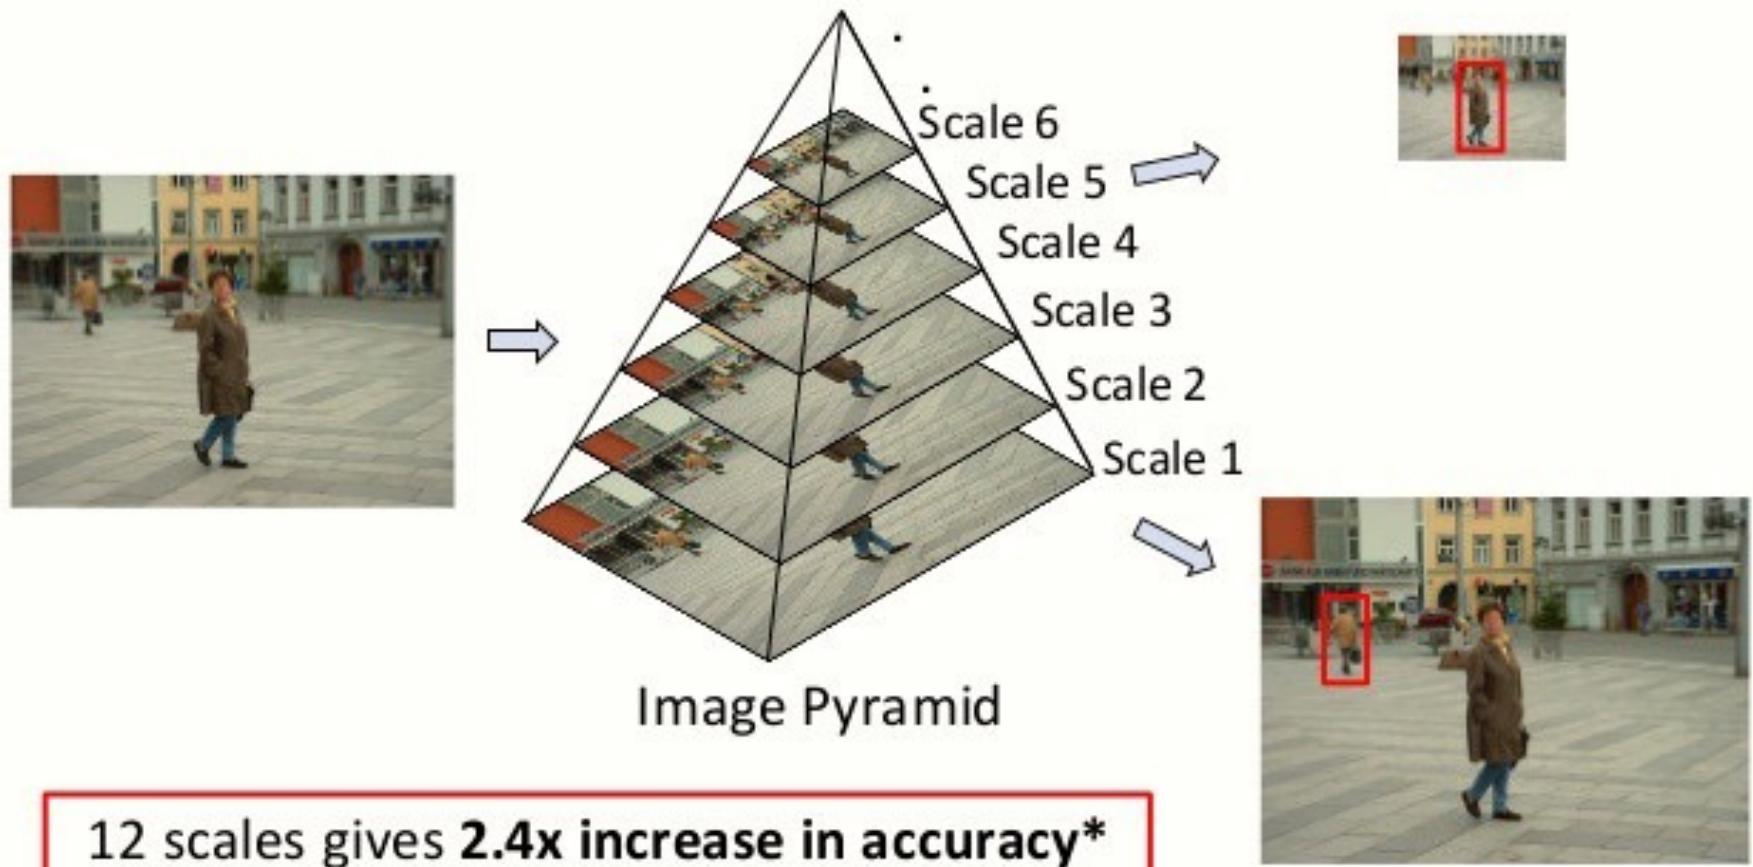


- ❖ Load the two images of apple and orange
- ❖ Find the Gaussian Pyramids for apple and orange
- ❖ From Gaussian Pyramids, find their Laplacian Pyramids
- ❖ Now join the left half of apple and right half of orange in each levels of Laplacian Pyramids
- ❖ From this joint image pyramids, reconstruct the original image.

https://docs.opencv.org/3.4/dc/dff/tutorial_py_pyramids.html

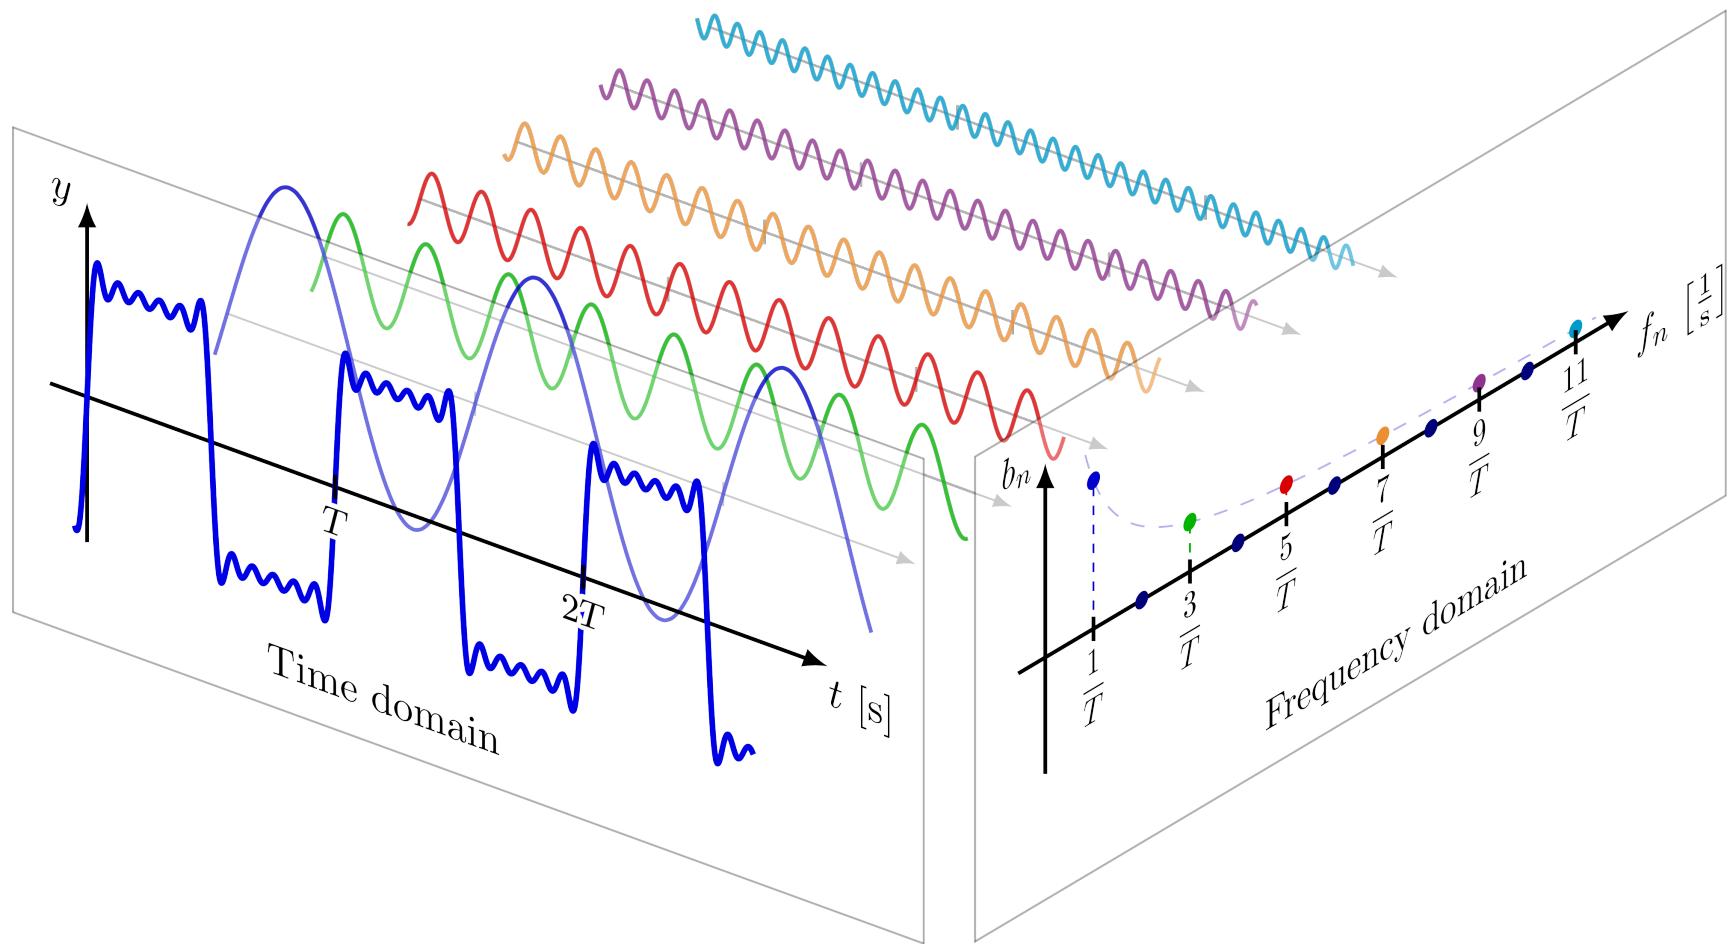


Multi-scale Object Detection



12 scales gives **2.4x increase in accuracy***
at the cost of **3.2x increase in processing**

Fourier Transform



Fourier transform

The connection to the ‘summation of sine waves’ idea

$$F(k) = \frac{1}{N} \sum_{x=0}^{N-1} f(x) e^{-j2\pi kx/N}$$

$$f(x) = \sum_{k=0}^{N-1} F(k) e^{j2\pi kx/N}$$

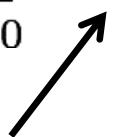
$$\text{Euler's formula } e^{j\theta} = \cos \theta + j \sin \theta$$

sum over frequencies

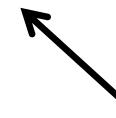


$$f(x) = \sum_{k=0}^{N-1} F(k) \left\{ \cos(2\pi kx/N) + j \sin(2\pi kx/N) \right\}$$

scaling parameter



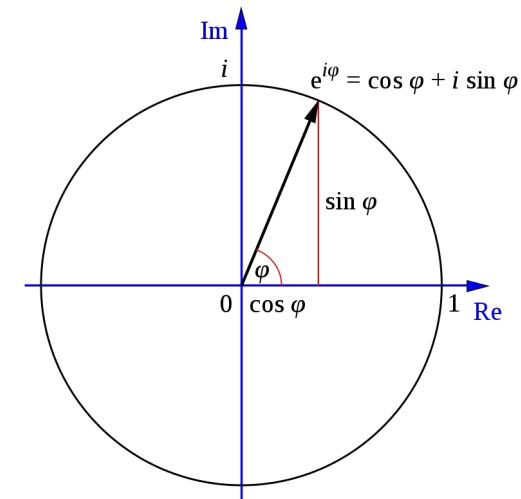
wave components



2D-Fourier Transform

$$F[k, l] = \frac{1}{MN} \sum_{m=0}^{M-1} \sum_{n=0}^{N-1} f[m, n] e^{-j2\pi \left(\frac{k}{M}m + \frac{l}{N}n \right)}$$

$$f[m, n] = \sum_{k=0}^{M-1} \sum_{l=0}^{N-1} F[k, l] e^{j2\pi \left(\frac{k}{M}m + \frac{l}{N}n \right)}$$



2D-Fourier Transform

There is the **fft2** function in the **numpy** toolbox that can do a two-dimensional fast Fourier transform on the image.

```
def my_fft(img):
    dft = np.fft.fft2(img)

    ## shift transform
    dft_shift = np.fft.fftshift(dft)

    ## calculate the magnitude
    magnitude_spectrum = np.abs(dft_shift)

    ## calculate the phase
    phase_spectrum = np.angle(dft_shift)

    return magnitude_spectrum, phase_spectrum
```

To better visualize the output spectrogram, we need to move the center points of the four corners to the center of the matrix and do logarithmic transformation

<https://numpy.org/doc/stable/reference/generated/numpy.fft.fft2.html>

2D-Fourier Transform

There is the **ifft2** function in the **numpy** toolbox that can do a two-dimensional inverse Fourier transform from magnitude and phase.

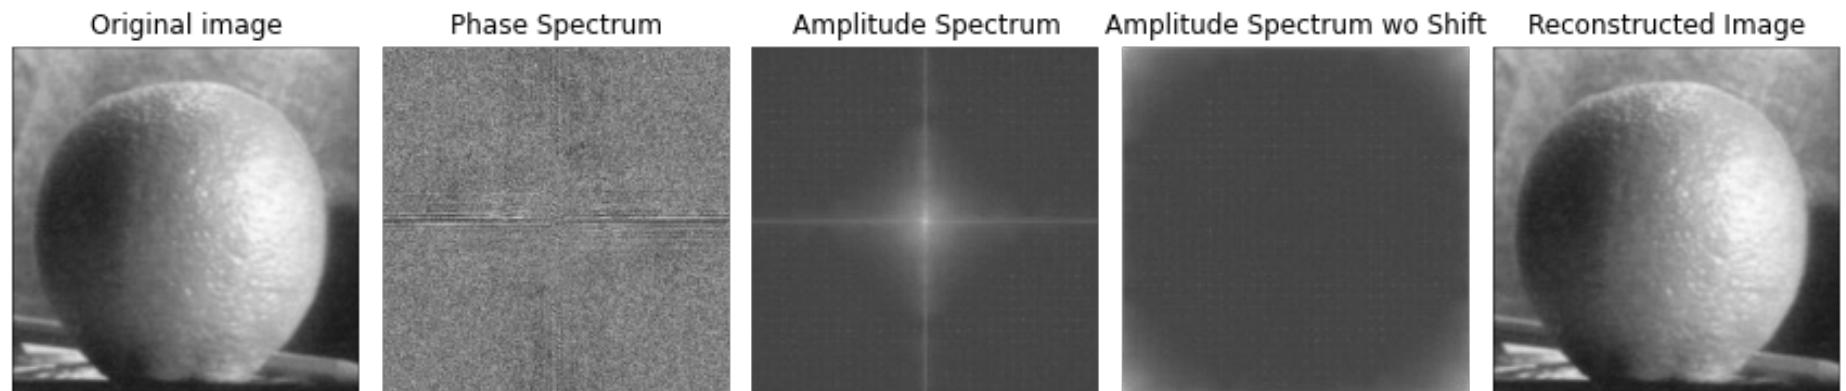
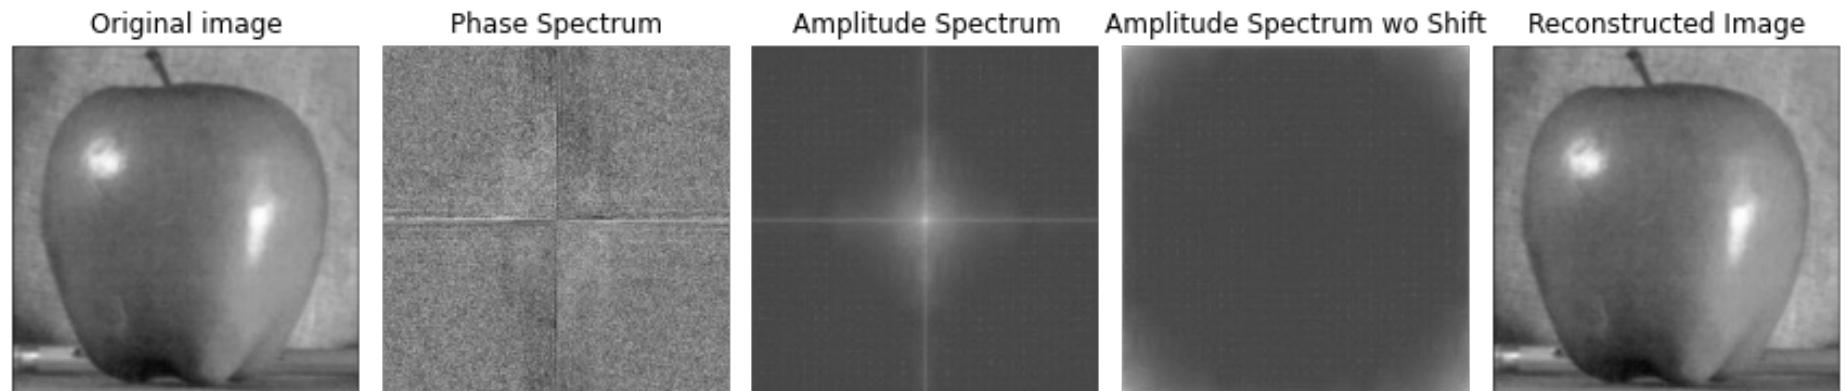
```
def my_ifft(img_m, img_p):
    ## integrate magnitude and phase
    img_mandp = img_m*np.e**(1j*img_p)

    img_mandp = np.uint8(np.abs(np.fft.ifft2(img_mandp)))

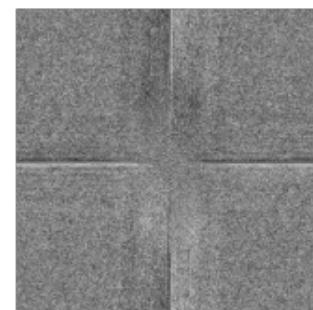
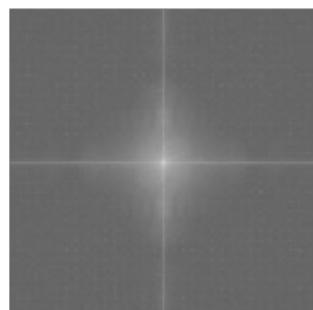
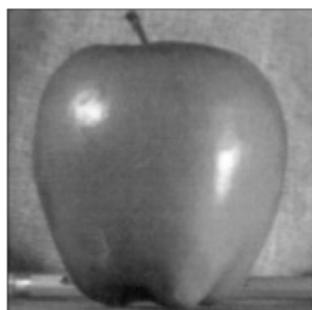
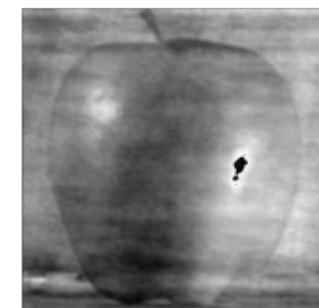
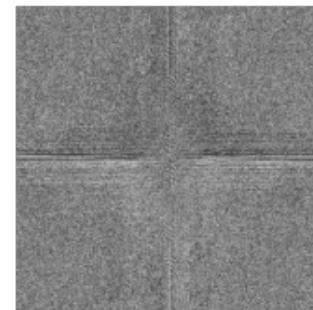
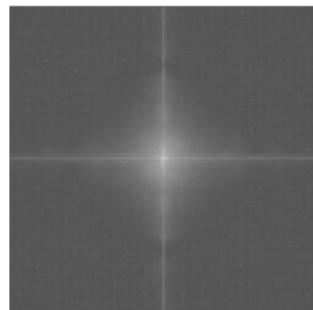
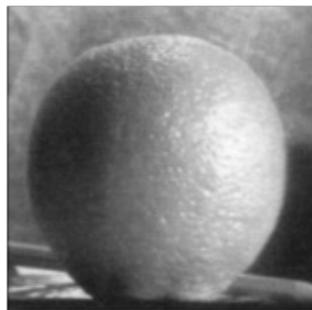
    img_mandp=img_mandp/np.max(img_mandp)*255
    return img_mandp
```

[https://numpy.org/doc/stable/reference
/generated/numpy.fft.ifft2.html](https://numpy.org/doc/stable/reference/generated/numpy.fft.ifft2.html)

2D-Fourier Transform



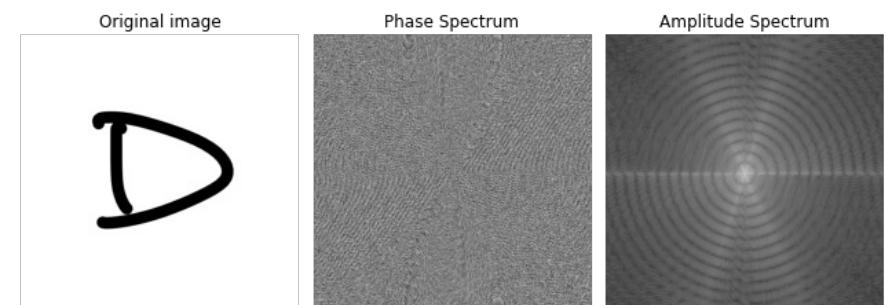
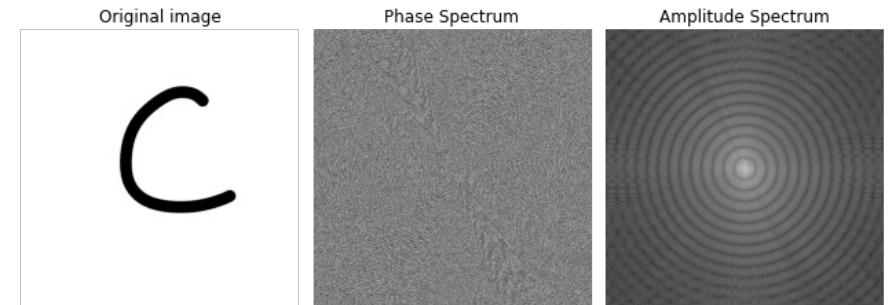
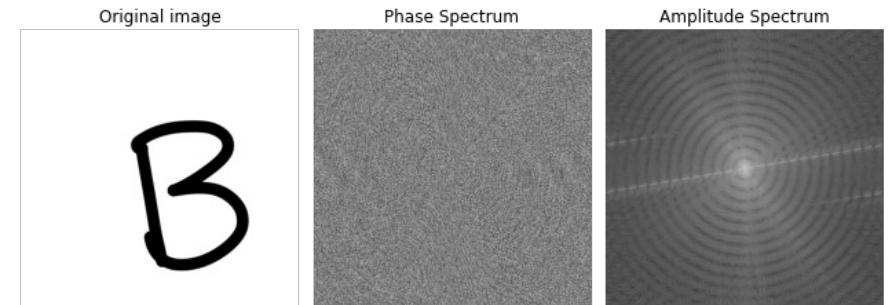
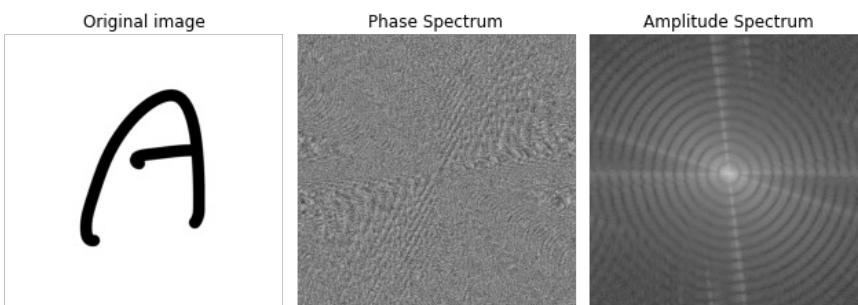
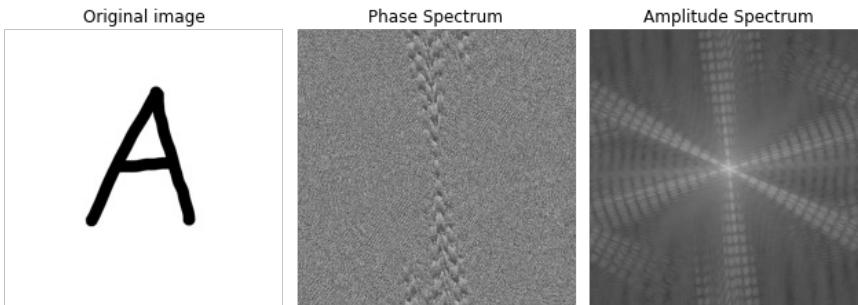
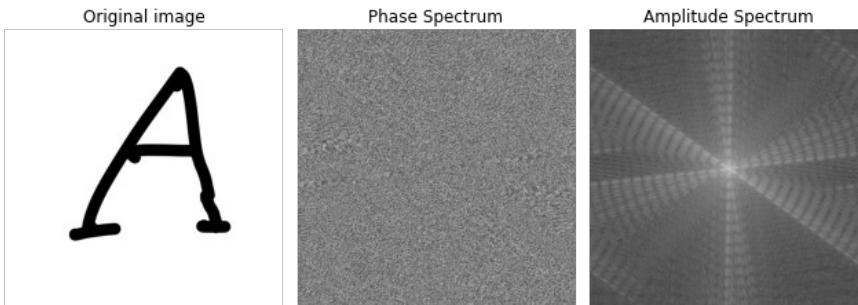
2D-Fourier Transform

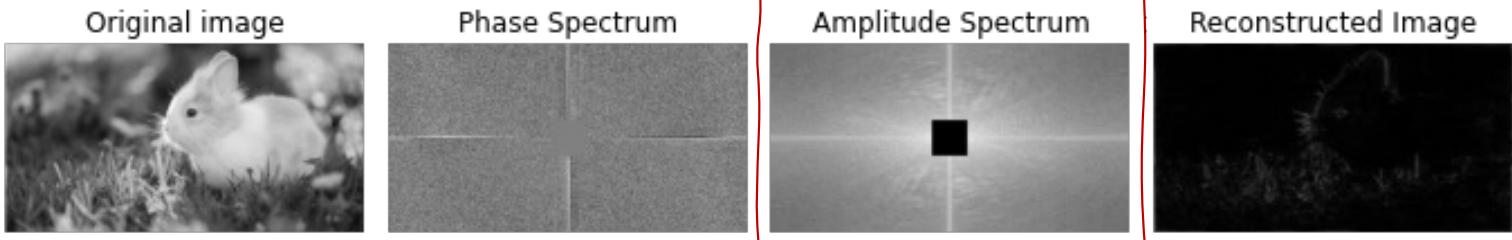
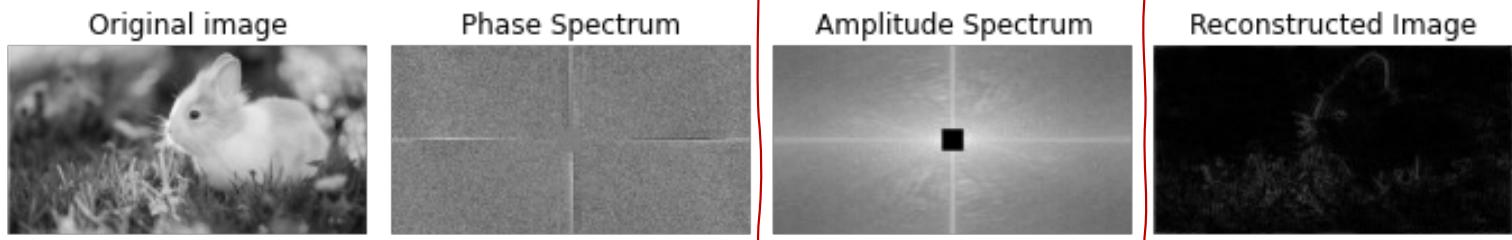
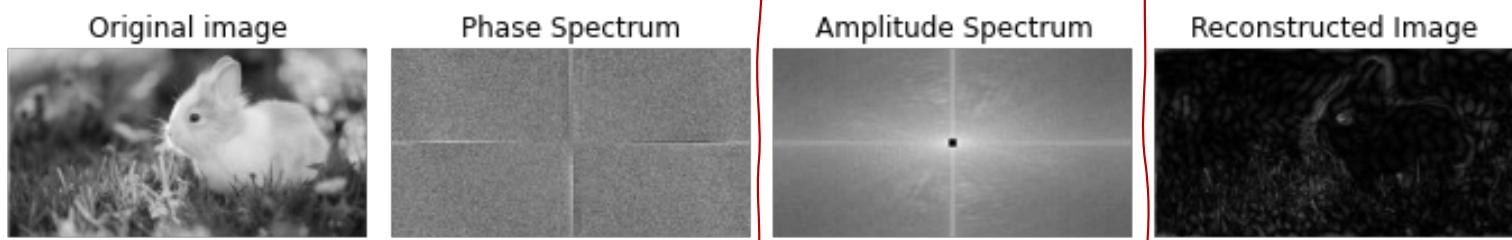
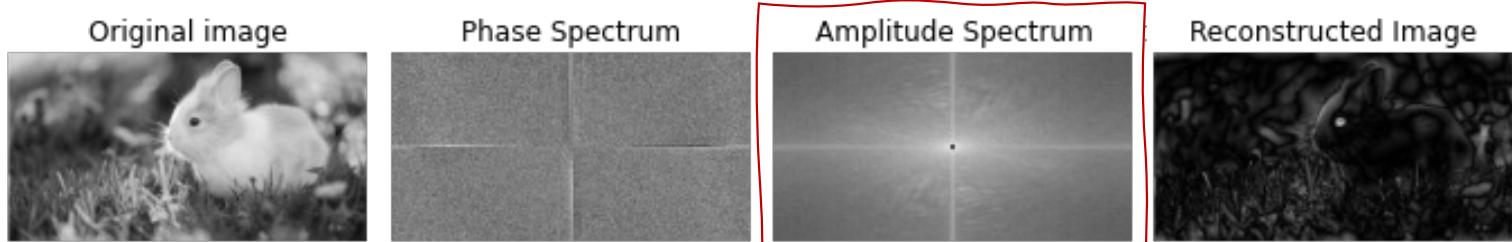


original

amplitude

phase

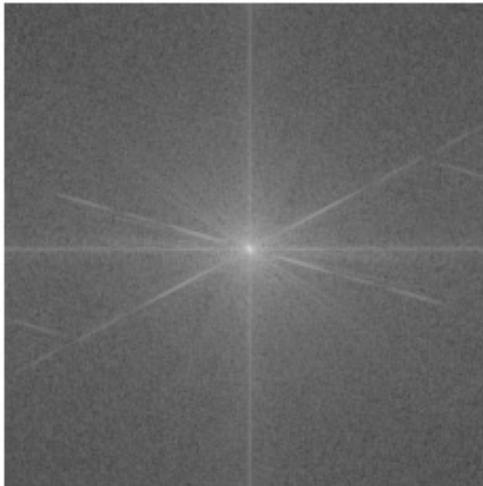




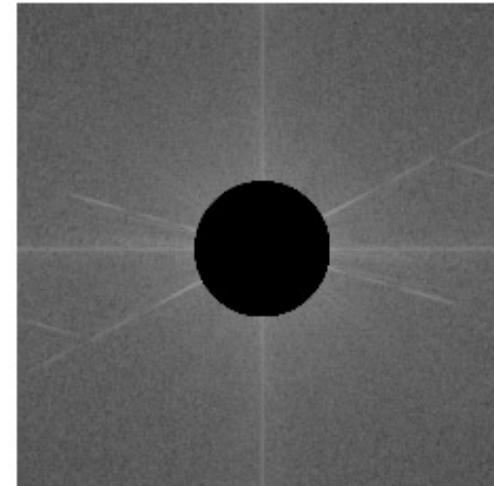
Original image



Log-scaled FFT magnitude



Low freqs removed



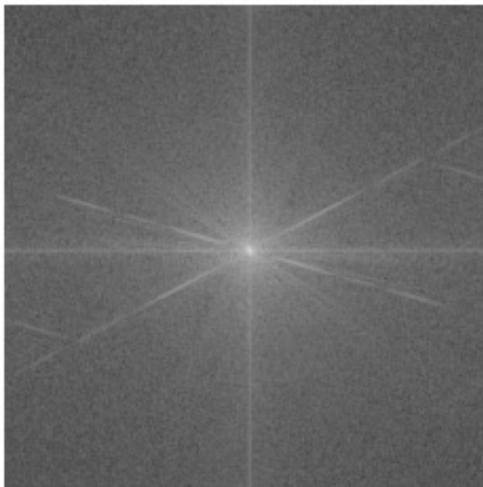
Visualized high frequencies



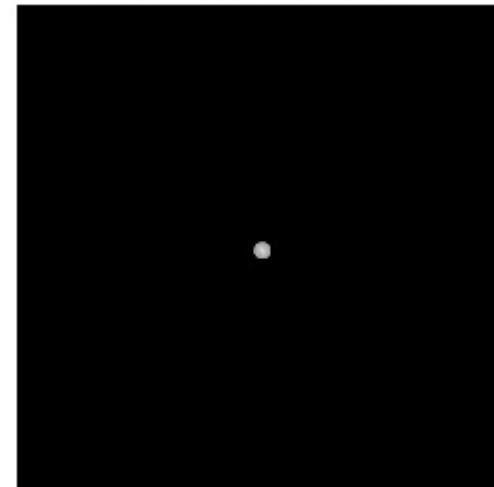
Original image



Log-scaled FFT magnitude



High freqs removed



Visualized low frequencies



Frequency-domain filtering

The convolution theorem

The Fourier transform of the convolution of two functions is the product of their Fourier transforms:

$$\mathcal{F}\{g * h\} = \mathcal{F}\{g\}\mathcal{F}\{h\}$$

The Fourier transform of the product of two functions is the convolution of the two Fourier transforms:

$$\mathcal{F}\{g \cdot h\} = \mathcal{F}\{g\} * \mathcal{F}\{h\}$$

Convolution in spatial domain is equivalent to multiplication in frequency domain!

$$g * h = \mathcal{F}^{-1} \{\mathcal{F}\{g\} \cdot \mathcal{F}\{h\}\}$$

$$g \cdot h = \mathcal{F}^{-1} \{\mathcal{F}\{g\} * \mathcal{F}\{h\}\}$$

What do we use convolution for?

Convolution for 1D continuous signals

Definition of linear shift-invariant filtering as convolution:

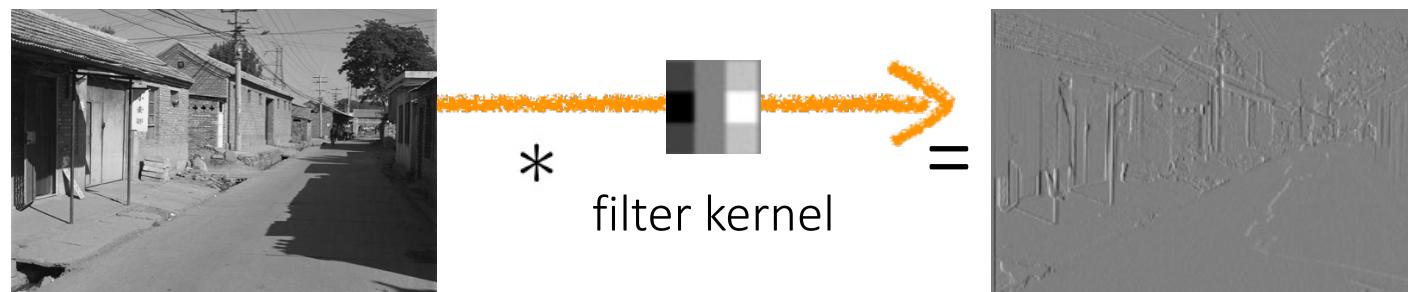
$$(f * g)(x) = \int_{-\infty}^{\infty} f(y)g(x - y)dy$$

filtered signal filter input signal

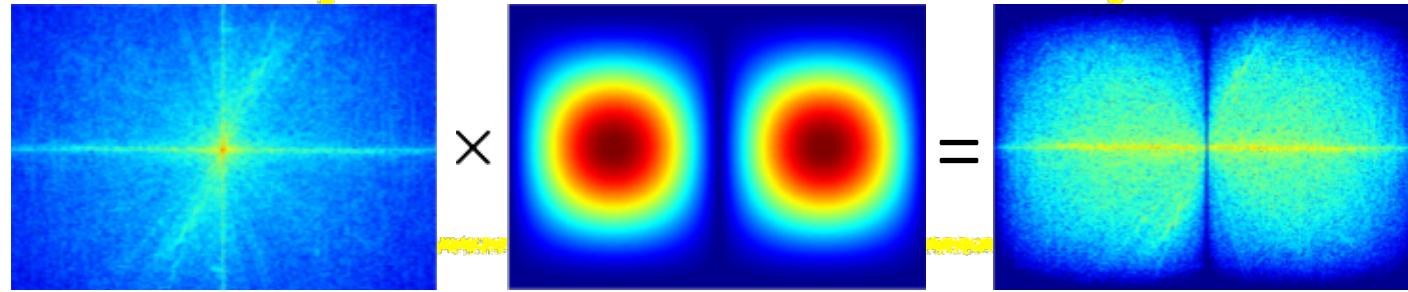
Using the convolution theorem, we can interpret and implement all types of linear shift-invariant filtering as multiplication in frequency domain.

Why implement convolution in frequency domain?

Spatial domain filtering



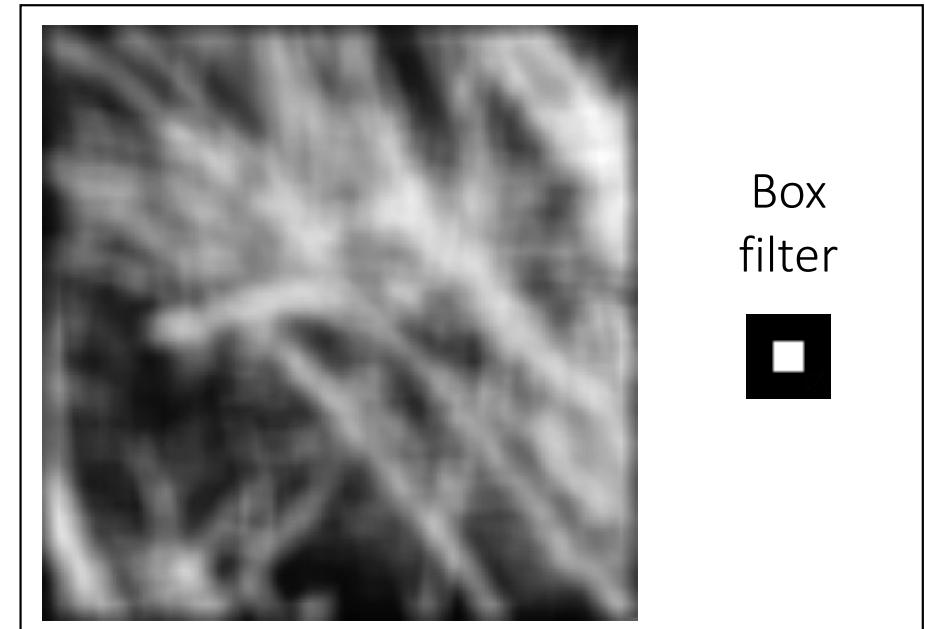
Fourier transform



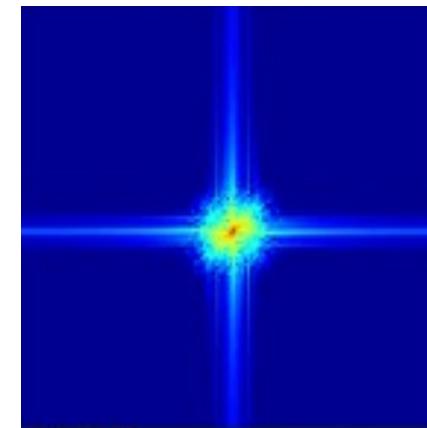
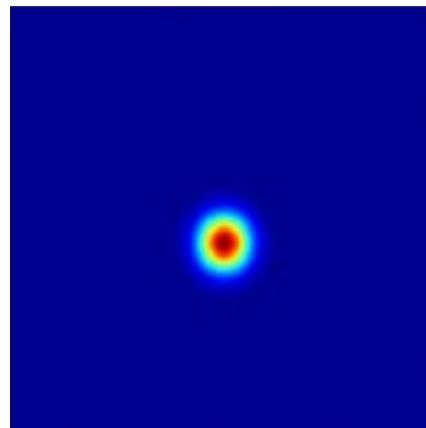
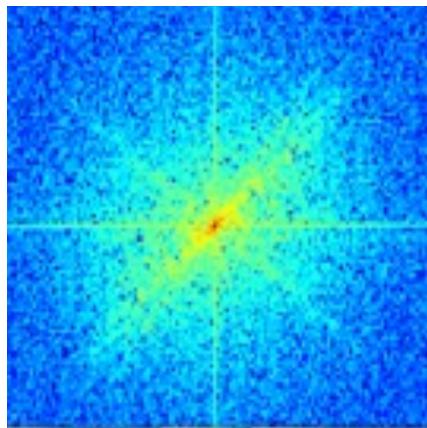
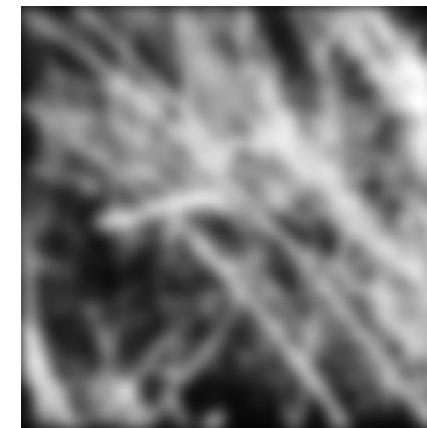
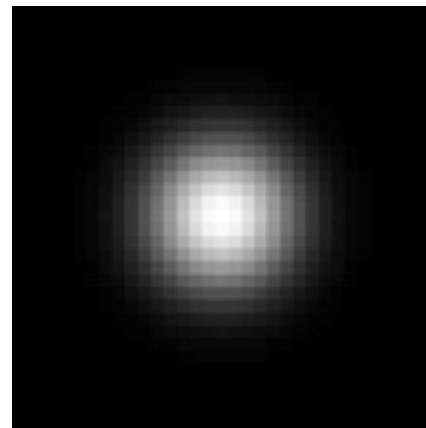
Frequency domain filtering

Revisiting blurring

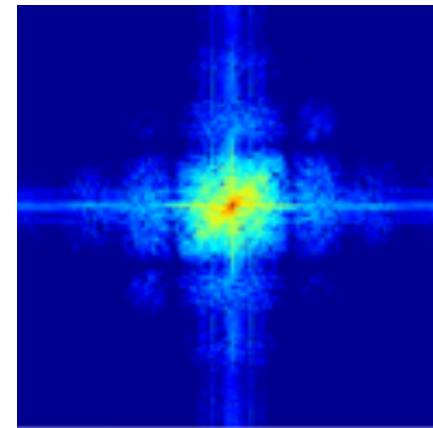
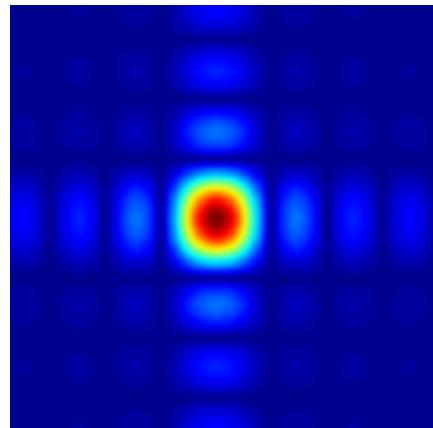
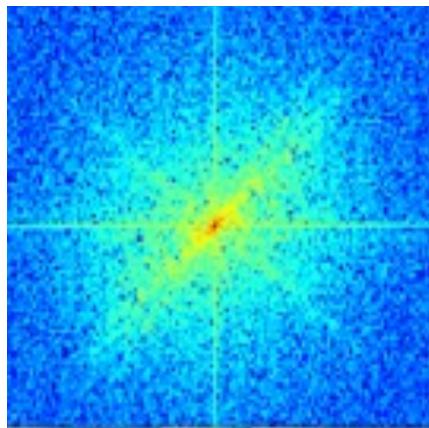
Why does the Gaussian give a nice smooth image, but the square filter give edgy artifacts?



Gaussian blur

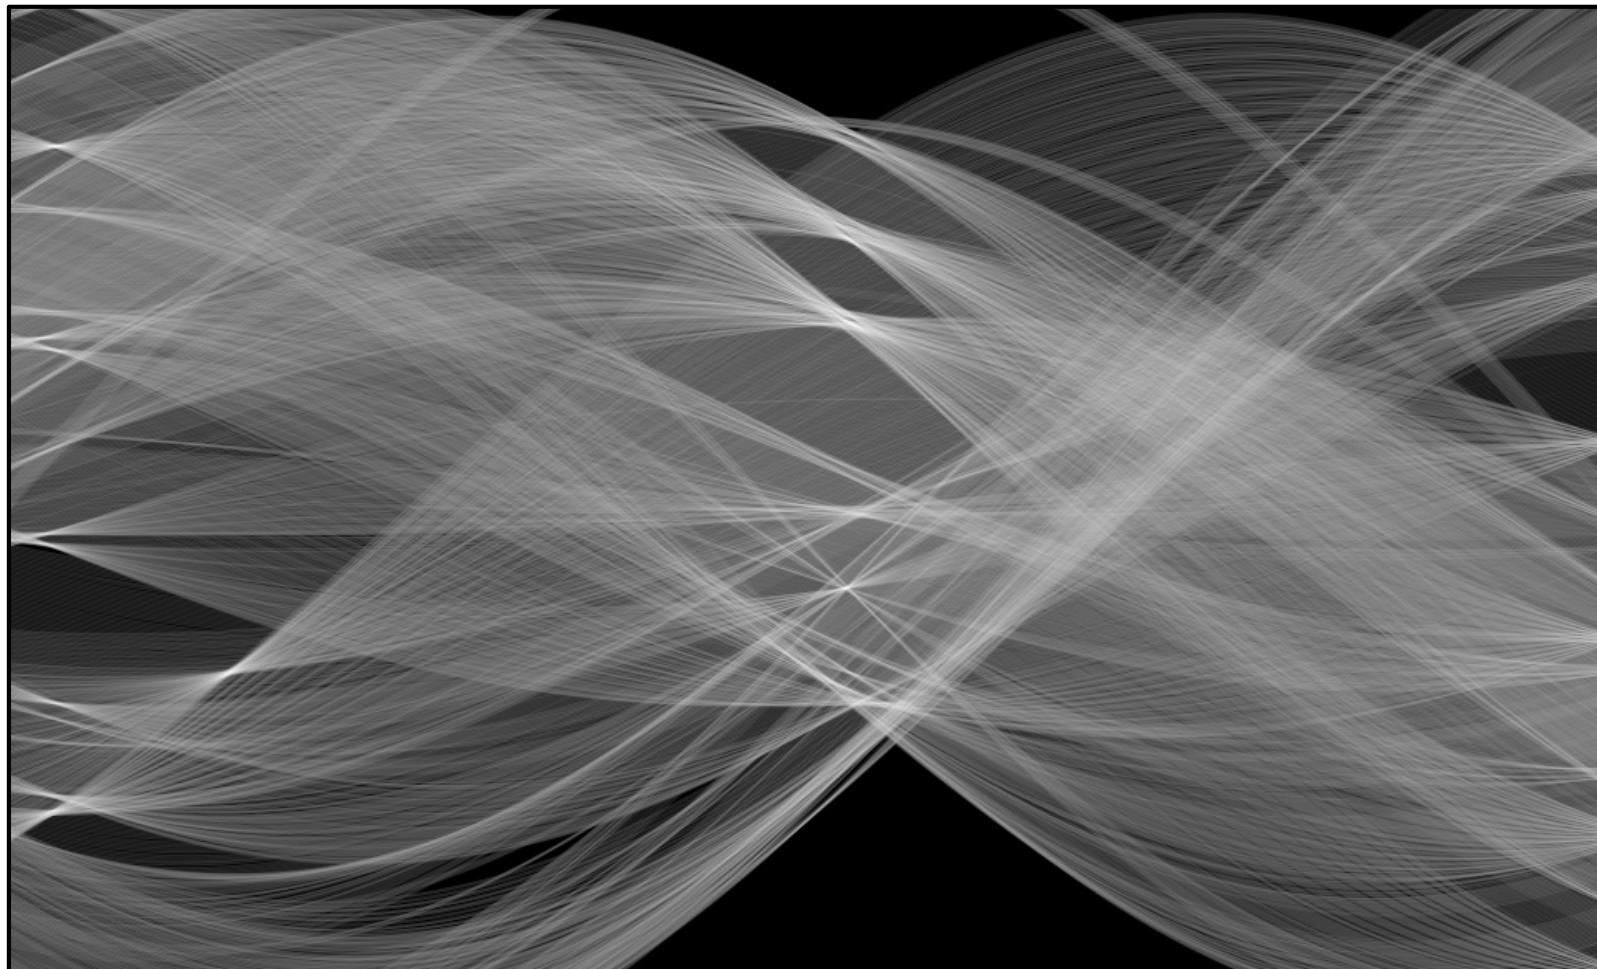


Box blur



Demo: Python Program

Hough transform

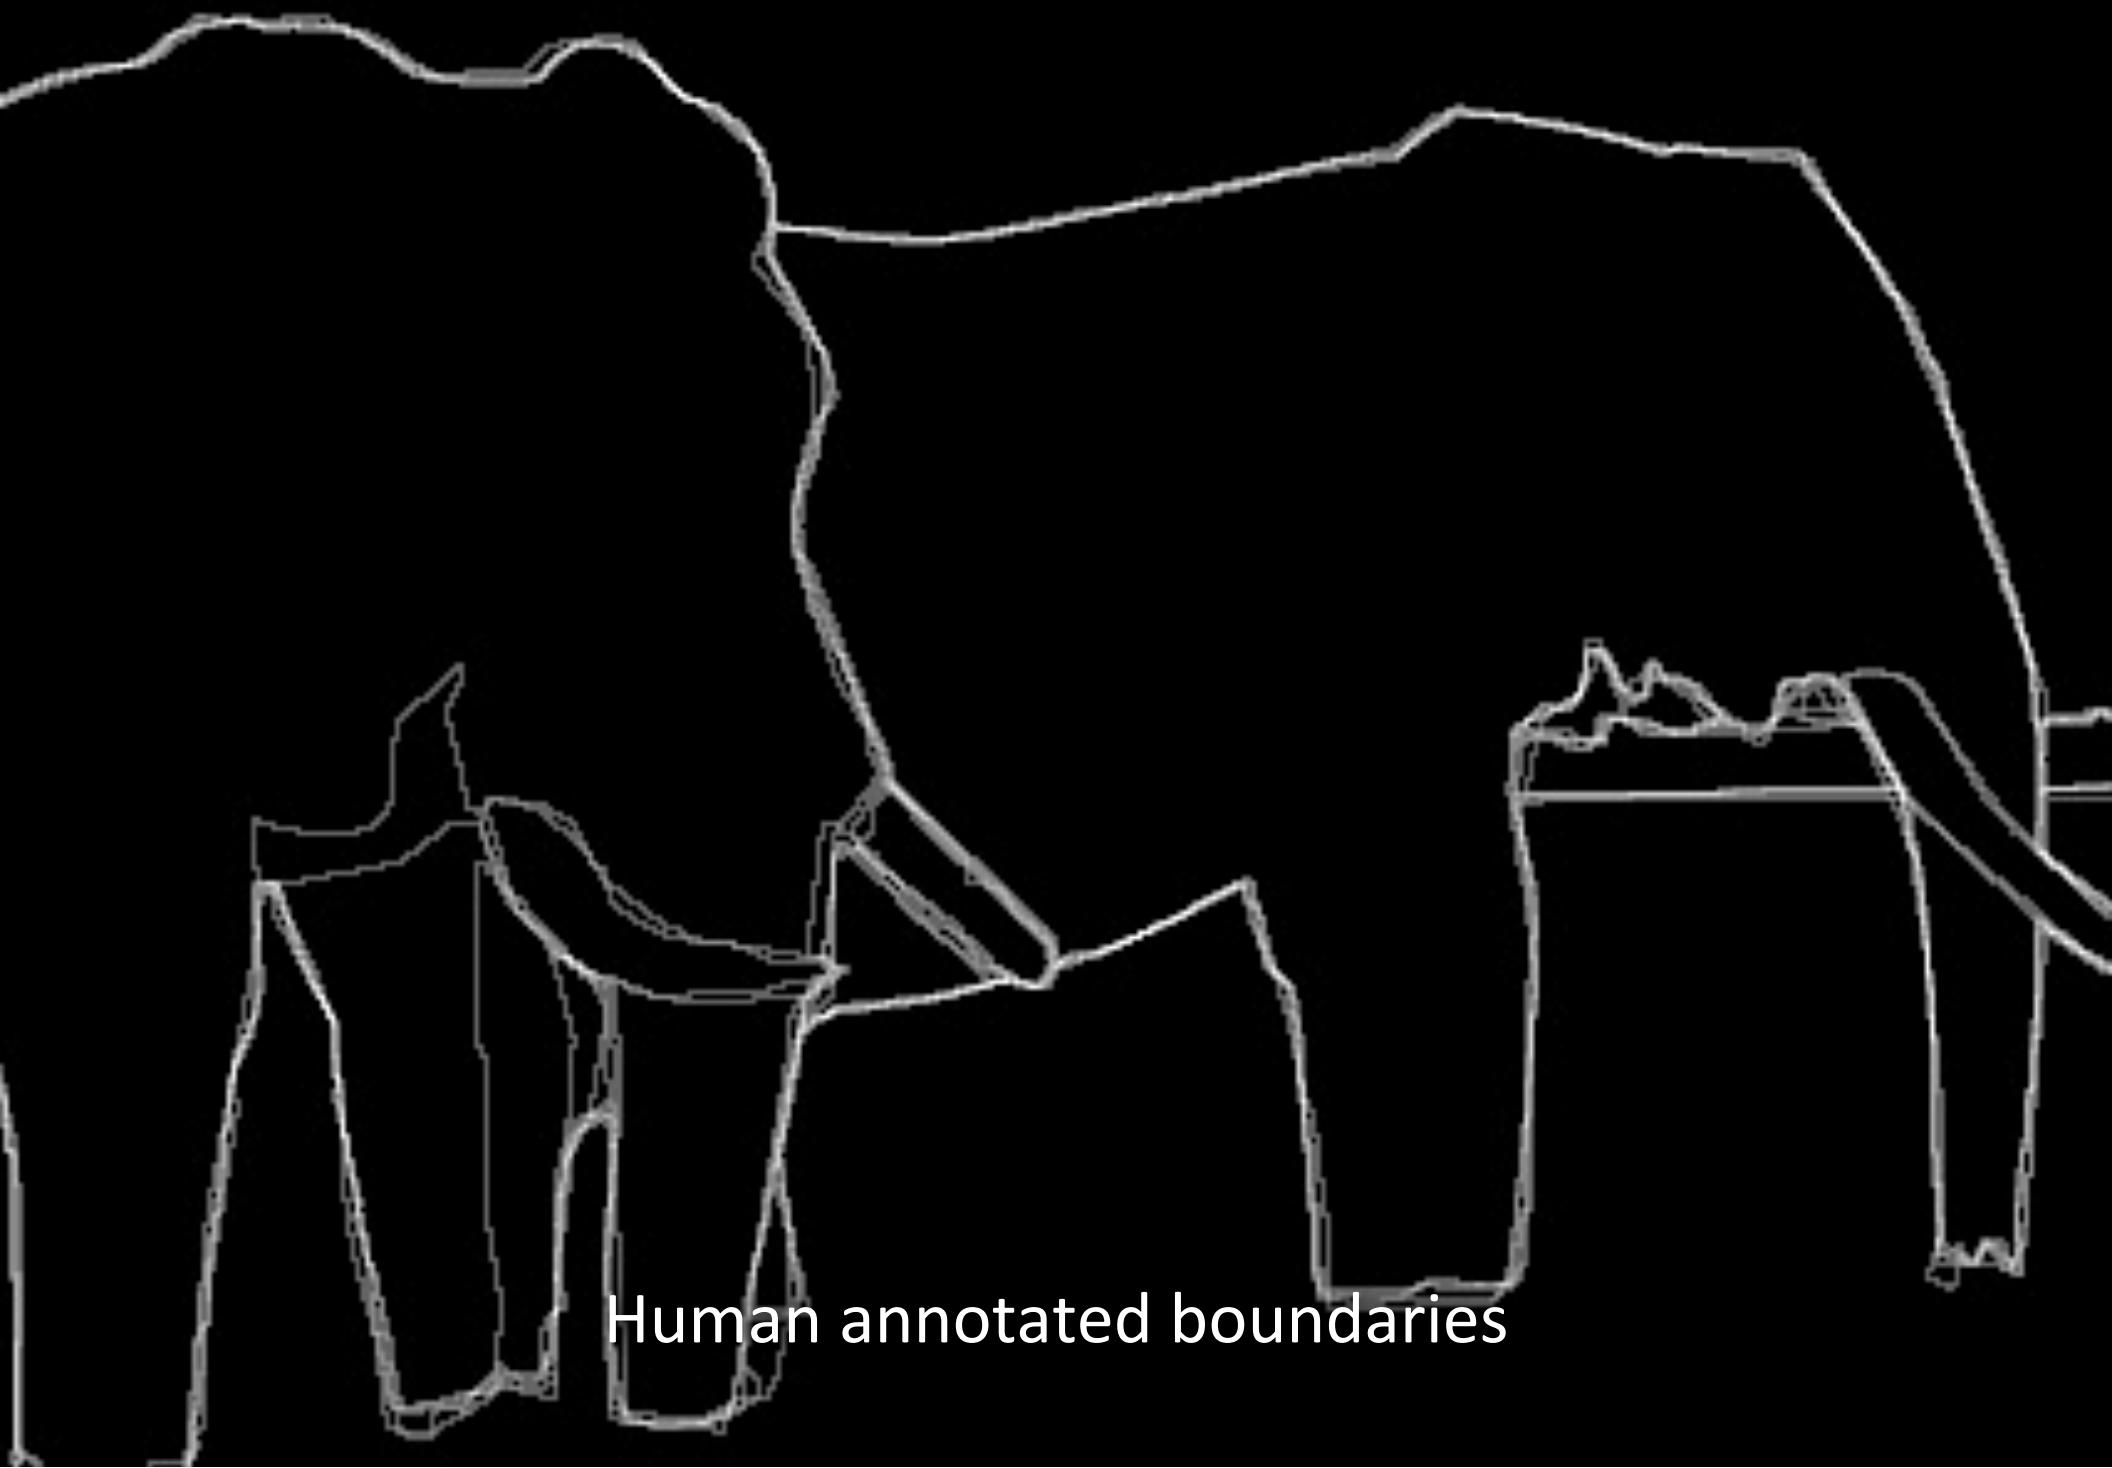


Computer Vision
Fall 2022, Lecture 4

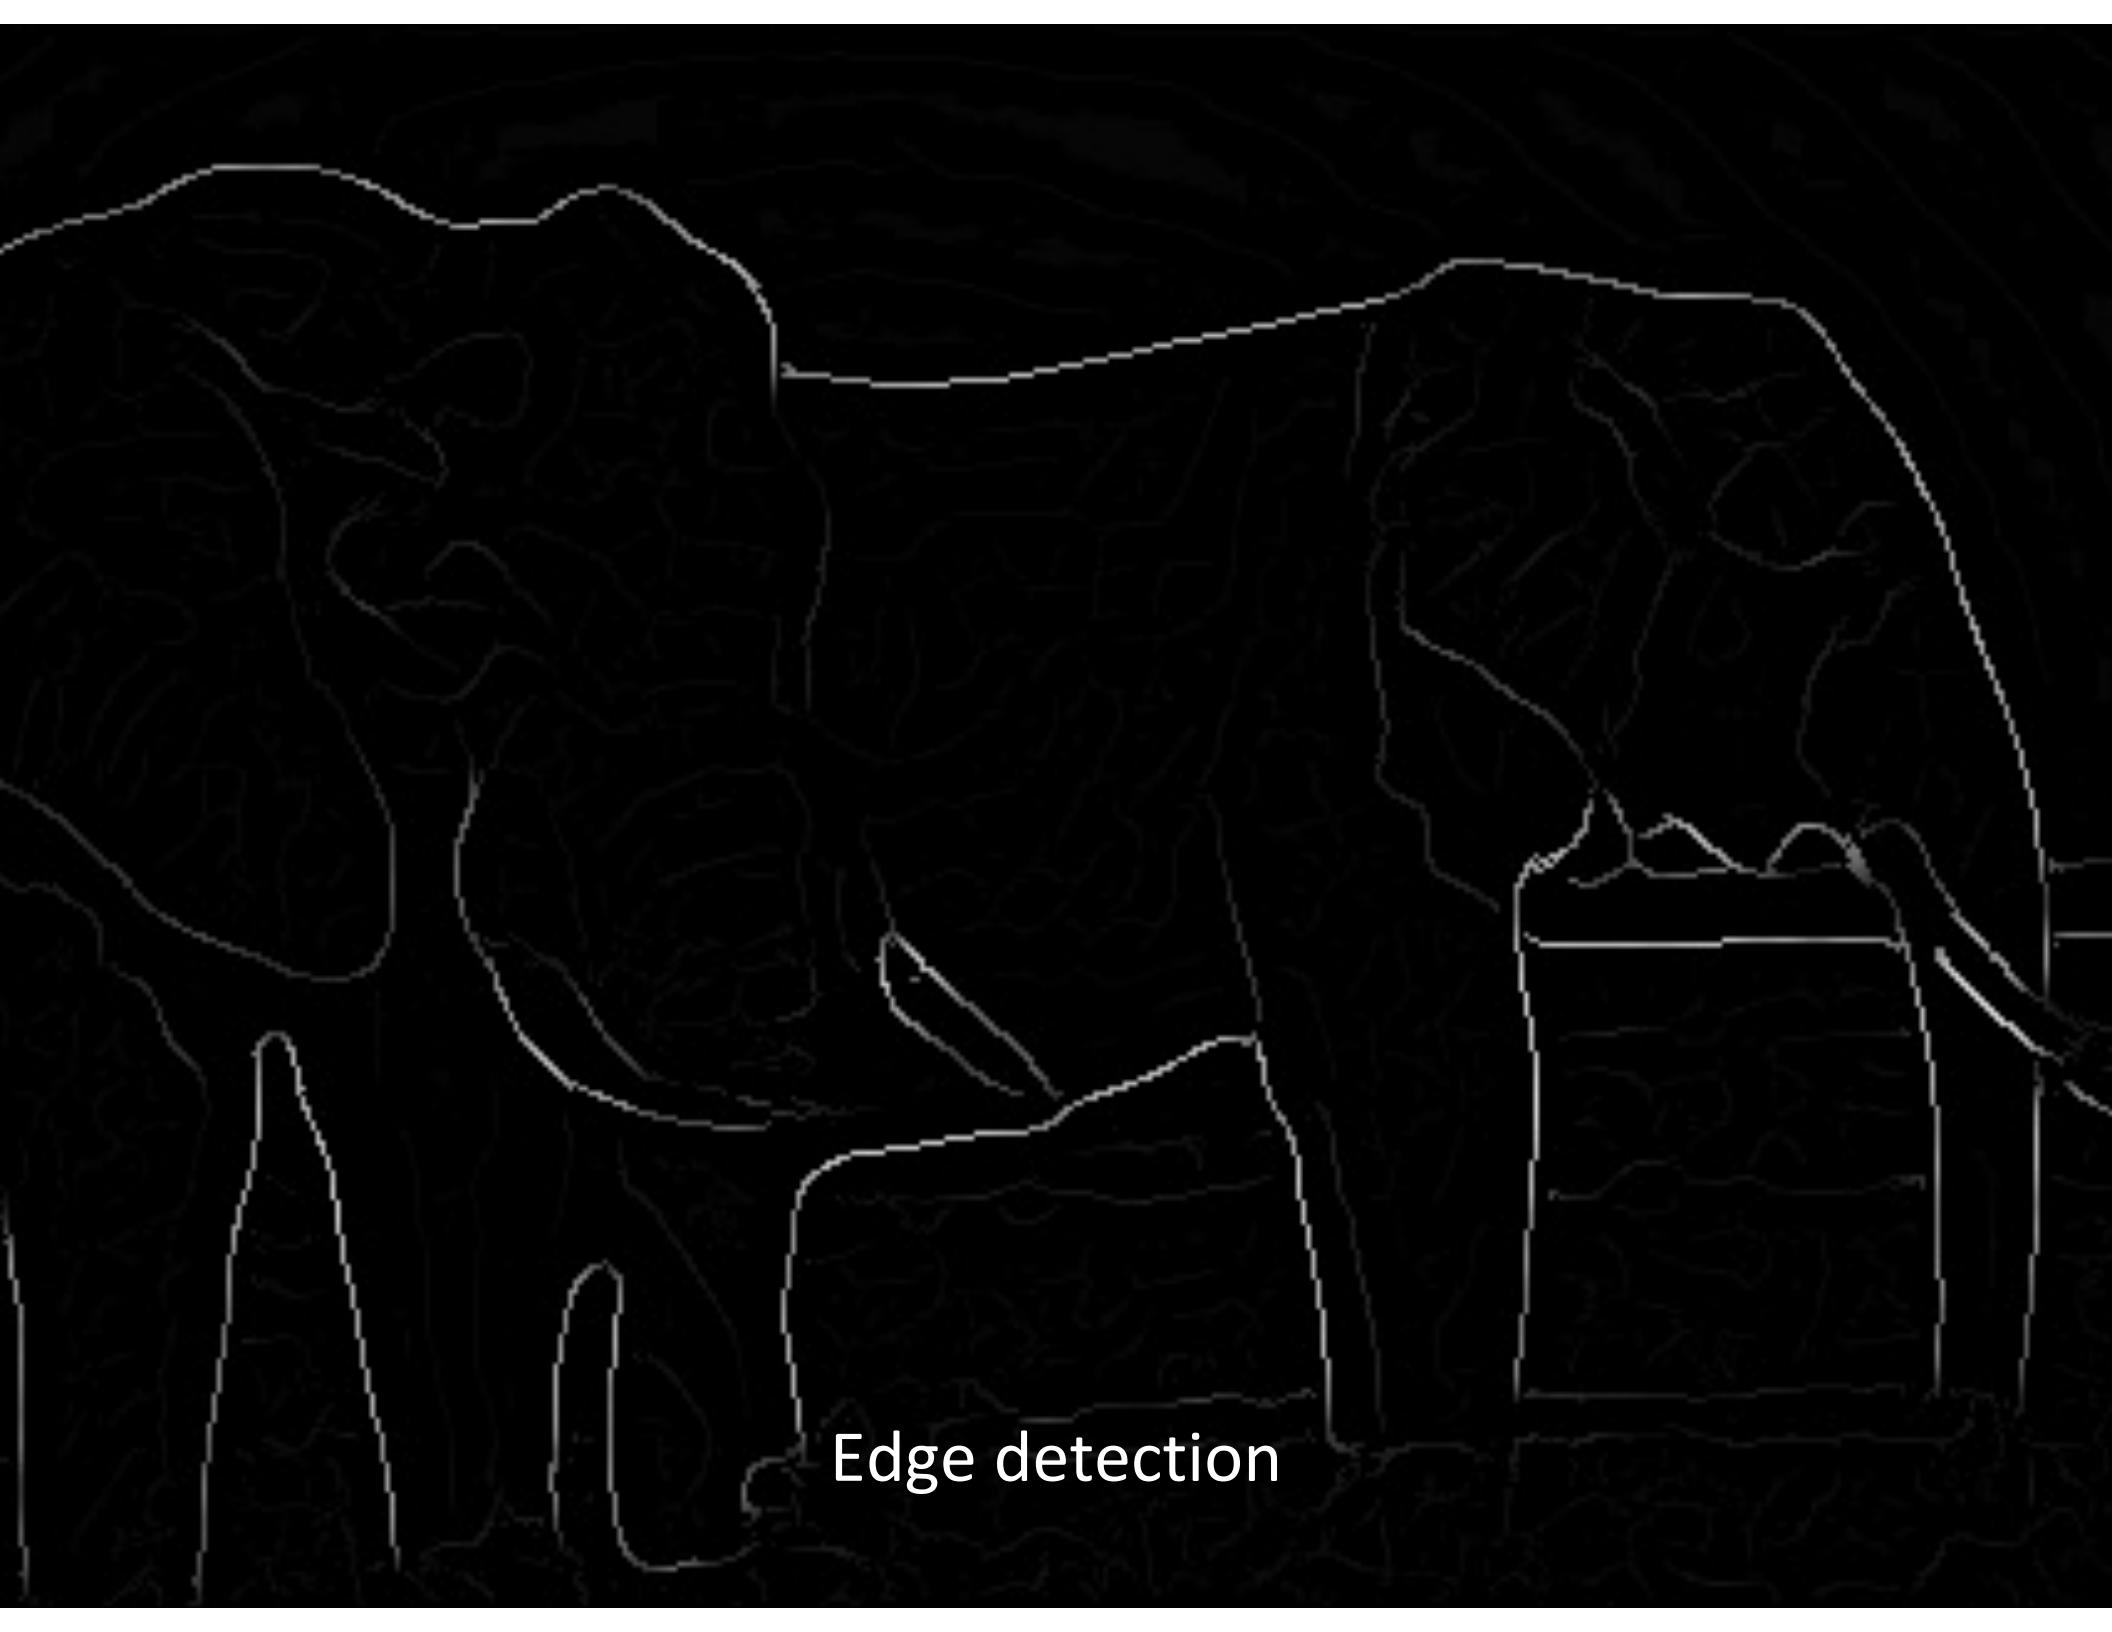
Finding boundaries



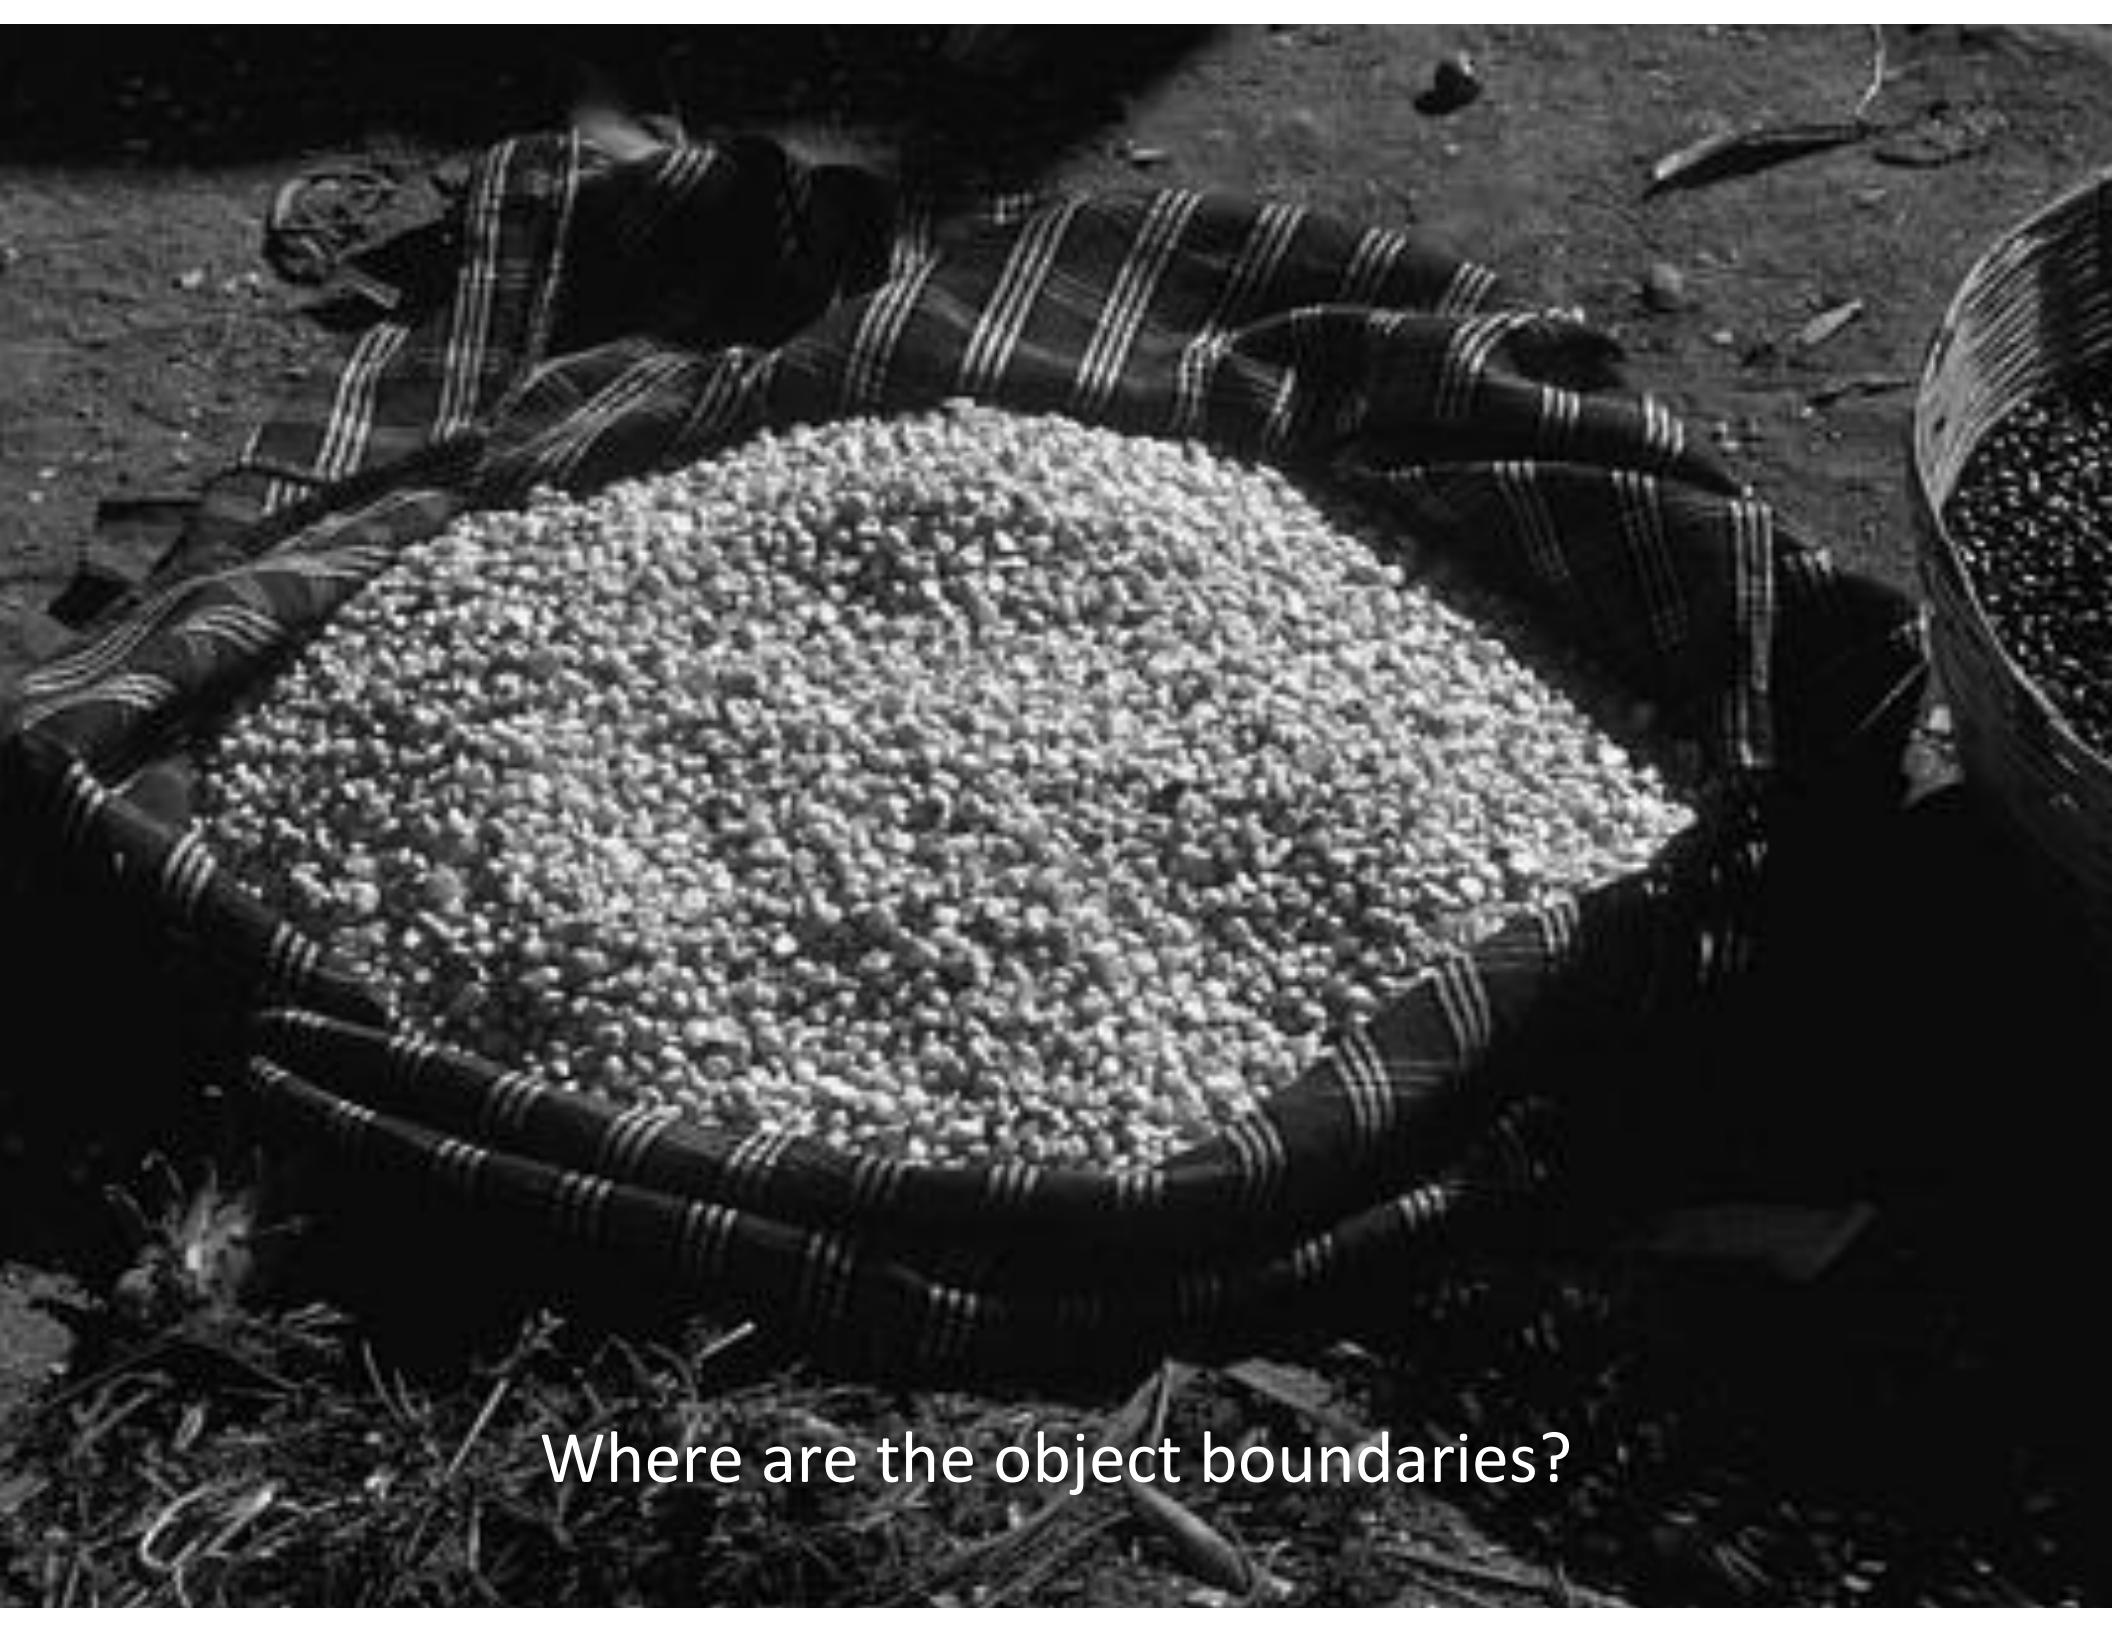
Where are the object boundaries?



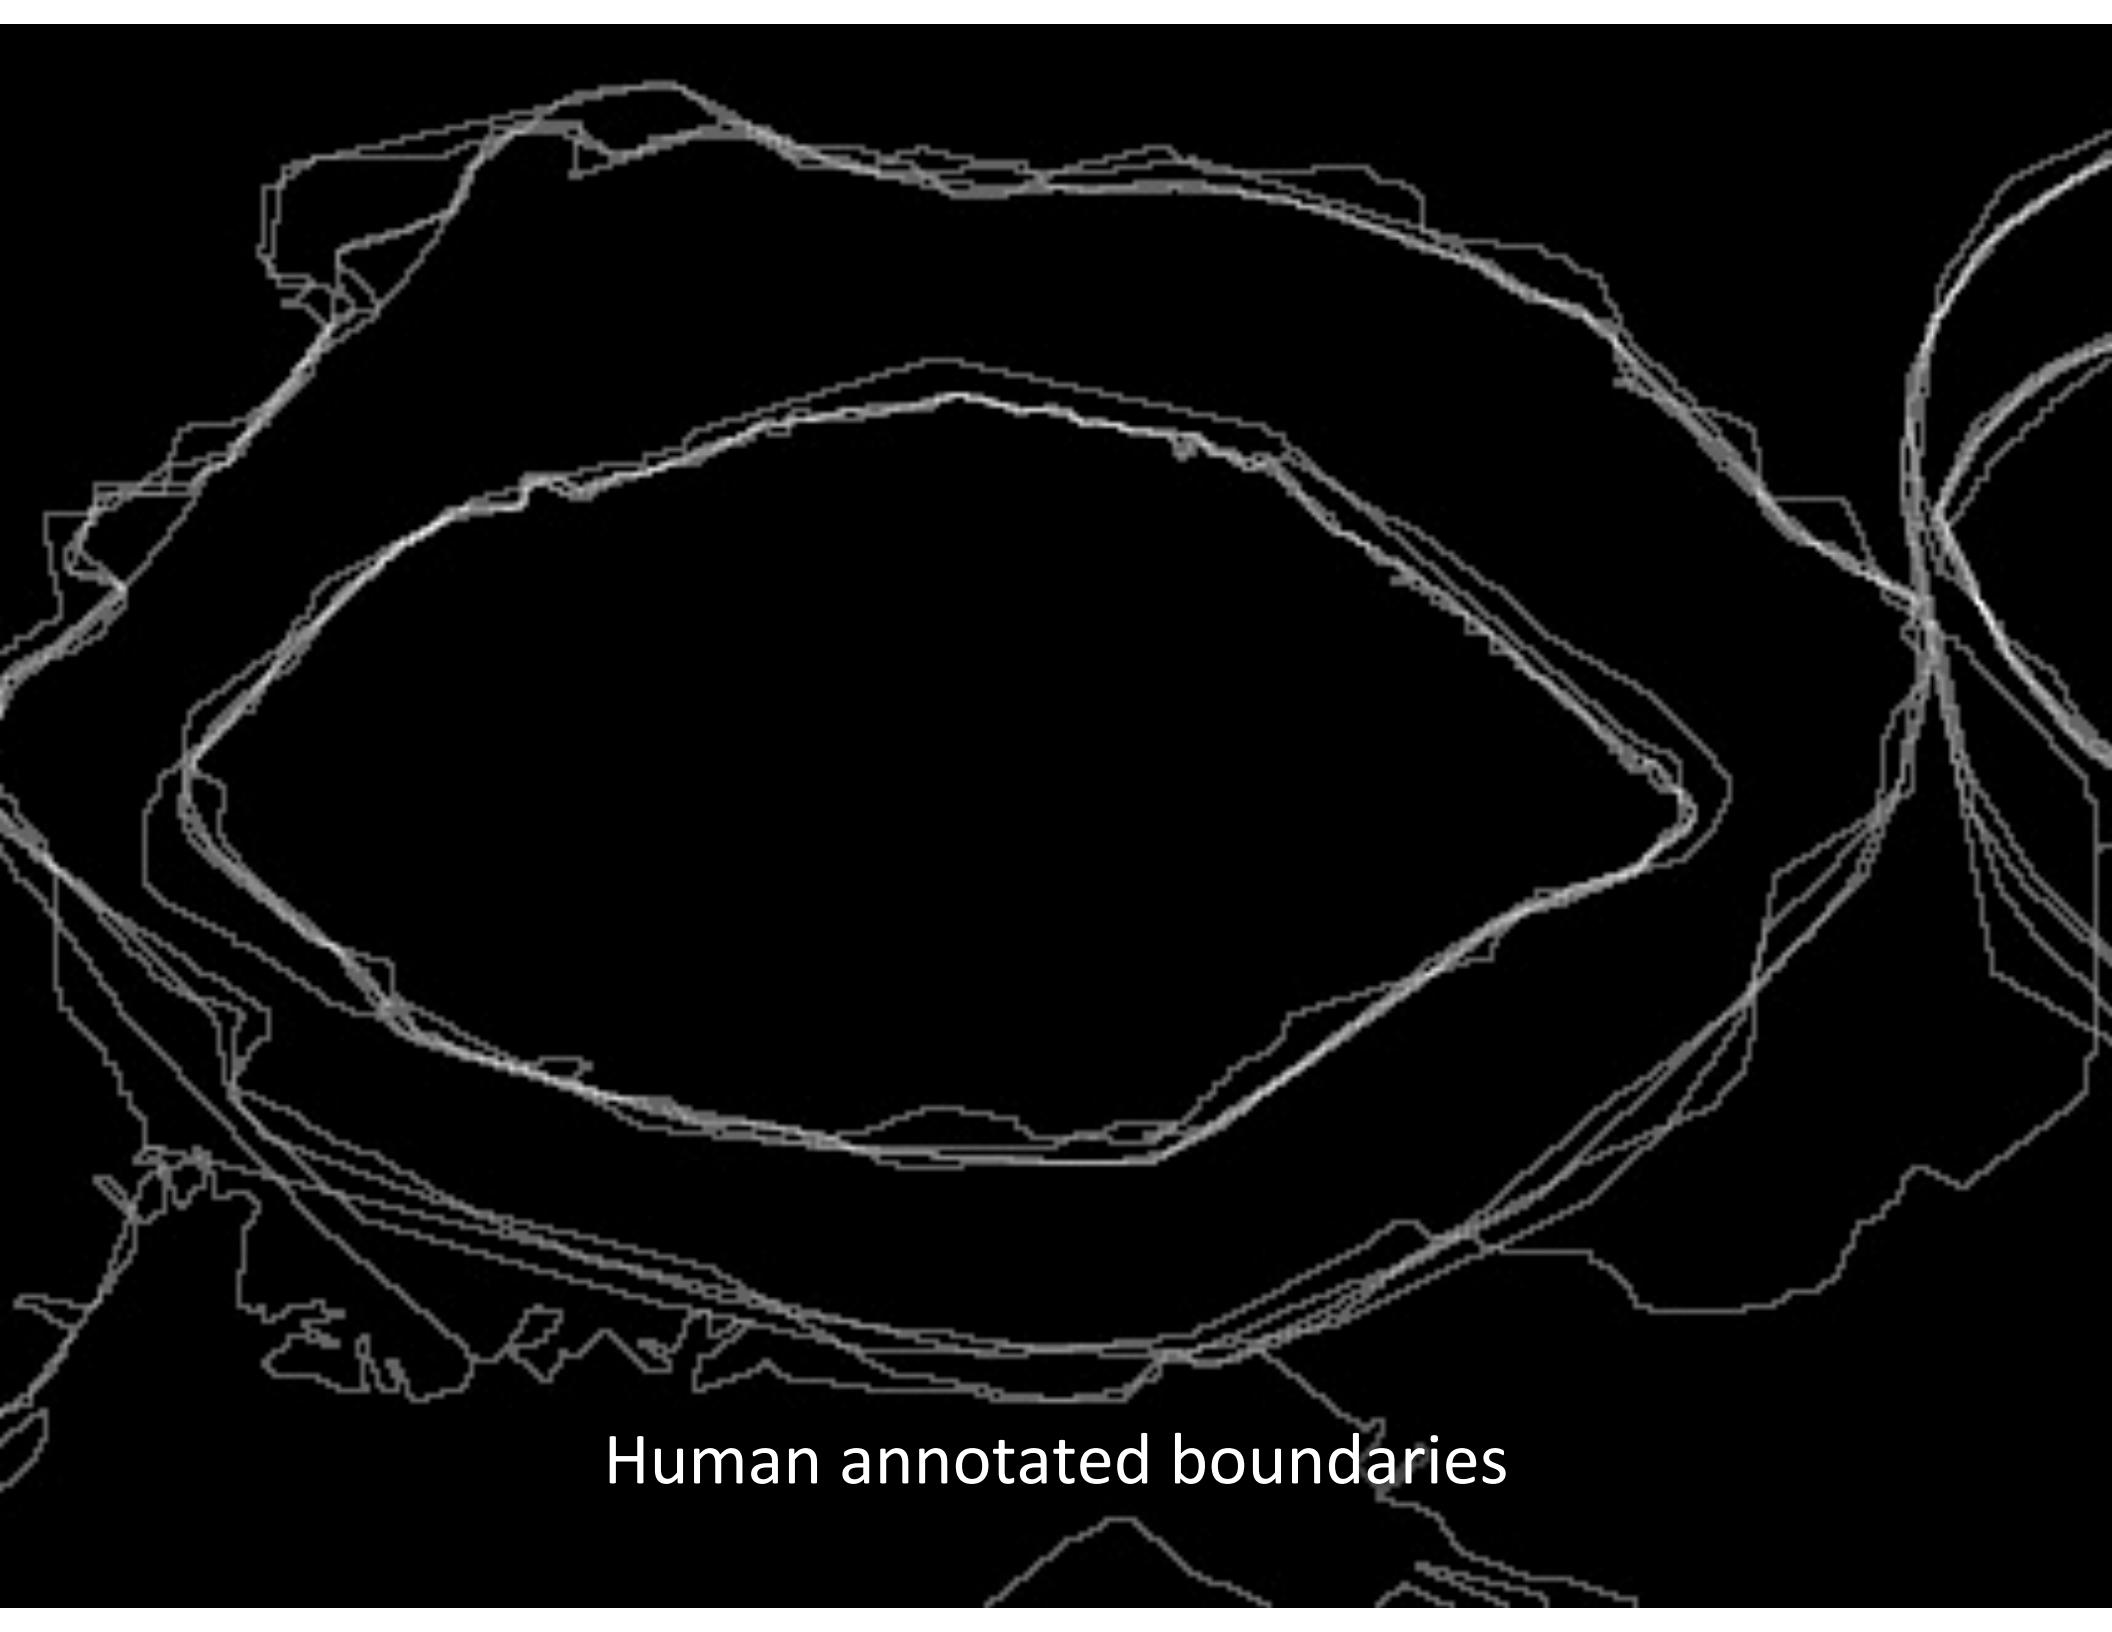
Human annotated boundaries



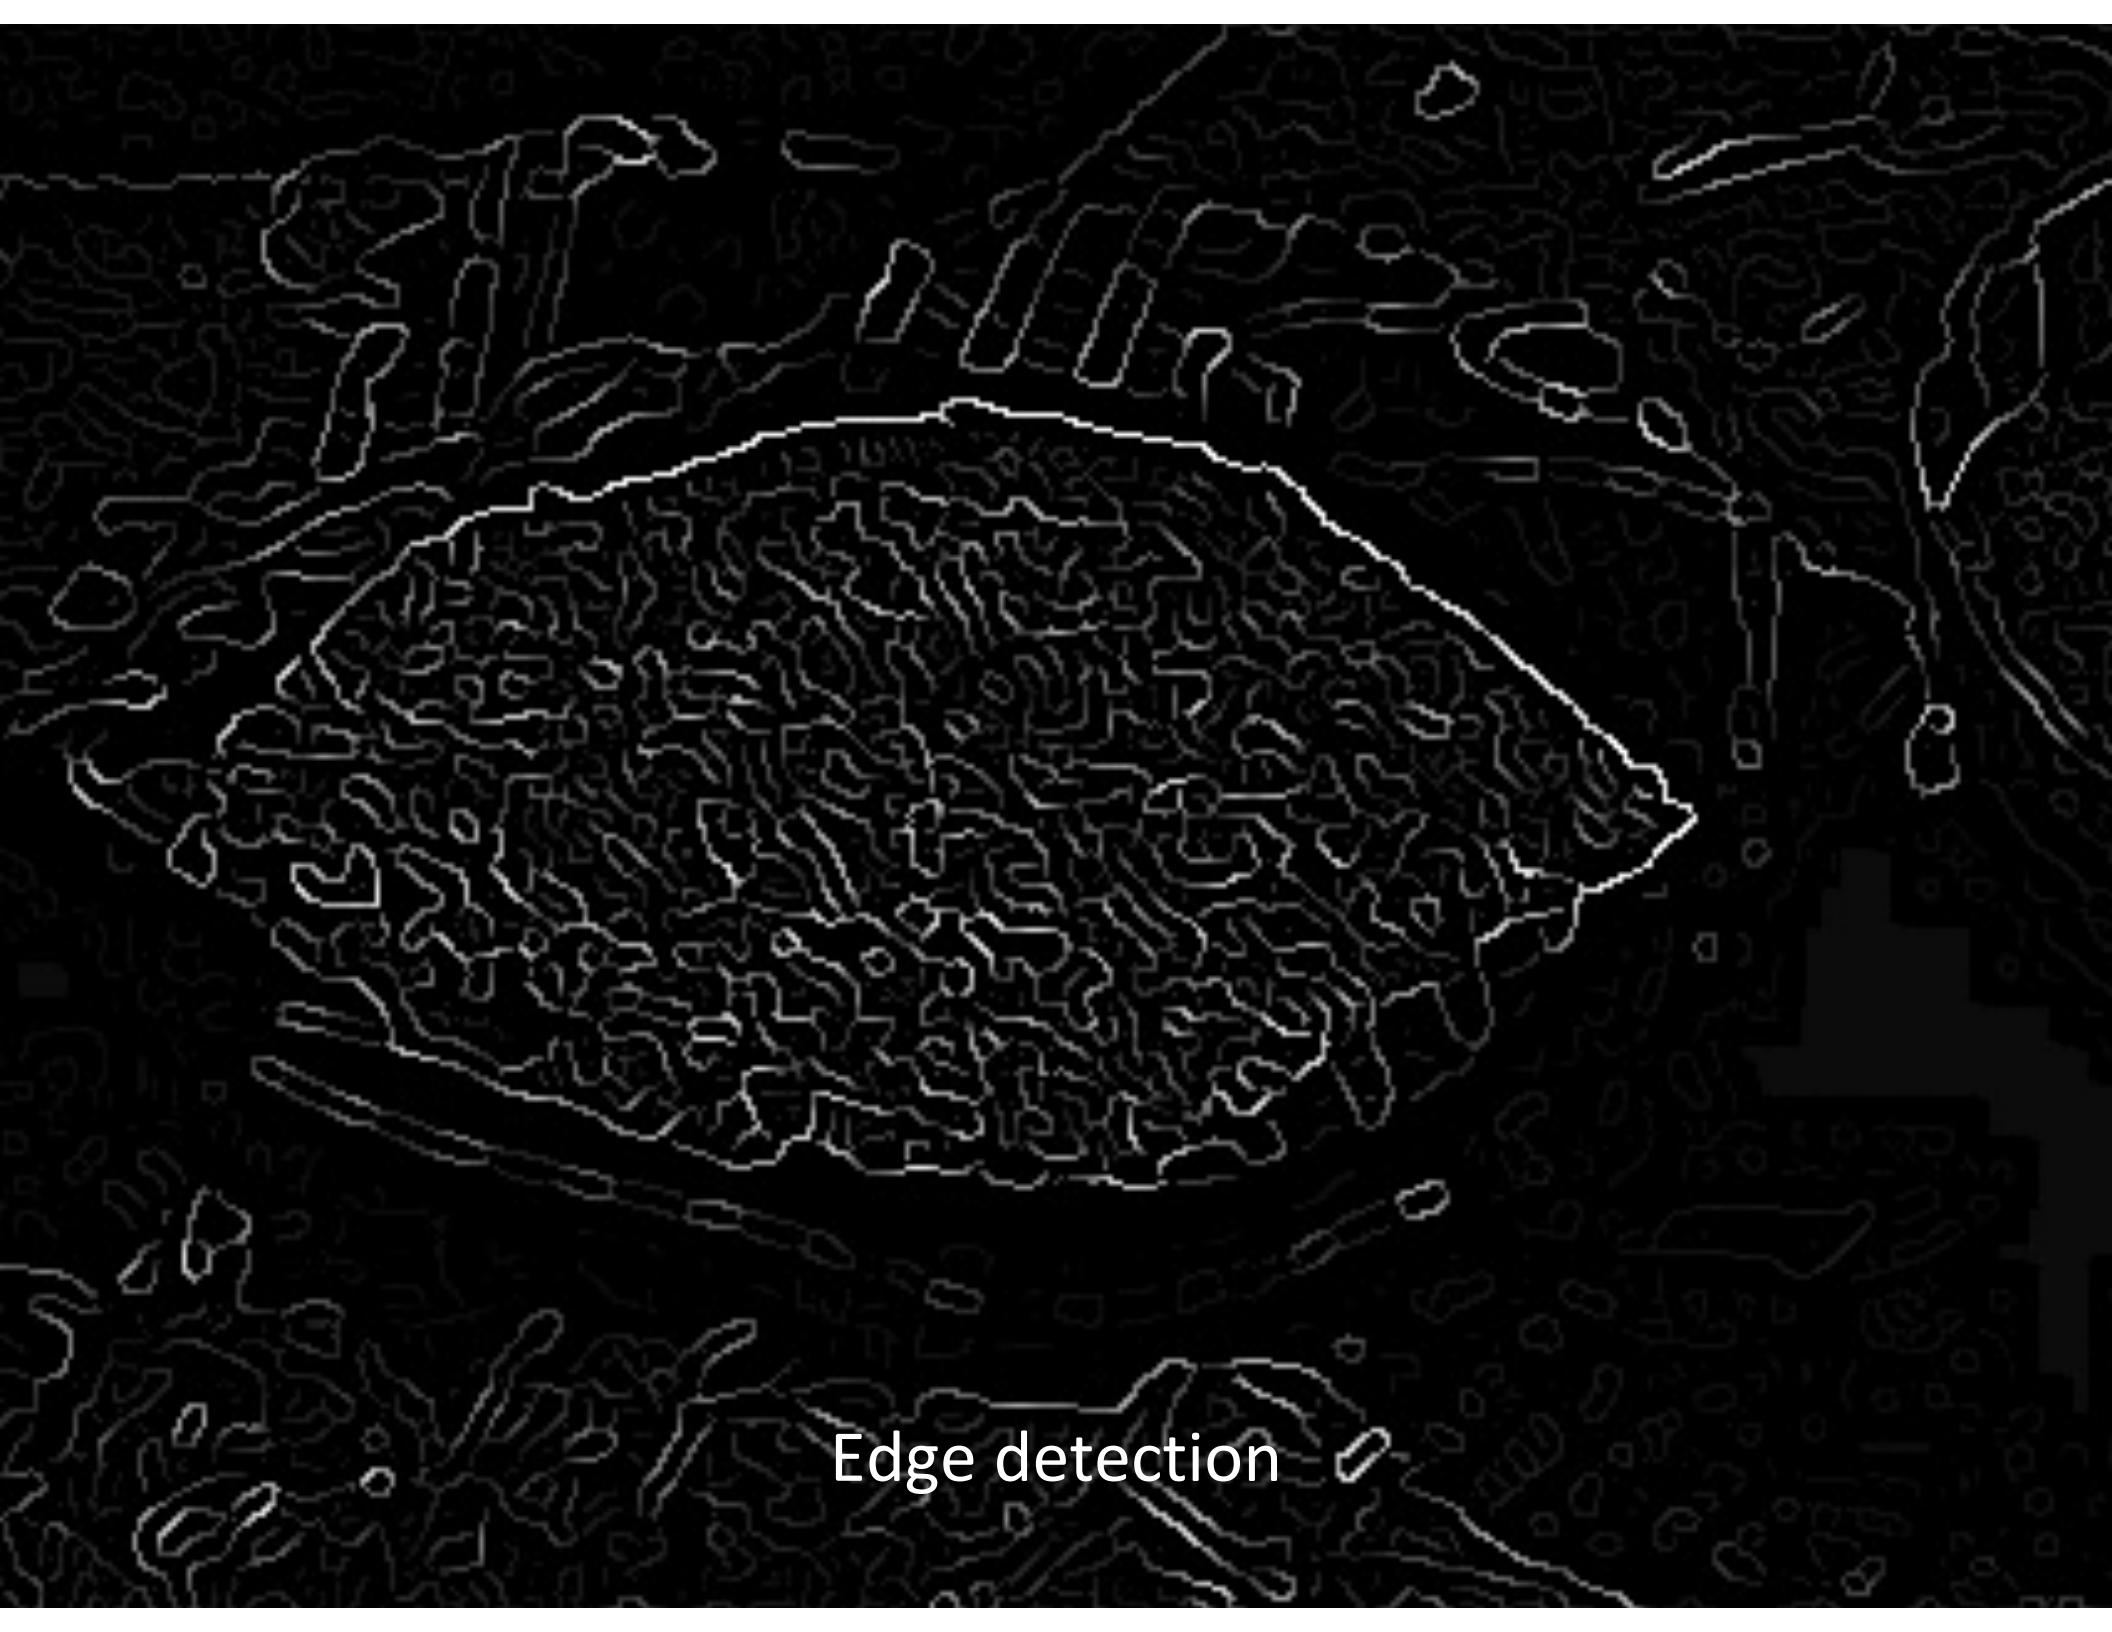
Edge detection



Where are the object boundaries?



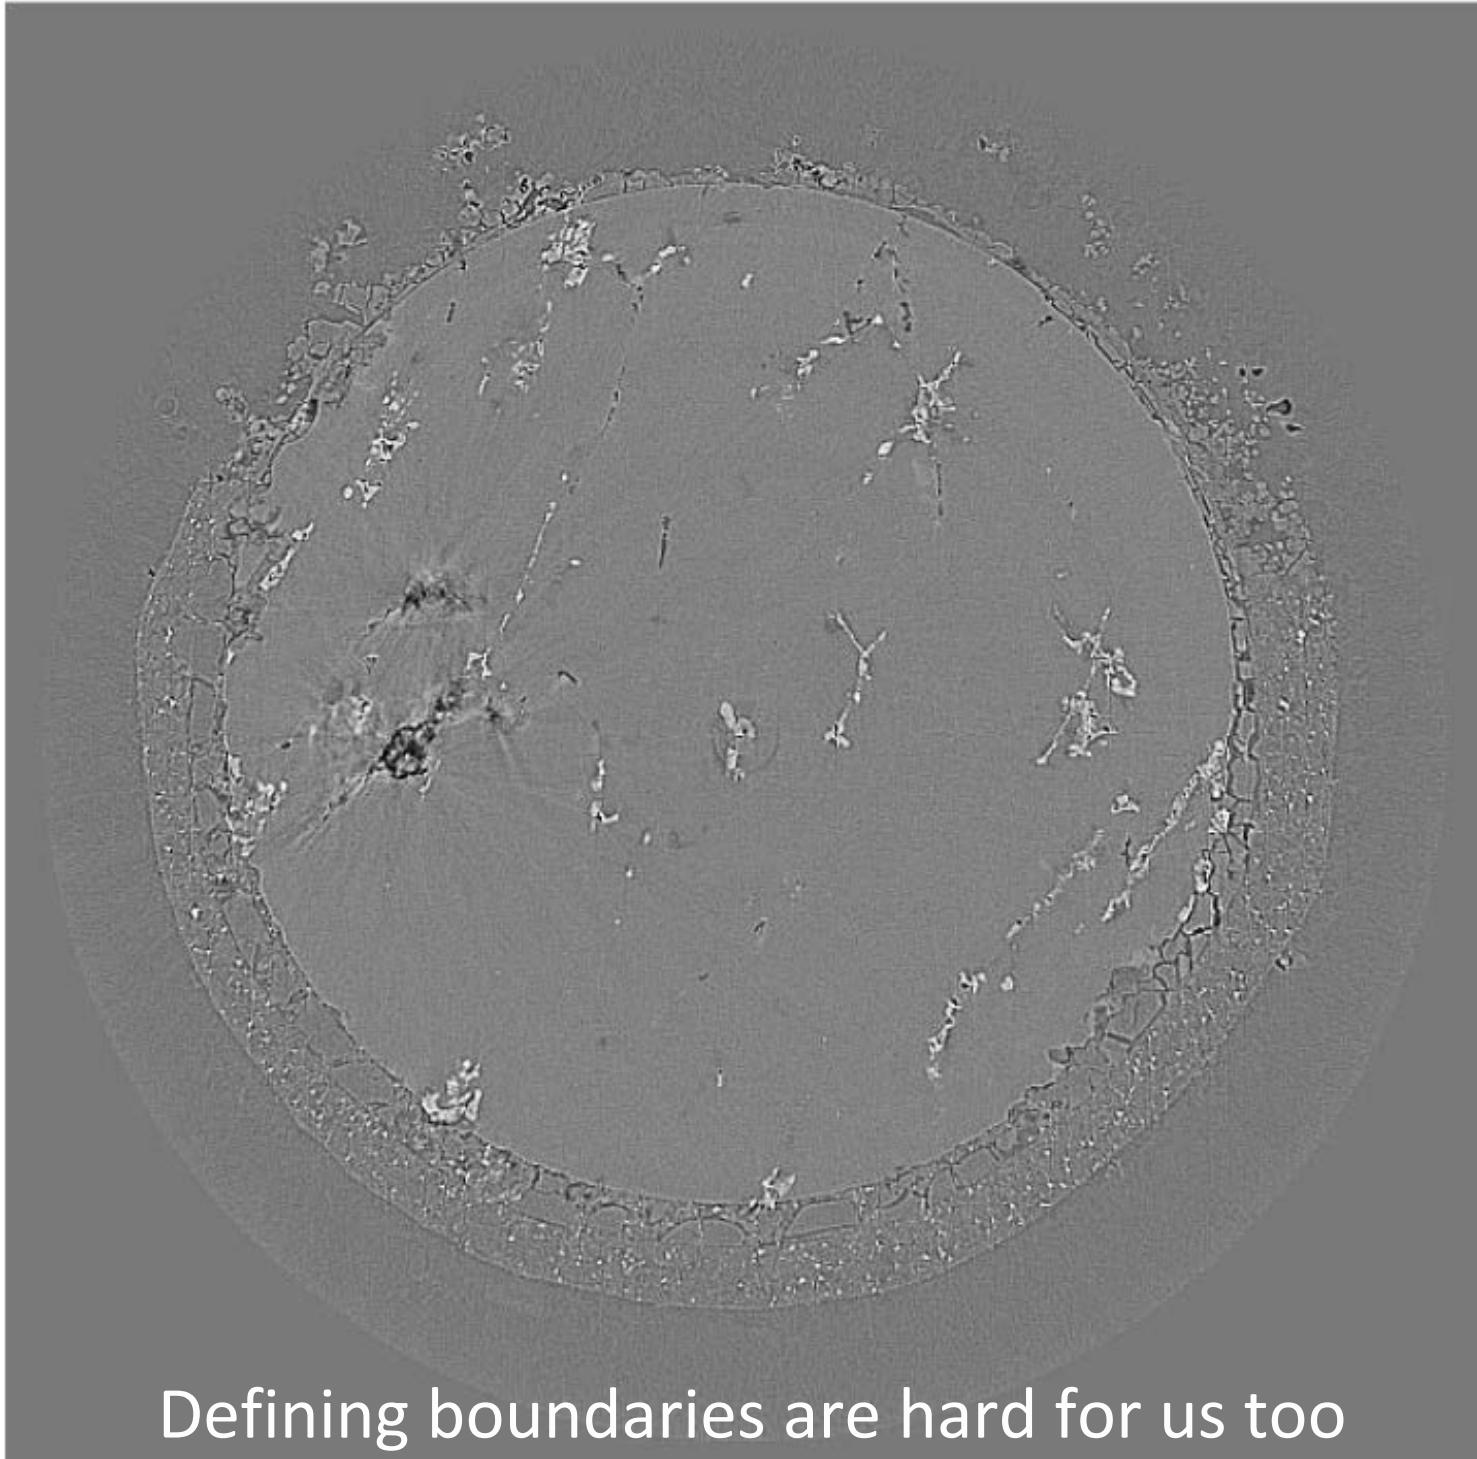
Human annotated boundaries



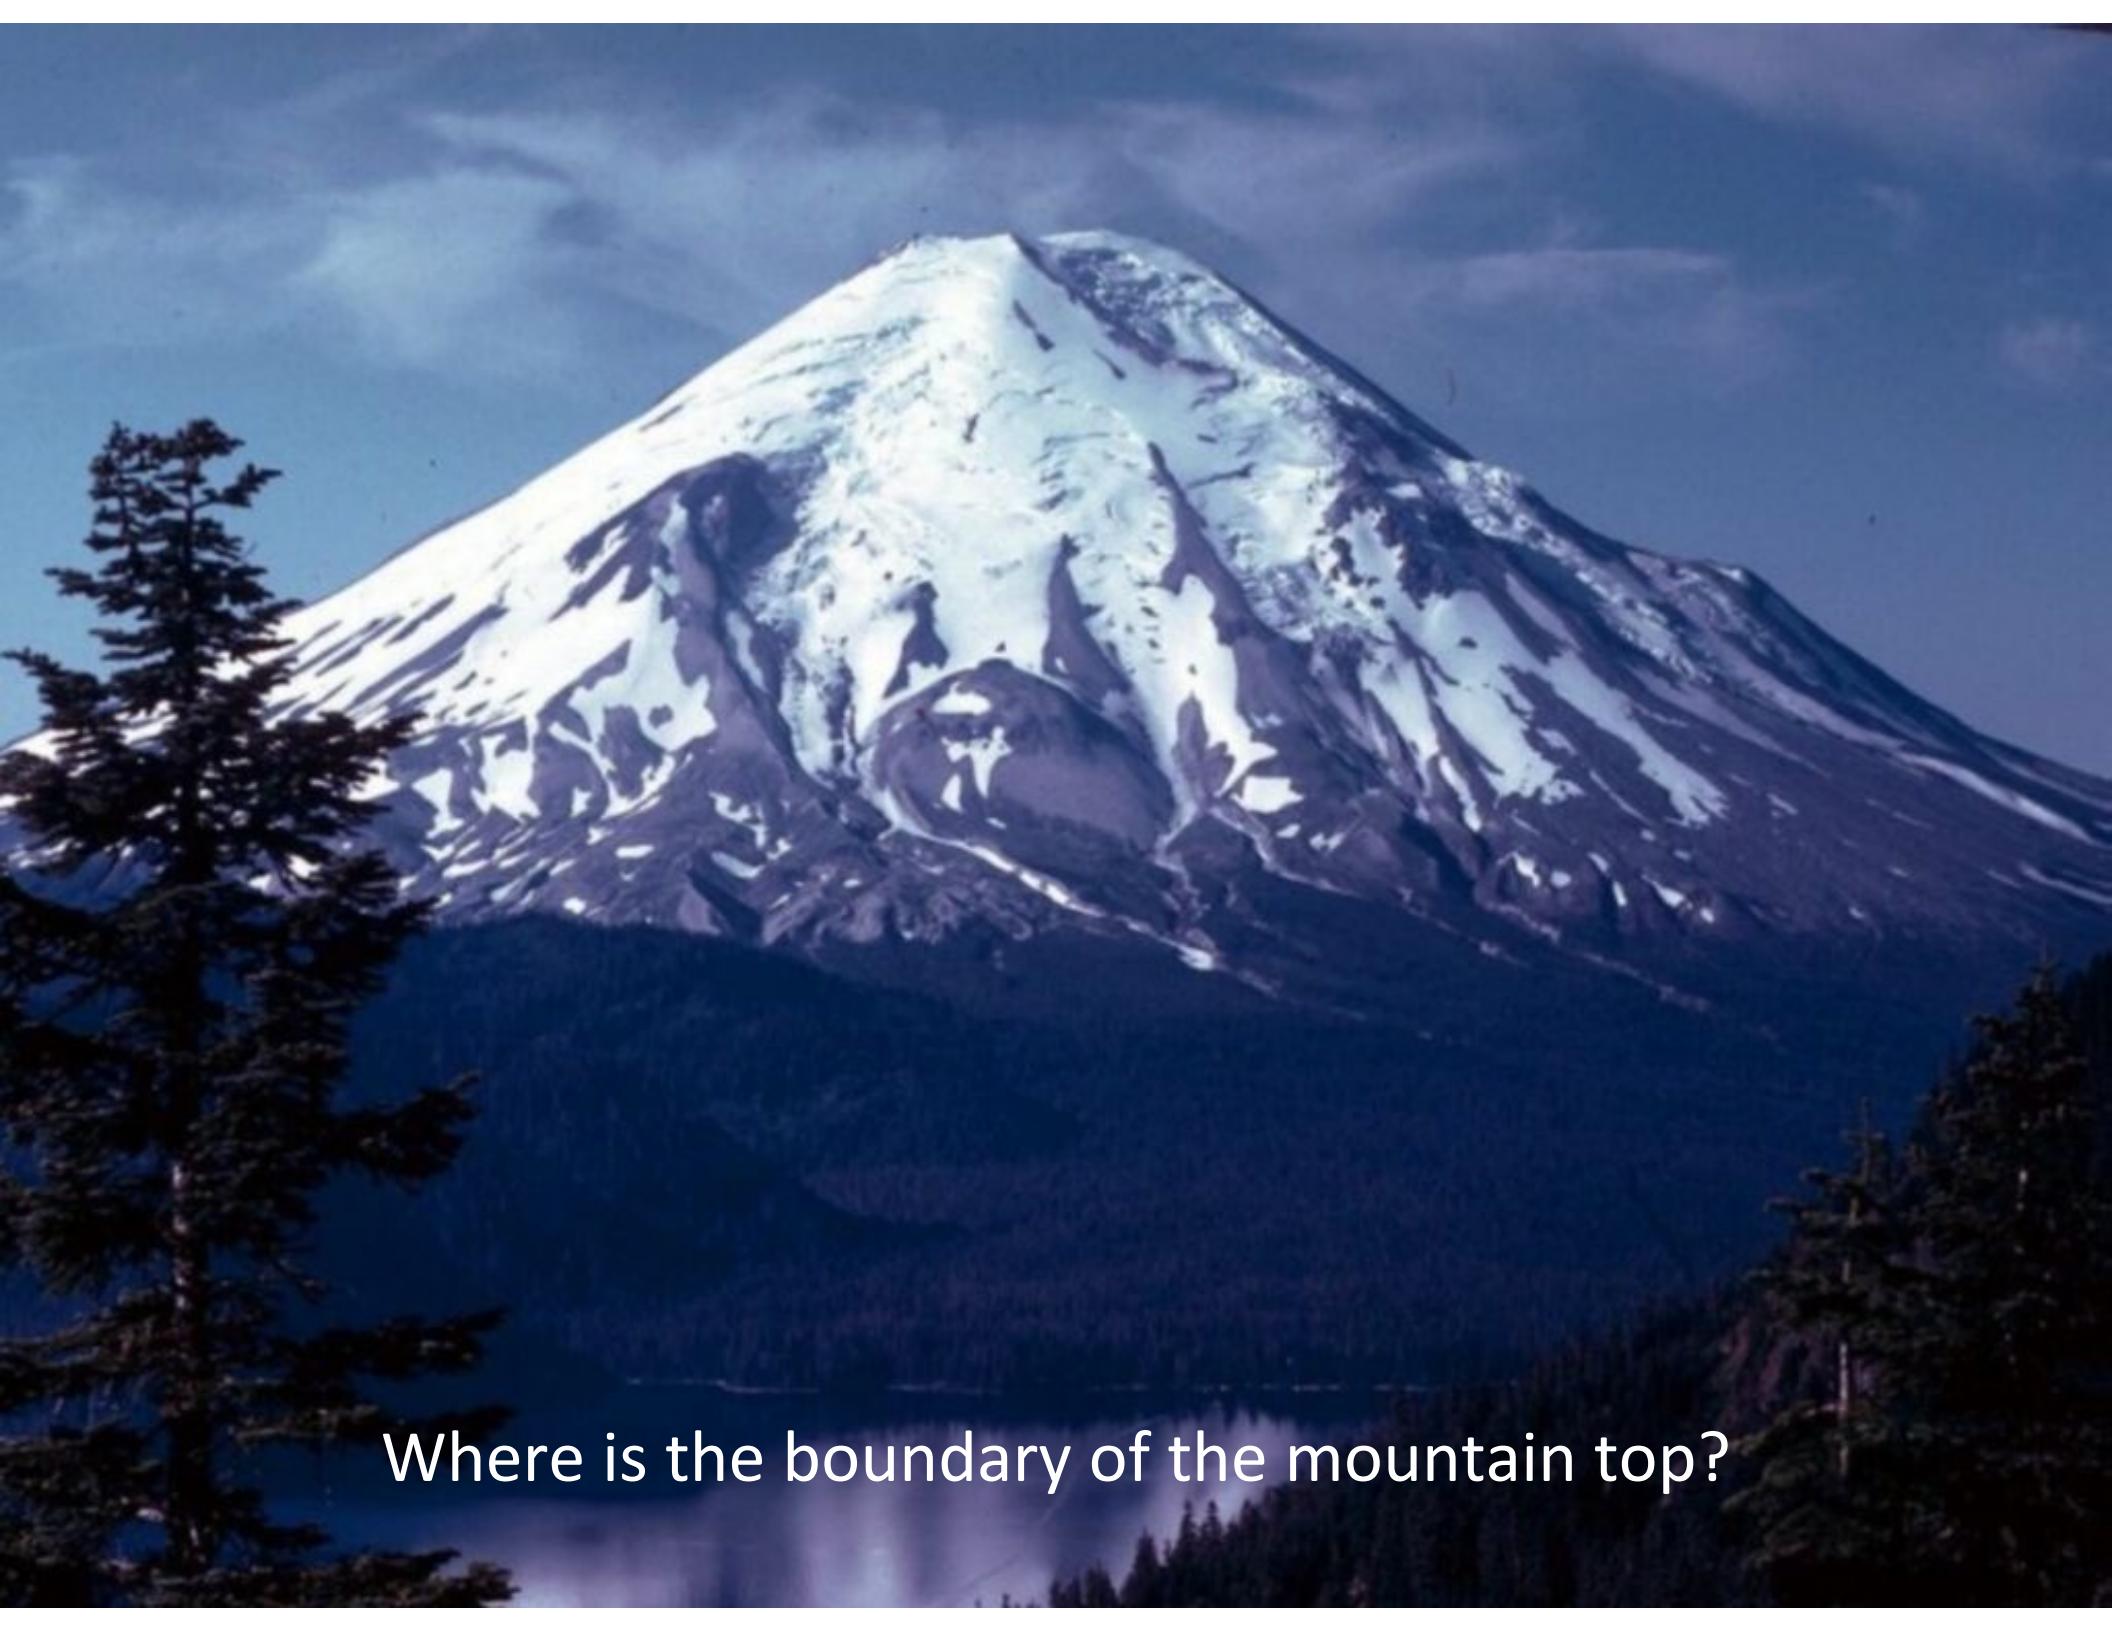
Edge detection

A high-magnification microscopic image of a tissue sample. The image shows a dense arrangement of cells with distinct nuclei. A prominent feature is a large, central area where the cellular structure is less dense, appearing more orange-red. This area contains many thin, radiating fibers that resemble collagen or other extracellular matrix components. The overall color palette includes shades of red, orange, green, and white.

Defining boundaries are hard for us too



Defining boundaries are hard for us too

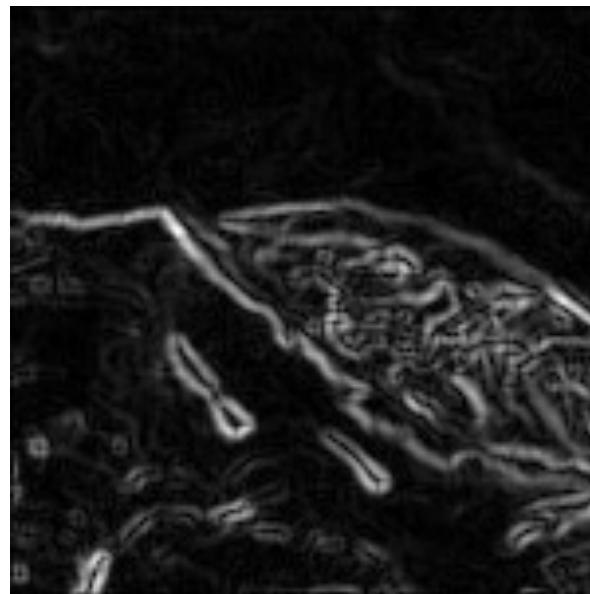


Where is the boundary of the mountain top?

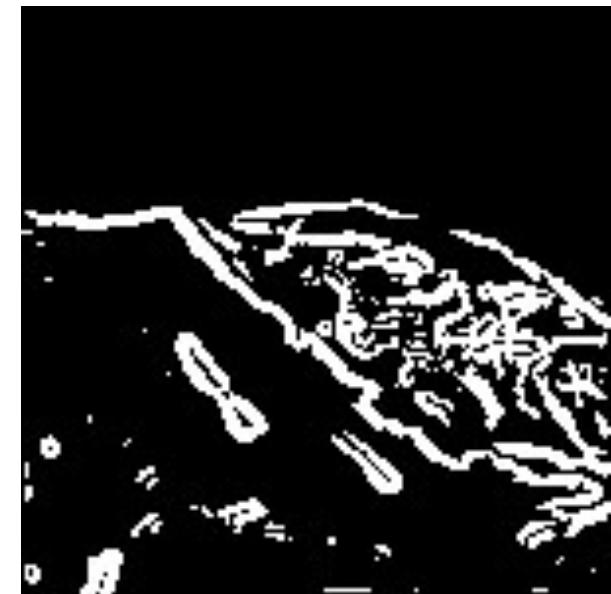
Lines are hard to find



Original image



Edge detection



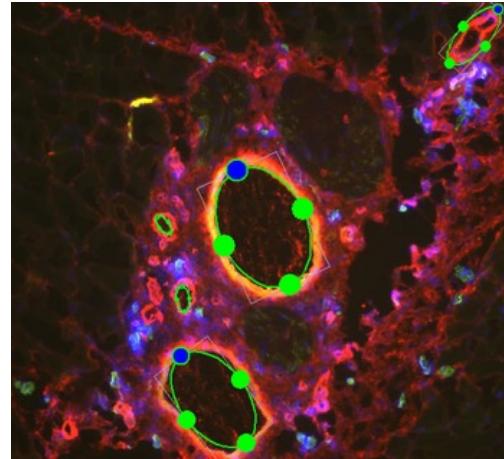
Thresholding

Noisy edge image
Incomplete boundaries

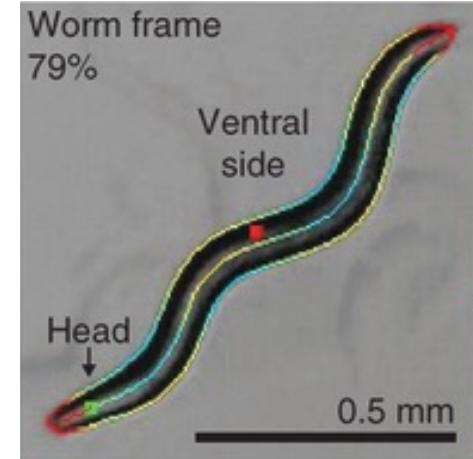
Applications



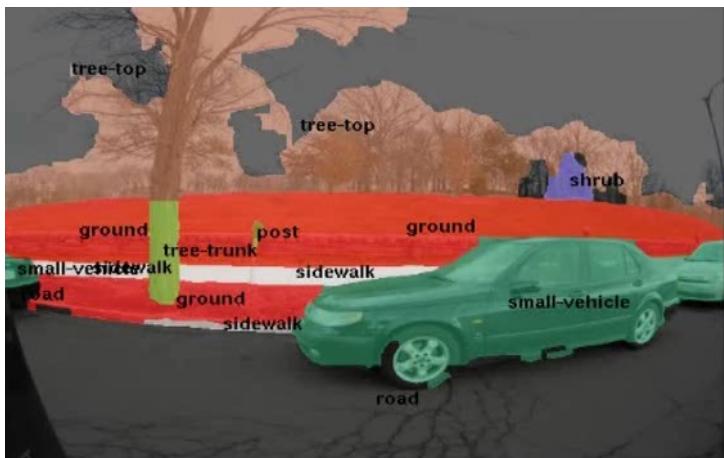
Autonomous Vehicles
(lane line detection)



tissue engineering
(blood vessel counting)



behavioral genetics
(earthworm contours)



Autonomous Vehicles
(semantic scene segmentation)



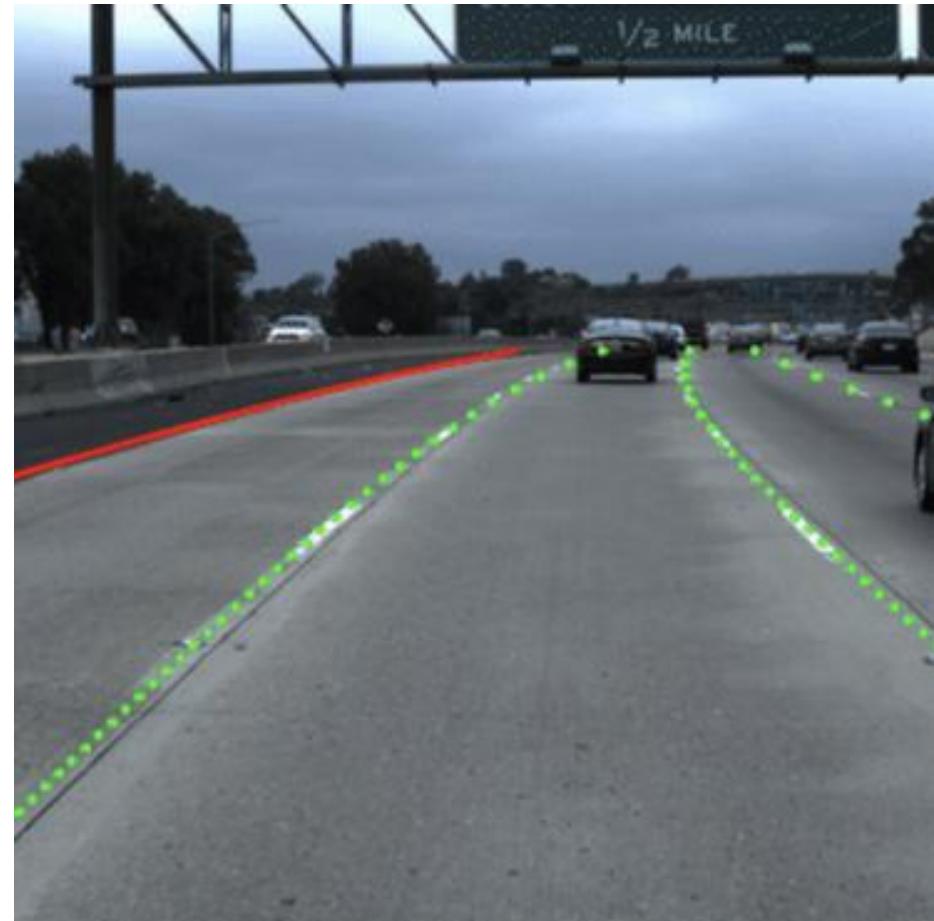
Computational Photography
(image inpainting)

Line fitting

Slope intercept form

$$y = mx + b$$

slope y-intercept

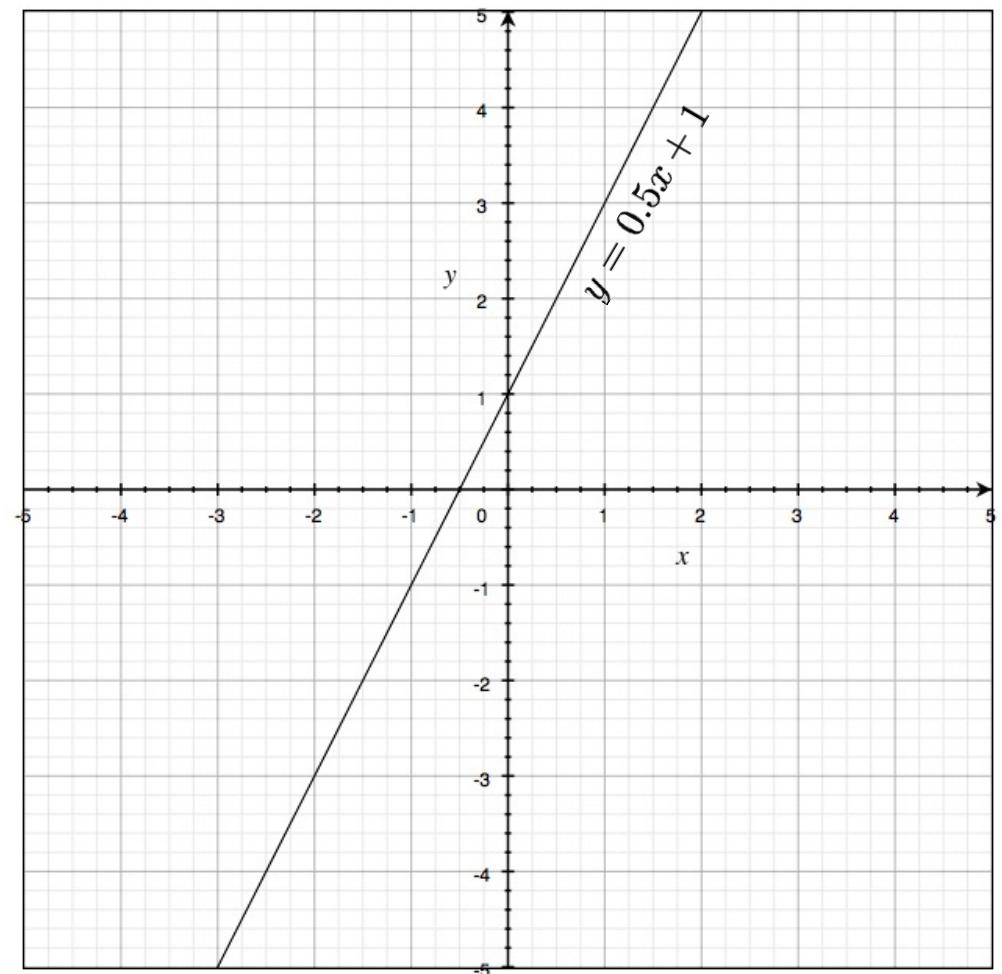


What are m and b ?

Slope intercept form

$$y = mx + b$$

slope y-intercept



Double intercept form

$$\frac{x}{a} + \frac{y}{b} = 1$$

x-intercept y-intercept

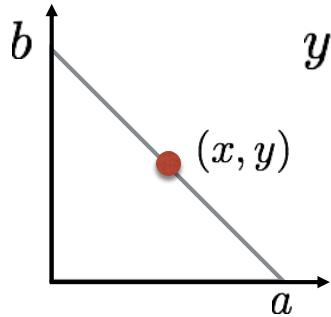
What are a and b?

Double intercept form

$$\frac{x}{a} + \frac{y}{b} = 1$$

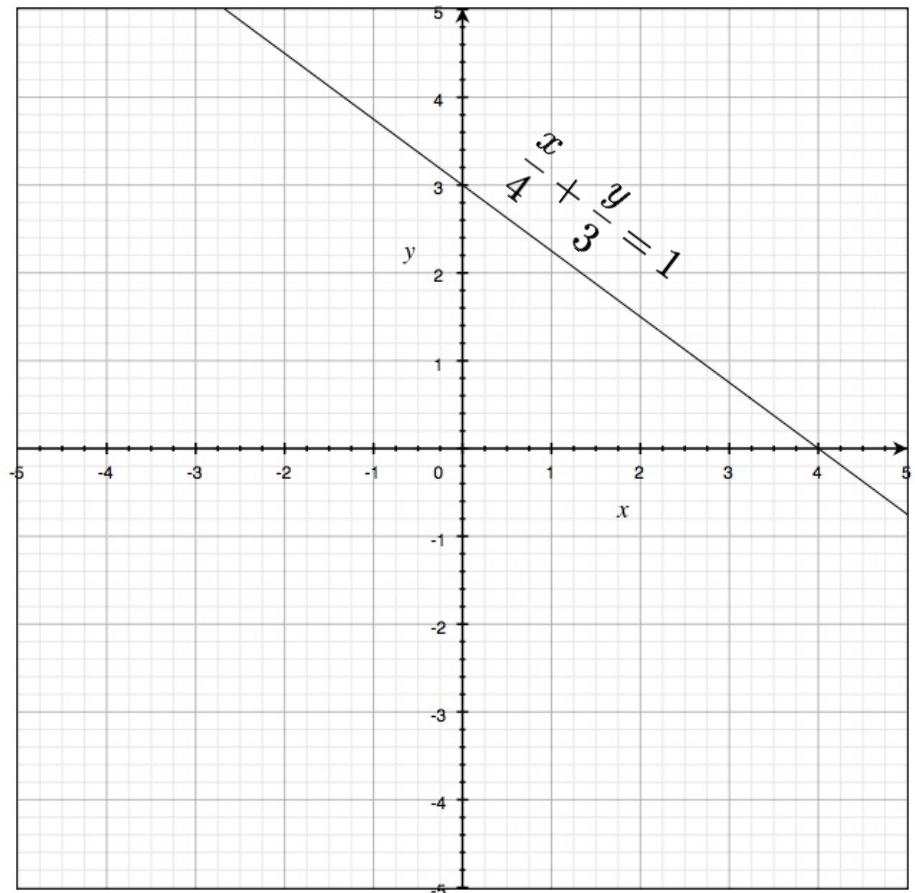
x-intercept  y-intercept 

Derivation:



(Similar slope)

$$\frac{y - b}{x - 0} = \frac{0 - y}{a - x}$$
$$ya + yx - ba + bx = -yx$$
$$ya + bx = ba$$
$$\frac{y}{b} + \frac{x}{a} = 1$$



Normal Form

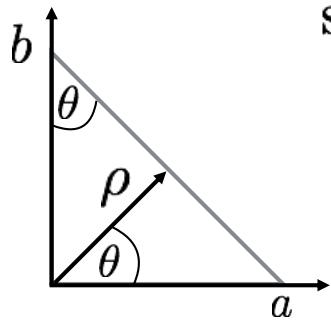
$$x \cos \theta + y \sin \theta = \rho$$

What are rho and theta?

Normal Form

$$x \cos \theta + y \sin \theta = \rho$$

Derivation:

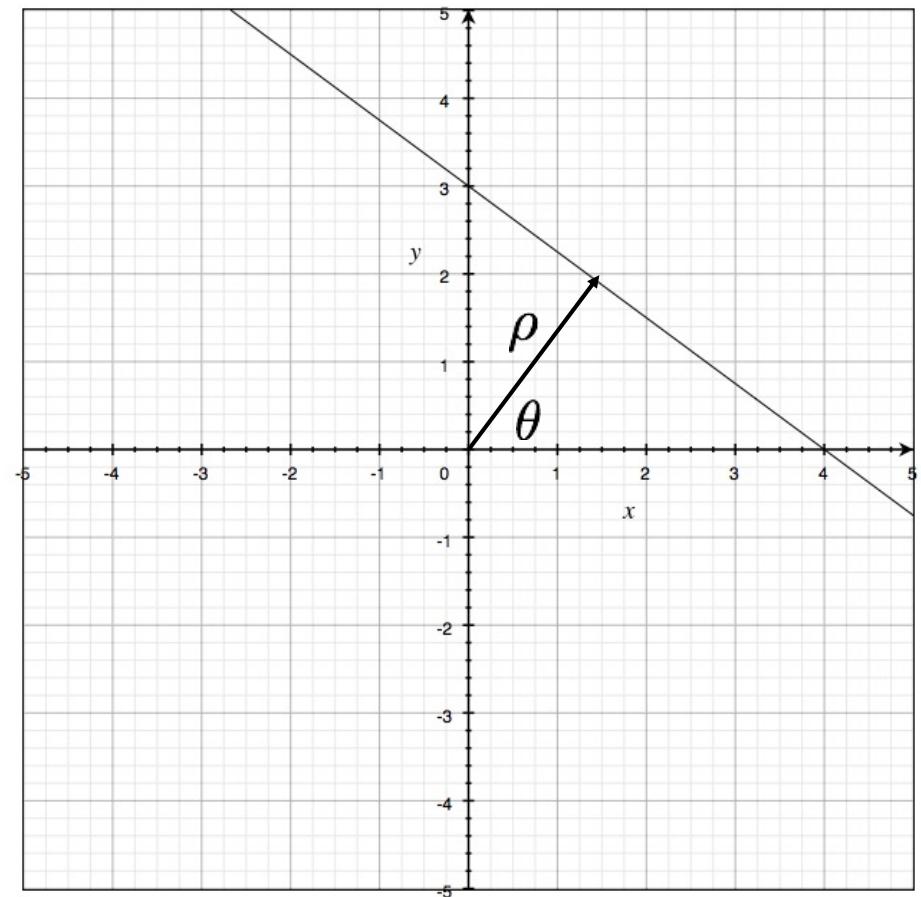


$$\cos \theta = \frac{\rho}{a} \rightarrow a = \frac{\rho}{\cos \theta}$$

$$\sin \theta = \frac{\rho}{b} \rightarrow b = \frac{\rho}{\sin \theta}$$

plug into: $\frac{x}{a} + \frac{y}{b} = 1$

$$x \cos \theta + y \sin \theta = \rho$$



Hough transform

Hough transform

- Generic framework for detecting a parametric model
- Edges don't have to be connected
- Lines can be occluded
- Key idea: edges **vote** for the possible models

Image and parameter space

variables
 $y = mx + b$
parameters

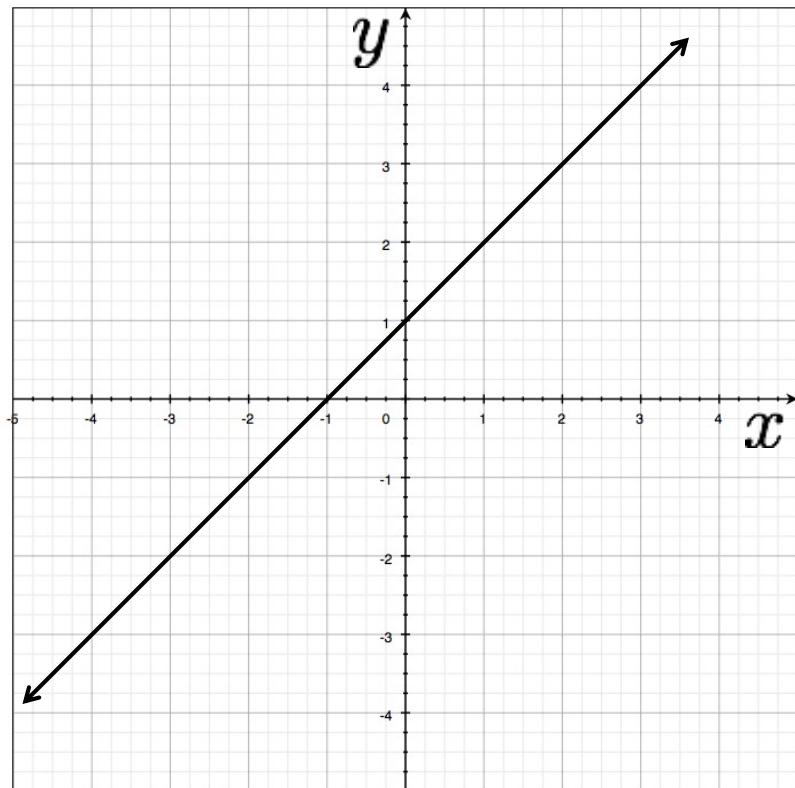
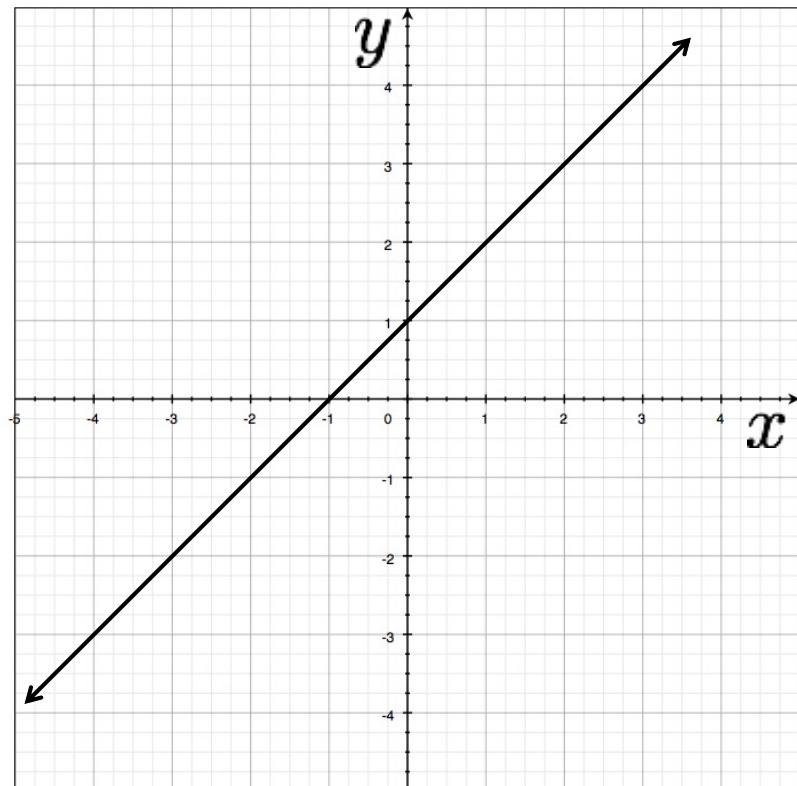


Image space

Image and parameter space

variables
 $y = mx + b$
parameters



a line
becomes a
point

variables
 $y - mx = b$
parameters

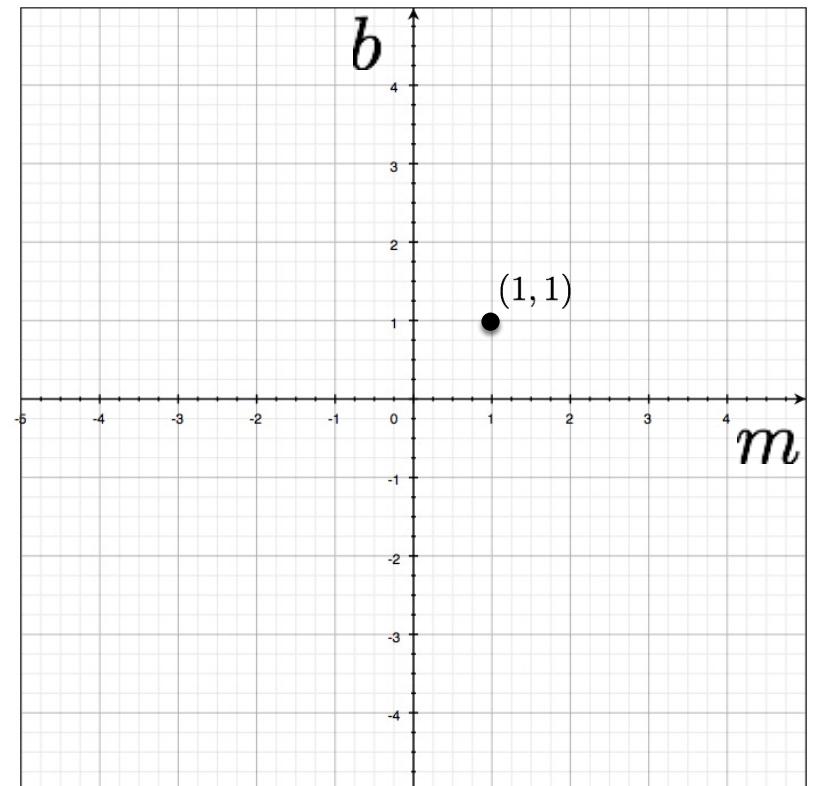


Image space

Parameter space

Image and parameter space

$$y = mx + b$$

variables
parameters

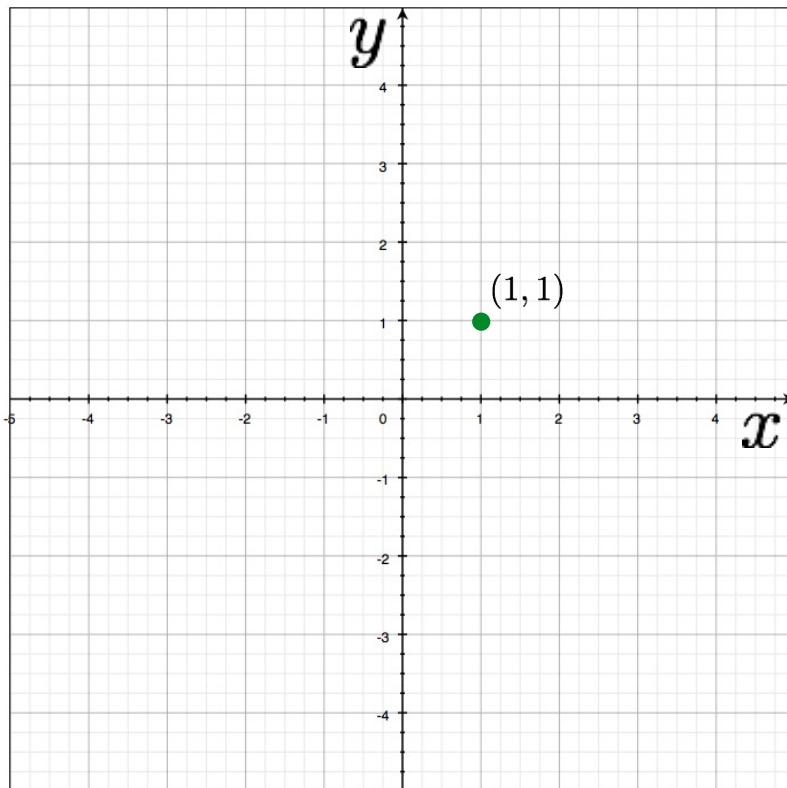


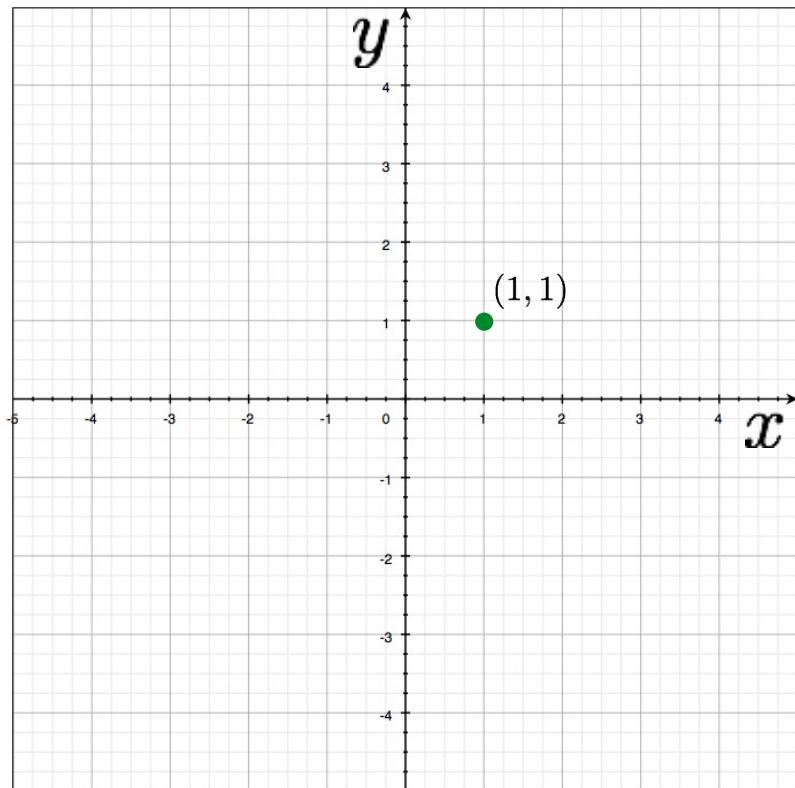
Image space

What would a point in image space become in parameter space?

Image and parameter space

$$y = mx + b$$

variables
parameters



a point becomes a line

$$y - mx = b$$

variables
parameters

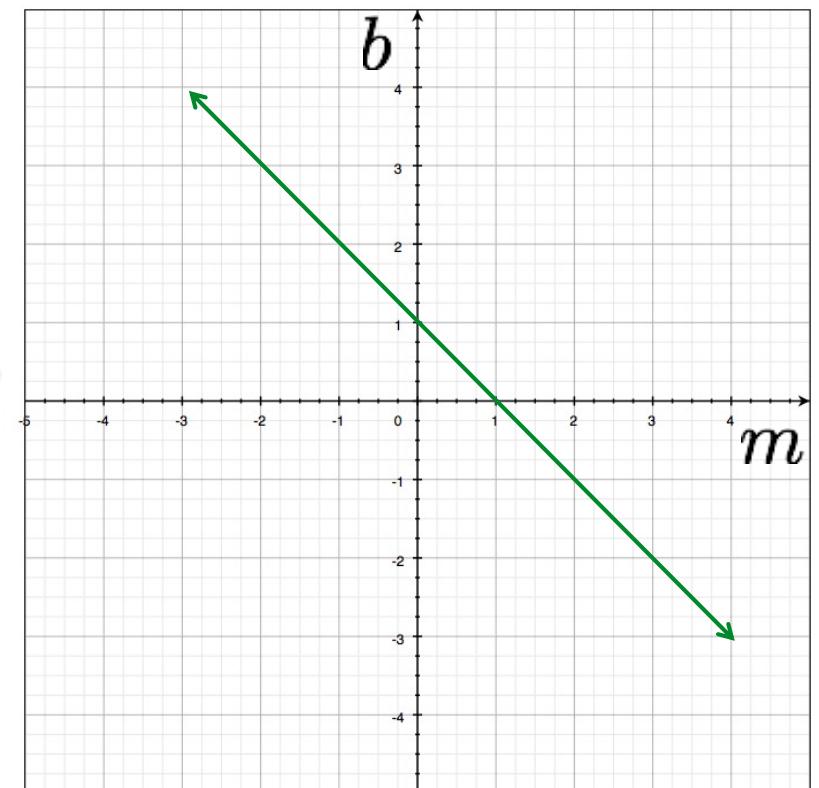
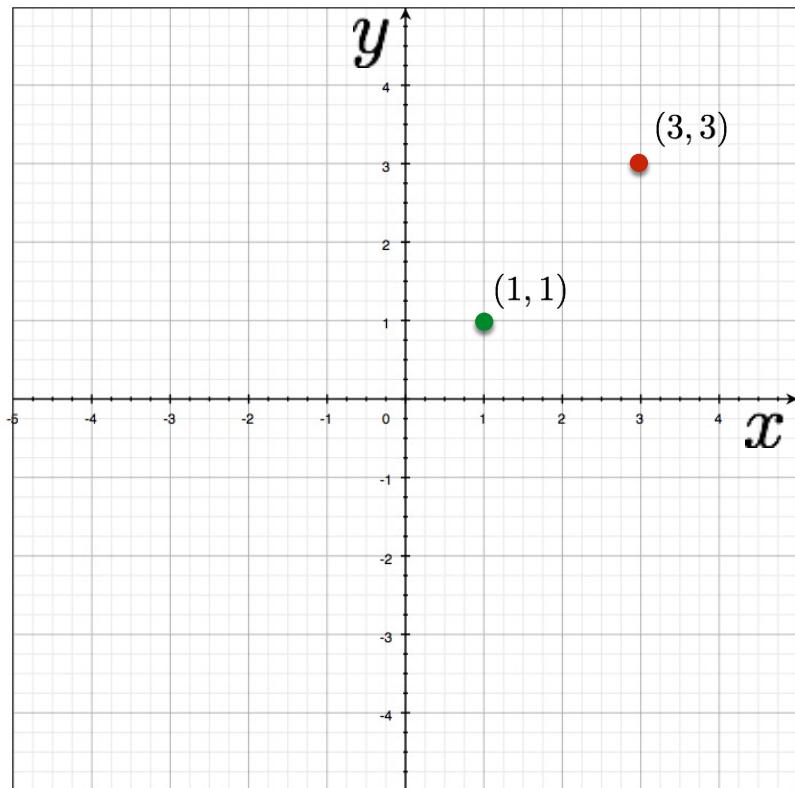


Image space

Parameter space

Image and parameter space

variables
 $y = mx + b$
parameters



two points
become
?

variables
 $y - mx = b$
parameters

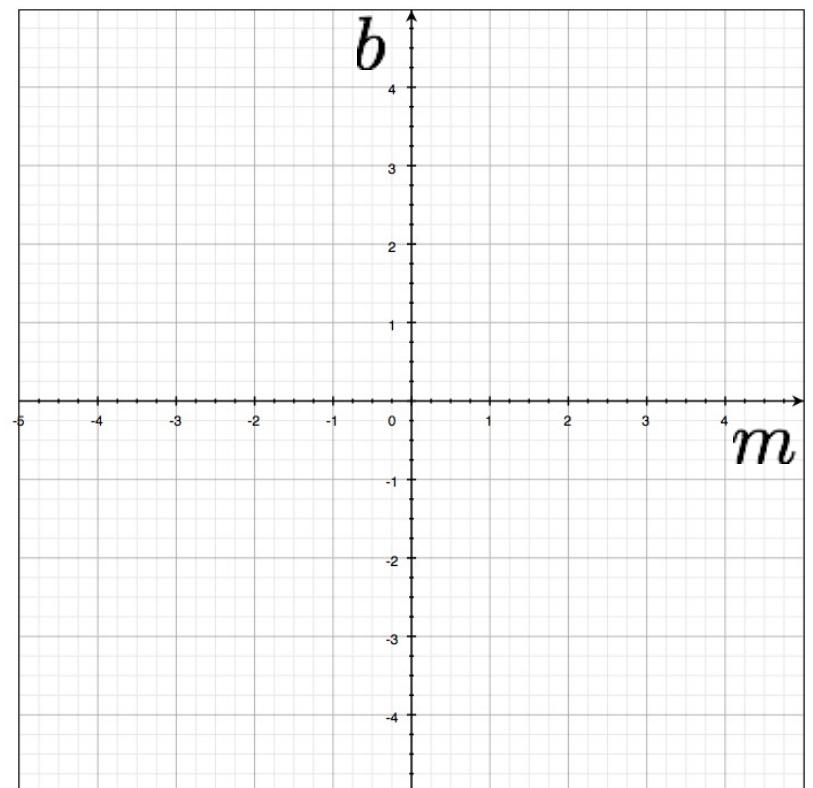


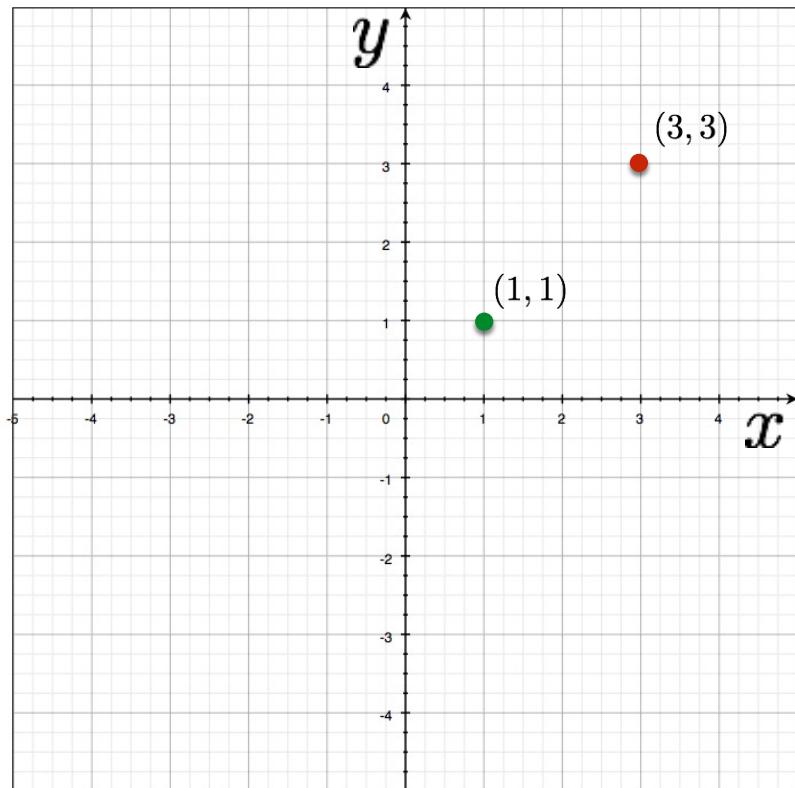
Image space

Parameter space

Image and parameter space

$$y = mx + b$$

variables
parameters



two points
become
?

$$y - mx = b$$

variables
parameters

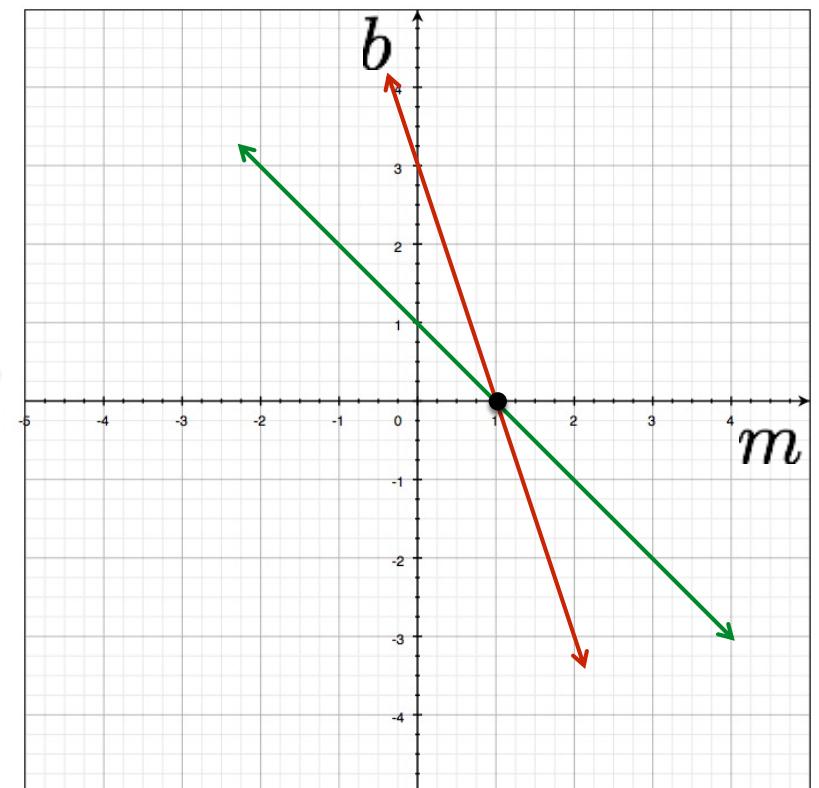
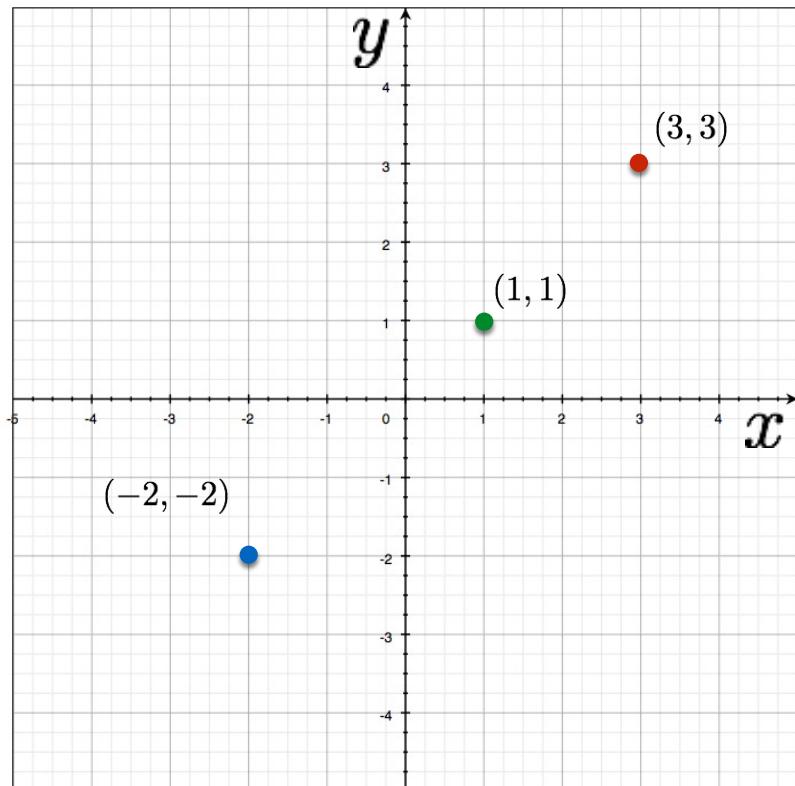


Image space

Parameter space

Image and parameter space

variables
 $y = mx + b$
parameters



three points
become
?

variables
 $y - mx = b$
parameters

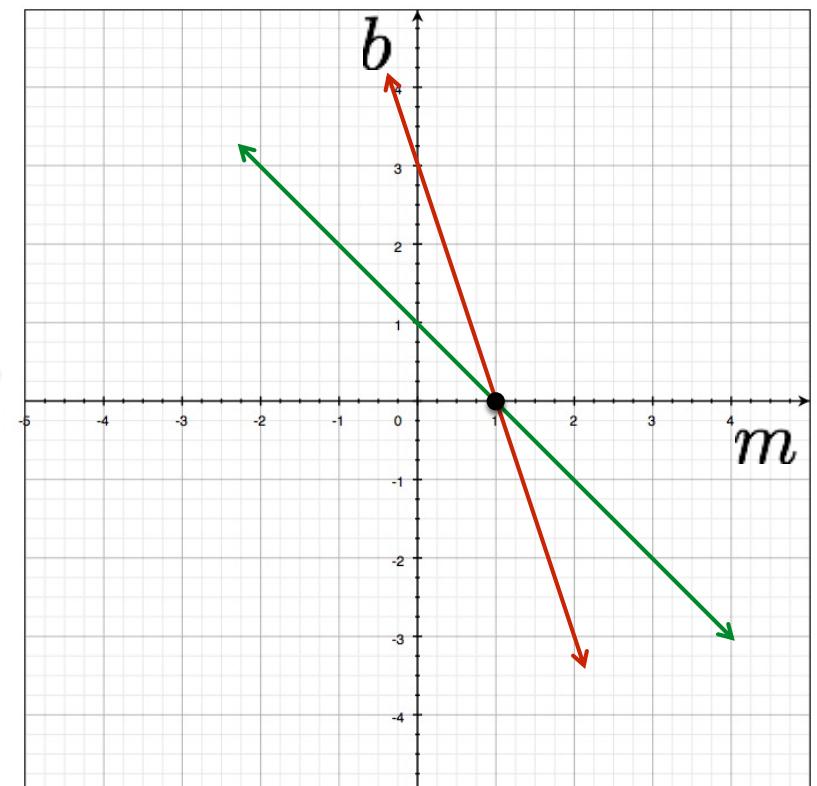


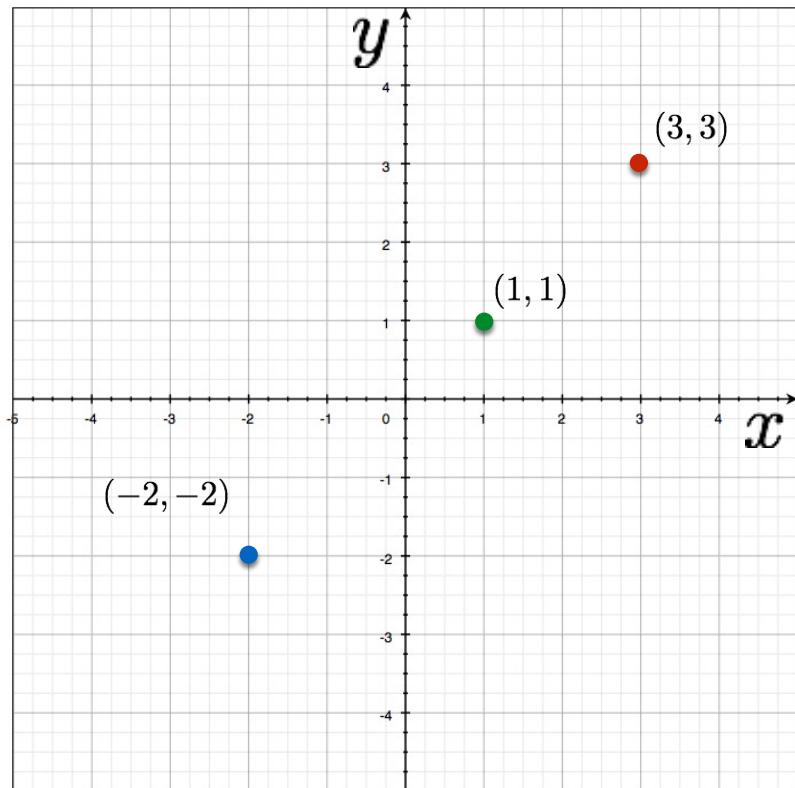
Image space

Parameter space

Image and parameter space

$$y = mx + b$$

variables
parameters



three points
become
?

$$y - mx = b$$

variables
parameters

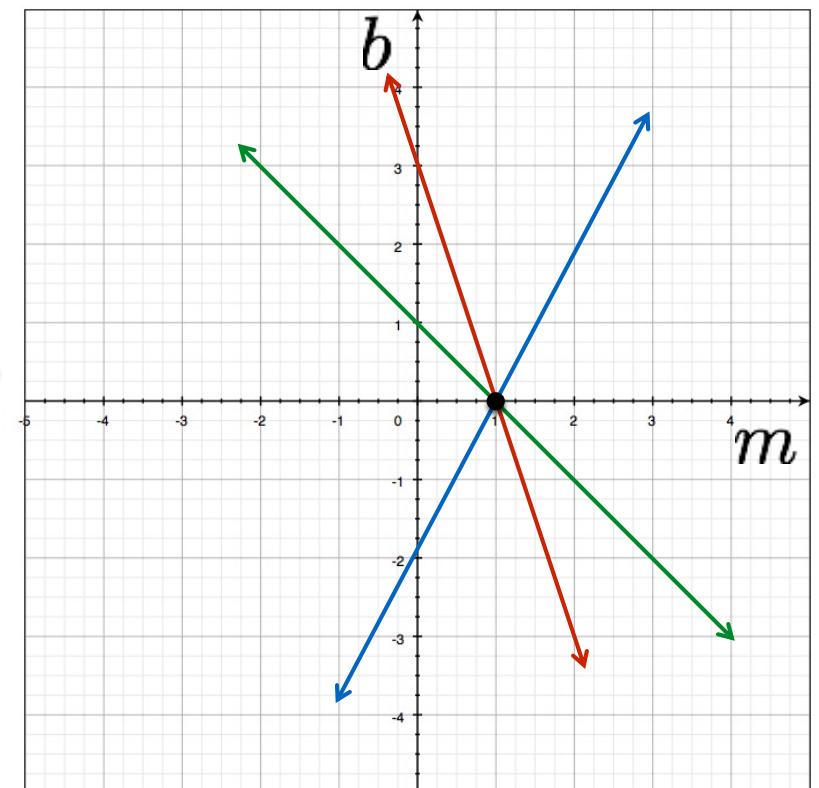
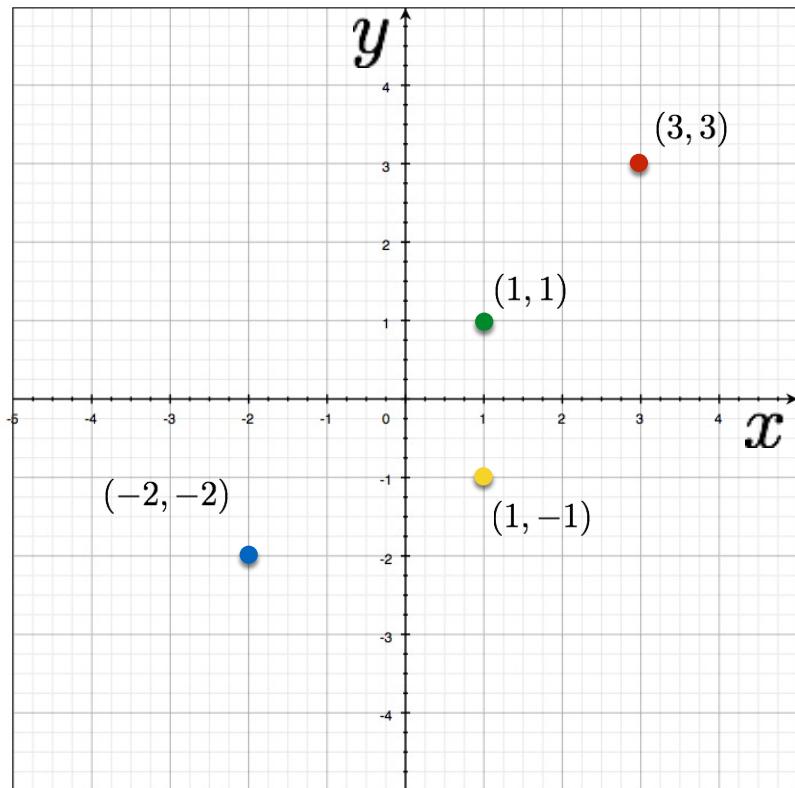


Image space

Parameter space

Image and parameter space

variables
 $y = mx + b$
parameters



four points
become
?

variables
 $y - mx = b$
parameters

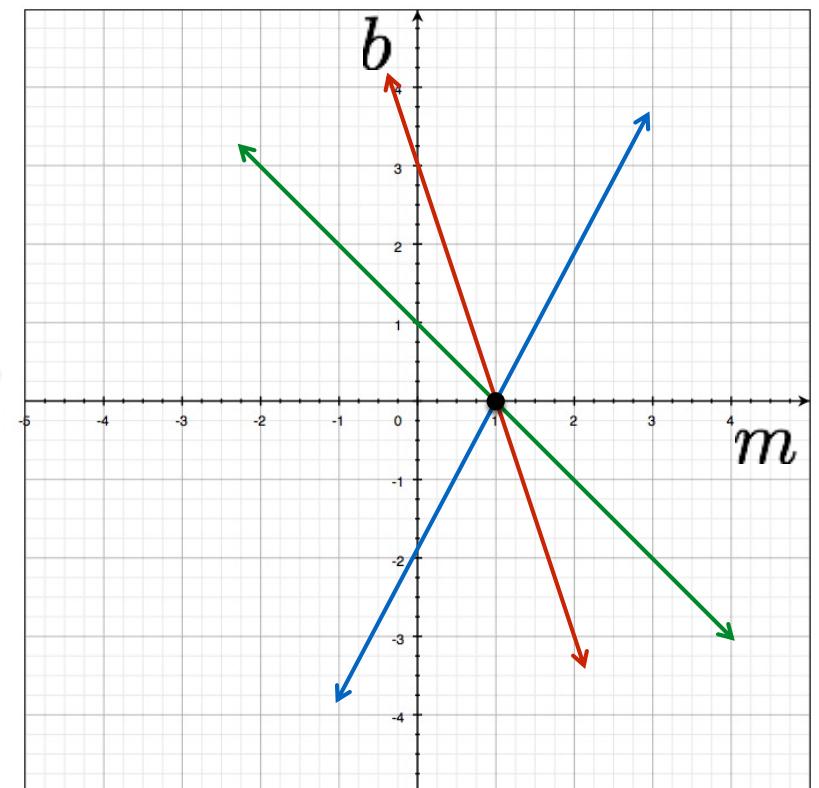
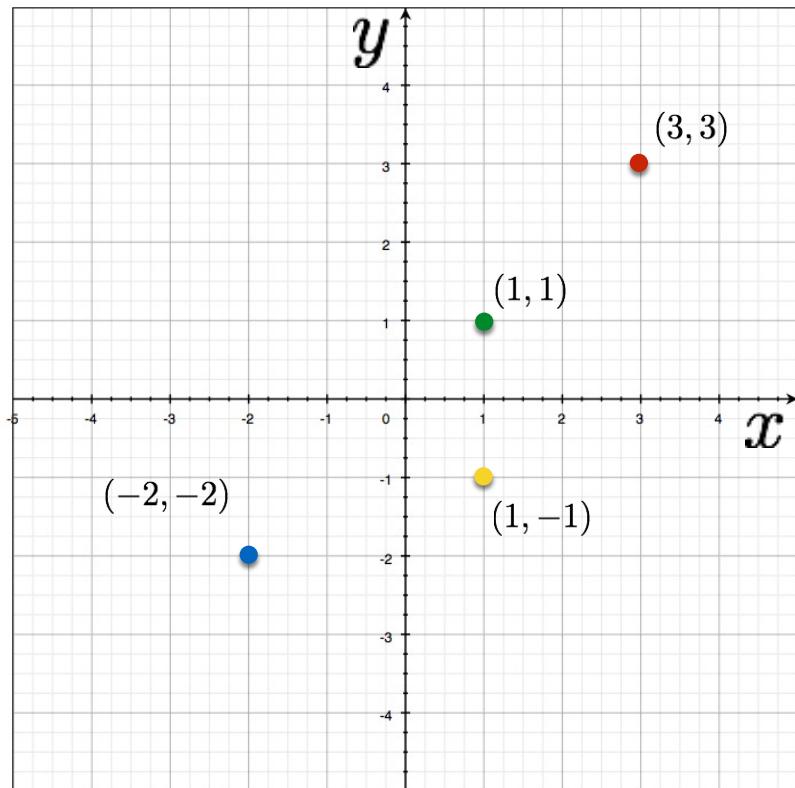


Image space

Parameter space

Image and parameter space

variables
 $y = mx + b$
parameters



four points
become
?

variables
 $y - mx = b$
parameters

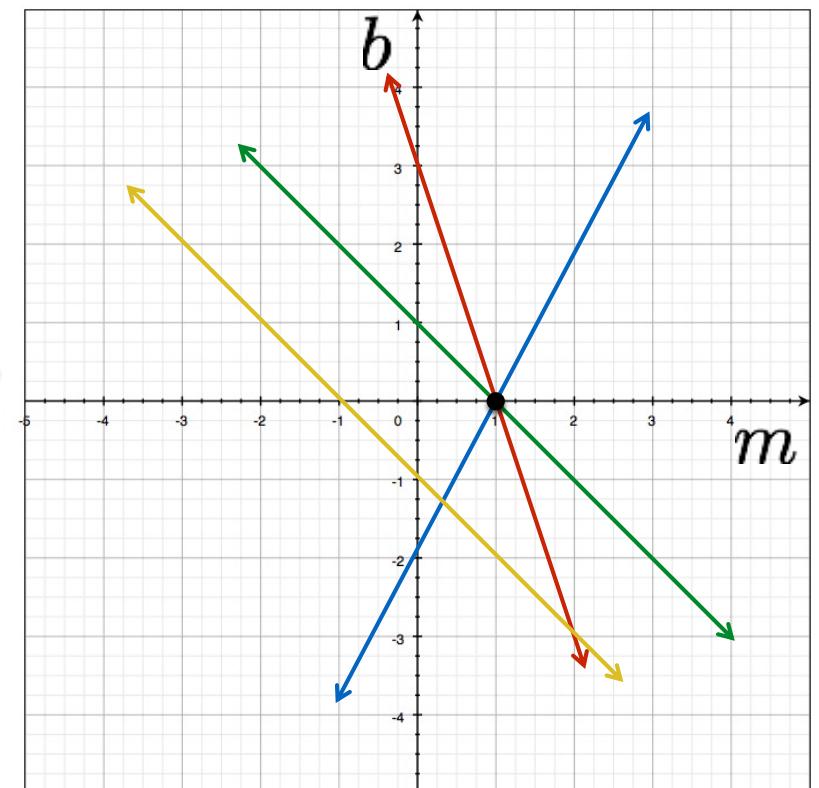


Image space

Parameter space

How would you find the best fitting line?

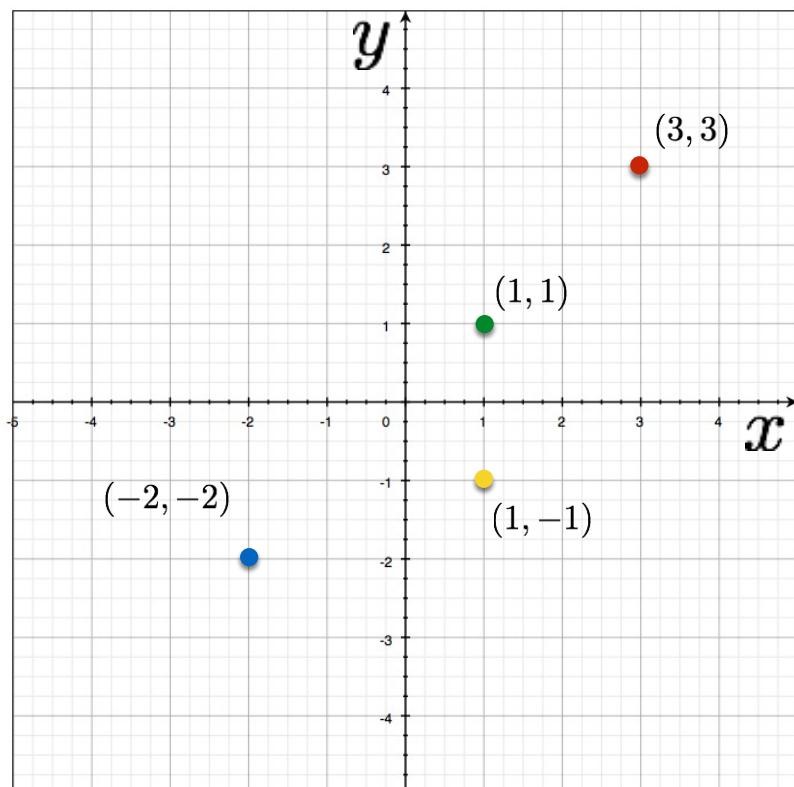
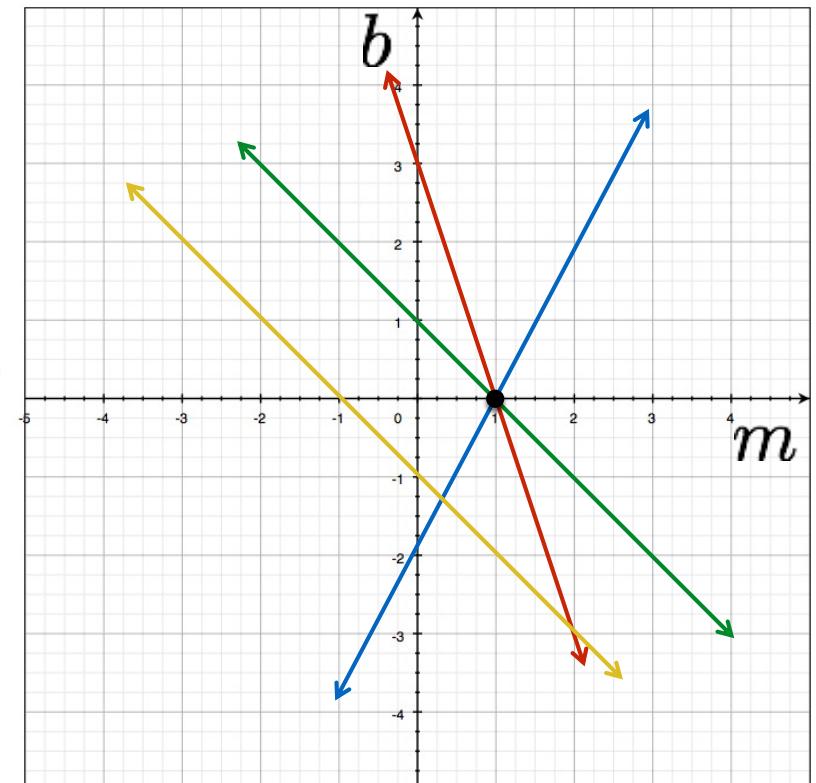


Image space



Parameter space

Is this method robust to measurement noise?

Is this method robust to outliers?

Line Detection by Hough Transform

Algorithm:

1. Quantize Parameter Space (m, c)

2. Create Accumulator Array $A(m, c)$

3. Set $A(m, c) = 0 \quad \forall m, c$

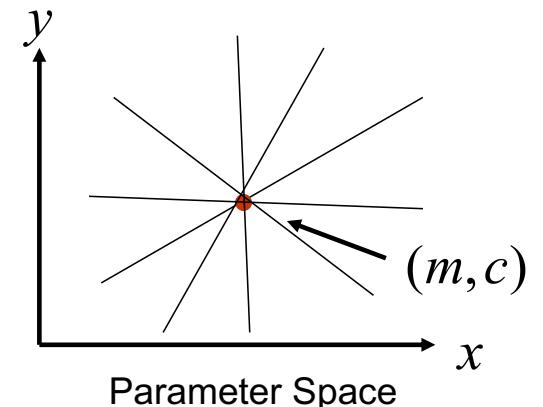
4. For each image edge (x_i, y_i)

For each element in $A(m, c)$

If (m, c) lies on the line: $c = -x_i m + y_i$

Increment $A(m, c) = A(m, c) + 1$

5. Find local maxima in $A(m, c)$



$A(m, c)$

1				1
	1			1
	1	1		
		2		
	1	1	1	
1			1	
1				1

Problems with parameterization

How big does the accumulator need to be for the parameterization (m, c) ?

$A(m, c)$

1					1		
	1				1		
		1		1			
			2				
		1		1			
	1				1		
1						1	

Problems with parameterization

How big does the accumulator need to be for the parameterization (m, c) ?

$$A(m, c)$$

1					1	
	1				1	
		1		1		
			2			
		1		1		
	1				1	
1						1

The space of m is huge!

$$-\infty \leq m \leq \infty$$

The space of c is huge!

$$-\infty \leq c \leq \infty$$

Better Parameterization

Use normal form:

$$x \cos \theta + y \sin \theta = \rho$$

Given points (x_i, y_i) find (ρ, θ)

Hough Space Sinusoid

$$0 \leq \theta \leq 2\pi$$

$$0 \leq \rho \leq \rho_{\max}$$

(Finite Accumulator Array Size)

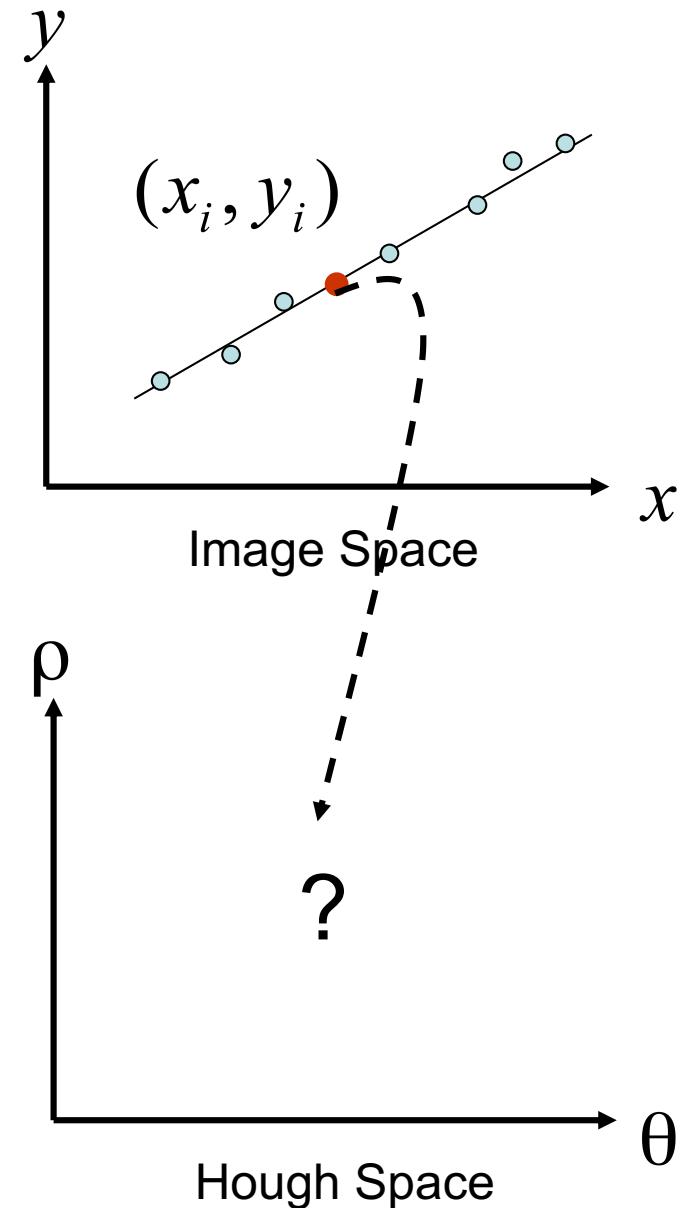
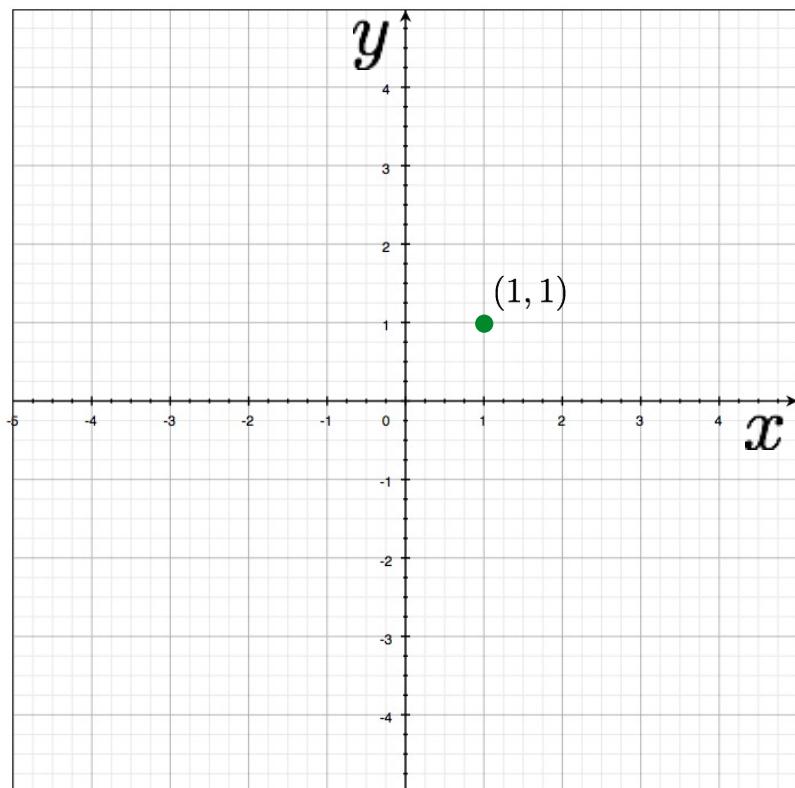


Image and parameter space

variables
 $y = mx + b$
parameters



a point becomes?

parameters
 $x \cos \theta + y \sin \theta = \rho$
variables

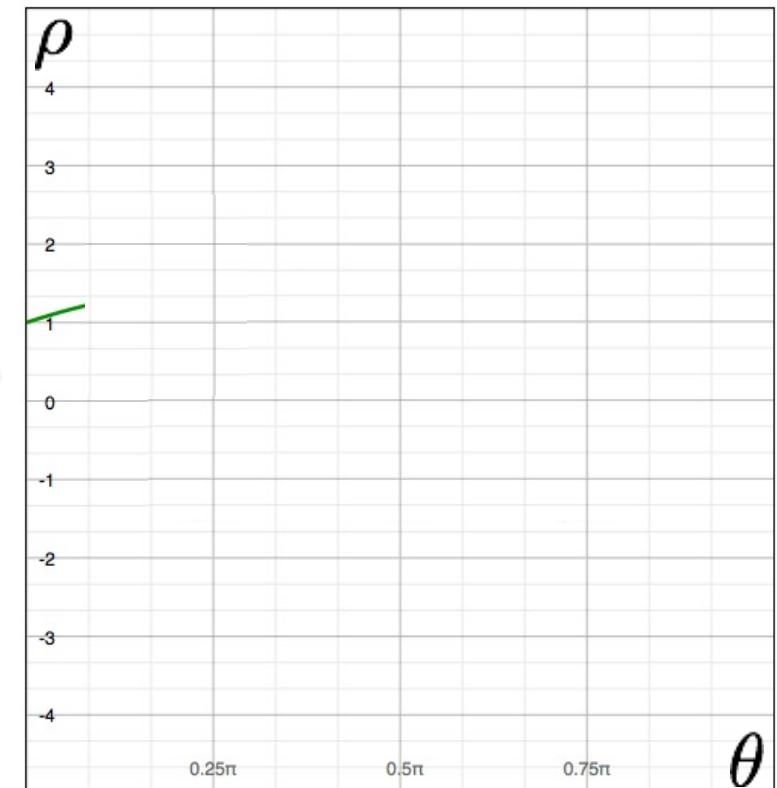


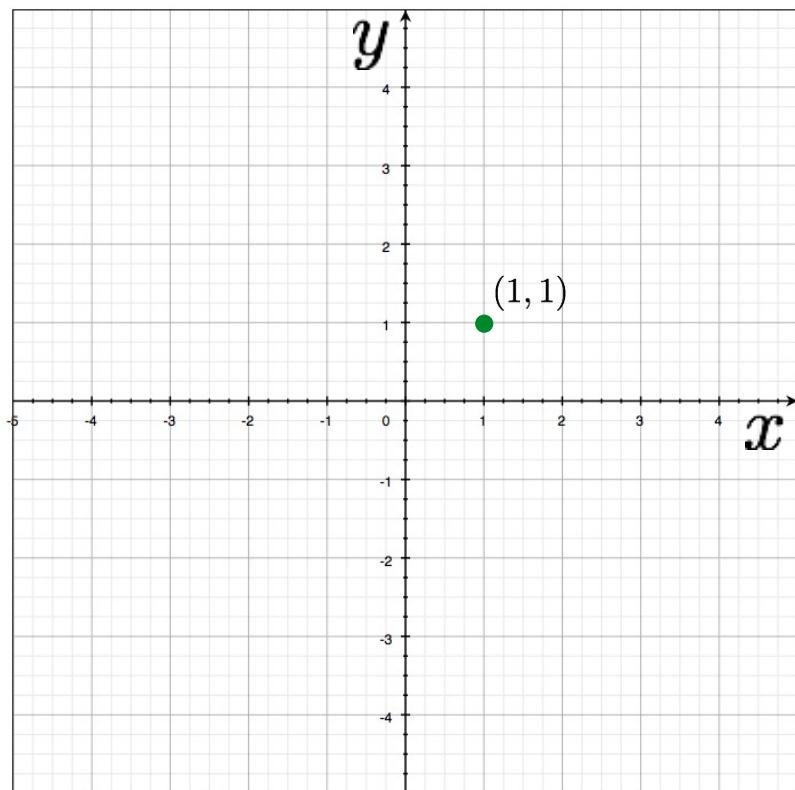
Image space

Parameter space

Image and parameter space

$$y = mx + b$$

variables
parameters



a point
becomes a
wave

$$x \cos \theta + y \sin \theta = \rho$$

parameters
variables



Image space

Parameter space

Image and parameter space

variables
 $y = mx + b$
parameters

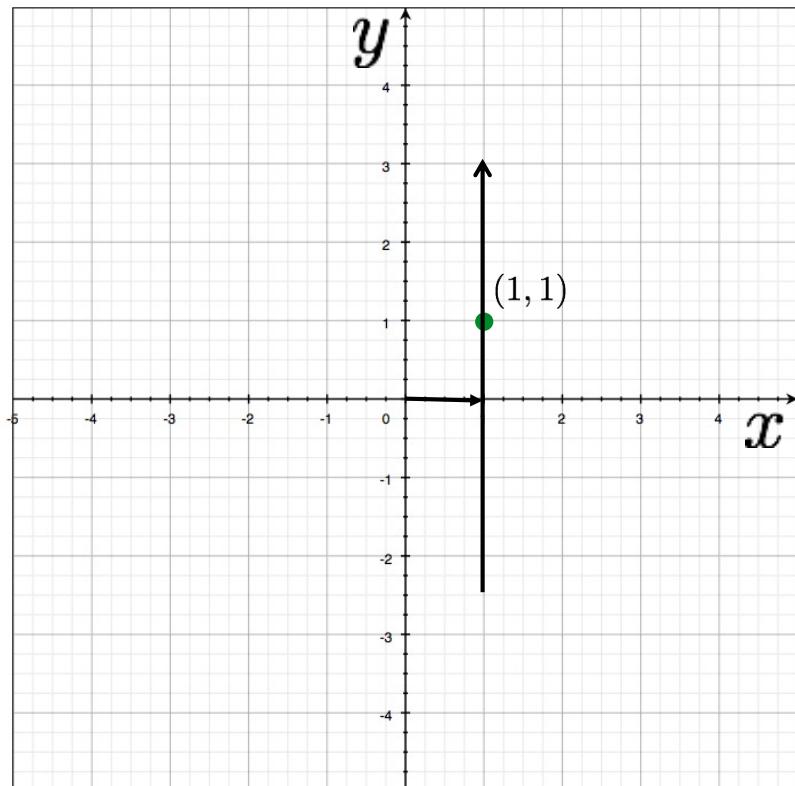
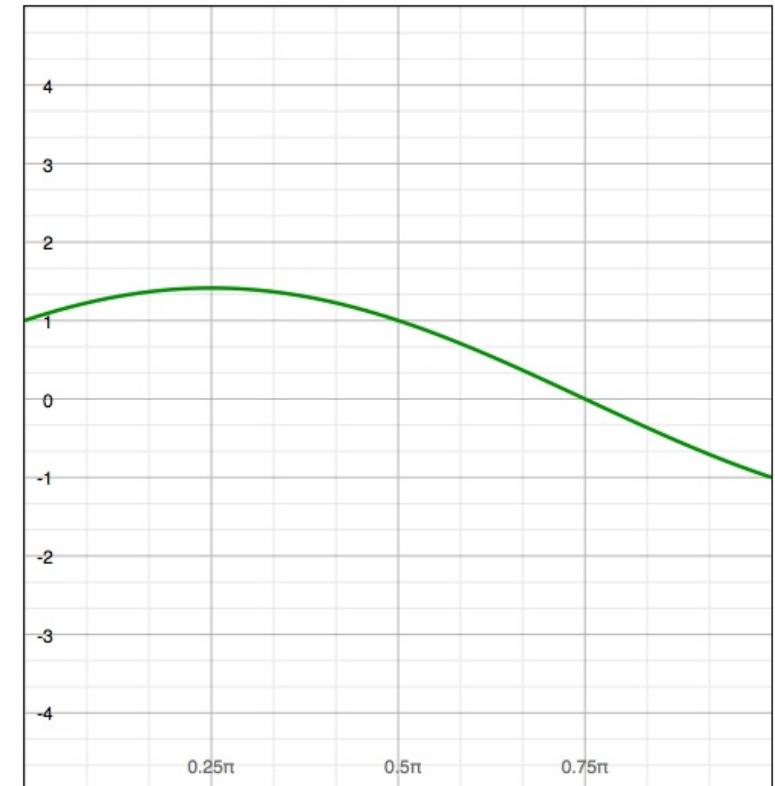


Image space

$$x \cos \theta + y \sin \theta = \rho$$

a line
becomes?



Parameter space

Image and parameter space

variables
 $y = mx + b$
parameters

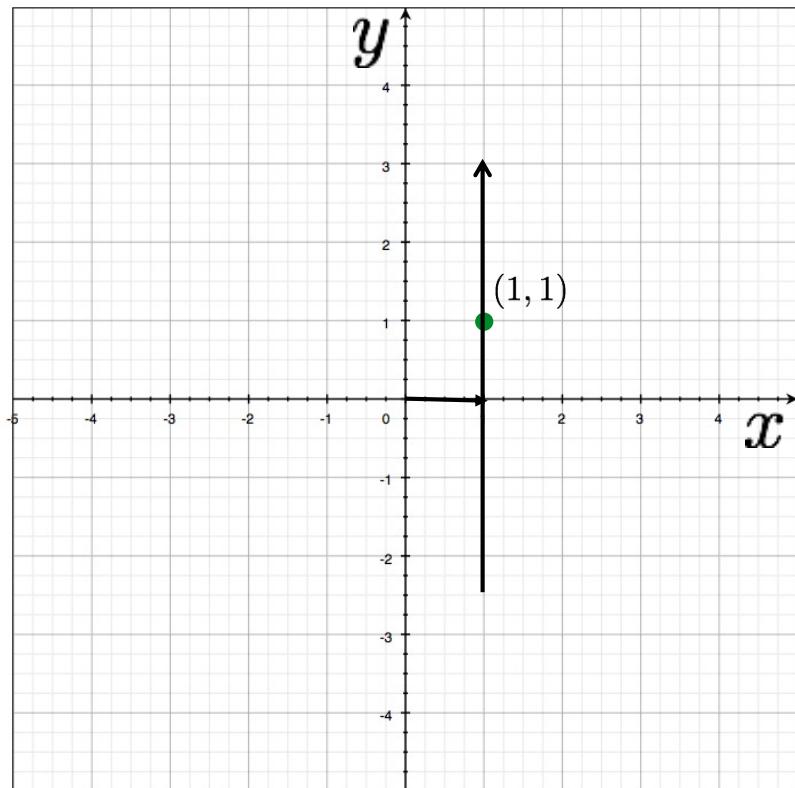
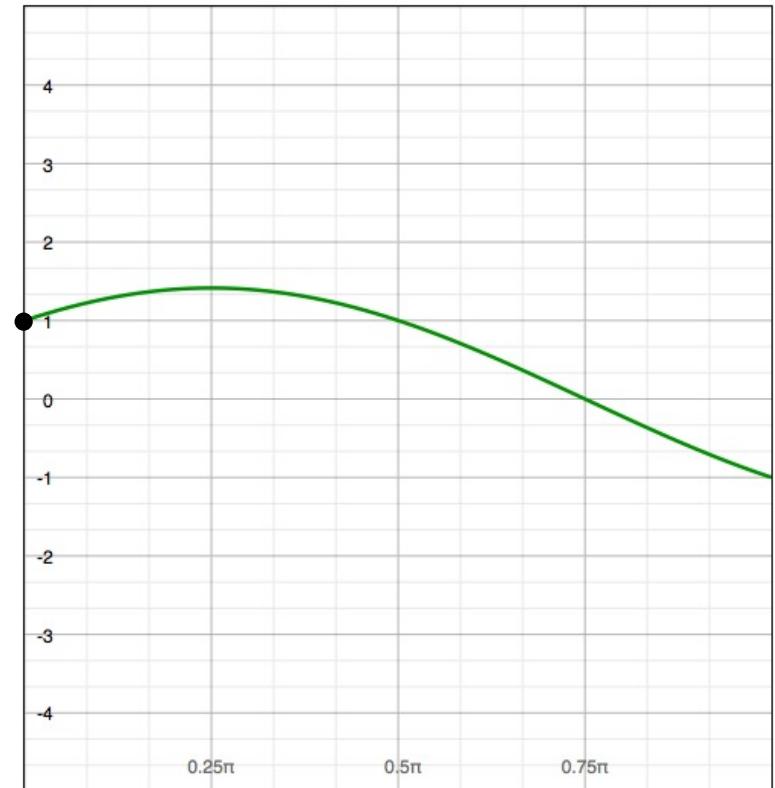


Image space

$$x \cos \theta + y \sin \theta = \rho$$

a line
becomes a
point



Parameter space

Image and parameter space

variables
 $y = mx + b$
parameters

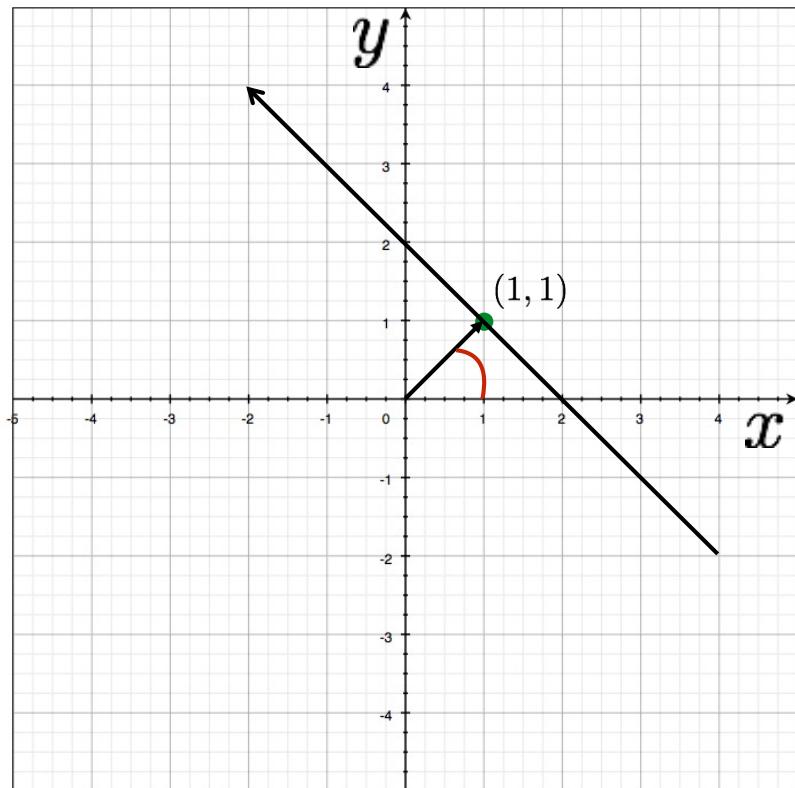
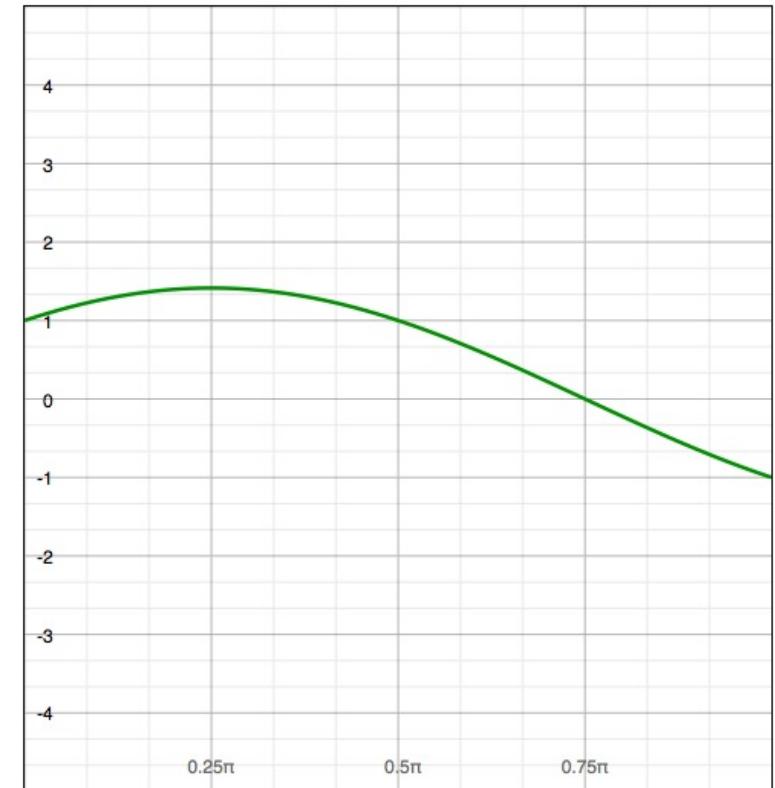


Image space

$$x \cos \theta + y \sin \theta = \rho$$

a line
becomes?



Parameter space

Image and parameter space

variables
 $y = mx + b$
parameters

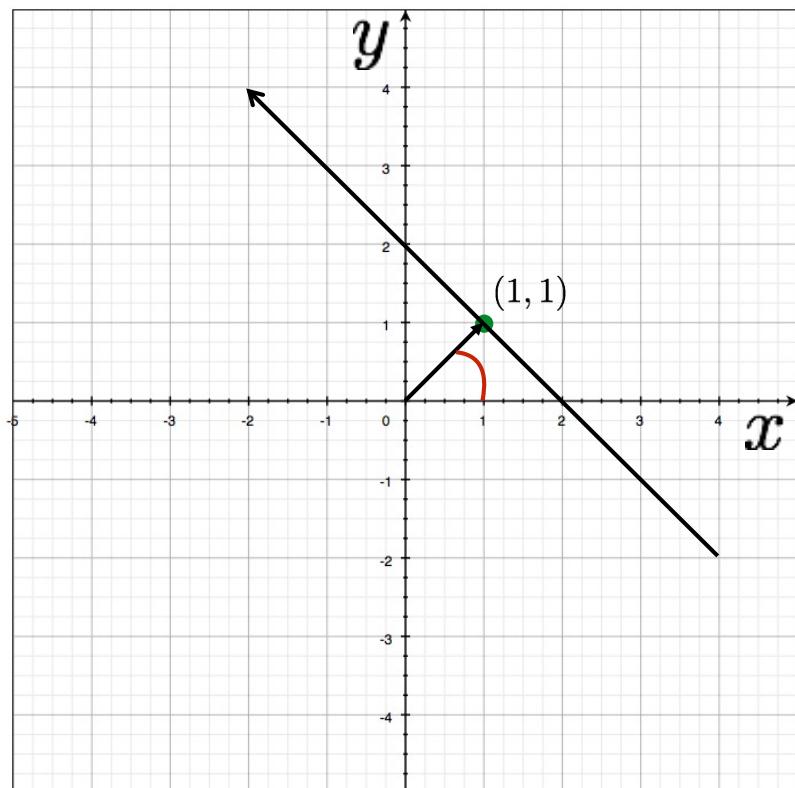
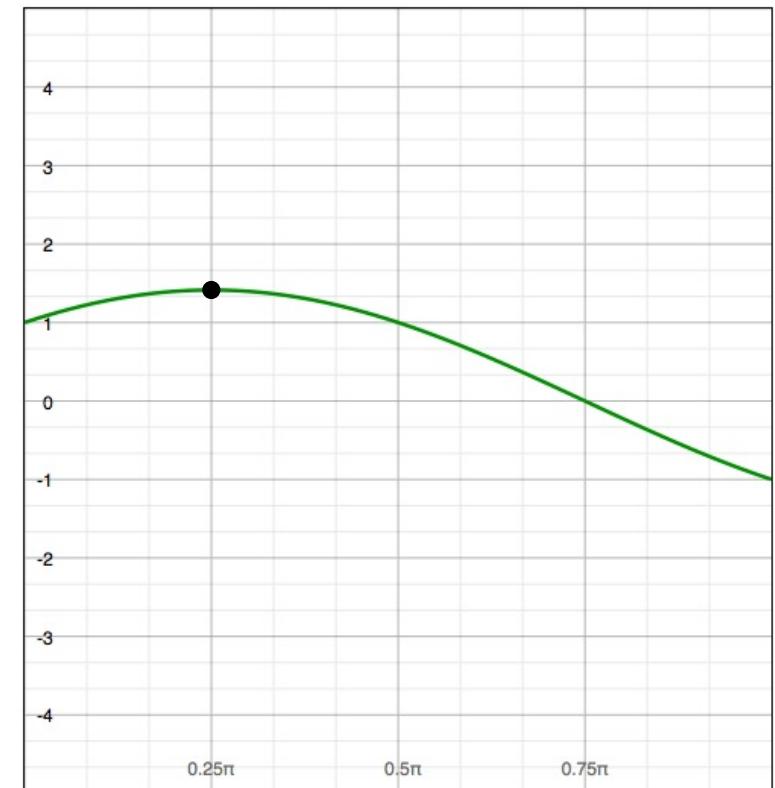


Image space

$$x \cos \theta + y \sin \theta = \rho$$

a line
becomes a
point



Parameter space

Image and parameter space

variables
 $y = mx + b$
parameters

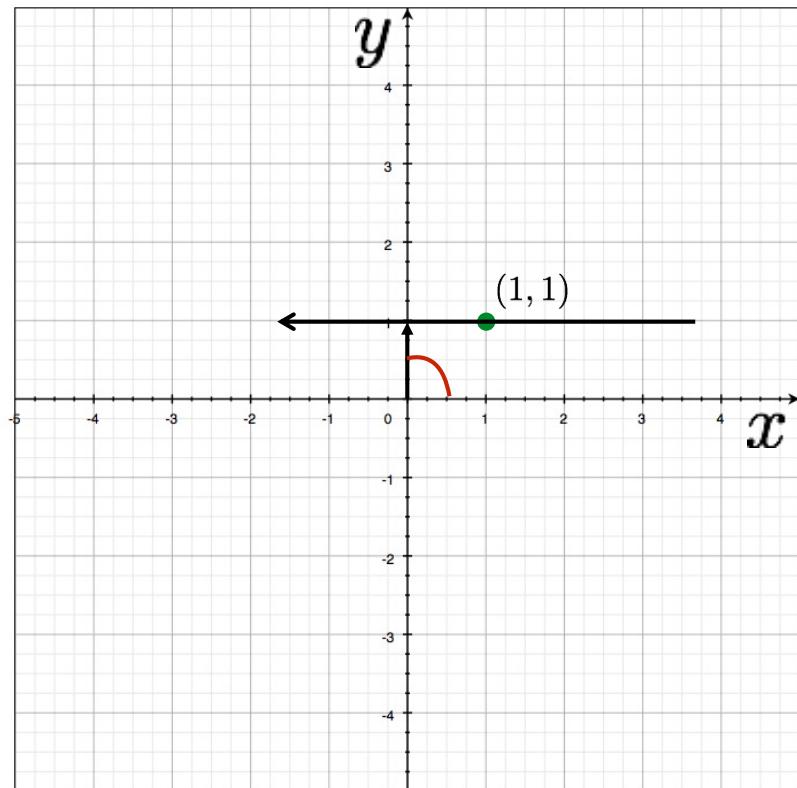
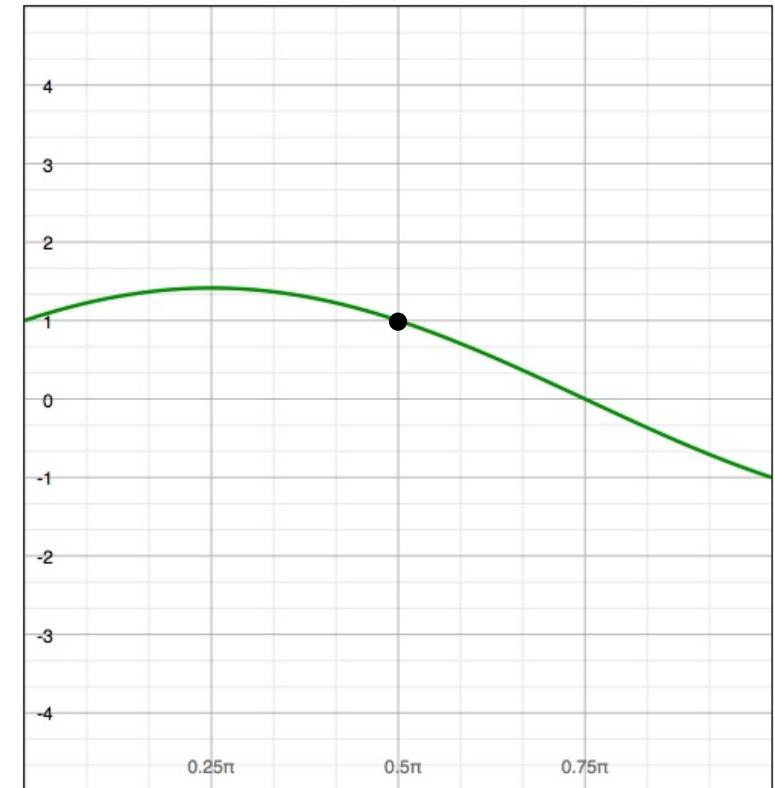


Image space

$$x \cos \theta + y \sin \theta = \rho$$

a line
becomes a
point



Parameter space

Image and parameter space

variables
 $y = mx + b$
parameters

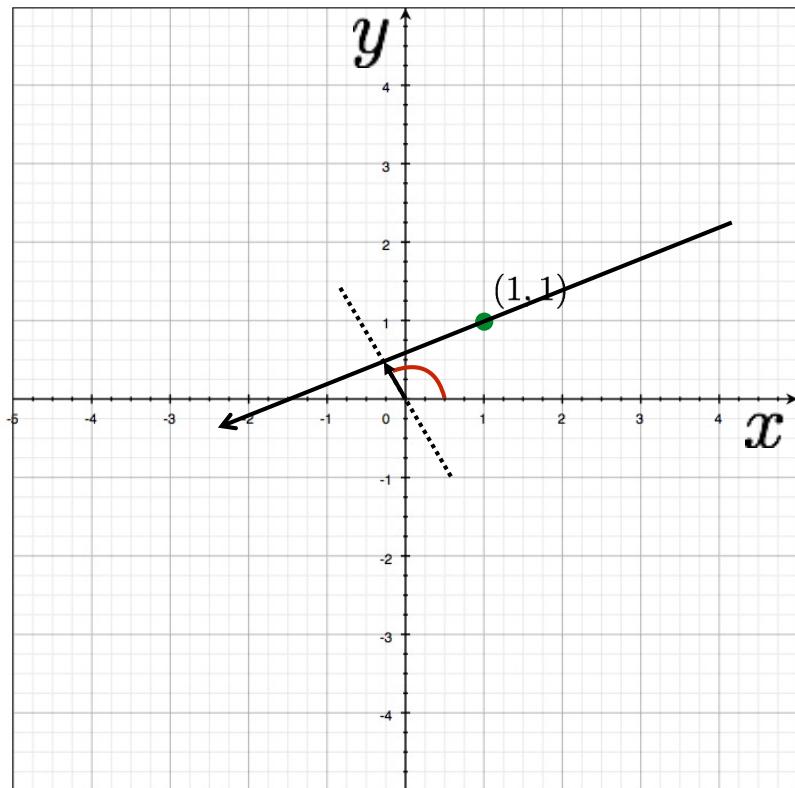
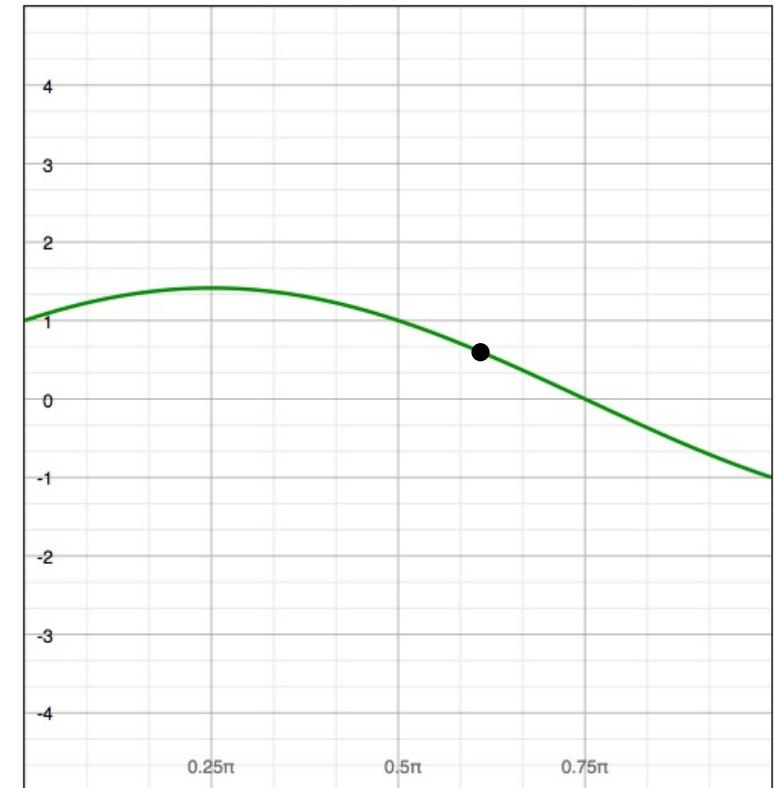


Image space

$$x \cos \theta + y \sin \theta = \rho$$

a line
becomes a
point



Parameter space

Image and parameter space

variables
 $y = mx + b$
parameters

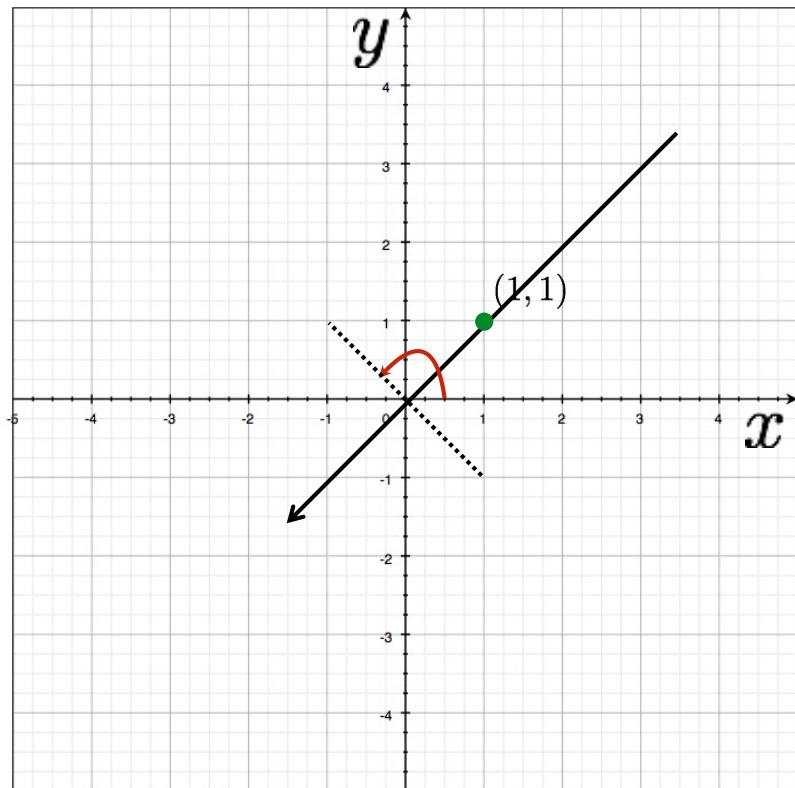
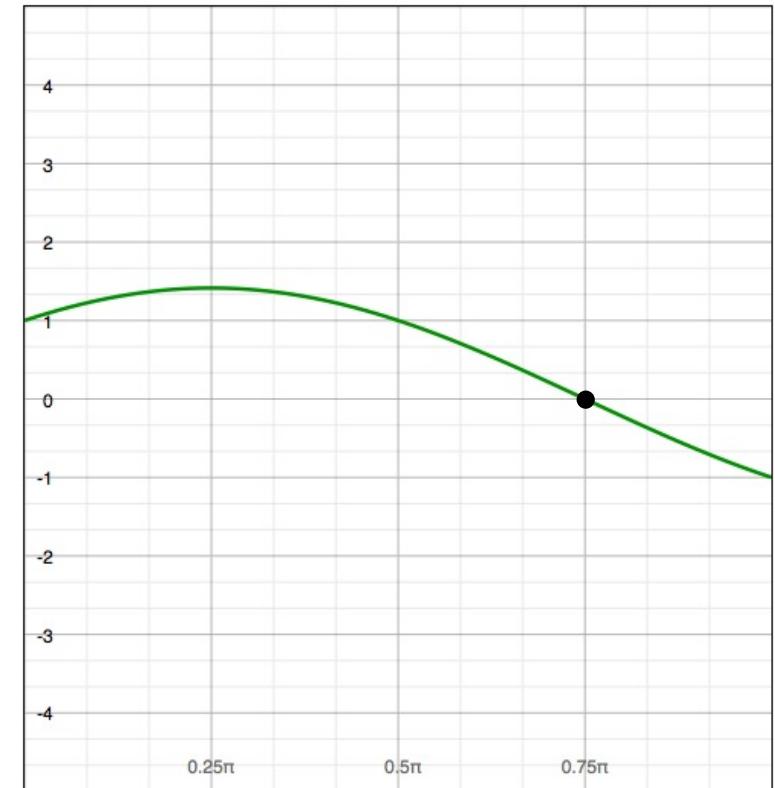


Image space

$$x \cos \theta + y \sin \theta = \rho$$

a line
becomes a
point



Parameter space

Image and parameter space

variables
 $y = mx + b$
parameters

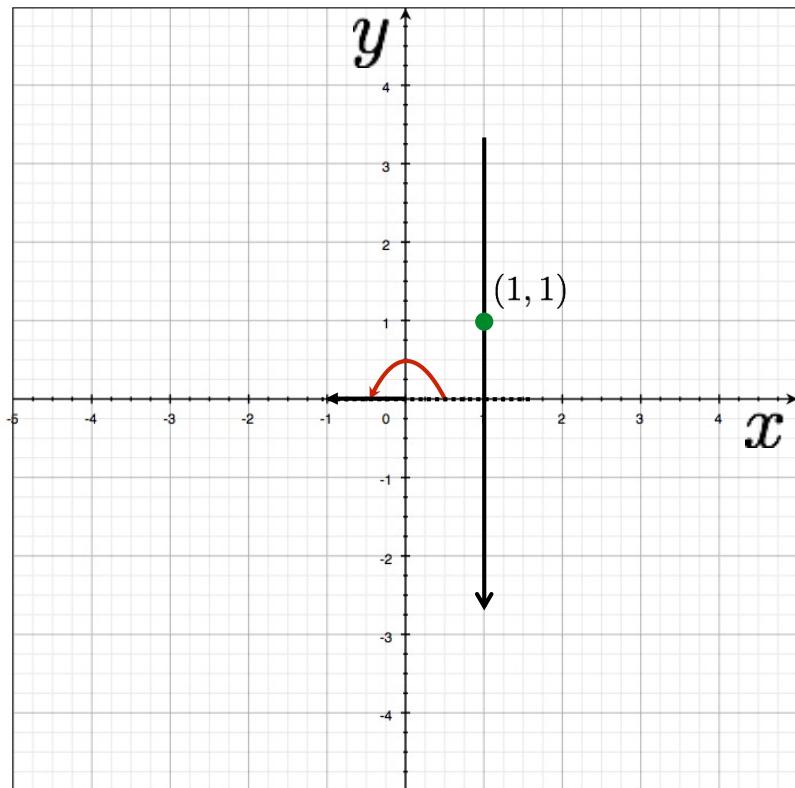
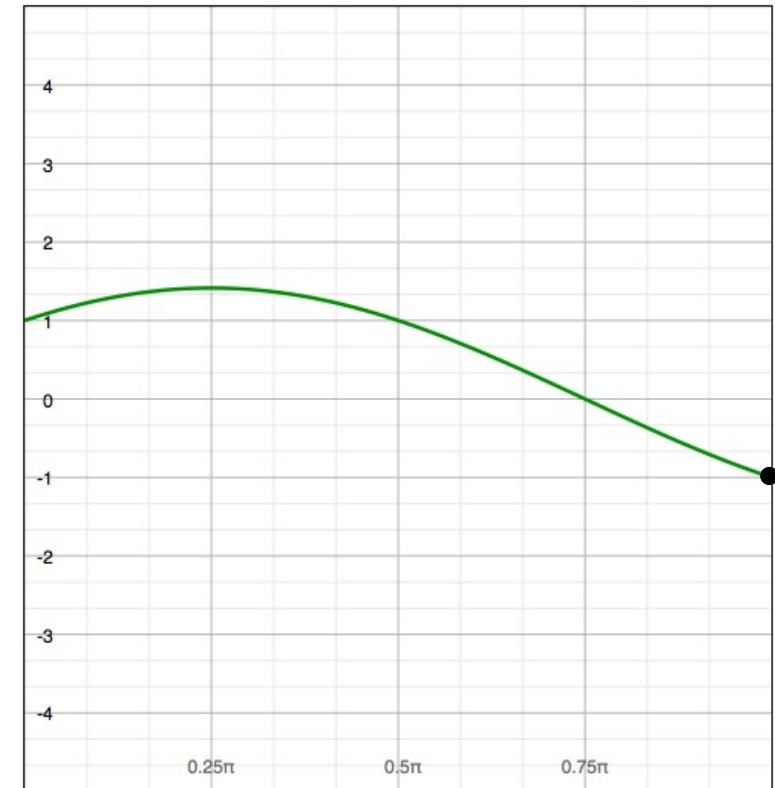


Image space

$$x \cos \theta + y \sin \theta = \rho$$

a line
becomes a
point



Parameter space

Image and parameter space

variables
 $y = mx + b$
parameters

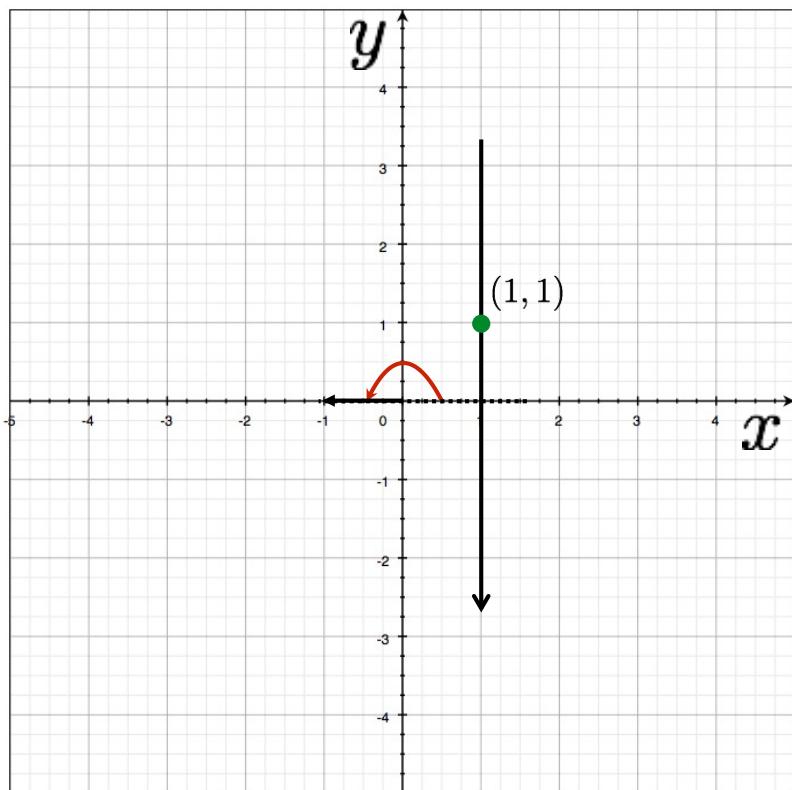
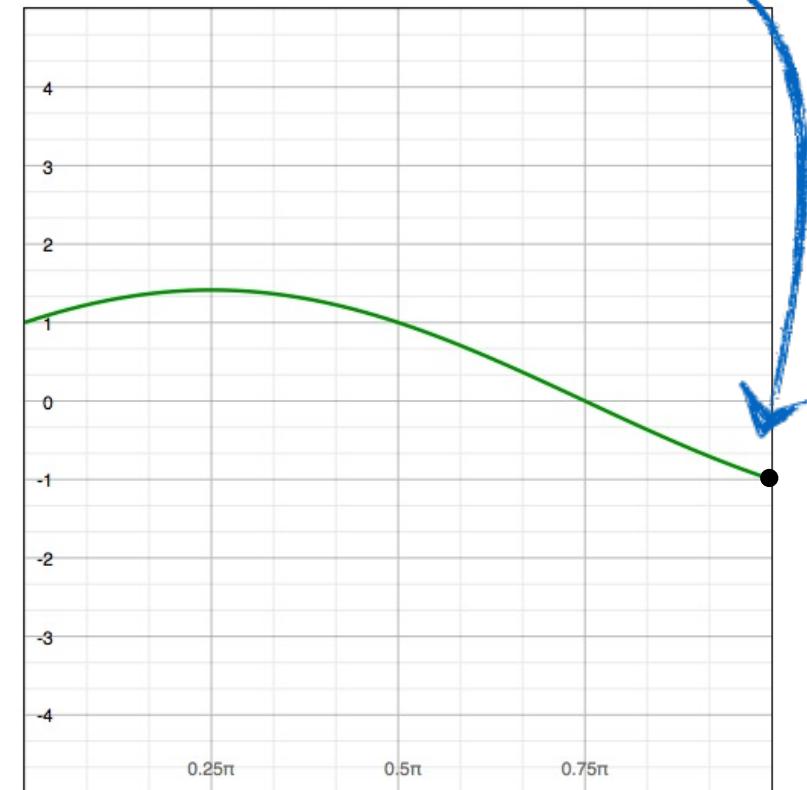


Image space

$$x \cos \theta + y \sin \theta = \rho$$

Wait ...why is rho negative?

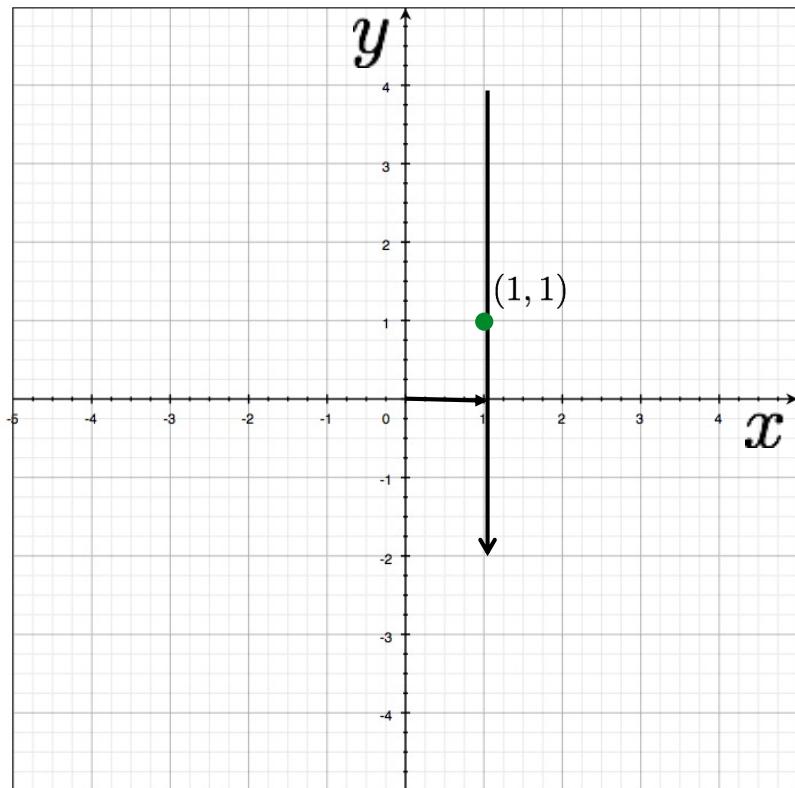


a line
becomes a
point

Parameter space

Image and parameter space

variables
 $y = mx + b$
parameters



a line
becomes a
point

$$x \cos \theta + y \sin \theta = \rho$$

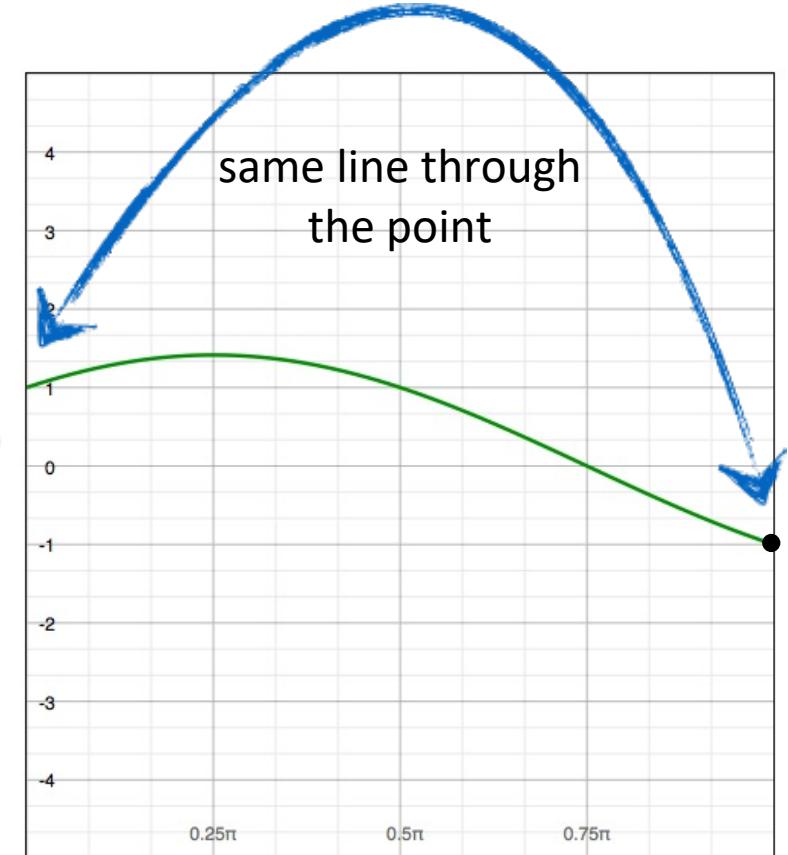


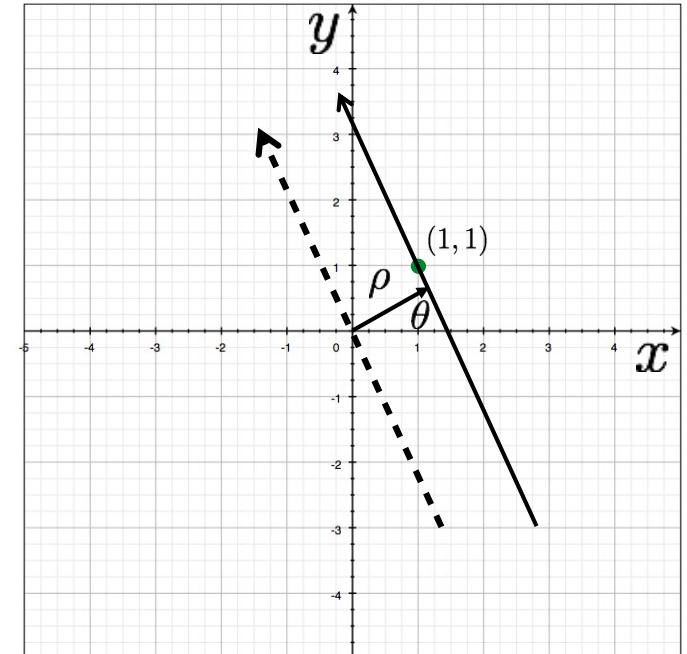
Image space

Parameter space

There are two ways to write the same line:

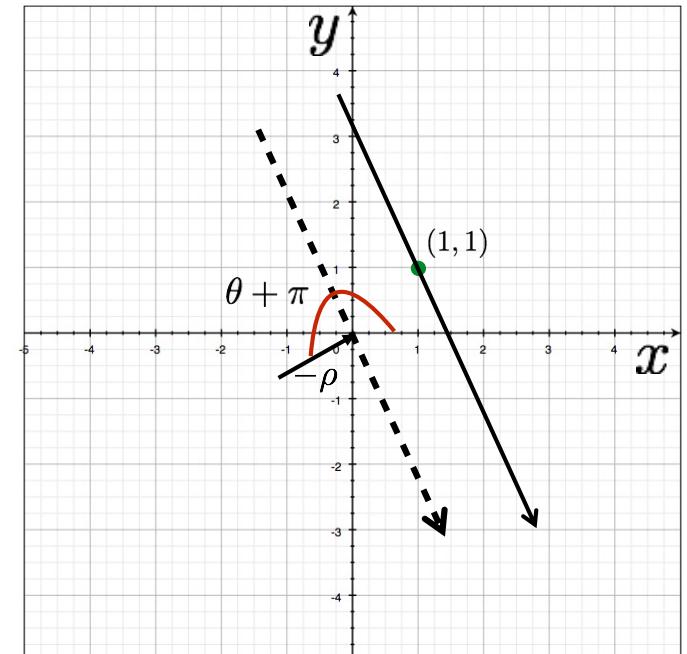
Positive rho version:

$$x \cos \theta + y \sin \theta = \rho$$



Negative rho version:

$$x \cos(\theta + \pi) + y \sin(\theta + \pi) = -\rho$$



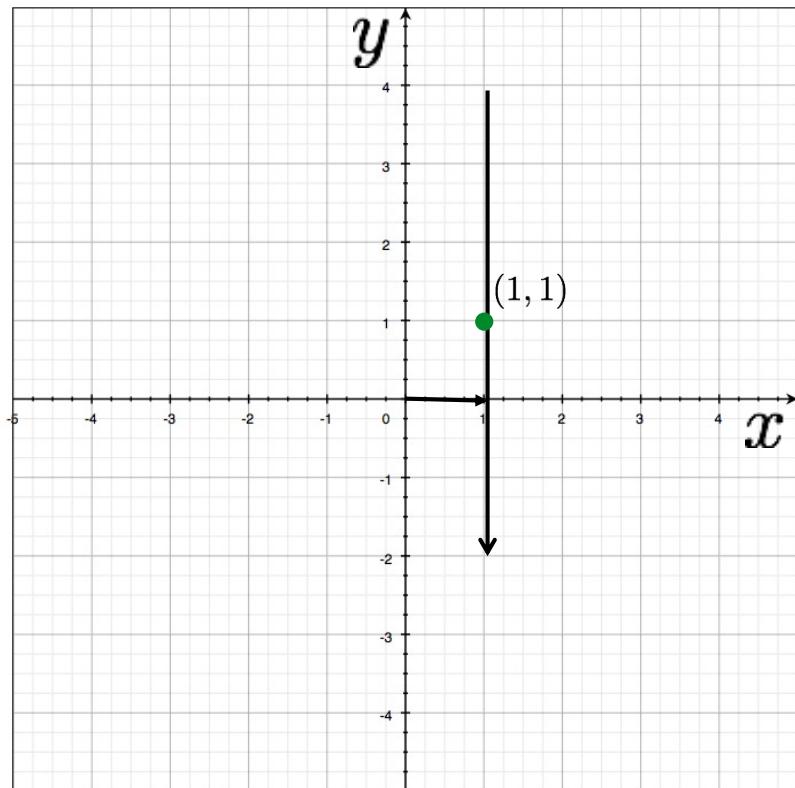
Recall:

$$\sin(\theta) = -\sin(\theta + \pi)$$

$$\cos(\theta) = -\cos(\theta + \pi)$$

Image and parameter space

variables
 $y = mx + b$
parameters



a line becomes a point

$$x \cos \theta + y \sin \theta = \rho$$



Image space

Parameter space

Image and parameter space

variables
 $y = mx + b$
parameters

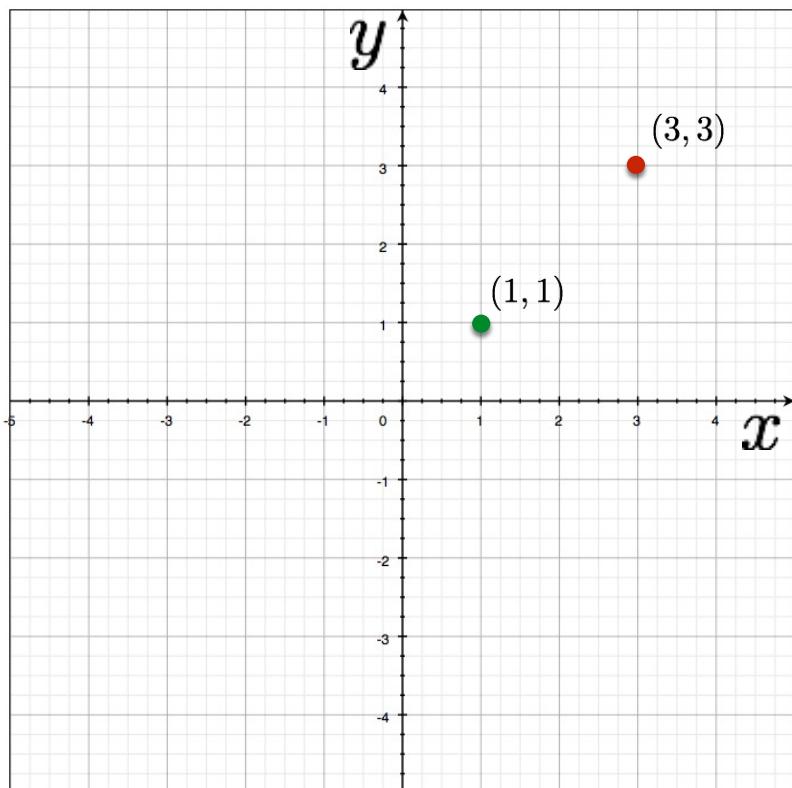
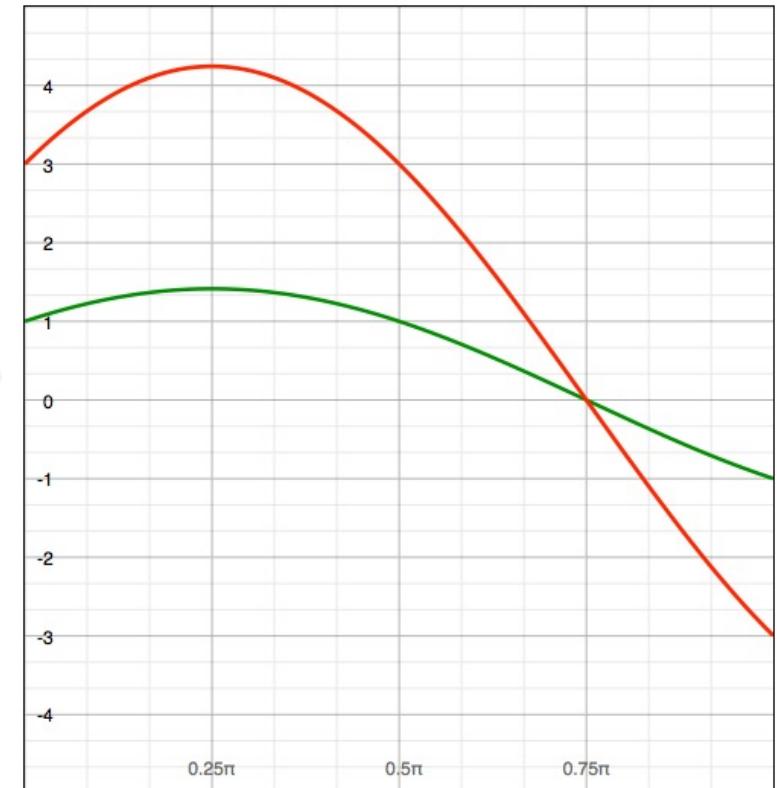


Image space

two points
become
?



Parameter space

Image and parameter space

variables
 $y = mx + b$
parameters

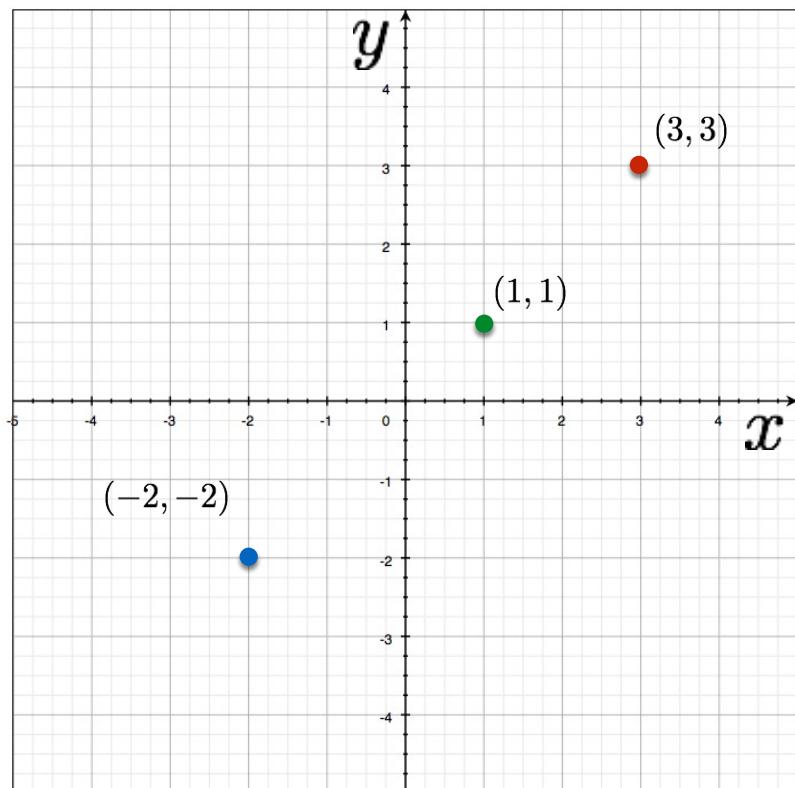
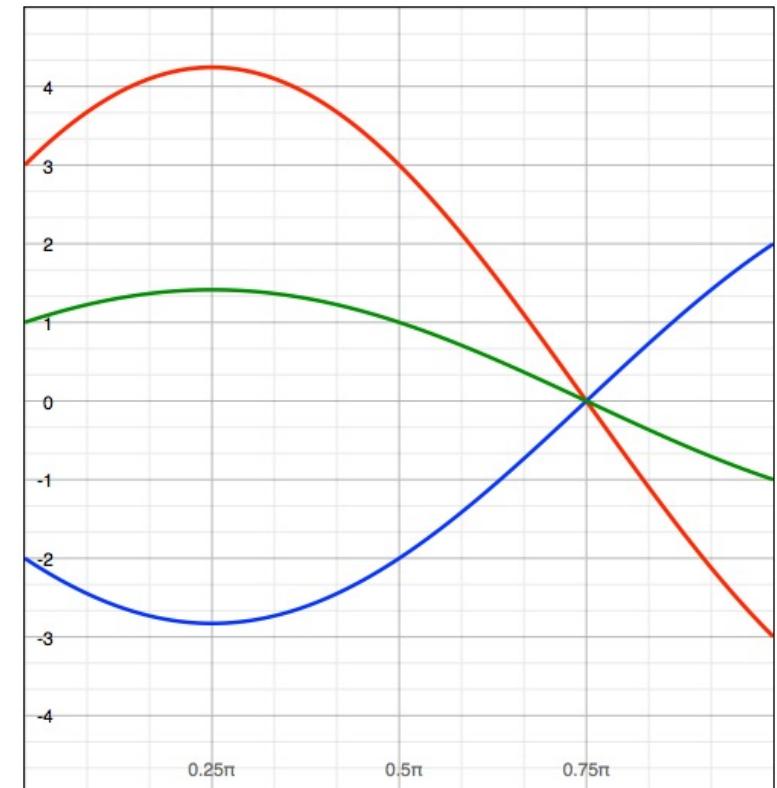


Image space

three points
become
?



Parameter space

Image and parameter space

variables
 $y = mx + b$
parameters

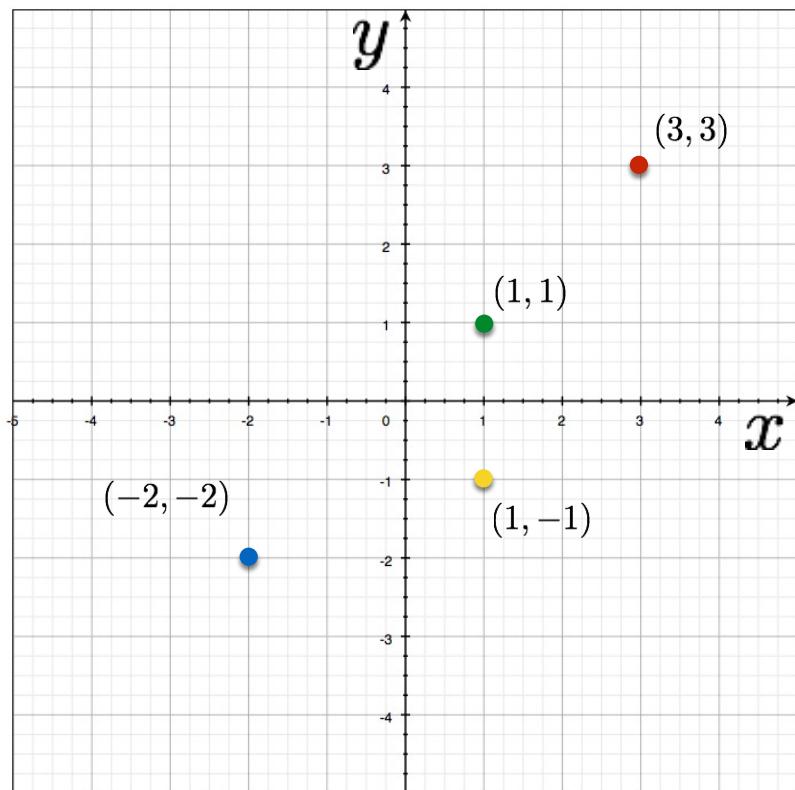
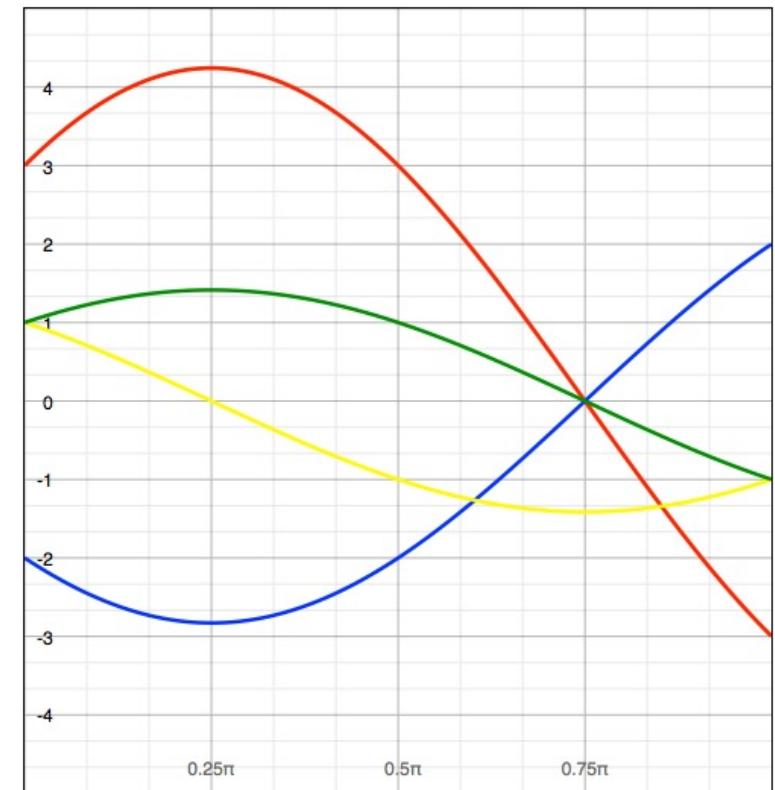


Image space

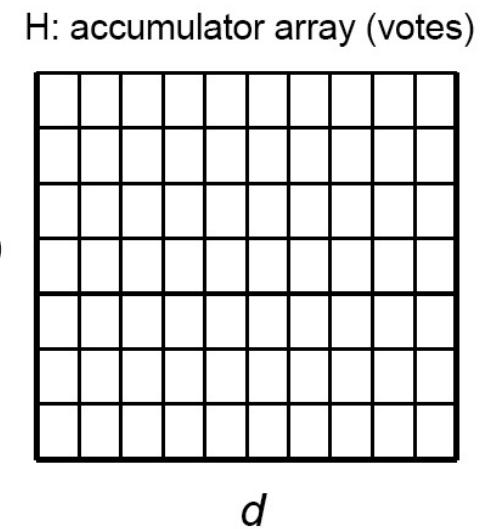
four points
become
?



Parameter space

Implementation

1. Initialize accumulator H to all zeros
2. For each edge point (x, y) in the image
 For $\theta = 0$ to 180
 $\rho = x \cos \theta + y \sin \theta$
 $H(\theta, \rho) = H(\theta, \rho) + 1$
 end
end
3. Find the value(s) of (θ, ρ) where $H(\theta, \rho)$ is a local maximum
4. The detected line in the image is given by
$$\rho = x \cos \theta + y \sin \theta$$



NOTE: Watch your coordinates. Image origin is top left!

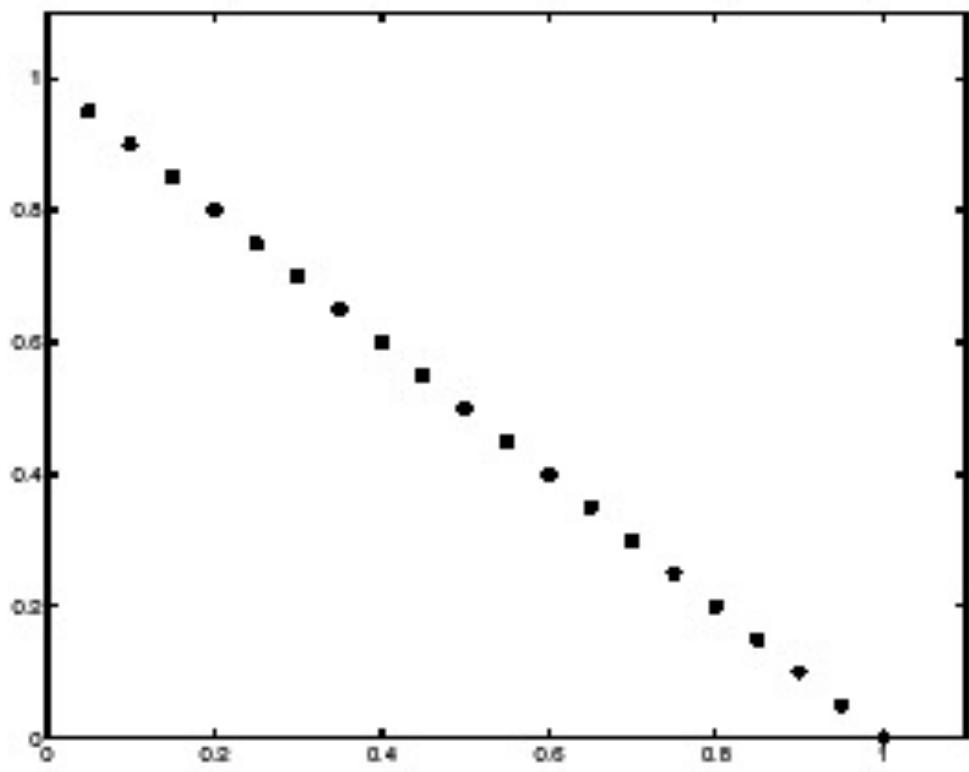
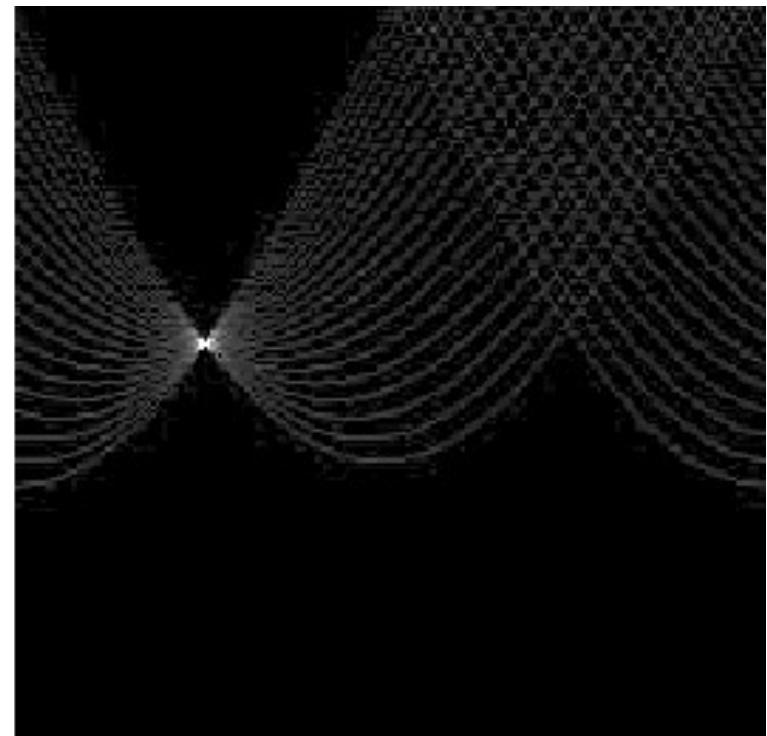


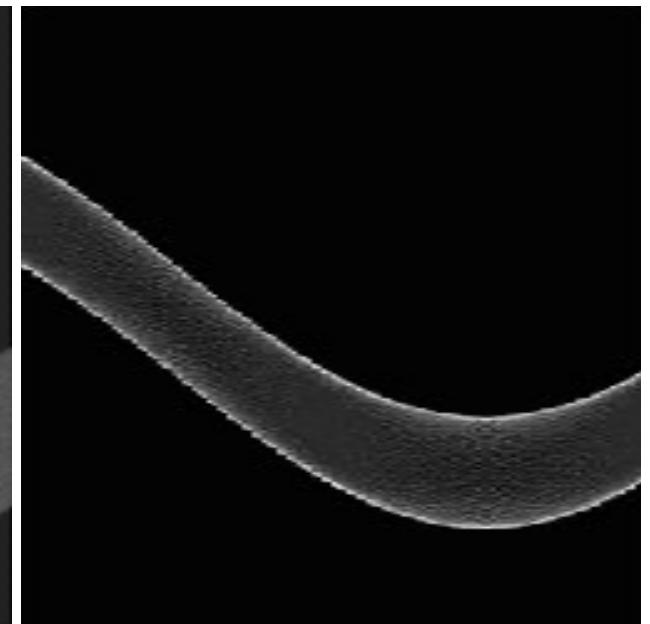
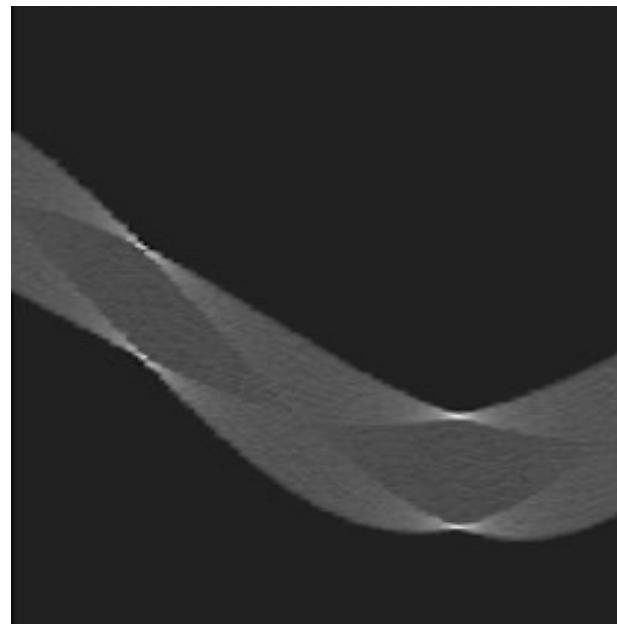
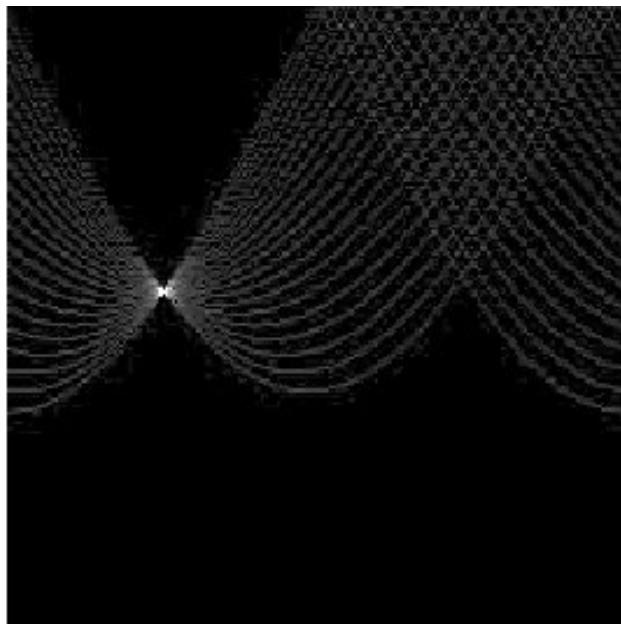
Image space



Votes

Basic shapes

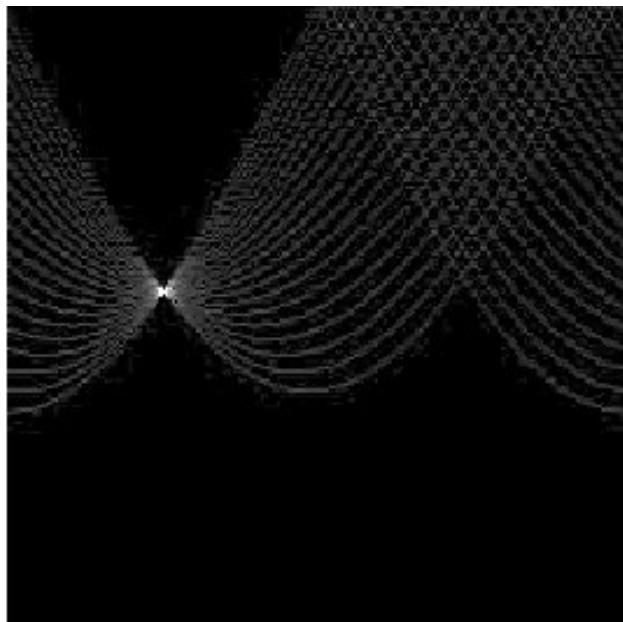
(in parameter space)



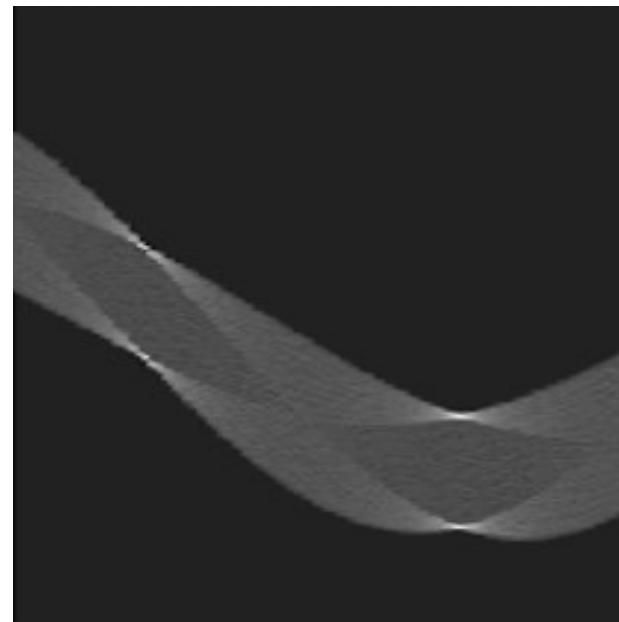
can you guess the shape?

Basic shapes

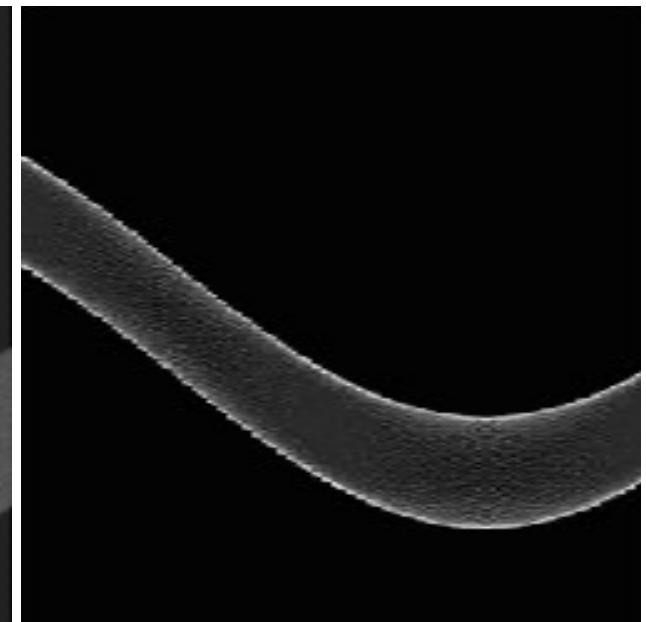
(in parameter space)



line

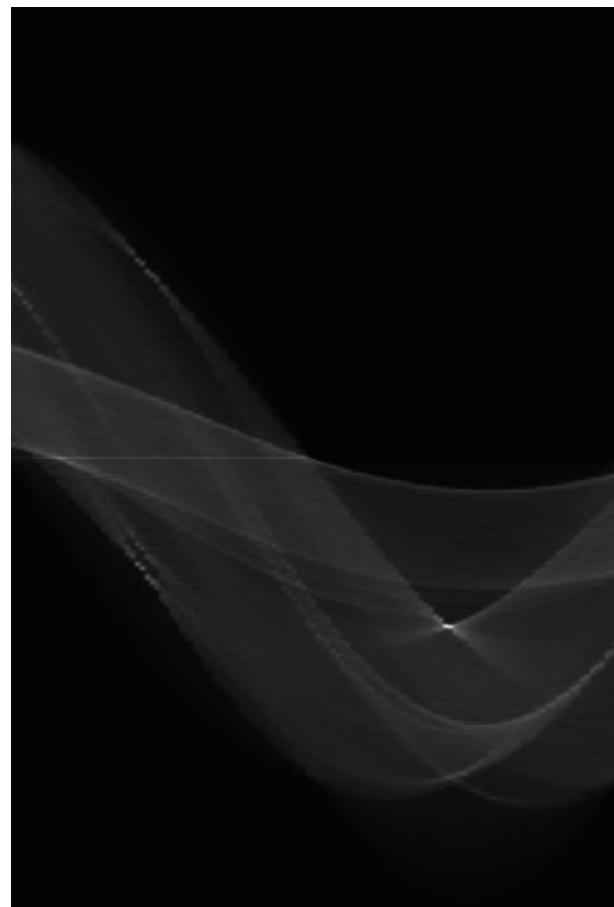
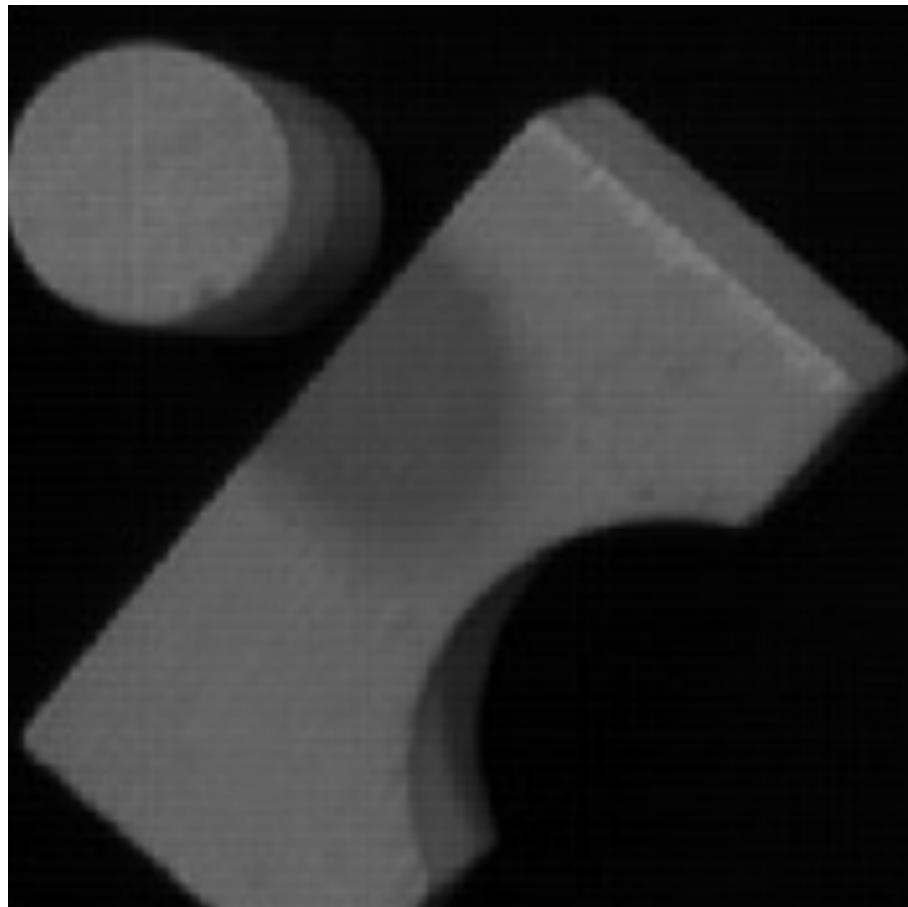


rectangle

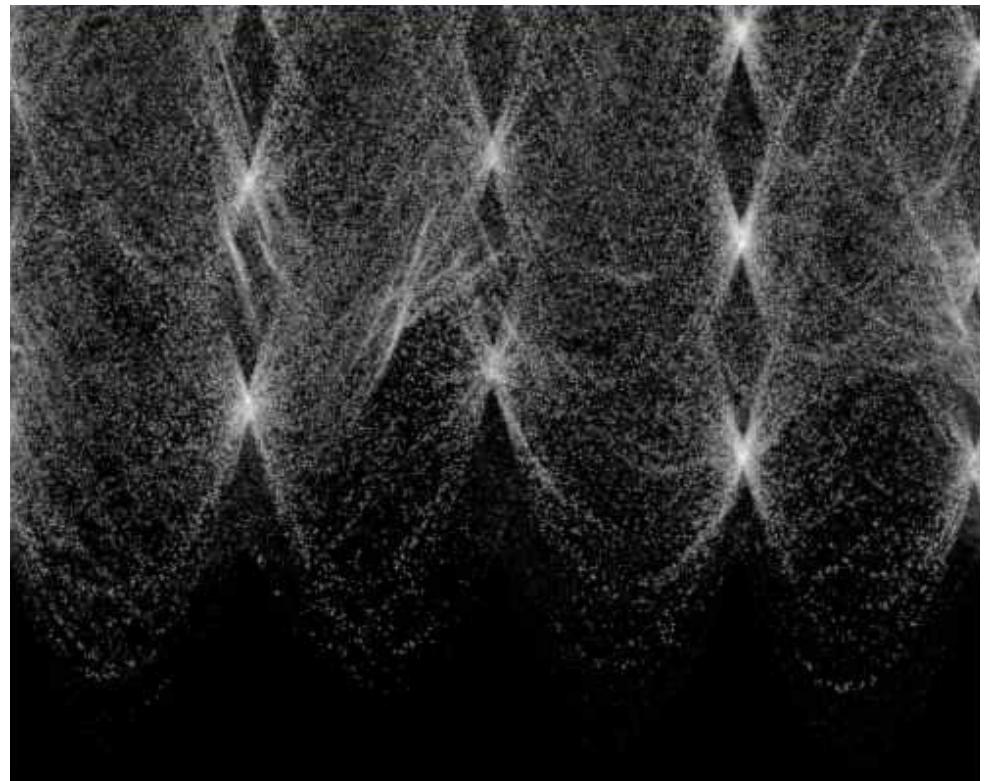


circle

Basic Shapes



More complex image



In practice, measurements are noisy...

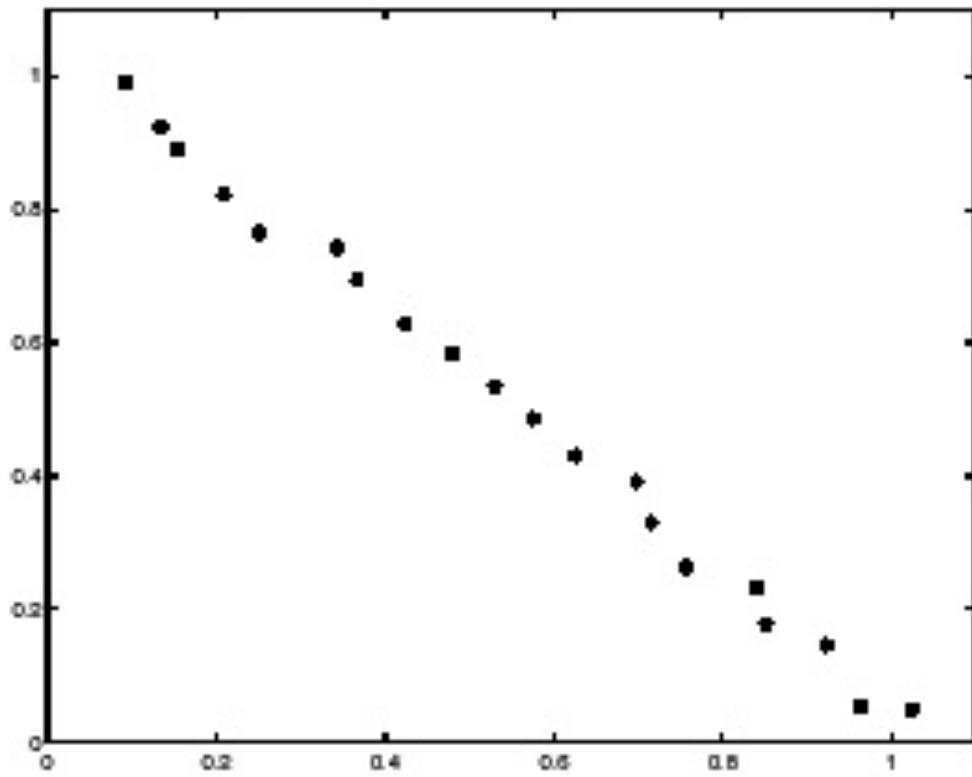
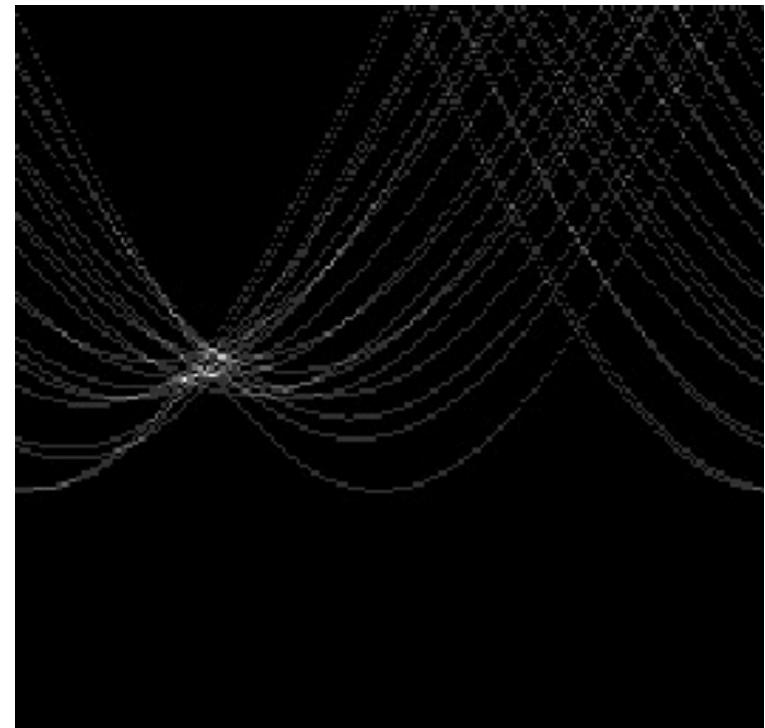


Image space



Votes

Too much noise ...

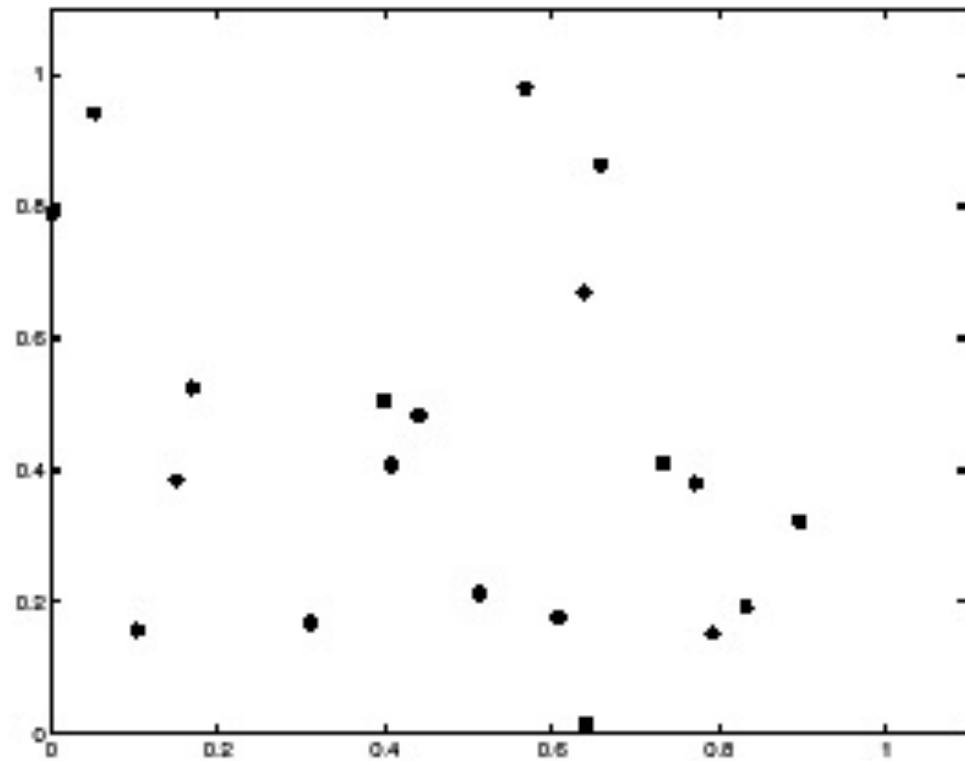
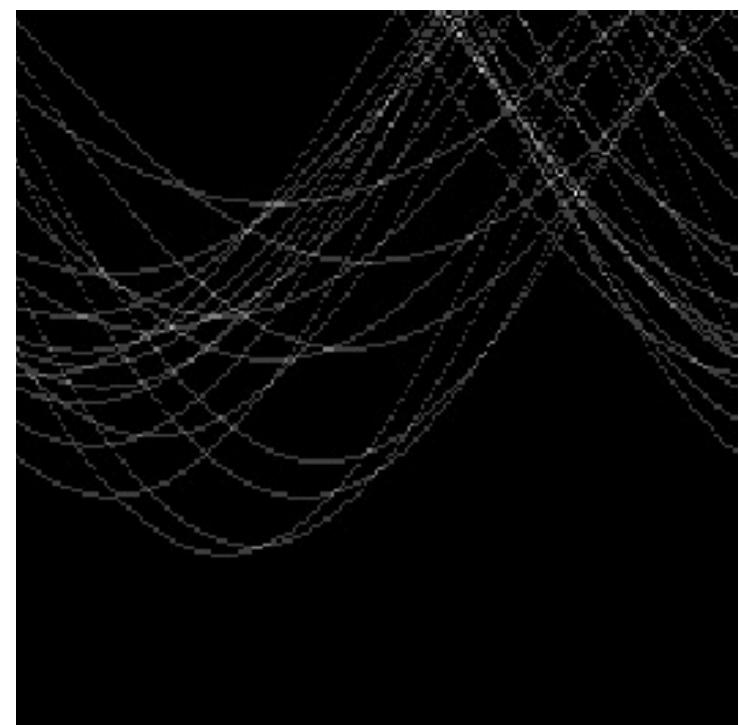


Image space

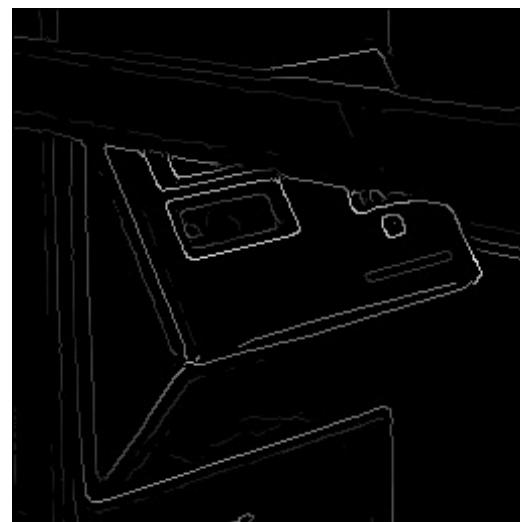


Votes

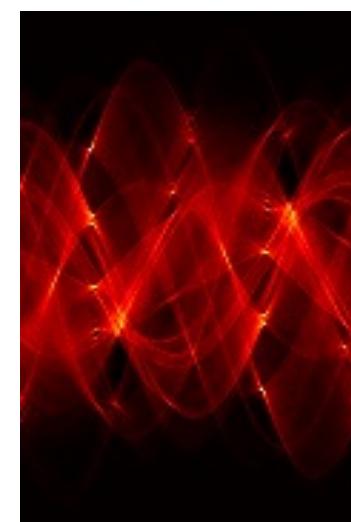
Real-world example



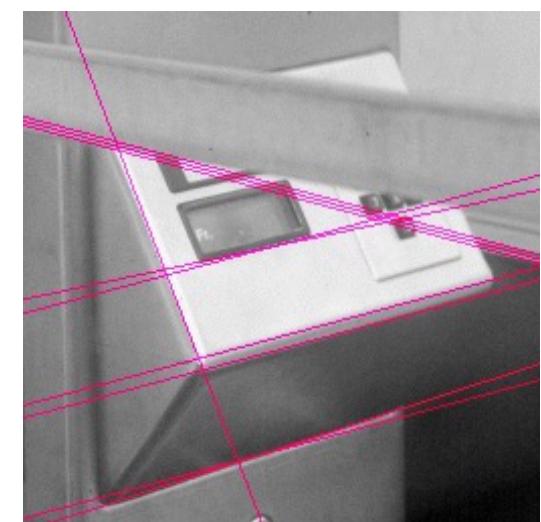
Original



Edges



parameter space



Hough Lines