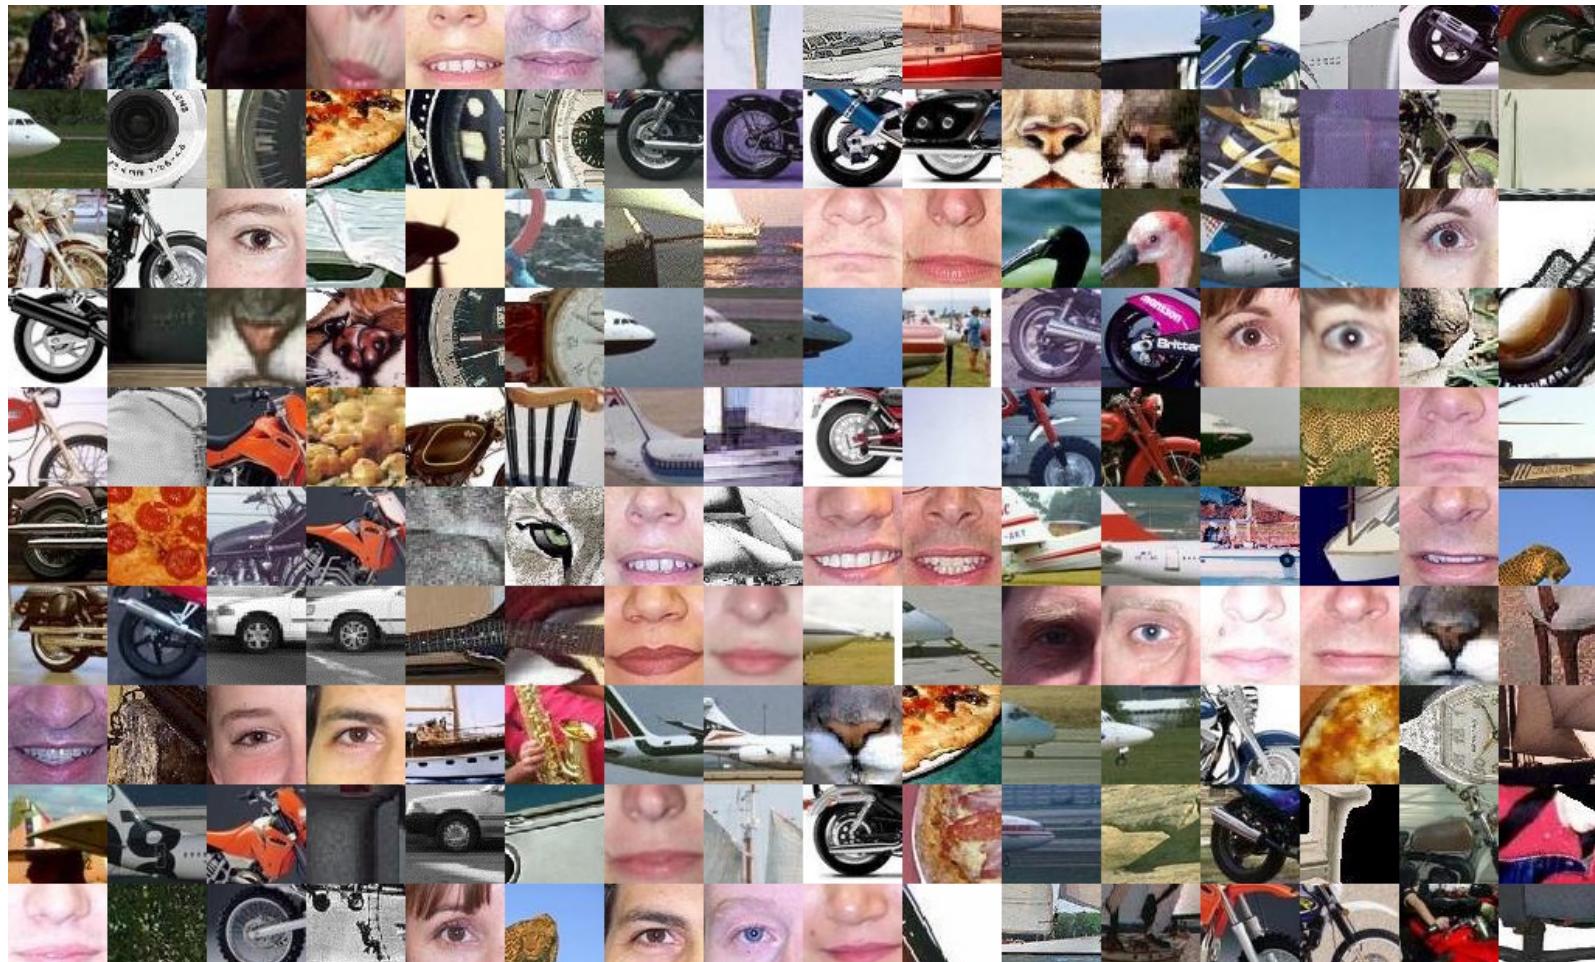


# Feature detectors and descriptors



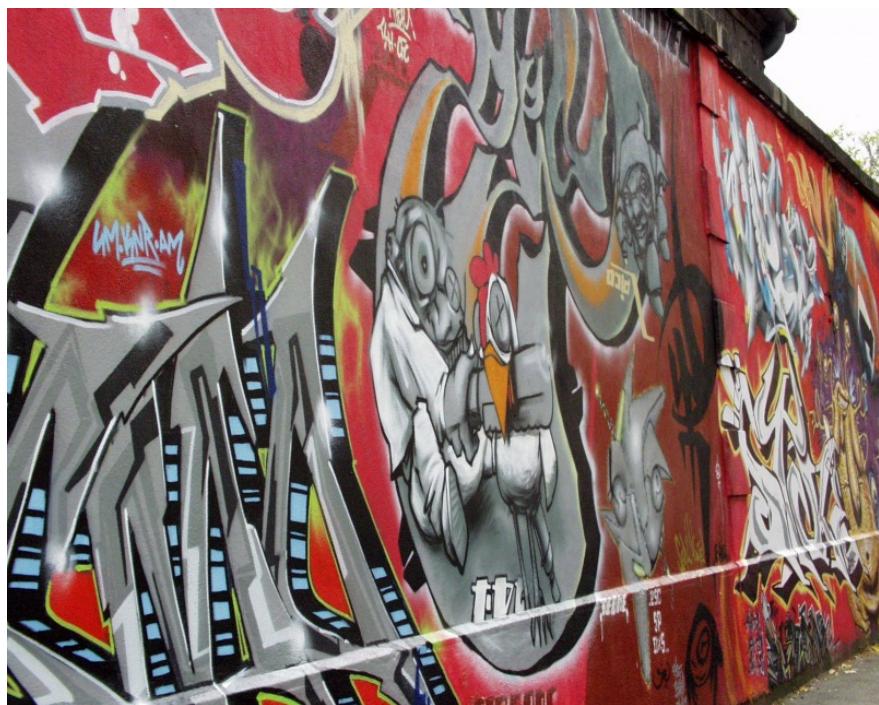
Computer Vision  
Fall 2022, Lecture 7

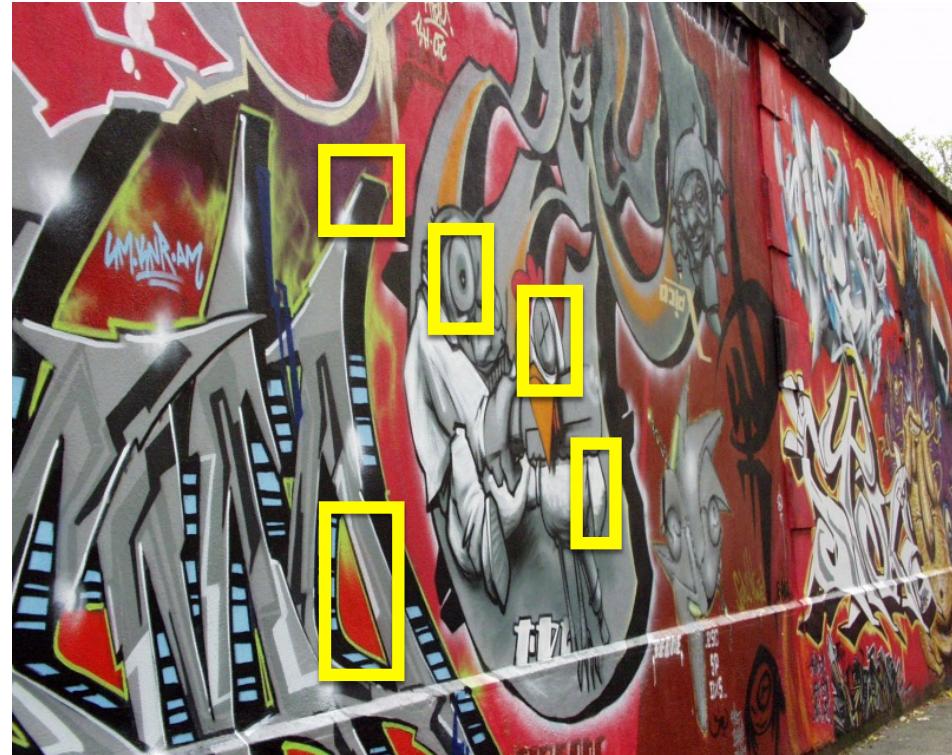
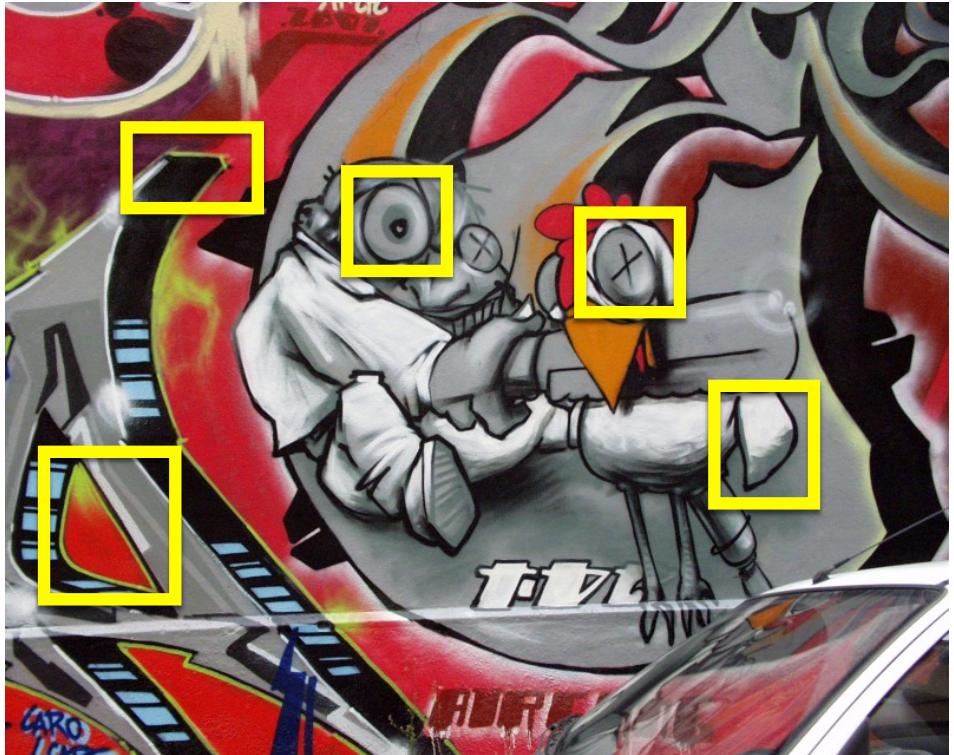
Why do we need feature  
descriptors?

A



B



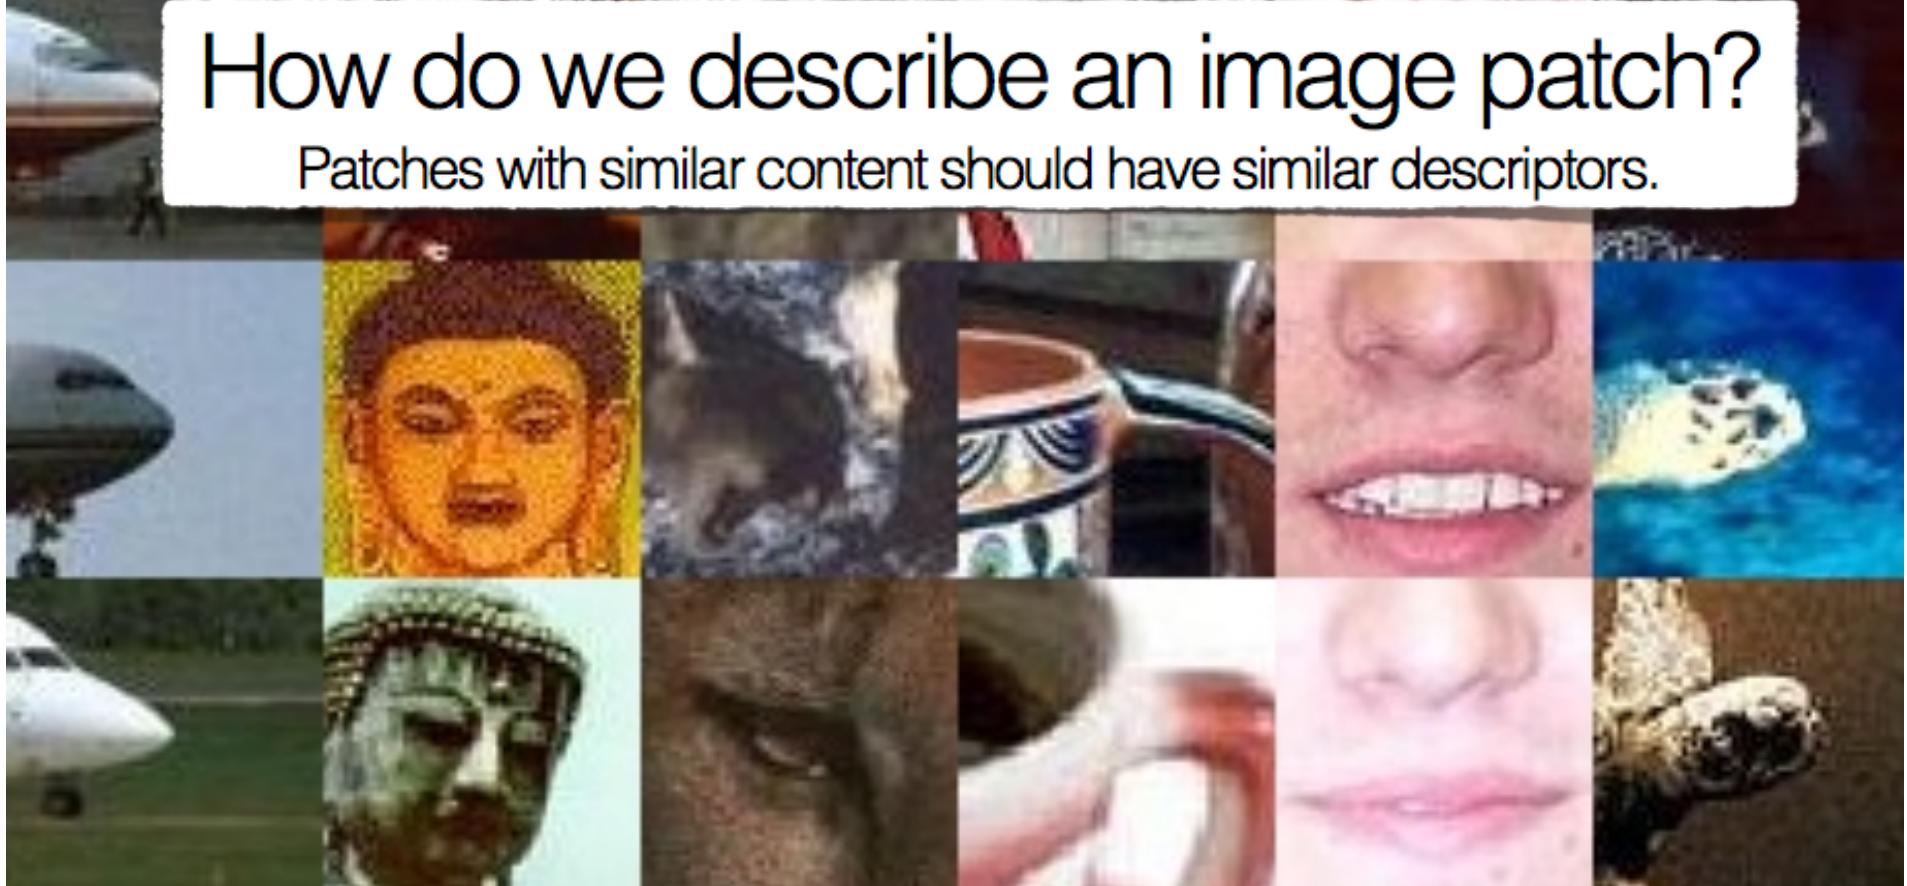


*If we know where the good features are,  
how do we match them?*



## How do we describe an image patch?

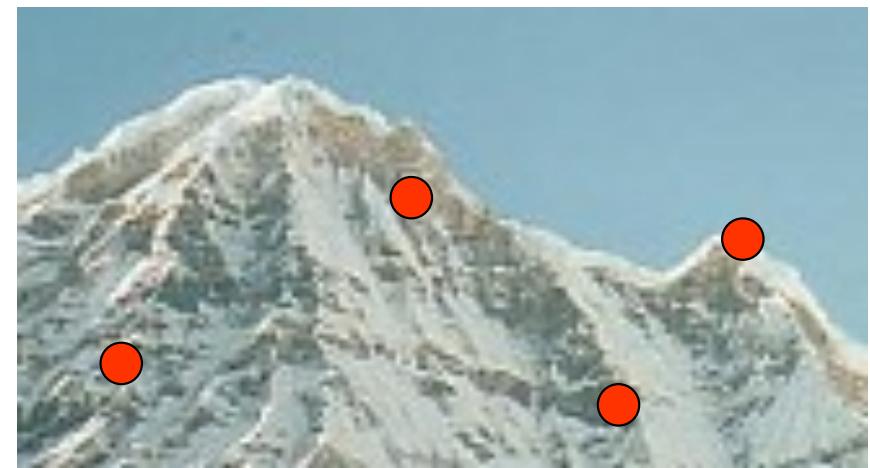
Patches with similar content should have similar descriptors.



# Designing feature descriptors

# Goal: interest operator repeatability

- We want to detect (at least some of) the same points in both images.

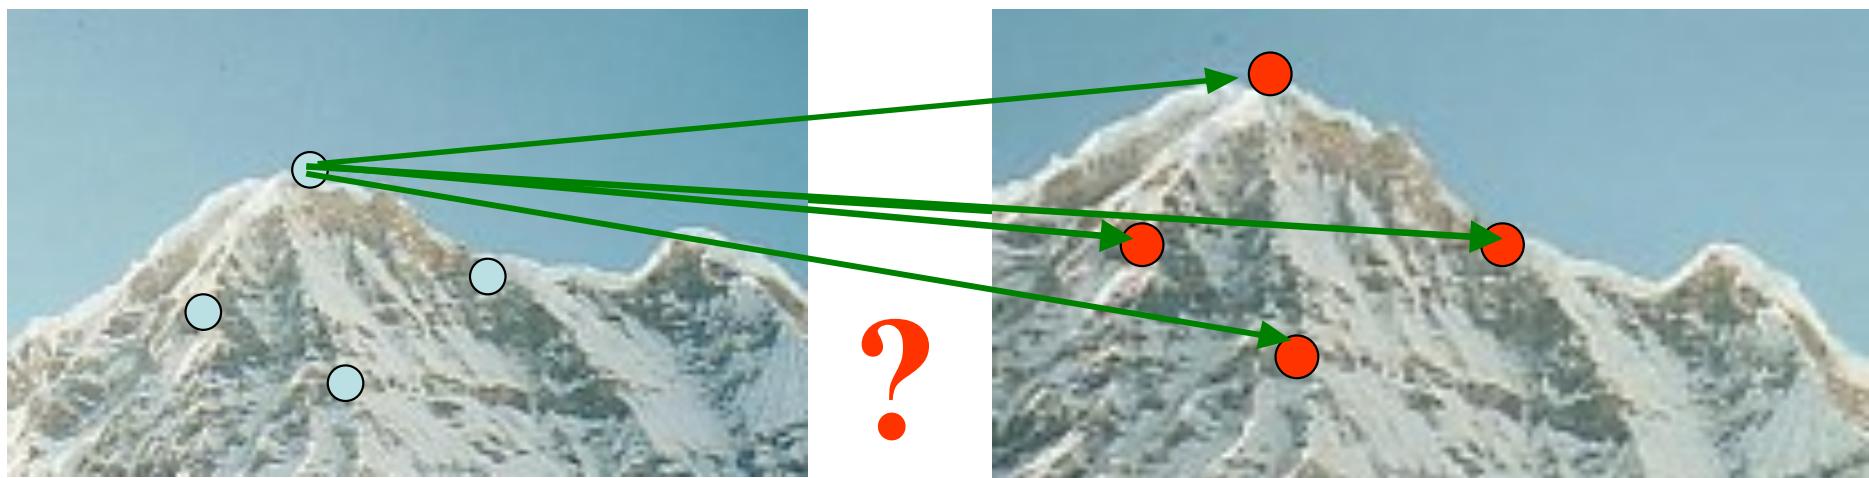


No chance to find true matches!

- Yet we have to be able to run the detection procedure independently per image.

# Goal: descriptor distinctiveness

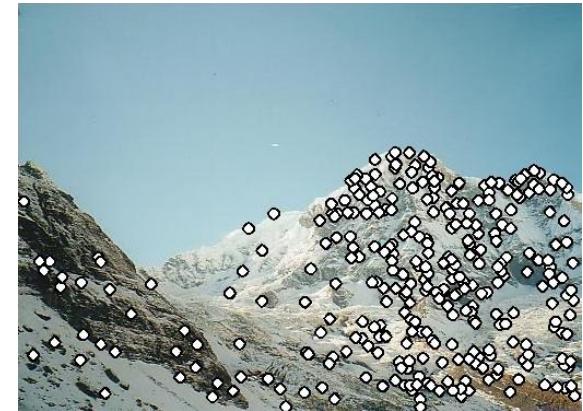
- We want to be able to reliably determine which point goes with which.



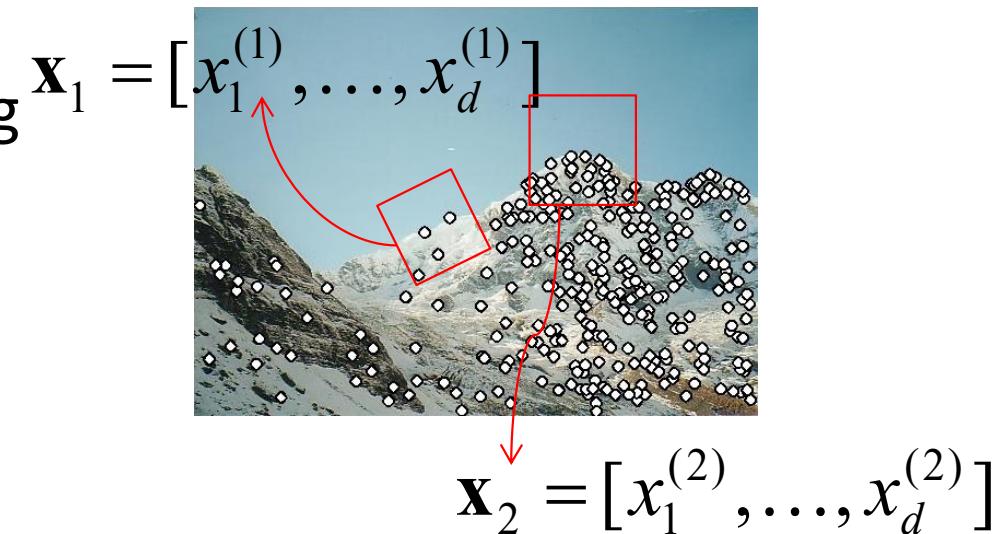
- Must provide some invariance to geometric and photometric differences between the two views.

# Local invariant features: Outline

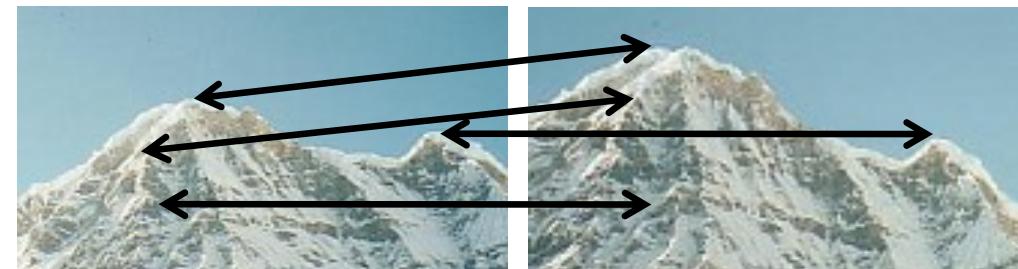
- 1) **Detection:** Identify the interest points



- 2) **Description:** Extract vector feature descriptor surrounding each interest point.

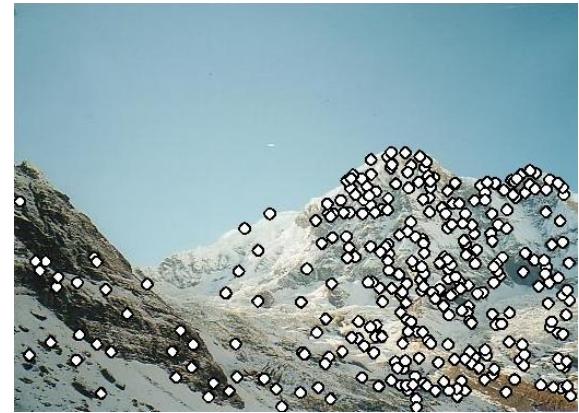


- 3) **Matching:** Determine correspondence between descriptors in two views



# Local features: main components

- 1) **Detection:** Identify the interest points



- 2) Description: Extract vector feature descriptor surrounding each interest point.
- 3) Matching: Determine correspondence between descriptors in two views

# We want invariance!!!

- To illumination
- To scale
- To rotation
- To affine
- To perspective projection

# Types of invariance

- Illumination



# Types of invariance

- Illumination
- Scale



# Types of invariance

- Illumination
- Scale
- Rotation



# Types of invariance

- Illumination
- Scale
- Rotation
- Affine (view point change)



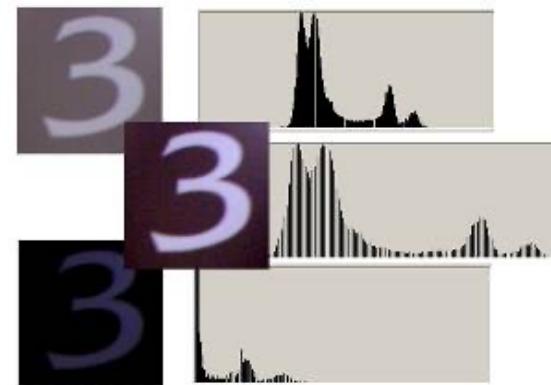
# Types of invariance

- Illumination
- Scale
- Rotation
- Affine
- Full Perspective



# How to achieve illumination invariance

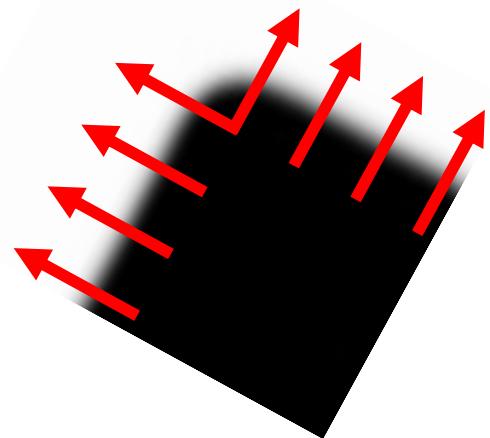
- The easy way  
(normalized)



## Recall: Corners as distinctive interest points

Since  $M$  is symmetric, we have

$$M = X \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} X^T$$

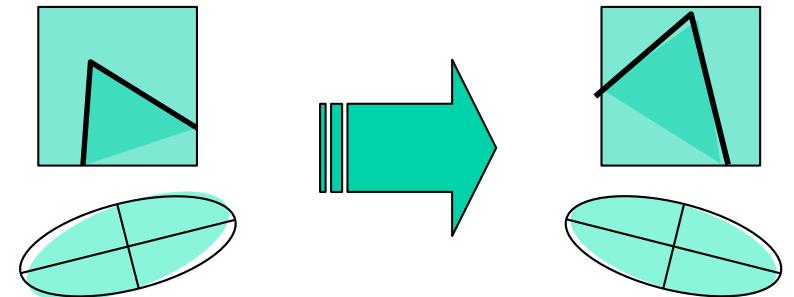


$$Mx_i = \lambda_i x_i$$

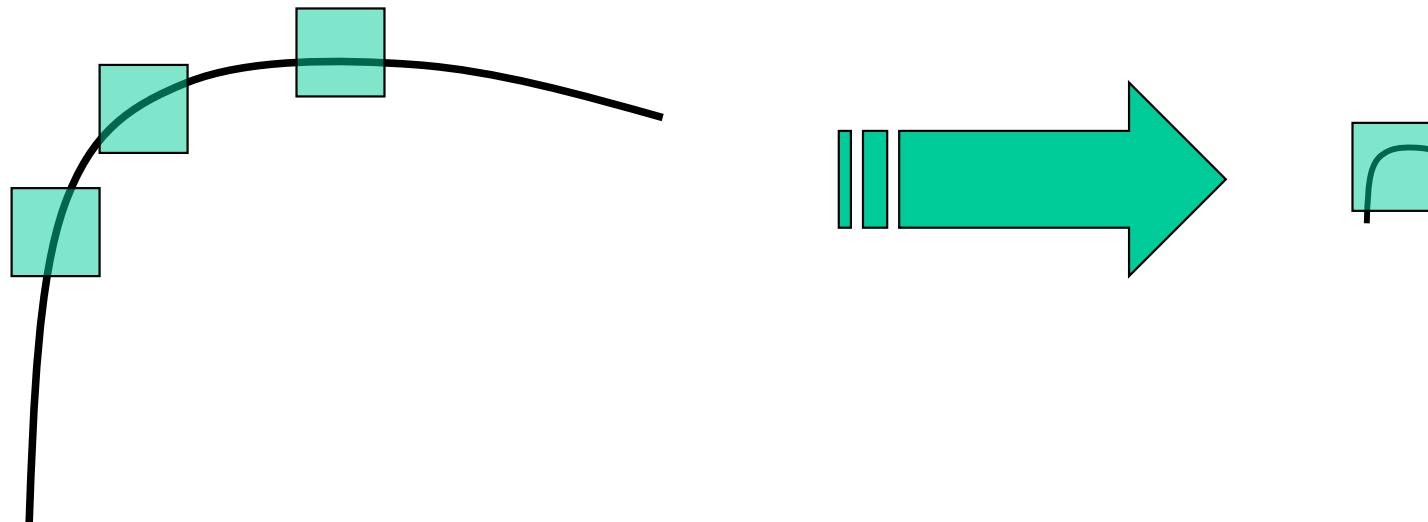
The eigenvalues of  $M$  reveal the amount of intensity change in the two principal orthogonal gradient directions in the window.

# Properties of the Harris corner detector

Rotation invariant? Yes



Scale invariant? No



All points will be  
classified as edges

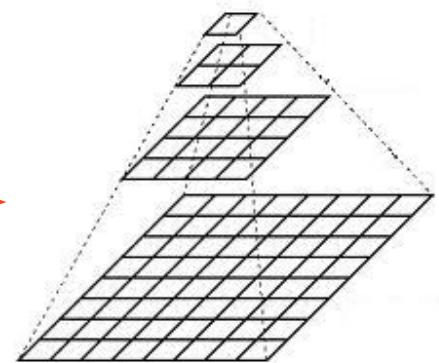
Corner !

# Scale invariant interest points

How can we independently select interest points in each image, such that the detections are repeatable across different scales?



 Image Pyramid -> Multi-Scale



Intuitively...

Find local maxima in both **position** and **scale**

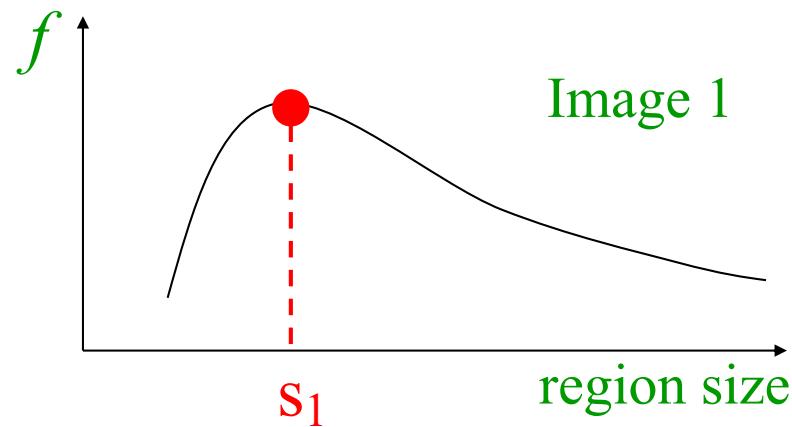
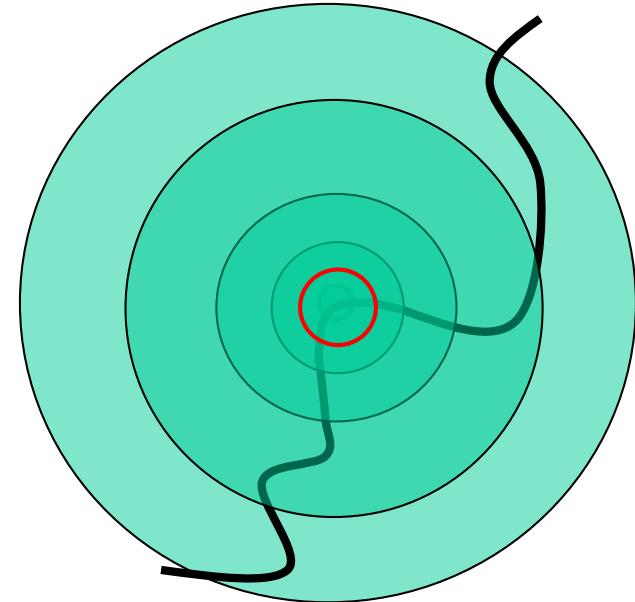
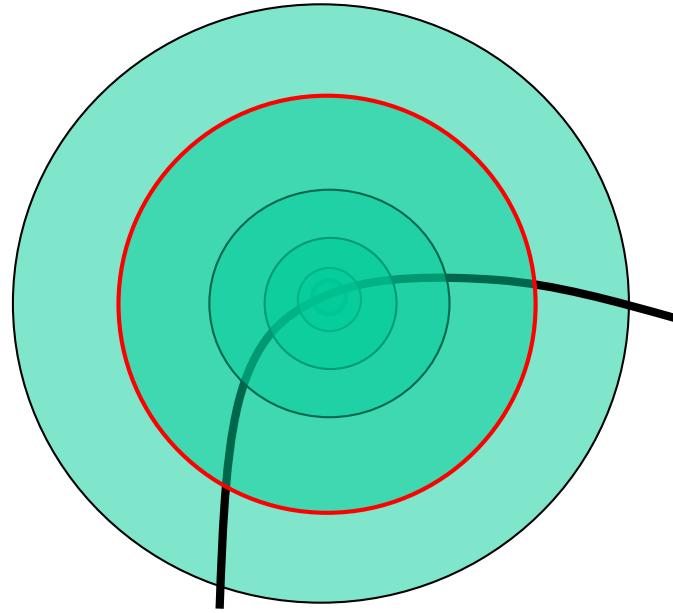


Image 1

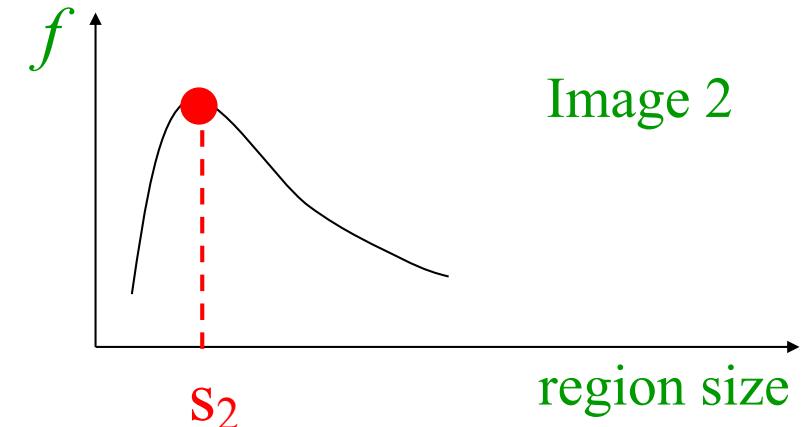
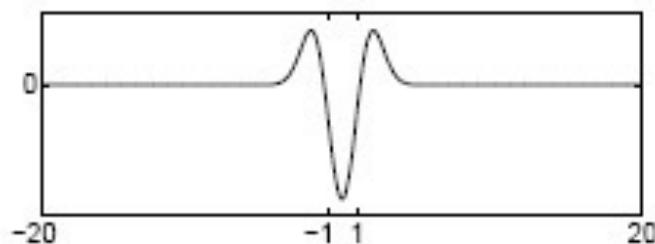


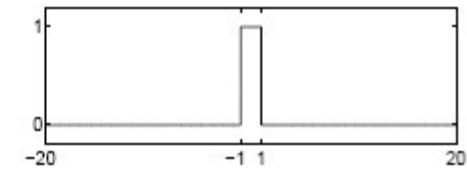
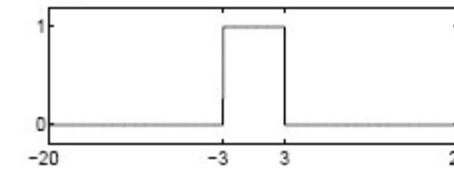
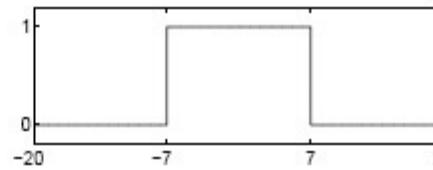
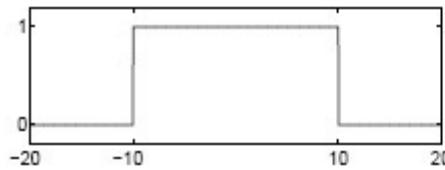
Image 2

Formally...

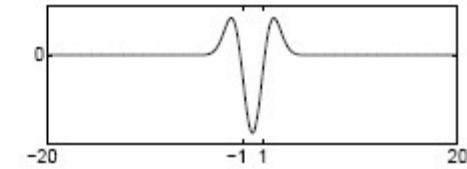
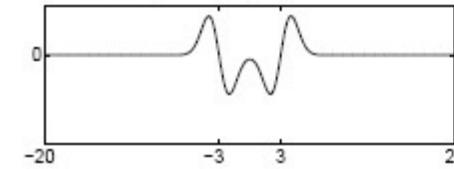
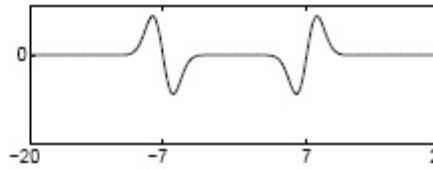
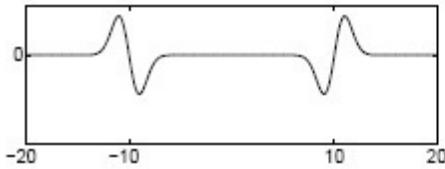
Laplacian filter



Original signal

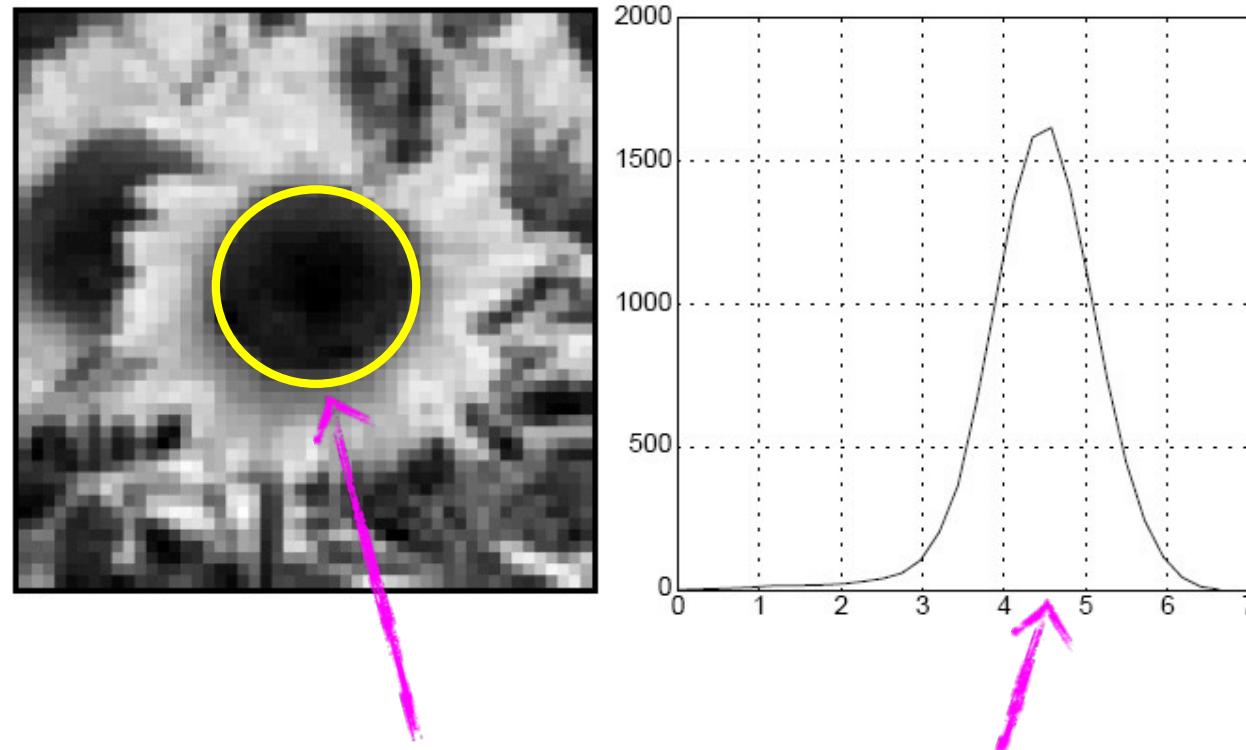


Convolved with Laplacian ( $\sigma = 1$ )



Highest response when the signal has the same **characteristic scale** as the filter

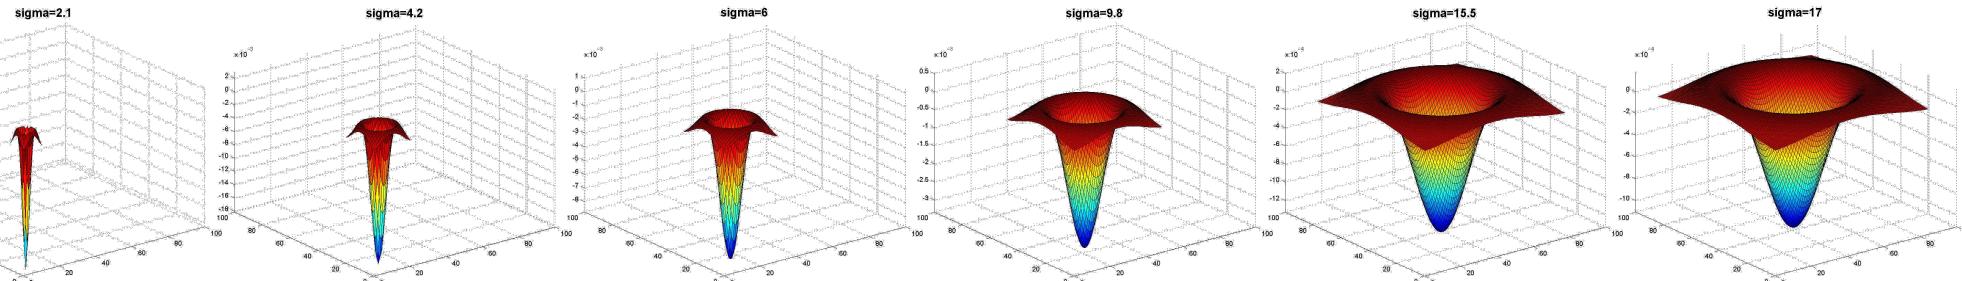
characteristic scale - the scale that produces peak filter response



characteristic scale

**we need to search over characteristic scales**

# What happens if you apply different Laplacian filters?



Full size



3/4 size



What happened when you applied different Laplacian filters?

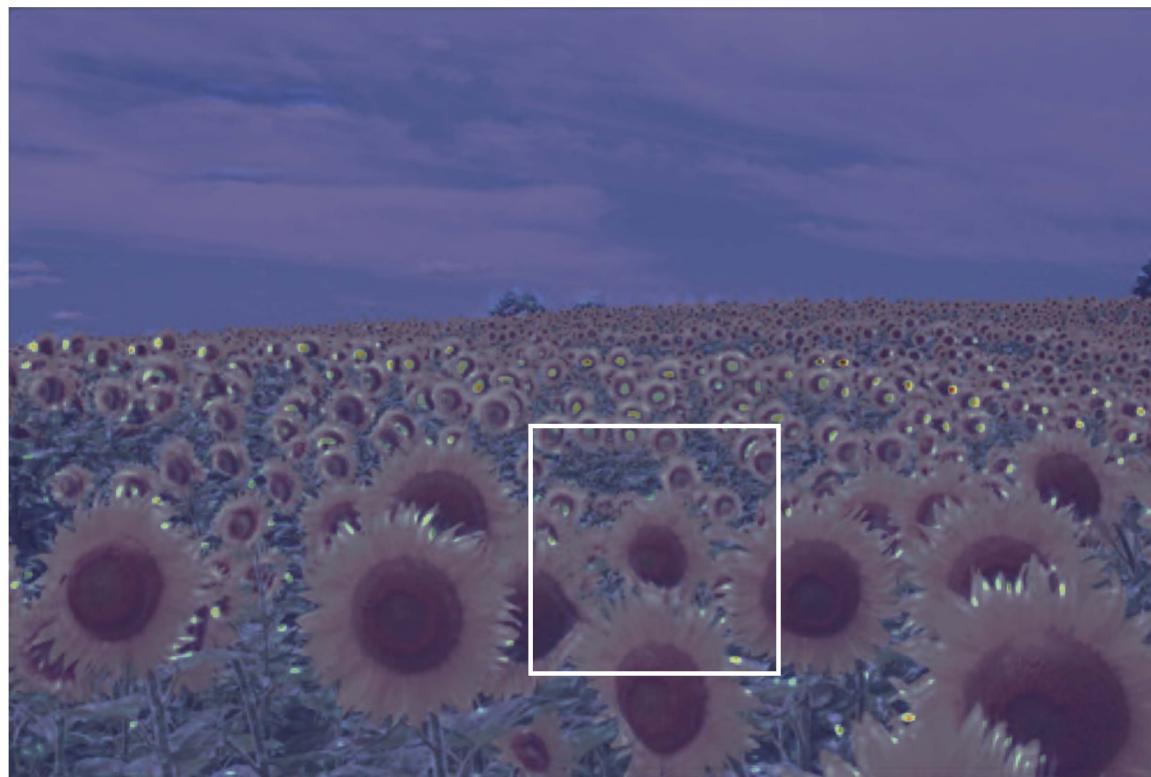
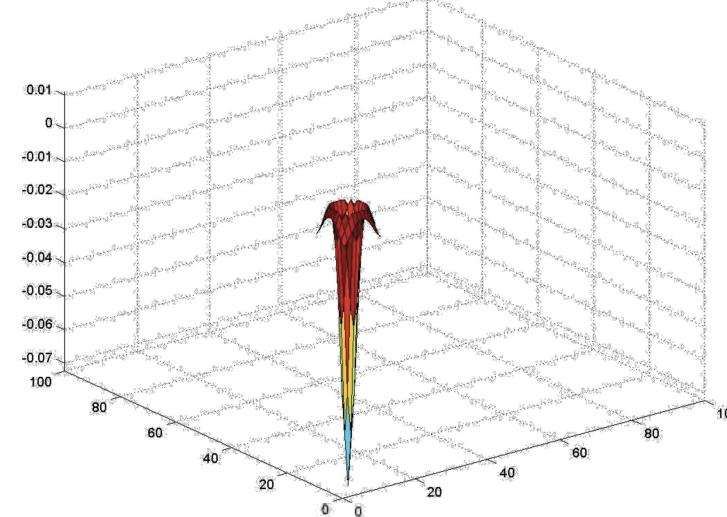
Full size

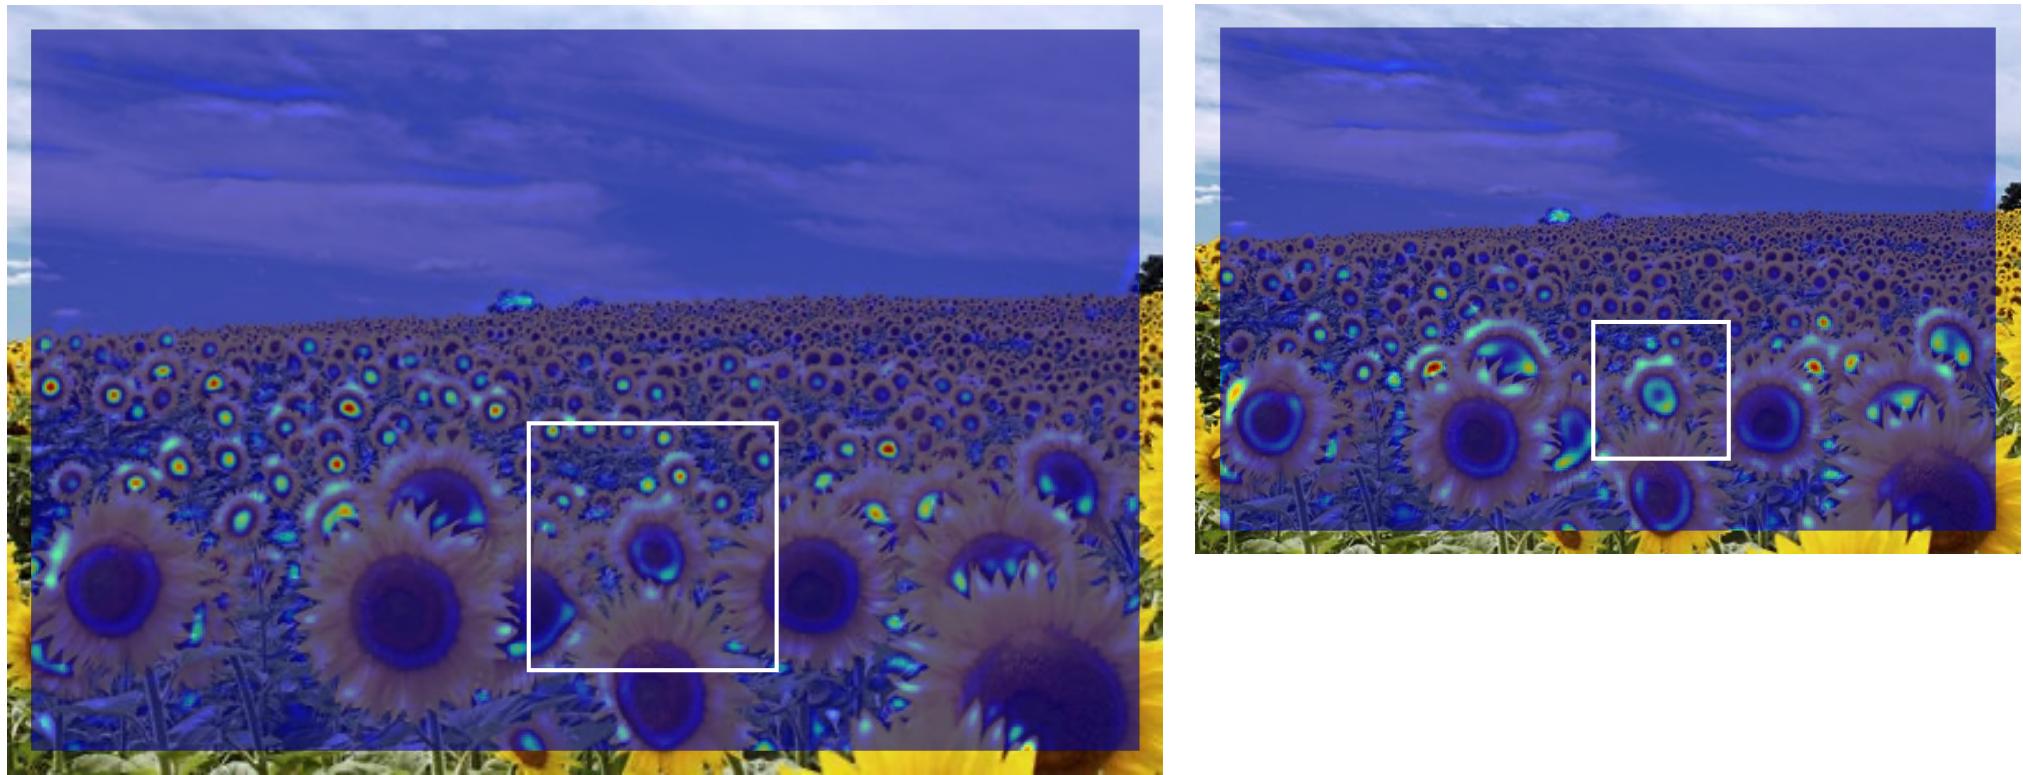
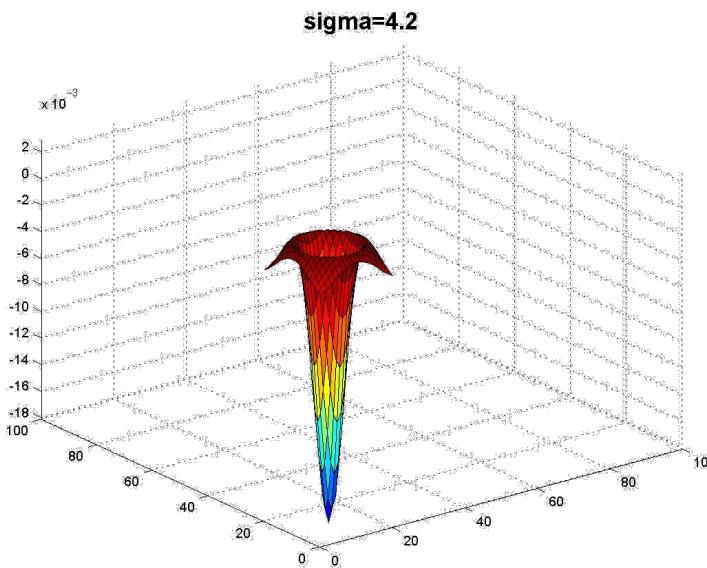


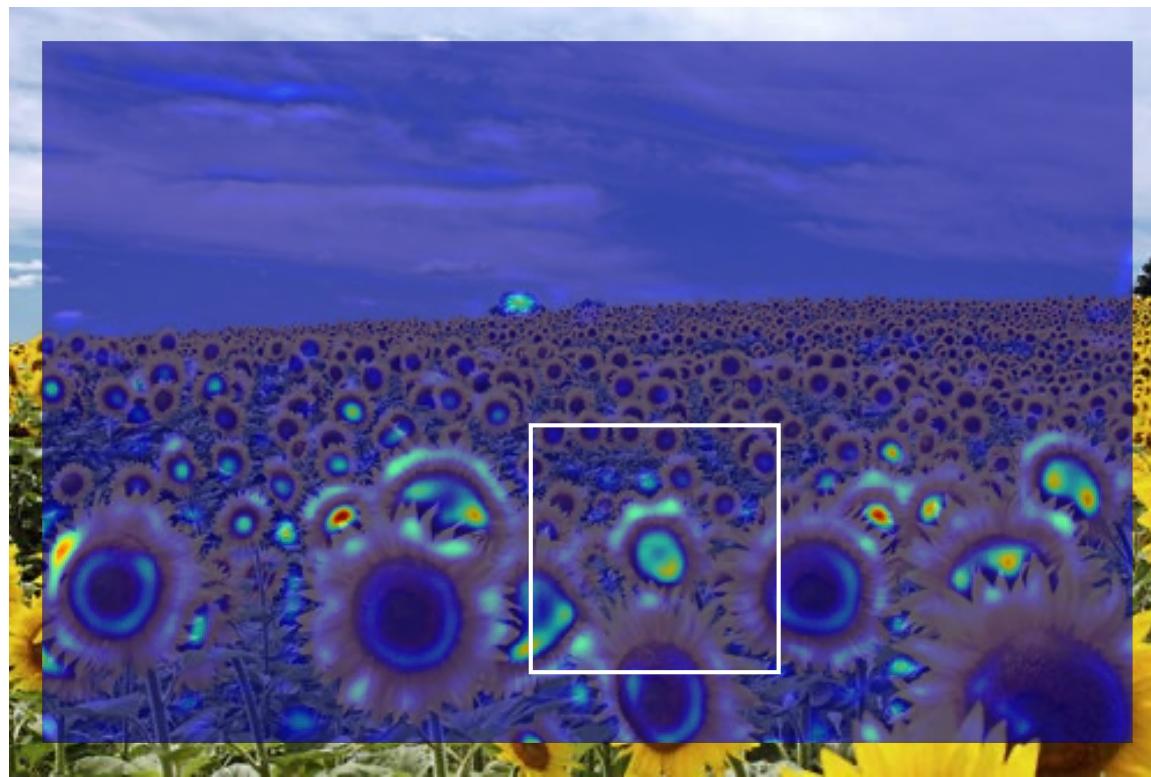
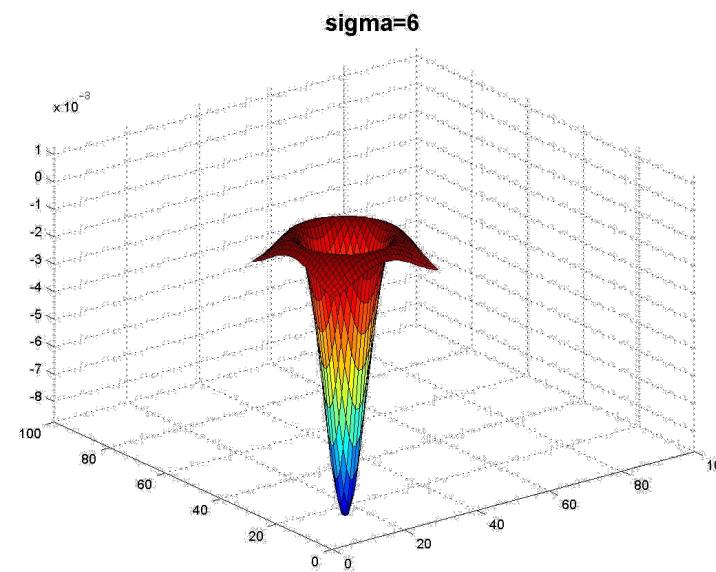
3/4 size

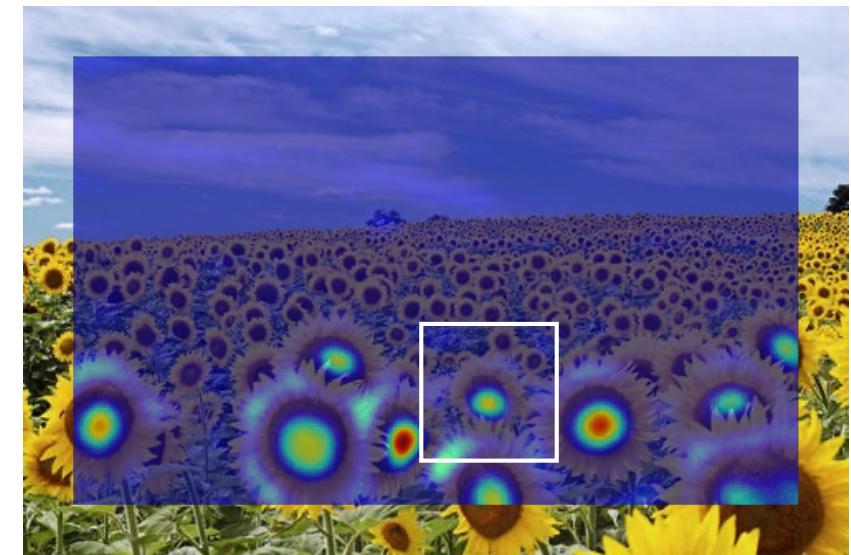
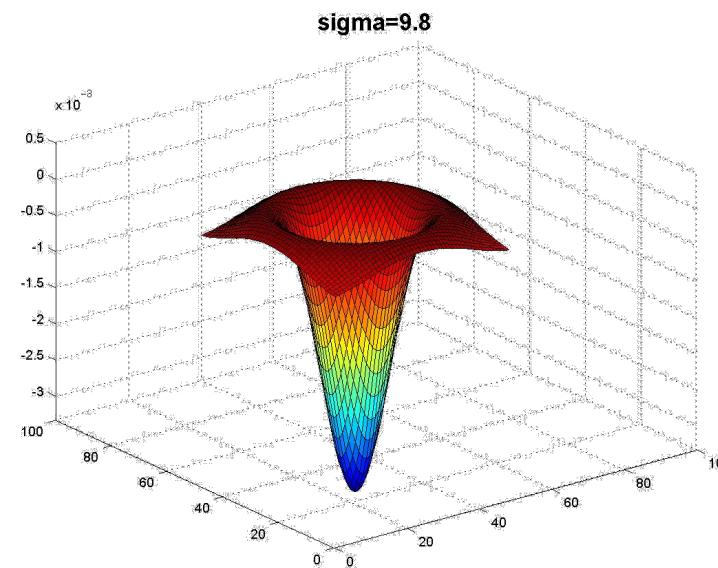


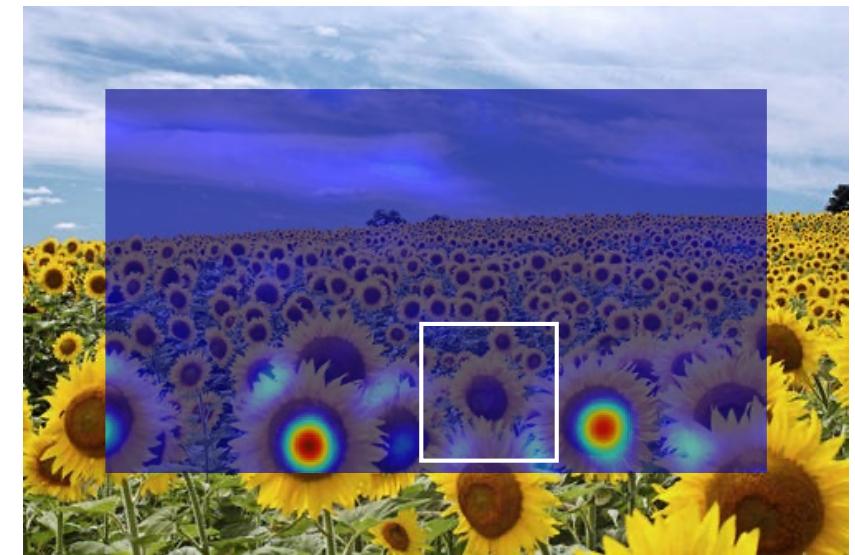
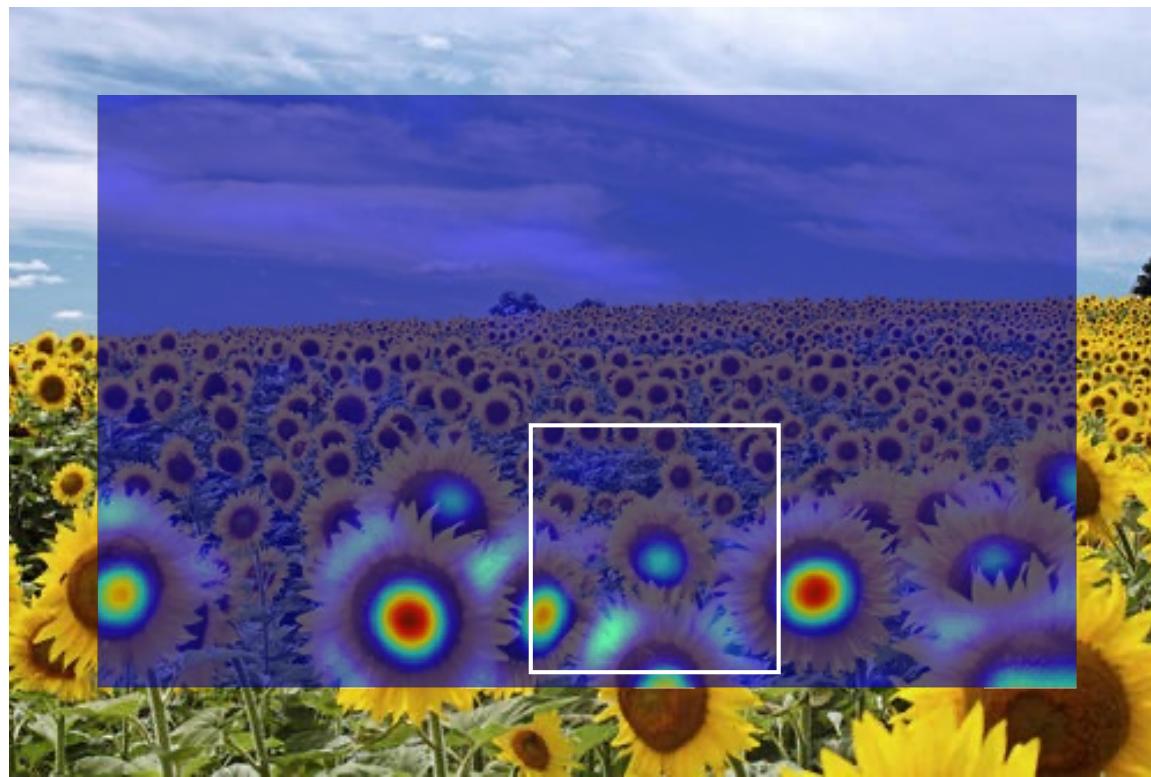
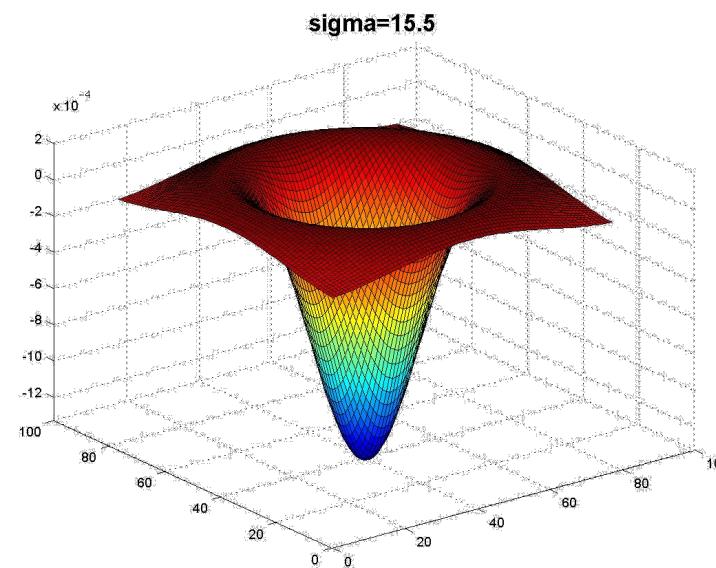
**sigma=2.1**



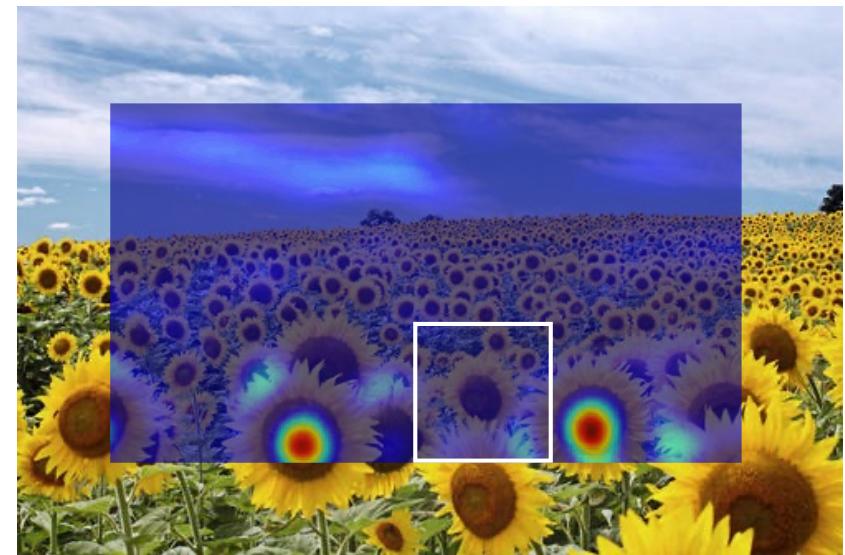
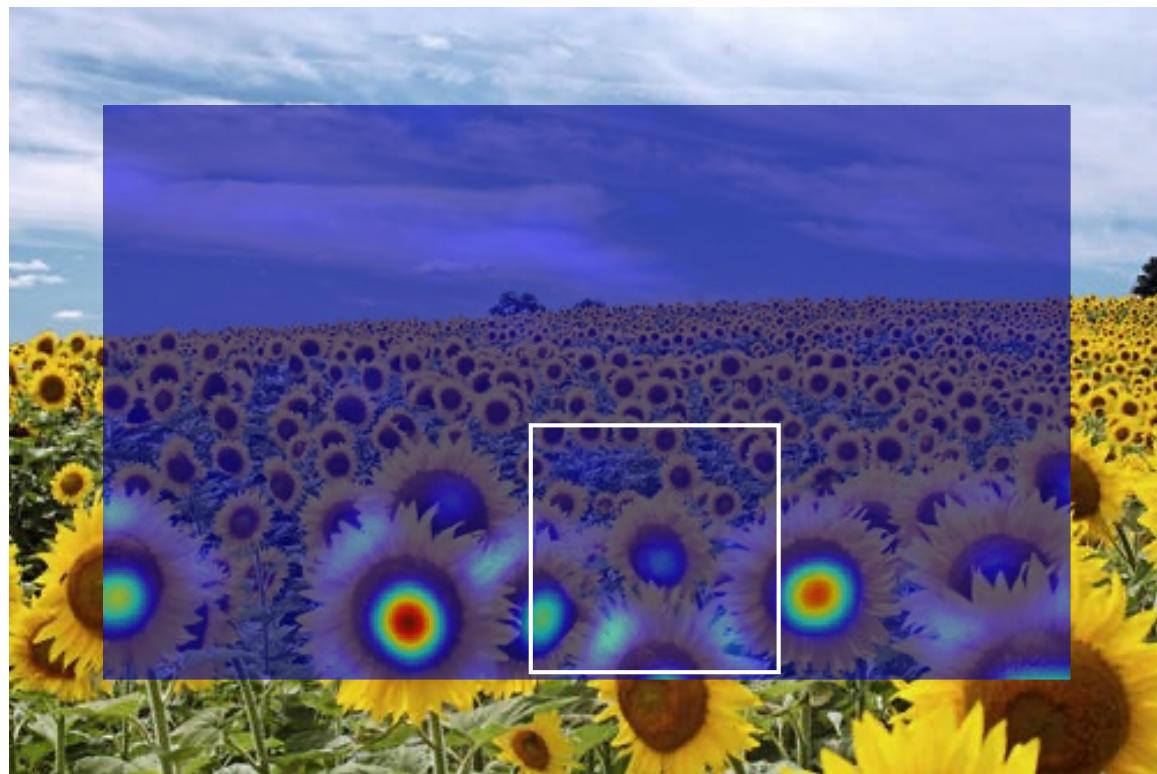
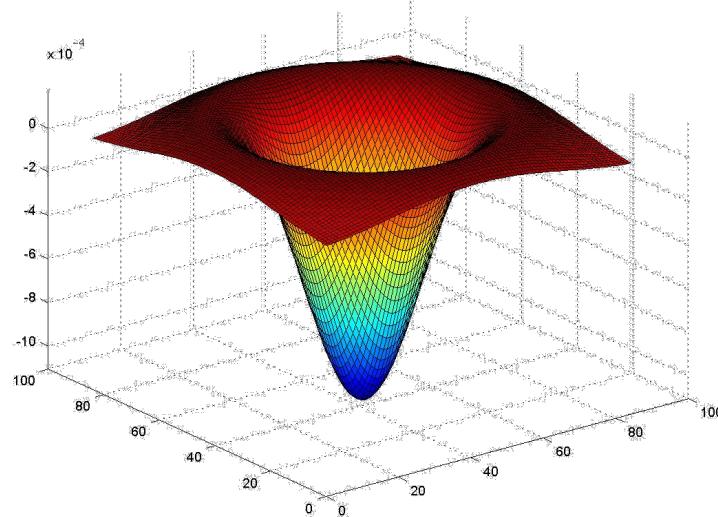








**sigma=17**

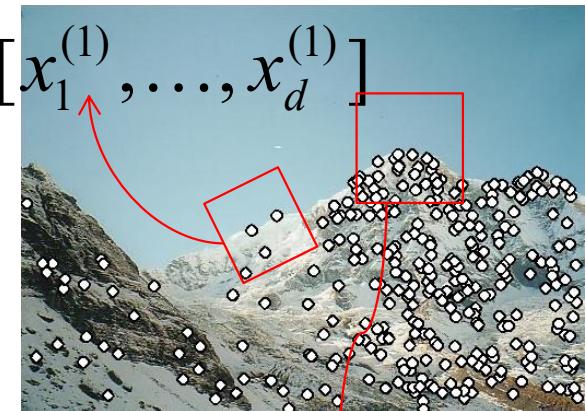


# Local features: main components

1) Detection: Identify the interest points

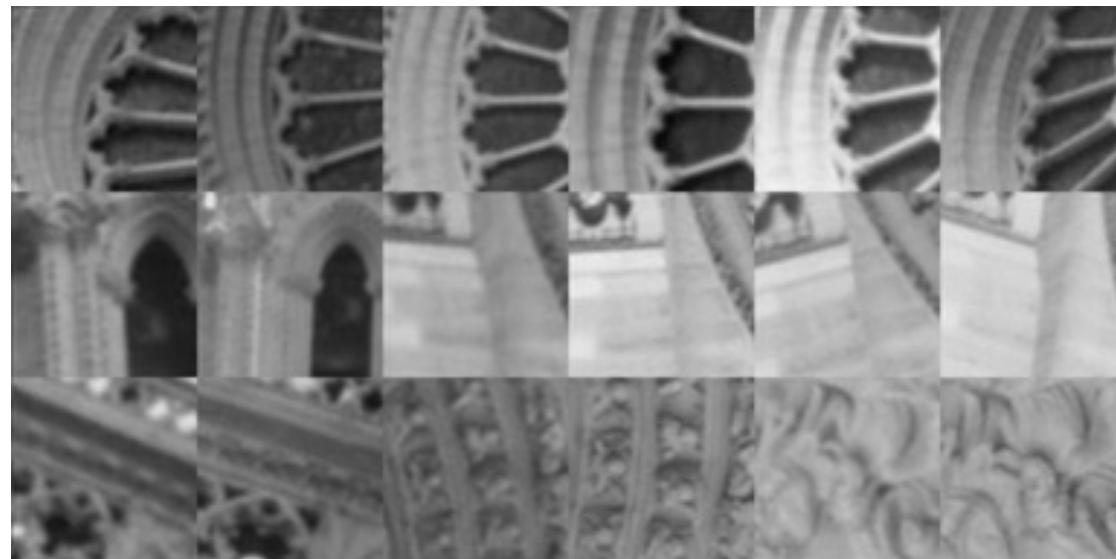
2) **Description**: Extract vector feature descriptor surrounding each interest point.

3) Matching: Determine correspondence between descriptors in two views

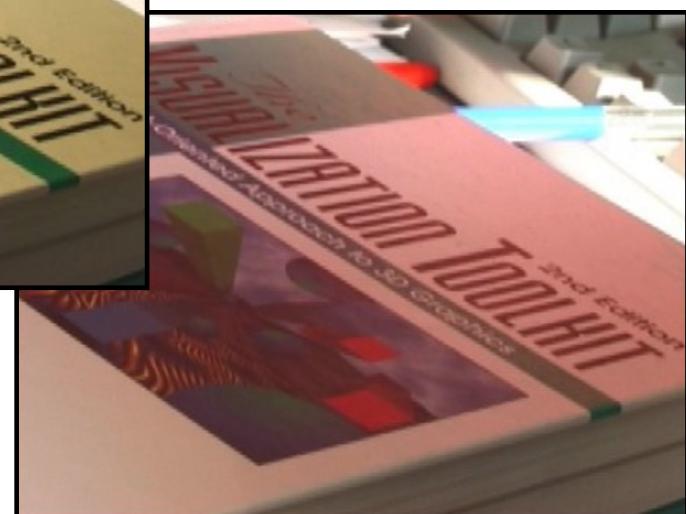
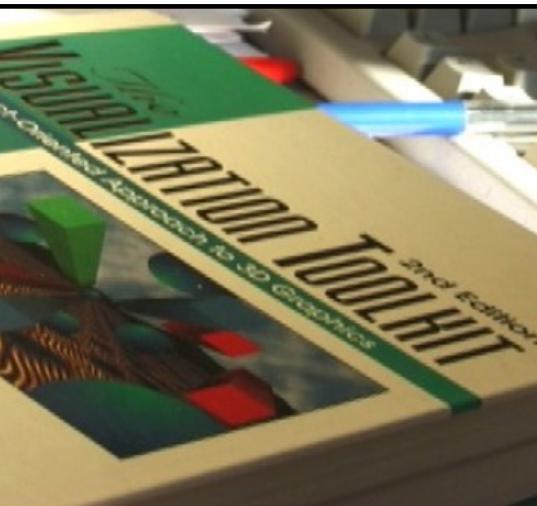
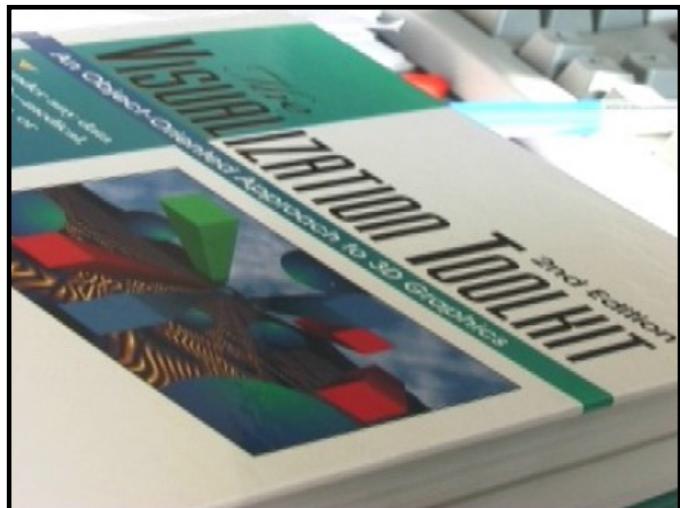




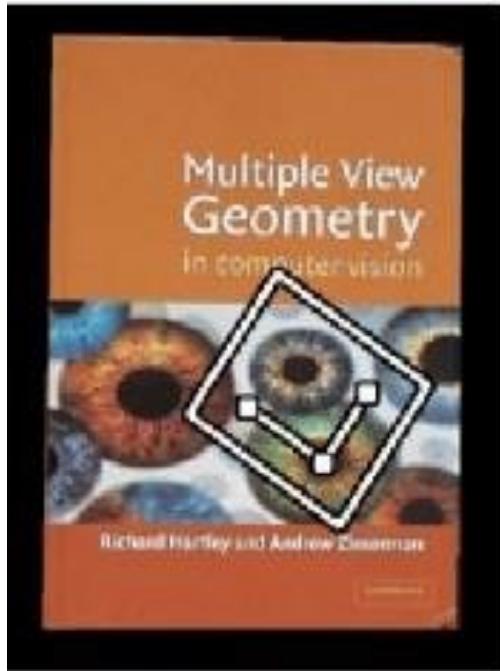
*What is the best descriptor for an image feature?*



# Photometric transformations



# Geometric transformations



objects will appear at different scales,  
translation and rotation

# Image patch

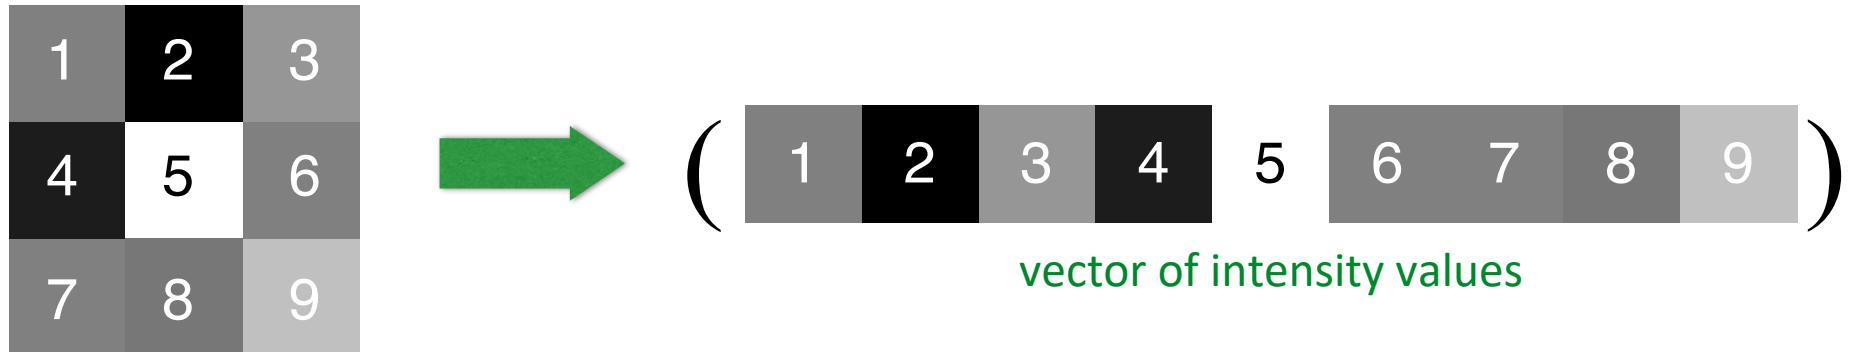
Just use the pixel values of the patch!



Perfectly fine if geometry and appearance is unchanged  
(a.k.a. template matching)

# Image patch

Just use the pixel values of the patch!

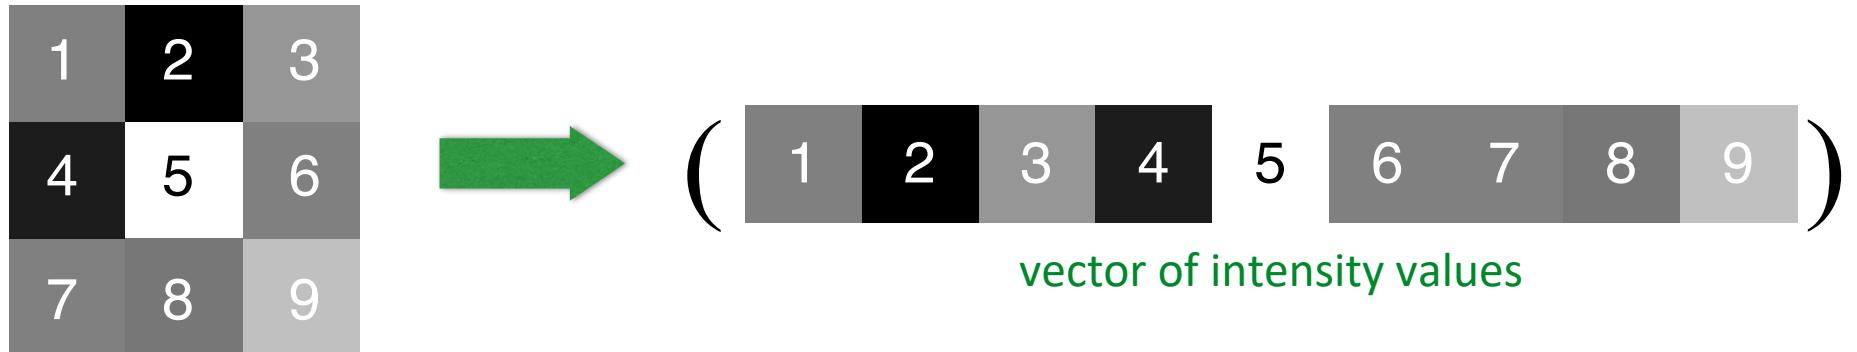


Perfectly fine if geometry and appearance is unchanged  
(a.k.a. template matching)

*What are the problems?*

# Image patch

Just use the pixel values of the patch!

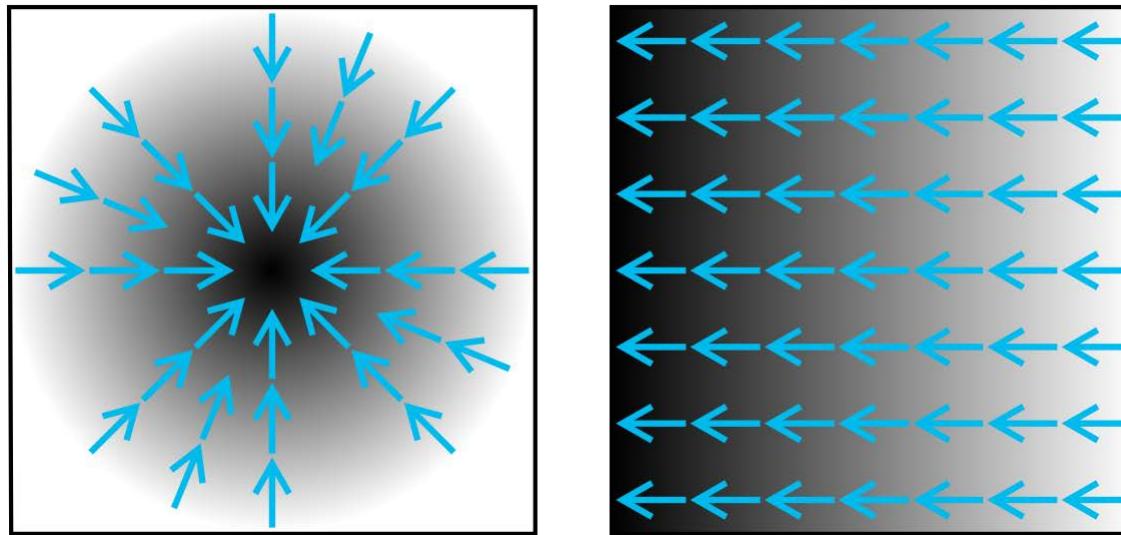


Perfectly fine if geometry and appearance is unchanged  
(a.k.a. template matching)

*What are the problems?*

*How can you be less sensitive to absolute intensity values?*

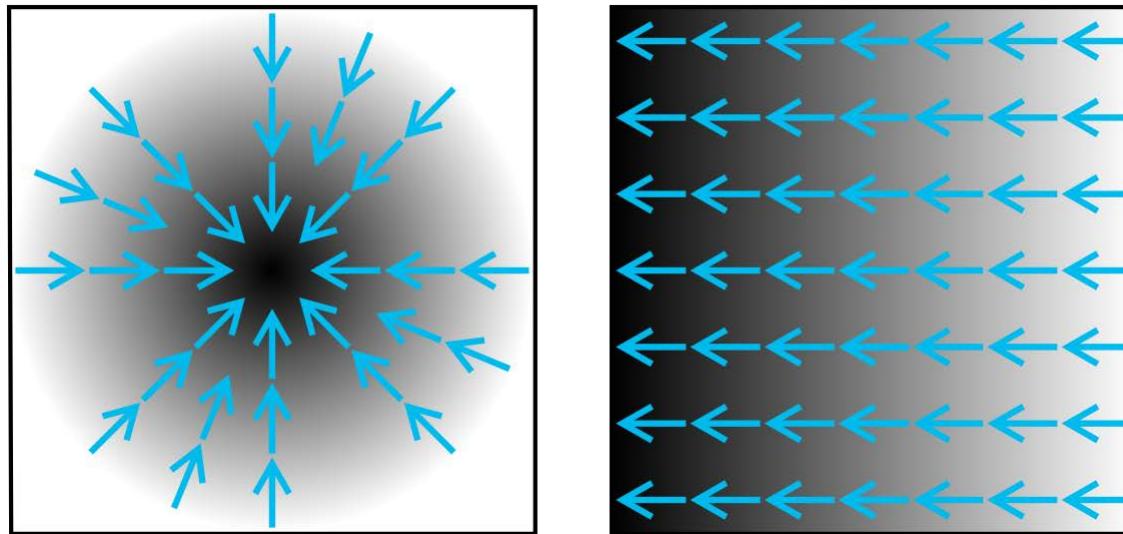
# Image gradients



Feature is invariant to absolute intensity values

*What are the problems?*

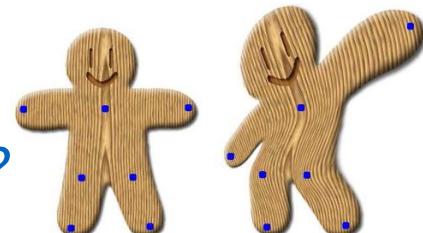
# Image gradients



Feature is invariant to absolute intensity values

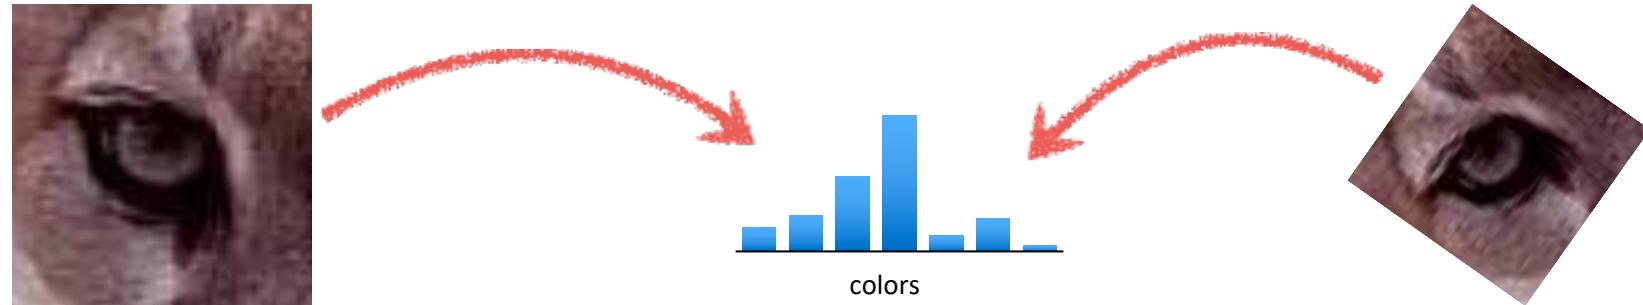
*What are the problems?*

*How can you be less sensitive to deformations?*



# Color histogram

Count the colors in the image using a histogram



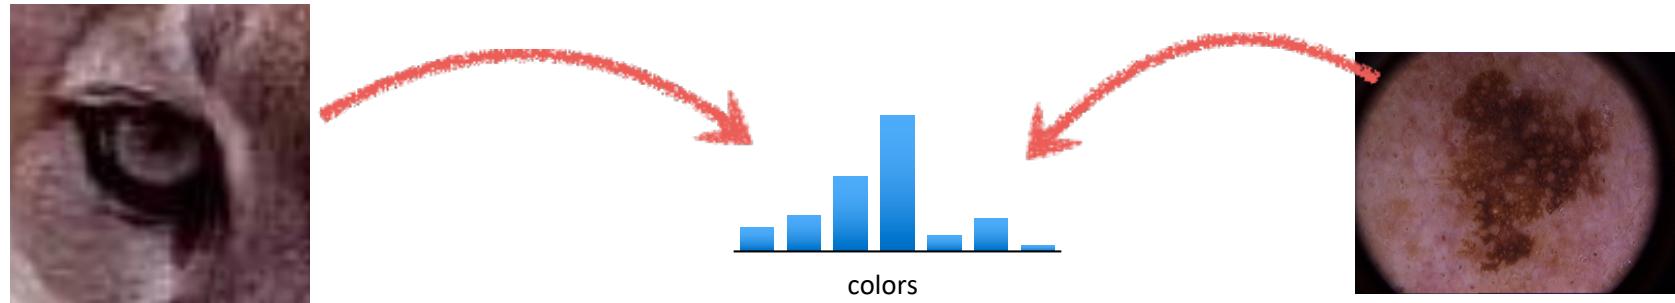
Invariant to changes in scale and rotation

*What are the problems?*

RGB  
HSV

# Color histogram

Count the colors in the image using a histogram



Invariant to changes in scale and rotation

*What are the problems?*

# Color histogram

Count the colors in the image using a histogram



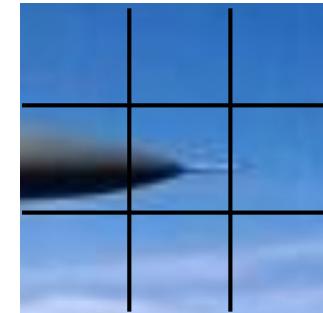
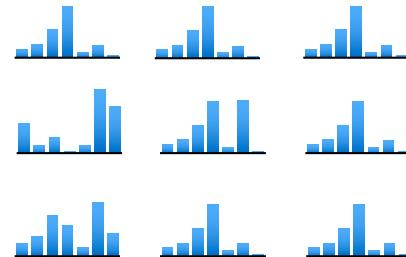
Invariant to changes in scale and rotation

*What are the problems?*

*How can you be more sensitive to spatial layout?*

# Spatial histograms

Compute histograms over spatial ‘cells’

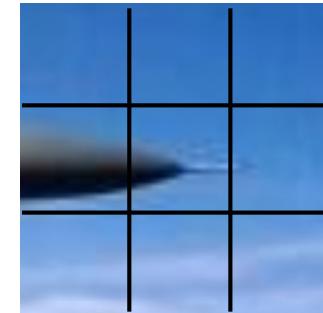
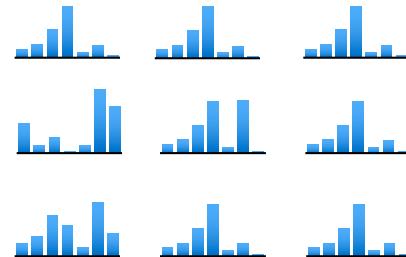


Retains rough spatial layout  
Some invariance to deformations

*What are the problems?*

# Spatial histograms

Compute histograms over spatial ‘cells’



Retains rough spatial layout

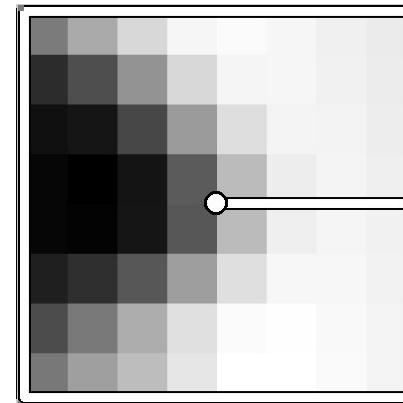
Some invariance to deformations

*What are the problems?*

*How can you be completely invariant to rotation?*

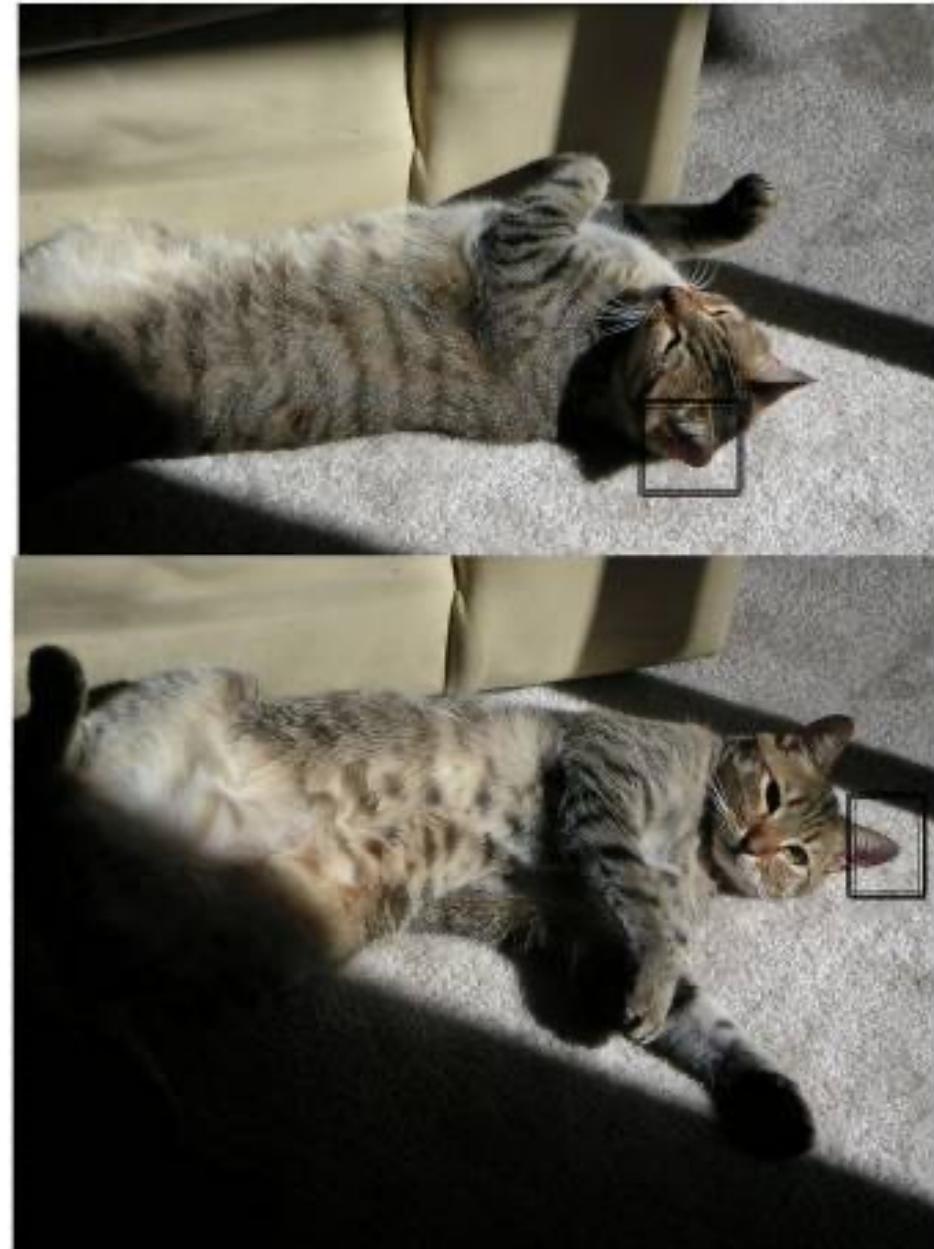
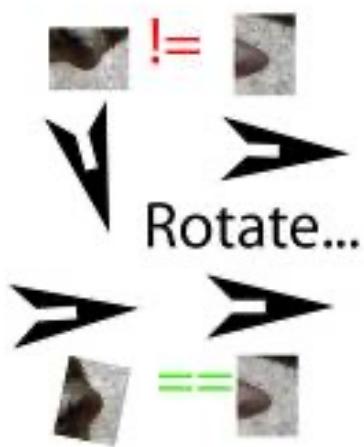
# Orientation normalization

Use the dominant image gradient direction to normalize the orientation of the patch



save the orientation angle  $\theta$  along with  $(x, y, s)$

# Rotation Invariance



# SIFT

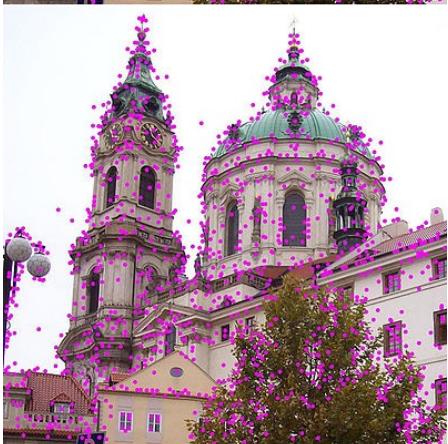
## (Scale Invariant Feature Transform)

Distinctive image features from scale-invariant keypoints

[DG Lowe - International journal of computer vision, 2004 - Springer](#)

This paper presents a method for extracting distinctive invariant features from images that can be used to perform reliable matching between different views of an object or scene. The features are invariant to image scale and rotation, and are shown to provide robust matching across a substantial range of affine distortion, change in 3D viewpoint, addition of noise, and change in illumination. The features are highly distinctive, in the sense that a single feature can be correctly matched with high probability against a large database of features from ...

[☆ Save](#) [✉ Cite](#) [Cited by 67710](#) [Related articles](#) [All 164 versions](#)



# SIFT

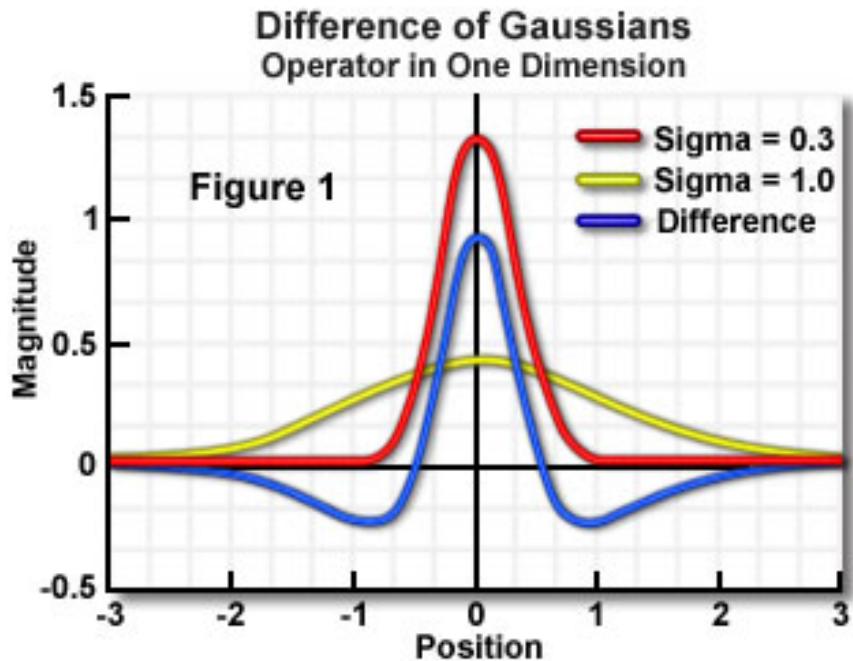
(Scale Invariant Feature Transform)

SIFT describes both a **detector** and **descriptor**

1. Multi-scale extrema detection
2. Keypoint localization
3. Orientation assignment
4. Keypoint descriptor

$$G(x, y) = \frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}}$$

# Difference of Gaussians



Laplacian  
Pyramid



$$DoG = G(x, y, k\sigma) - G(x, y, \sigma)$$

$I(k\sigma)$



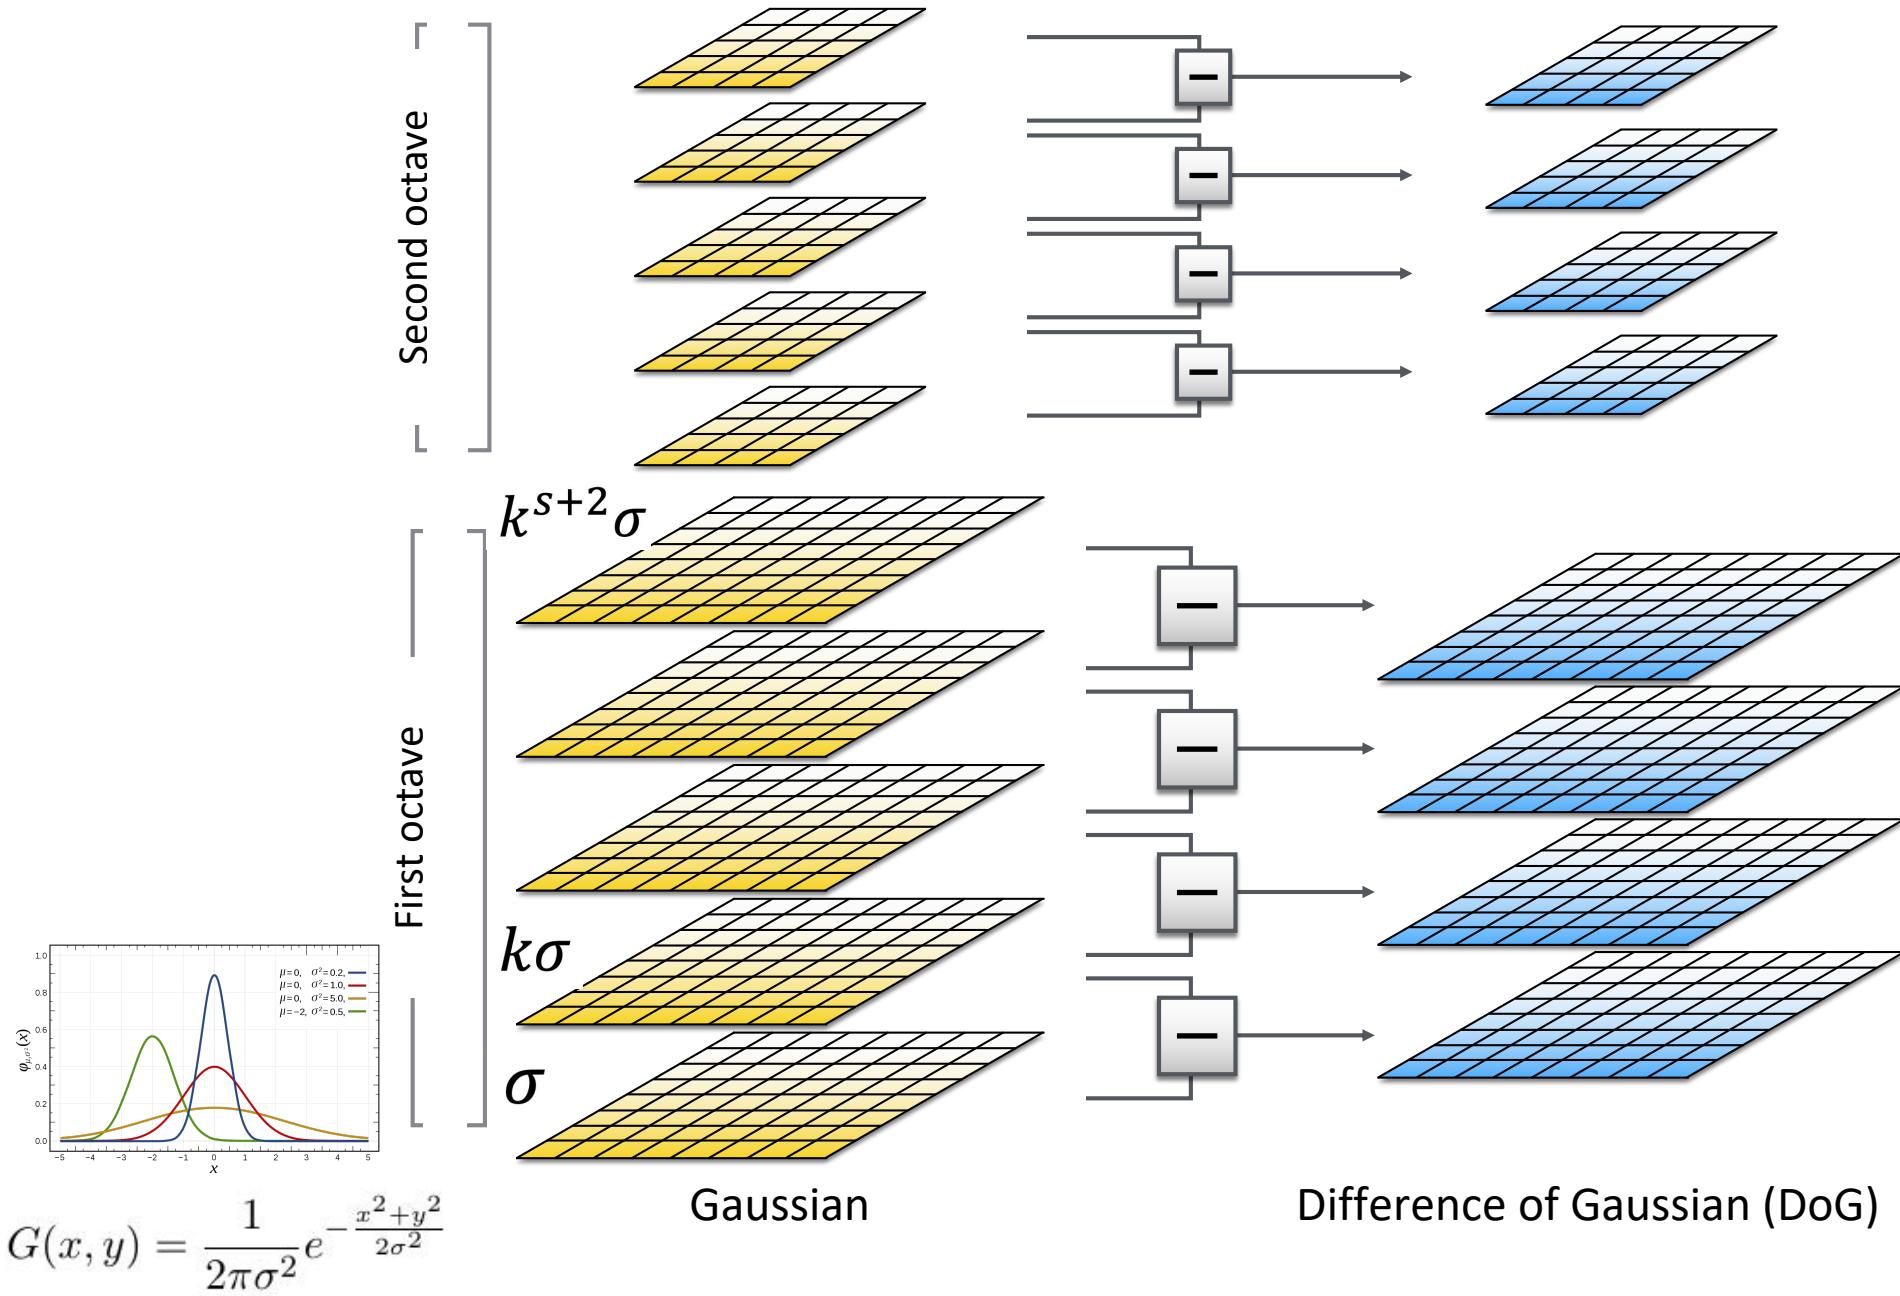
$I(\sigma)$



$I(k\sigma) - I(\sigma)$



# 1. Multi-scale extrema detection



# Difference of Gaussians

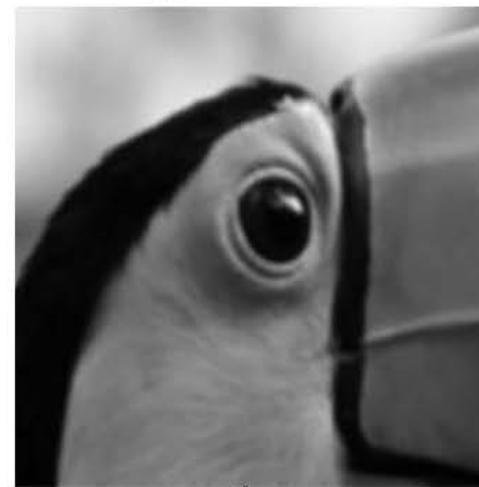
Sigma 0.7



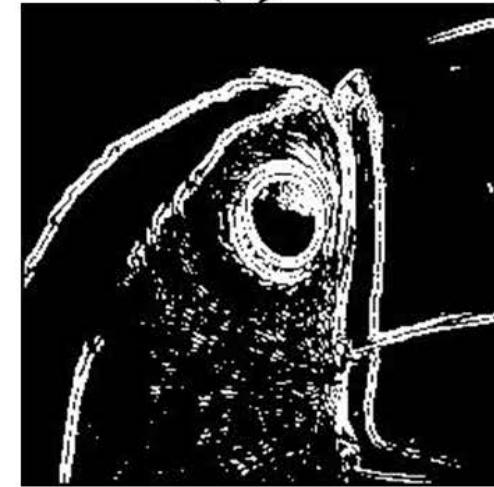
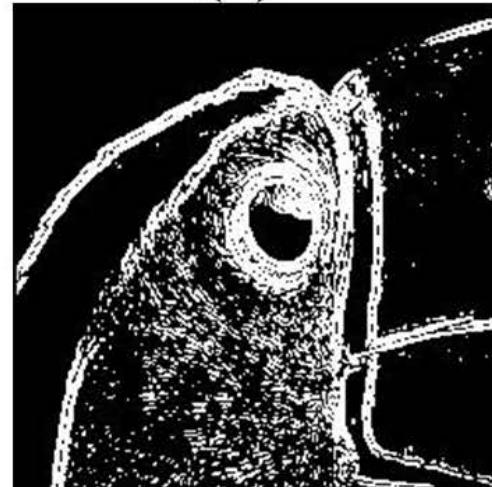
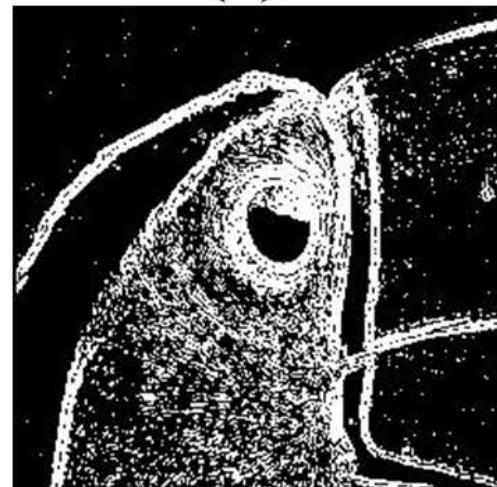
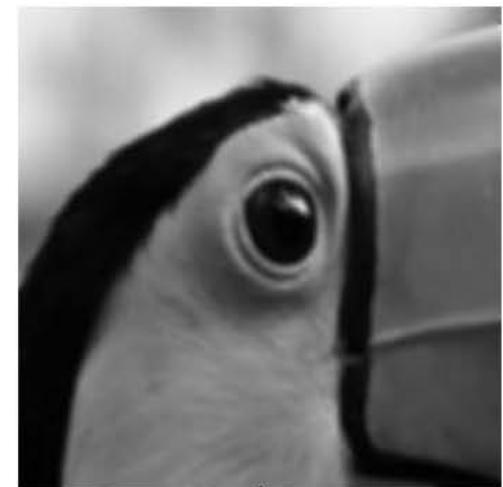
Sigma 1



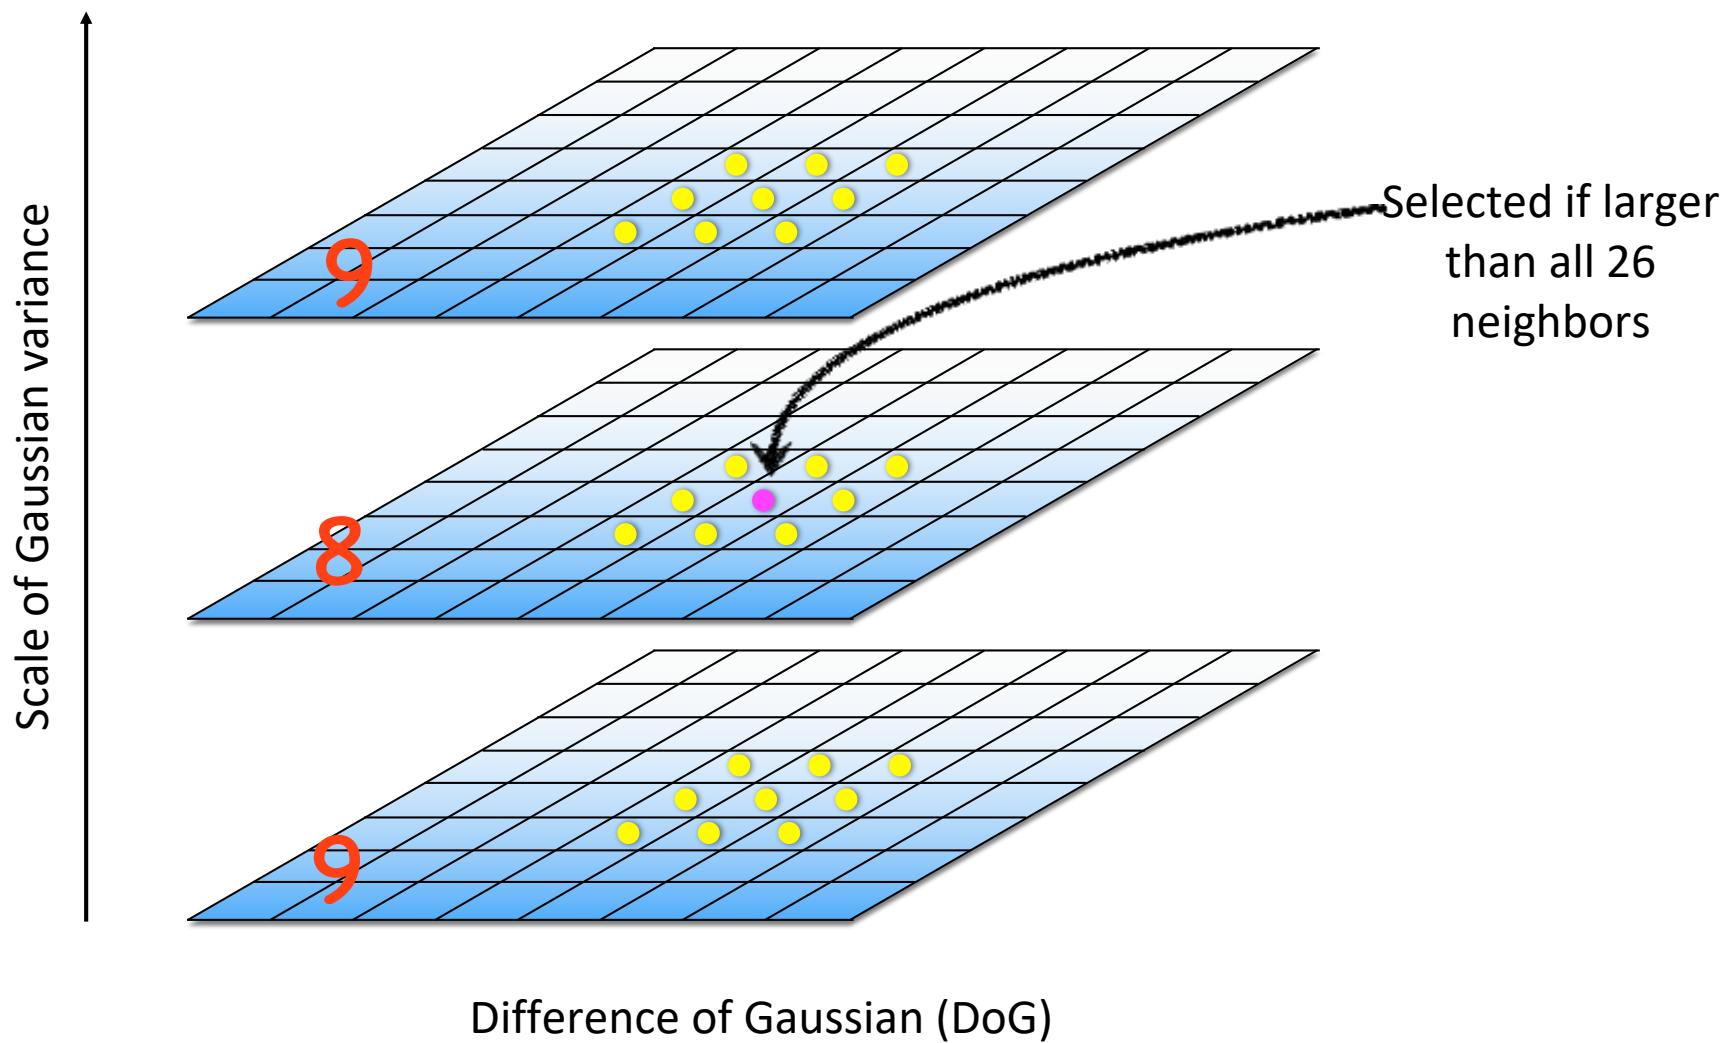
Sigma 1.4



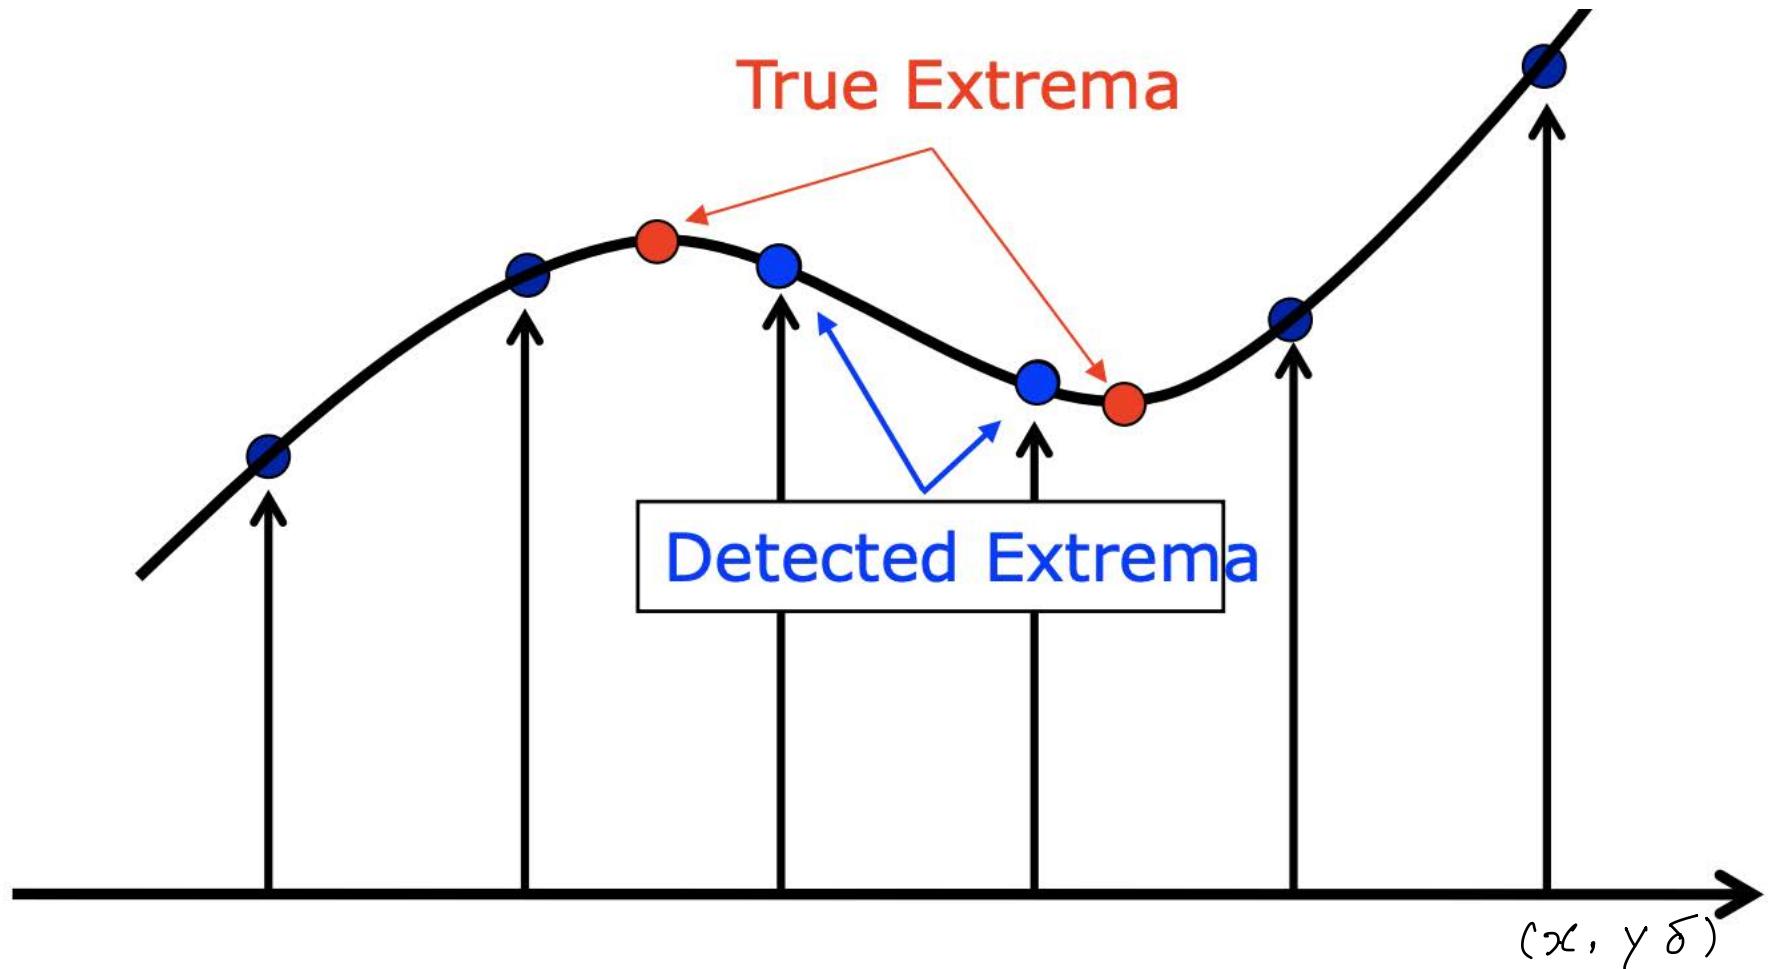
Sigma 2



# Scale-space extrema



# 2. Keypoint localization



Once a keypoint candidate has been found by comparing a pixel to its neighbors, the next step is to perform a detailed fit to the nearby data for location, scale, and ratio of principal curvatures.

# 2. Keypoint localization

2nd order Taylor series approximation of DoG scale-space

$$f(\mathbf{x}) = f + \frac{\partial f}{\partial \mathbf{x}}^T \mathbf{x} + \frac{1}{2} \mathbf{x}^T \frac{\partial^2 f}{\partial \mathbf{x}^2} \mathbf{x}$$

$$\mathbf{x} = \{x, y, \sigma\}$$

Take the derivative and solve for extrema

$$\mathbf{x}_m = - \frac{\partial^2 f}{\partial \mathbf{x}^2}^{-1} \frac{\partial f}{\partial \mathbf{x}}$$

Additional tests to retain only strong features

# 3. Orientation assignment

For a keypoint,  $L$  is the **Gaussian-smoothed** image with the closest scale,

$$m(x, y) = \sqrt{(L(x + 1, y) - L(x - 1, y))^2 + (L(x, y + 1) - L(x, y - 1))^2}$$

x-derivative    y-derivative

$$\theta(x, y) = \tan^{-1}((L(x, y + 1) - L(x, y - 1)) / (L(x + 1, y) - L(x - 1, y)))$$

Detection process returns

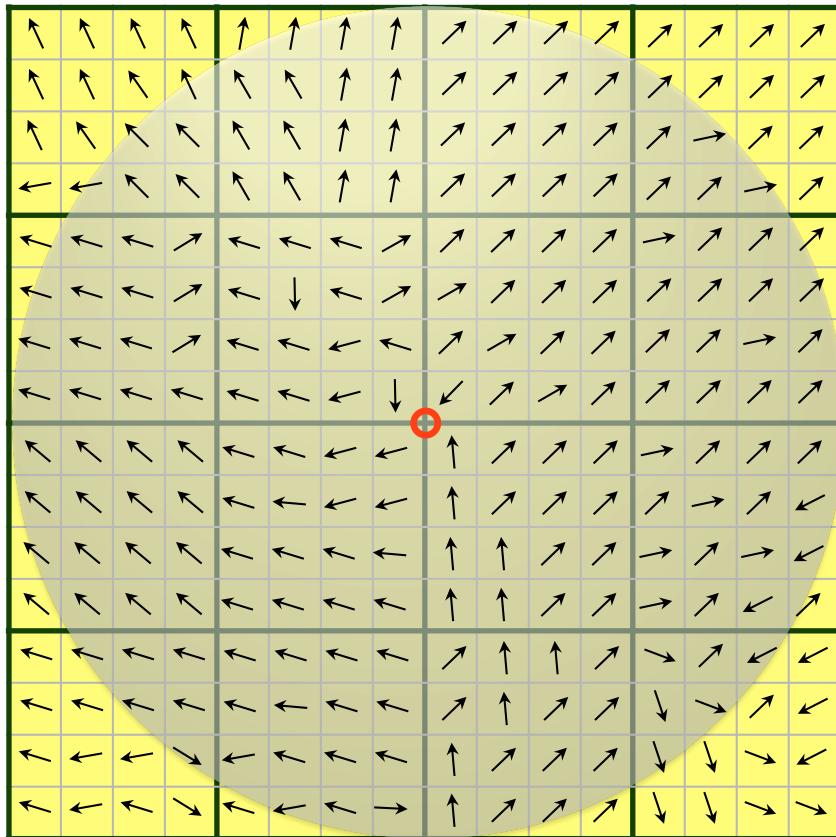
$$\{x, y, \sigma, \theta\}$$

location    scale    orientation

# 4. Keypoint descriptor

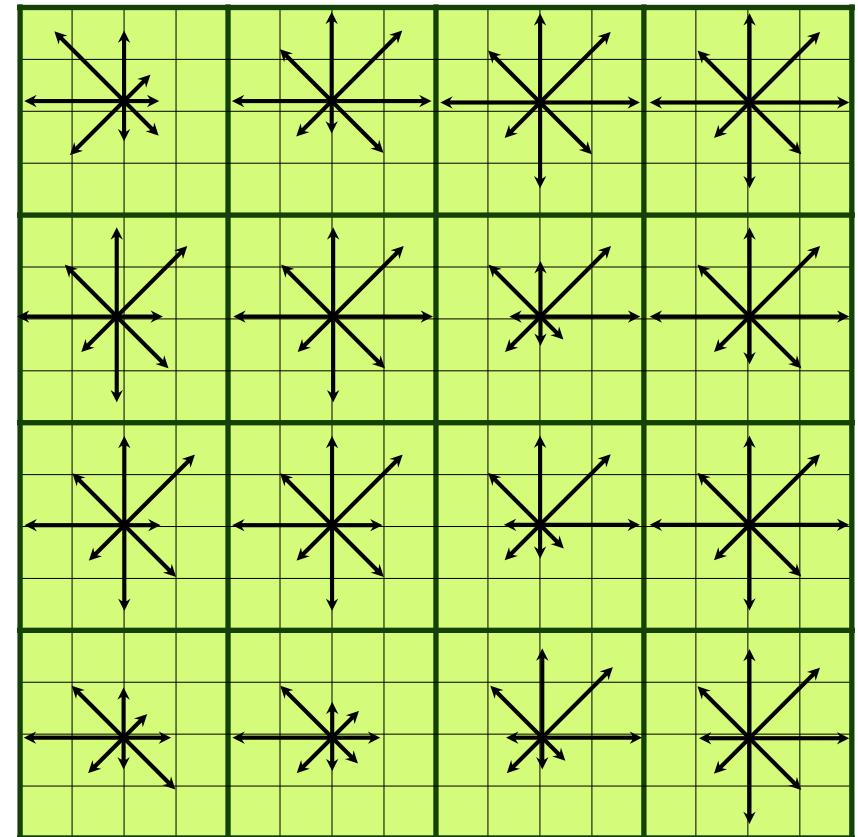
Image Gradients

(4 x 4 pixel per cell, 4 x 4 cells)



SIFT descriptor

(16 cells x 8 directions = 128 dims)

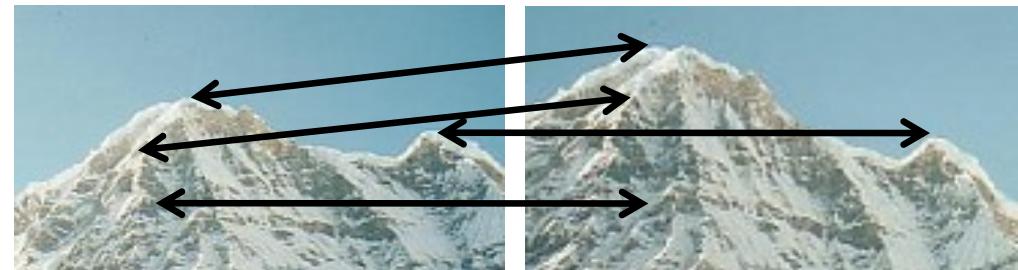


# SIFT properties

- Invariant to
  - Scale
  - Rotation
- Partially invariant to
  - Illumination changes
  - Camera viewpoint
  - Occlusion, clutter

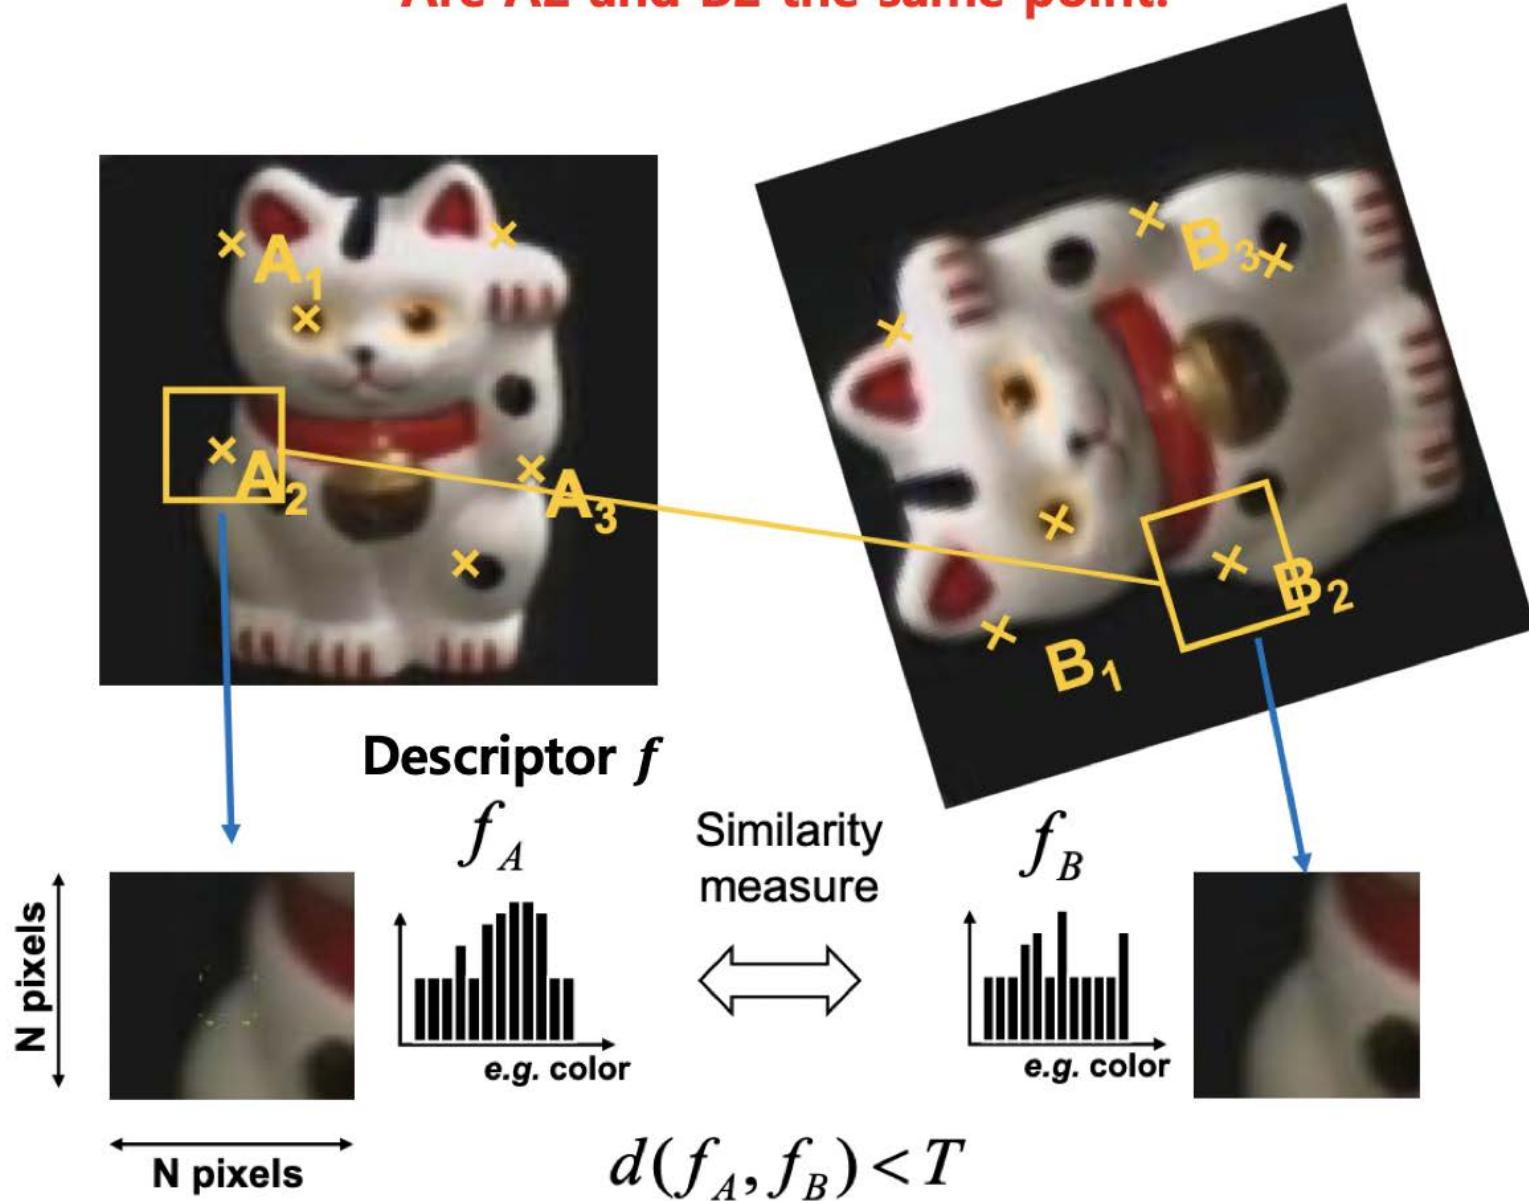
# Local features: main components

- 1) Detection: Identify the interest points
- 2) Description: Extract vector feature descriptor surrounding each interest point.
- 3) Matching: Determine correspondence between descriptors in two views

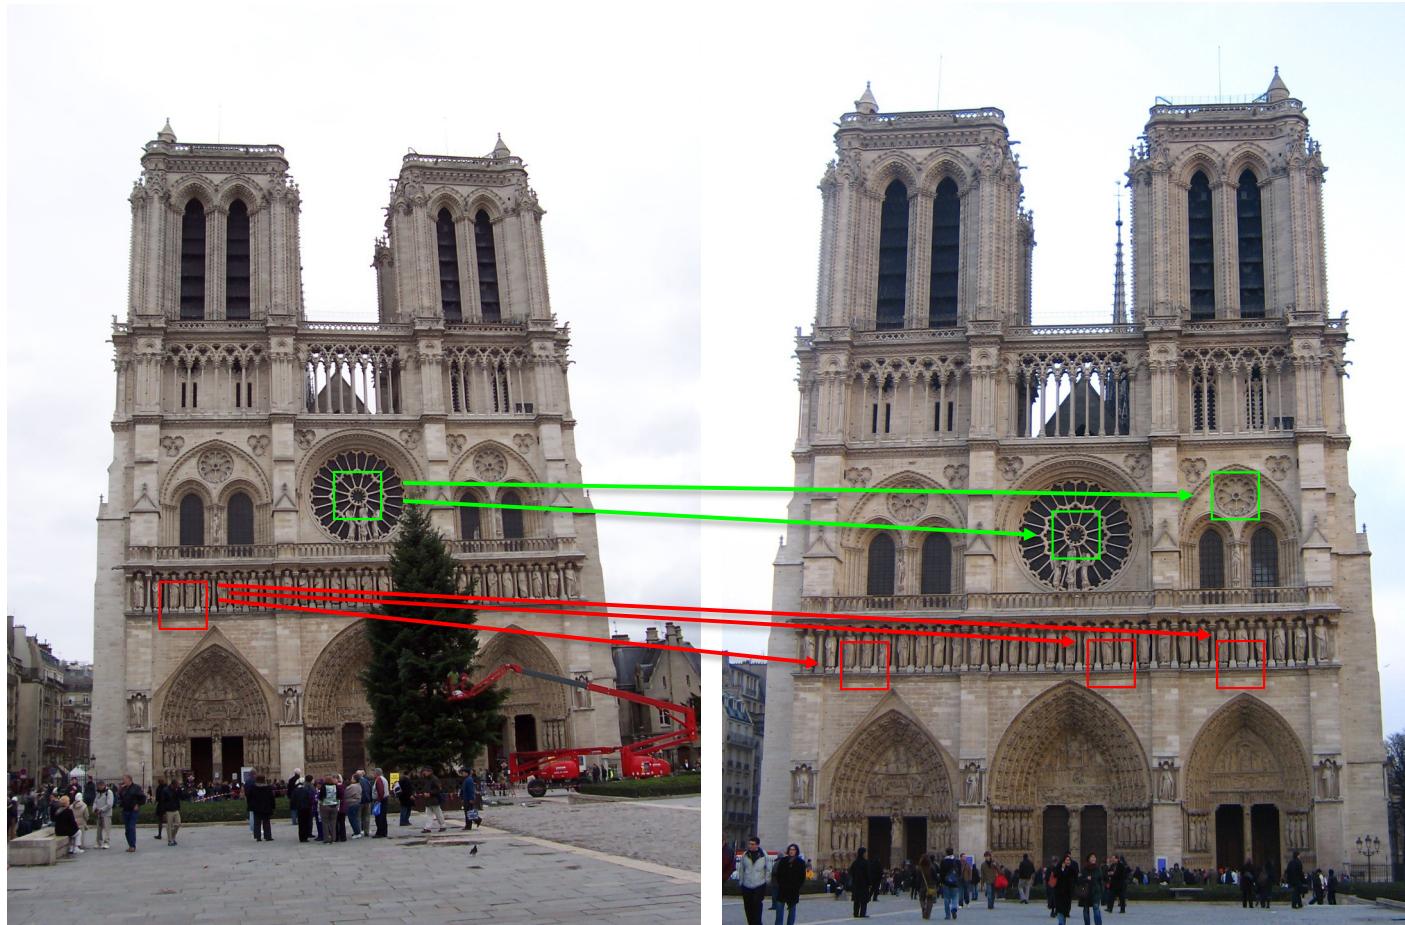


# Matching local features

Are A2 and B2 the same point?



# How do we decide which features match?



Distance: 0.34, 0.30, 0.40 Distance: 0.61, 1.22

# Euclidean distance vs. Cosine Similarity

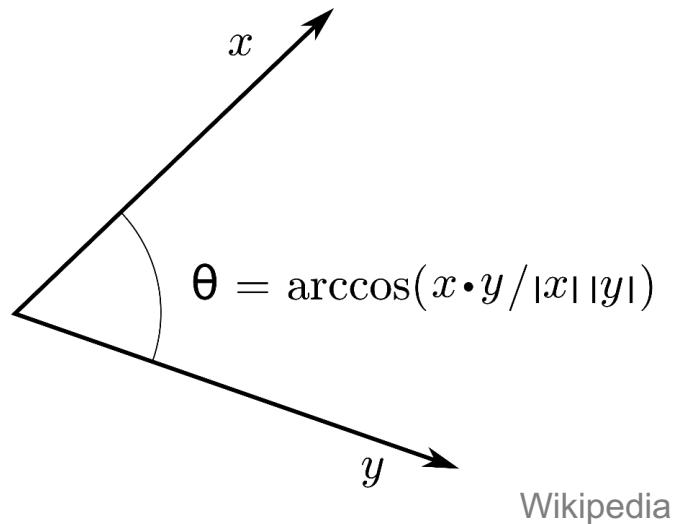
- Euclidean  $d(\mathbf{p}, \mathbf{q}) = d(\mathbf{q}, \mathbf{p}) = \sqrt{(q_1 - p_1)^2 + (q_2 - p_2)^2 + \cdots + (q_n - p_n)^2}$

$$= \sqrt{\sum_{i=1}^n (q_i - p_i)^2}.$$

$$\|\mathbf{q} - \mathbf{p}\| = \sqrt{(\mathbf{q} - \mathbf{p}) \cdot (\mathbf{q} - \mathbf{p})}.$$

- Cosine similarity:  $\mathbf{a} \cdot \mathbf{b} = \|\mathbf{a}\|_2 \|\mathbf{b}\|_2 \cos \theta$

$$\text{similarity} = \cos(\theta) = \frac{\mathbf{A} \cdot \mathbf{B}}{\|\mathbf{A}\|_2 \|\mathbf{B}\|_2}$$



Wikipedia

# Feature Matching

- Criteria 1:
  - Compute distance in feature space, e.g., Euclidean distance between 128-dim SIFT descriptors
  - Match point to lowest distance (nearest neighbor)
- Problems:
  - Does everything have a match?

# Feature Matching

- Criteria 2:
  - Compute distance in feature space, e.g., Euclidean distance between 128-dim SIFT descriptors
  - Match point to lowest distance (nearest neighbor)
  - Ignore anything higher than threshold (no match!)
- Problems:
  - Threshold is hard to pick
  - Non-distinctive features could have lots of close matches, only one of which is correct

# Evaluation

- Suppose our matcher computes 1,000 matches between two images
  - 800 are correct matches, 200 are incorrect (according to an oracle that gives us ground truth matches)
  - A given threshold (e.g., ratio distance = 0.6) gives us 600 correct matches and 100 incorrect matches that survive the threshold
  - True positive rate =  $600 / 800 = \frac{3}{4}$
  - False positive rate =  $100 / 200 = \frac{1}{2}$

# Evaluation

Precision: How many retrieved items are relevant?

Recall: How many relevant items are retrieved?

$$\text{Precision} = \frac{tp}{tp + fp}$$

$$\text{Recall} = \frac{tp}{tp + fn}$$

		Actual	
		Positive	Negative
Prediction	Positive	TP	FP
	Negative	FN	TN

